


Heap


2019 Spring





Lecture Outline

- Heap-allocated Memory
 - `malloc()` and `free()`
 - Memory leaks
 - Sample codes




Memory Allocation So Far (1/2)

- So far, we have seen two kinds of memory allocation:

```
int counter = 0;      // global var

int main(int argc, char** argv) {
    counter++;
    printf("count = %d\n", counter);
    return 0;
}
```

- counter is *statically*-allocated
 - Allocated when program is loaded
 - Deallocated when program exits




Memory Allocation So Far (2/2)

- So far, we have seen two kinds of memory allocation:

```
int foo(int a) {
    int x = a + 1;      // local var
    return x;
}

int main(int argc, char** argv) {
    int y = foo(10);   // local var
    printf("y = %d\n", y);
    return 0;
}
```

- a , x , y are *automatically*-allocated
 - Allocated when function is called
 - Deallocated when function returns




Dynamic Allocation

- Situations where static and automatic allocation aren't sufficient:
 - We need memory that persists across multiple function calls but not the whole lifetime of the program
 - We need more memory than can fit on the Stack
 - We need memory whose size is not known in advance to the caller


```
// this is pseudo-C code
char* ReadFile(char* filename) {
    int size = GetFileSize(filename);
    char* buffer = AllocateMem(size);

    ReadFileIntoBuffer(filename, buffer);
    return buffer;
}
```



Dynamic Allocation

- What we want is ***dynamically***-allocated memory
 - Your program explicitly requests a new block of memory
 - The language allocates it at runtime, perhaps with help from OS
 - Dynamically-allocated memory persists until either:
 - Your code explicitly deallocated it (*manual memory management*)
 - A garbage collector collects it (*automatic memory management*)
- C requires you to manually manage memory
 - Gives you more control, but causes headaches




Aside: NULL

- NULL is a memory location that is **guaranteed to be invalid**
 - In C on Linux, NULL is 0x0 and an attempt to dereference NULL *causes a segmentation fault*
- Useful as an indicator of an uninitialized (or currently unused) pointer or allocation error
 - It's better to cause a segfault than to allow the corruption of memory!

segfault.c

```
int main(int argc, char** argv) {
    int* p = NULL;
    *p = 1; // causes a segmentation fault
    return 0;
}
```



malloc()

- General usage:

```
var = (type*) malloc (size in bytes)
```

- malloc allocates a block of memory of the requested size
 - Returns a pointer to the first byte of that memory
 - And returns NULL if the memory allocation failed!
 - You should assume that the memory initially contains garbage
 - You'll typically use sizeof to calculate the size you need

```
// allocate a 10-float array
float* arr = (float*) malloc(10*sizeof(float));
if (arr == NULL) {
    return errcode;
}
...    // do stuff with arr
```



calloc()

- General usage:

```
var = (type*) calloc(num, bytes per element)
```

- Like **malloc**, but also zeros out the block of memory
 - Helpful for shaking out bugs
 - Slightly slower; preferred for non-performance-critical code
 - malloc** and **calloc** are found in **stdlib.h**

```
// allocate a 10-double array
double* arr = (double*) calloc(10, sizeof(double));
if (arr == NULL) {
    return errcode;
}
...    // do stuff with arr
```

free()




- Usage: `free(pointer);`
- Deallocates the memory pointed-to by the pointer
 - Pointer *must* point to the first byte of heap-allocated memory (*i.e.* something previously returned by `malloc` or `calloc`)
 - Freed memory becomes eligible for future allocation
 - Pointer is unaffected by call to free
 - Defensive programming: can set pointer to `NULL` after freeing it

```
float* arr = (float*) malloc(10*sizeof(float));
if (arr == NULL)
    return errcode;
...
// do stuff with arr
free(arr);
arr = NULL;    // OPTIONAL
```

The Heap



- The Heap is a large pool of unused memory that is used for dynamically-allocated data
 - `malloc` allocates chunks of data in the Heap; `free` deallocates those chunks
 - `malloc` maintains bookkeeping data in the Heap to track allocated blocks



Heap and Stack Example (1/11)



arraycopy.c

```
#include <stdlib.h>


int* copy(int a[], int size) {
    int i, *a2;

    a2 = malloc(size*sizeof(int));
    if (a2 == NULL)
        return NULL;

    for (i = 0; i < size; i++)
        a2[i] = a[i];

    return a2;
}

int main(int argc, char** argv) {
    int nums[4] = {1, 2, 3, 4};
    int* ncopy = copy(nums, 4);
    // .. do stuff with the array ..
    free(ncopy);
    return 0;
}
```



Heap and Stack Example (2/11)



arraycopy.c

```
#include <stdlib.h>


int* copy(int a[], int size) {
    int i, *a2;

    a2 = malloc(size*sizeof(int));
    if (a2 == NULL)
        return NULL;

    for (i = 0; i < size; i++)
        a2[i] = a[i];

    return a2;
}

int main(int argc, char** argv) {
    int nums[4] = {1, 2, 3, 4};
    int* ncopy = copy(nums, 4);
    // .. do stuff with the array ..
    free(ncopy);
    return 0;
}
```



Heap and Stack Example (3/11)



arraycopy.c

```
#include <stdlib.h>


int* copy(int a[], int size) {
    int i, *a2;

    a2 = malloc(size*sizeof(int));
    if (a2 == NULL)
        return NULL;

    for (i = 0; i < size; i++)
        a2[i] = a[i];

    return a2;
}

int main(int argc, char** argv) {
    int nums[4] = {1, 2, 3, 4};
    int* ncopy = copy(nums, 4);
    // .. do stuff with the array ..
    free(ncopy);
    return 0;
}
```



Heap and Stack Example (4/11)



arraycopy.c

```
#include <stdlib.h>


int* copy(int a[], int size) {
    int i, *a2;

    a2 = malloc(size*sizeof(int));
    if (a2 == NULL)
        return NULL;

    for (i = 0; i < size; i++)
        a2[i] = a[i];

    return a2;
}

int main(int argc, char** argv) {
    int nums[4] = {1, 2, 3, 4};
    int* ncopy = copy(nums, 4);
    // .. do stuff with the array ..
    free(ncopy);
    return 0;
}
```



Heap and Stack Example (5/11)



arraycopy.c

```
#include <stdlib.h>


int* copy(int a[], int size) {
    int i, *a2;

    a2 = malloc(size*sizeof(int));
    if (a2 == NULL)
        return NULL;

    for (i = 0; i < size; i++)
        a2[i] = a[i];

    return a2;
}

int main(int argc, char** argv) {
    int nums[4] = {1, 2, 3, 4};
    int* ncopy = copy(nums, 4);
    // .. do stuff with the array ..
    free(ncopy);
    return 0;
}
```



Heap and Stack Example (6/11)



```
#include <stdlib.h>


int* copy(int a[], int size) {
    int i, *a2;

    a2 = malloc(size*sizeof(int));
    if (a2 == NULL)
        return NULL;

    for (i = 0; i < size; i++)
        a2[i] = a[i];

    return a2;
}

int main(int argc, char** argv) {
    int nums[4] = {1, 2, 3, 4};
    int* ncopy = copy(nums, 4);
    // .. do stuff with the array ..
    free(ncopy);
    return 0;
}
```



Heap and Stack Example (7/11)



arraycopy.c

```
#include <stdlib.h>


int* copy(int a[], int size) {
    int i, *a2;

    a2 = malloc(size*sizeof(int));
    if (a2 == NULL)
        return NULL;

    for (i = 0; i < size; i++)
        a2[i] = a[i];

    return a2;
}

int main(int argc, char** argv) {
    int nums[4] = {1, 2, 3, 4};
    int* ncopy = copy(nums, 4);
    // .. do stuff with the array ..
    free(ncopy);
    return 0;
}
```



Heap and Stack Example (8/11)



arraycopy.c

```
#include <stdlib.h>


int* copy(int a[], int size) {
    int i, *a2;

    a2 = malloc(size*sizeof(int));
    if (a2 == NULL)
        return NULL;

    for (i = 0; i < size; i++)
        a2[i] = a[i];

    return a2;
}

int main(int argc, char** argv) {
    int nums[4] = {1, 2, 3, 4};
    int* ncopy = copy(nums, 4);
    // .. do stuff with the array ..
    free(ncopy);
    return 0;
}
```



Heap and Stack Example (9/11)



arraycopy.c

```
#include <stdlib.h>


int* copy(int a[], int size) {
    int i, *a2;

    a2 = malloc(size*sizeof(int));
    if (a2 == NULL)
        return NULL;

    for (i = 0; i < size; i++)
        a2[i] = a[i];

    return a2;
}

int main(int argc, char** argv) {
    int nums[4] = {1, 2, 3, 4};
    int* ncopy = copy(nums, 4);
    // .. do stuff with the array ..
    free(ncopy);
    return 0;
}
```



Heap and Stack Example (10/11)



arraycopy.c

```
#include <stdlib.h>


int* copy(int a[], int size) {
    int i, *a2;

    a2 = malloc(size*sizeof(int));
    if (a2 == NULL)
        return NULL;

    for (i = 0; i < size; i++)
        a2[i] = a[i];

    return a2;
}

int main(int argc, char** argv) {
    int nums[4] = {1, 2, 3, 4};
    int* ncopy = copy(nums, 4);
    // .. do stuff with the array ..
    free(ncopy);
    return 0;
}
```



Heap and Stack Example (11/11)



arraycopy.c

```
#include <stdlib.h>


int* copy(int a[], int size) {
    int i, *a2;


    a2 = malloc(size*sizeof(int));
    if (a2 == NULL)
        return NULL;

    for (i = 0; i < size; i++)
        a2[i] = a[i];

    return a2;
}

int main(int argc, char** argv) {
    int nums[4] = {1, 2, 3, 4};
    int* ncopy = copy(nums, 4);
    // .. do stuff with the array ..
    free(ncopy);
    return 0;
}
```





Memory Corruption

- There are all sorts of ways to corrupt memory in C

memcorrupt.c

```
#include <stdio.h>
#include <stdlib.h>

int main(int argc, char** argv) {
    int a[2];
    int* b = malloc(2*sizeof(int));
    int* c;


    a[2] = 5;    // assign past the end of an array
    b[0] += 2;   // assume malloc zeros out memory
    c = b+3;     // mess up your pointer arithmetic
    free(&(a[0])); // free something not malloc'ed
    free(b);
    free(b);     // double-free the same block
    b[0] = 5;    // use a freed pointer

    // any many more!
    return 0;
}
```



Memory Leak (1/2)

- A **memory leak** occurs when
 - code fails to deallocate dynamically-allocated memory that is no longer used
 - *e.g.* forget to **free** malloc-ed block, lose/change pointer to malloc-ed block



Memory Leak (2/2)

- Implication: program's VM footprint will keep growing
 - This might be OK for *short-lived* program, since memory deallocated when program ends
 - Usually has bad repercussions for *long-lived* programs
 - Might slow down over time (e.g. lead to VM thrashing)
 - Might exhaust all available memory and crash
 - Other programs might get starved of memory

프로그램 11-1




```
9  #include <stdio.h>
10 #include <stdlib.h>
11
12 int main()
13 {
14     int *p;      // 동적 할당된 블록을 가리킬 포인터
15     int n, i;
16
17     printf("몇 개의 정수를 입력하고 싶으십니까? ");
18     scanf("%d", &n);
19     if (n <= 0) {
20         printf("오류: 정수 개수를 잘못 입력하였습니다. ");
21         printf("프로그램을 종료합니다.\n");
22         return -1;
23     }
24
25     p = (int *) malloc(n * sizeof(int)); ← 입력할 정수를 저장할 변수를 동적으로 할당함
26     if (p == NULL) {
27         printf("오류: 메모리가 부족합니다. ");
28         printf("프로그램을 종료합니다.\n");
29         return -1;
30     }
31
32     printf("%d 개의 정수를 입력해 주세요: ", n);
33     for (i = 0; i < n; ++i)
34         scanf("%d", &p[i]);
35     printf("당신이 입력한 정수를 역순으로 출력합니다: ");
36     for (i = n-1; i >= 0; --i)
37         printf("%d ", p[i]);
38     printf("\n");
39
40     return 0;
41 }
42
```

} 입력한 n이 올바른 값인지 검사

} 입력할 정수를 저장할 변수를 동적으로 할당함

} malloc이 성공적으로 메모리를 할당하였는지 검사

assert 메크로



- assert
 - 디버깅할 때에 특정 조건을 검사하기 위한 매크로
 - **assert(수식)** ; 형태로 사용함
 - 프로그램 수행 중 '수식'이 참이면 오류 메시지를 출력하고 프로그램을 종료시킴
- assert 활용 시 주의사항
 - assert 매크로는 전처리기를 통해 무시하도록 조정할 수 있음
 - NDEBUG(no debug)를 선언하면 assert 매크로를 무시함
 - IDE 환경에서 **Release 모드**로 프로젝트 설정
 - 명령줄 환경에서: **c1 -DNDEBUG 파일이름**
 - 요즘엔 소프트웨어 안정성을 위해, 디버깅이 끝난 후에도 assert를 생략하지 않도록 권유하고 있음

프로그램 11-2



```
10 #include <stdio.h>
11 #include <stdlib.h>
12 #include <assert.h>
13
14 int main()
15 {
16     int *p;      // 동적 할당된 블록을 가리킬 포인터
17     int n, i;
18
19     printf("몇 개의 정수를 입력하고 싶으십니까? ");
20     scanf("%d", &n);
21     if (n <= 0) {
22         printf("오류: 정수 개수를 잘못 입력하였습니다. ");
23         printf("프로그램을 ");
24         return -1;
25     }
26
27     p = (int *) malloc(n * sizeof(int)); } malloc 리턴값을 assert로 검사함
28     assert(p != NULL); } (오류 발생을 위한 표준 입력 값 예시: 1234567890123)
29
30     printf("%d 개의 정수를 입력해 주세요: ", n);
31     for (i = 0; i < n; ++i)
32         scanf("%d", &p[i]);
33     printf("당신이 입력한 정수를 역순으로 출력합니다: ");
34     for (i = n-1; i >= 0; --i)
35         printf("%d ", p[i]);
36     printf("\n");
37
38     return 0;
39 }
40
```

Questions?

