

# System Programming

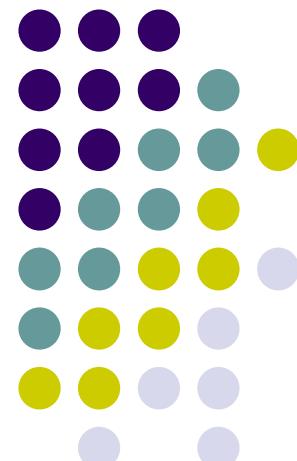
## 05. Floating Point (ch 2.4)

2019. Fall

Instructor: Joonho Kwon

[jhwon@pusan.ac.kr](mailto:jhwon@pusan.ac.kr)

Data Science Lab @ PNU





# Roadmap

C:

```
car *c = malloc(sizeof(car));  
c->miles = 100;  
c->gals = 17;  
float mpg = get_mpg(c);  
free(c);
```

Java:

```
Car c = new Car();  
c.setMiles(100);  
c.setGals(17);  
float mpg =  
    c.getMPG();
```

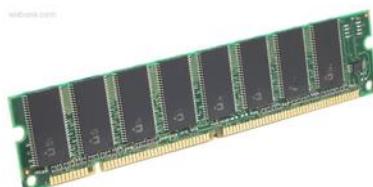
Assembly  
language:

```
get_mpg:  
    pushq  %rbp  
    movq   %rsp, %rbp  
    ...  
    popq   %rbp  
    ret
```

Machine  
code:

```
0111010000011000  
100011010000010000000010  
1000100111000010  
11000001111101000011111
```

Computer  
system:



- Memory & data
- Integers & floats
- x86 assembly
- Procedures & stacks
- Executables
- Arrays & structs
- Memory & caches
- Processes
- Virtual memory
- Memory allocation
- Java vs. C

OS:



# Number Representation Revisited



- What can we represent so far?
  - Signed and Unsigned Integers
  - Characters (ASCII)
  - Addresses
- How do we encode the following:
  - Real numbers (e.g. 3.14159)
  - Very large numbers (e.g.  $6.02 \times 10^{23}$ )
  - Very small numbers (e.g.  $6.626 \times 10^{-34}$ )
  - Special numbers (e.g.  $\infty$ , NaN)

Floating Point

# Floating point topics

- Fractional binary numbers
- IEEE floating-point standard
- Floating-point operations and roundir
- Floating-point in C
- There are many more details that we won't cover
  - It's a 58-page standard...

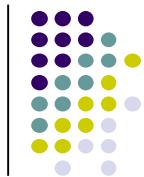


# Floating Point Summary



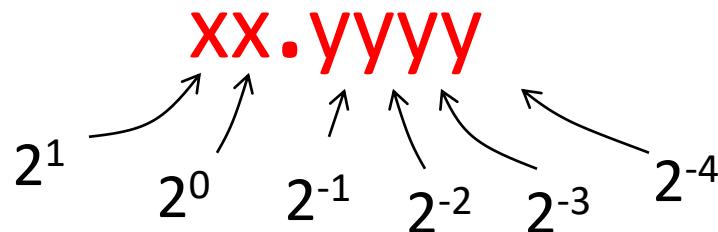
- As with integers, floats suffer from the fixed number of bits available to represent them
  - Can get overflow/underflow, just like ints
  - Some “simple fractions” have no exact representation (e.g., 0.2)
  - Can also lose precision, unlike ints
    - “Every operation gets a slightly wrong result”
- Floating point arithmetic not associative or distributive
  - Mathematically equivalent ways of writing an expression may compute different results
- **Never** test floating point values for equality!
- **Careful** when converting between ints and floats!

# Representation of Fractions



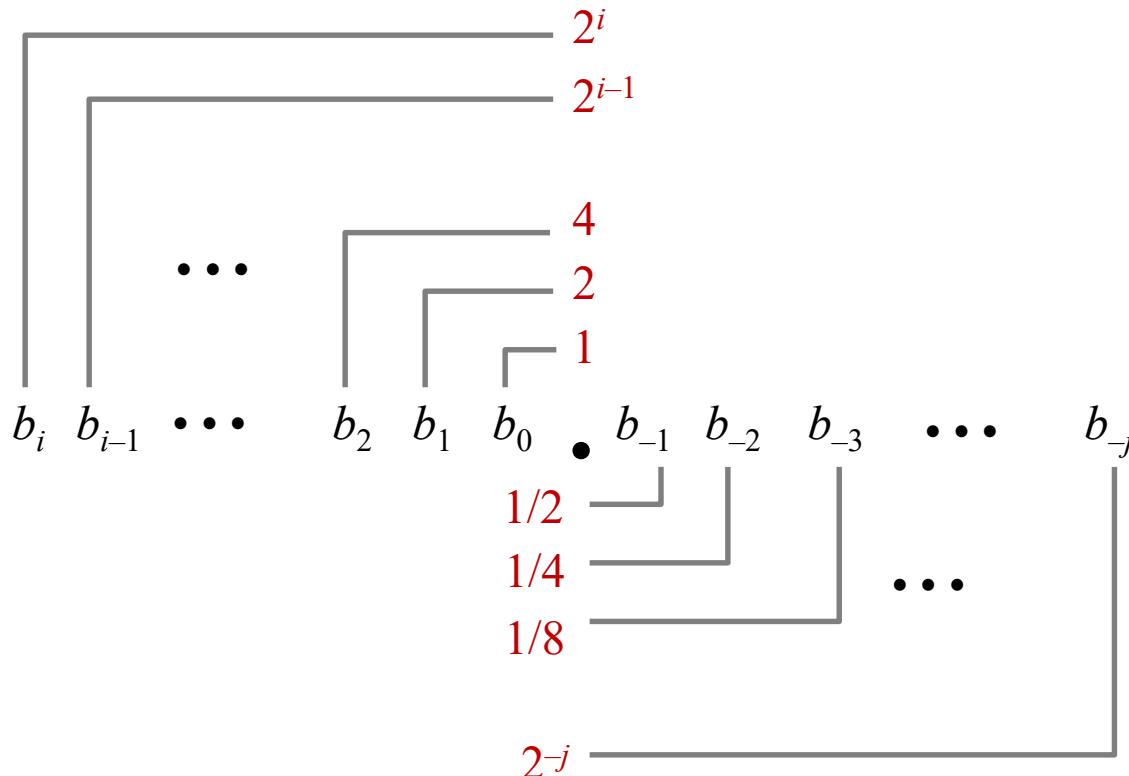
- “Binary Point,” like decimal point, signifies boundary between integer and fractional parts:

Example 6-bit representation:



- Example:  $10.1010_2 = 1 \times 2^1 + 1 \times 2^{-1} + 1 \times 2^{-3} = 2.625_{10}$
- Binary point numbers that match the 6-bit format above range from  $0 (00.0000_2)$  to  $3.9375 (11.1111_2)$

# Fractional Binary Numbers



- Representation
  - Bits to right of “binary point” represent fractional powers of 2
  - Represents rational number:

$$\sum_{k=-j}^i b_k \cdot 2^k$$

# Fractional Binary Numbers



- Value              Representation
  - 5 and 3/4         $101.11_2$
  - 2 and 7/8         $10.111_2$
  - $47/64$              $0.101111_2$
- Observations
  - Shift left = multiply by power of 2
  - Shift right = divide by power of 2
  - Numbers of the form  $0.\mathbf{111111\dots}_2$  are just below 1.0
    - $1/2 + 1/4 + 1/8 + \dots + 1/2^i + \dots \rightarrow 1.0$
    - Use notation  $1.0 - \varepsilon$

# Limits of Representation



- Limitations:
  - Even given an arbitrary number of bits, can only exactly represent numbers of the form  $x * 2^y$  ( $y$  can be negative)
  - Other rational numbers have repeating bit representations

**Value:**

- $1/3 = 0.\overline{3}_{10} = 0.01010101[01]..._2$
- $1/5 = 0.\overline{02}_{10} = 0.001100110011[0011]..._2$
- $1/10 = 0.\overline{01}_{10} = 0.0001100110011[0011]..._2$

**Binary Representation:**

# Fixed Point Representation



- Implied binary point. Two example schemes:
  - #1: the binary point is between bits 2 and 3  
 $b_7 \ b_6 \ b_5 \ b_4 \ b_3 \ [.] \ b_2 \ b_1 \ b_0$
  - #2: the binary point is between bits 4 and 5  
 $b_7 \ b_6 \ b_5 \ [.] \ b_4 \ b_3 \ b_2 \ b_1 \ b_0$
- Wherever we put the binary point, with fixed point representations there is a trade off between the amount of range and precision we have
- Fixed point = fixed *range* and fixed *precision*
  - range: difference between largest and smallest numbers possible
  - precision: smallest possible difference between any two numbers
- Hard to pick how much you need of each!

# Floating Point Representation



- Analogous to scientific notation
  - In Decimal:
    - Not 12000000, but  $1.2 \times 10^7$       In C: 1.2e7
    - Not 0.0000012, but  $1.2 \times 10^{-6}$       In C: 1.2e-6
  - In Binary:
    - Not 11000.000, but  $1.1 \times 2^4$
    - Not 0.000101, but  $1.01 \times 2^{-4}$
- We have to divvy up the bits we have (e.g., 32) among:
  - the sign (1 bit)
  - the significand
  - the exponent

# Scientific Notation Translation



- Convert from scientific notation to binary point
  - Perform the multiplication by shifting the decimal until the exponent disappears
    - Example:  $1.011_2 \times 2^4 = 10110_2 = 22_{10}$
    - Example:  $1.011_2 \times 2^{-2} = 0.01011_2 = 0.34375_{10}$
- Convert from binary point to *normalized* scientific notation
  - Distribute out exponents until binary point is to the right of a single digit
    - Example:  $1101.001_2 = 1.101001_2 \times 2^3$
- **Practice:** Convert  $11.375_{10}$  to binary scientific notation
- **Practice:** Convert  $1/5$  to binary

# Floating point topics



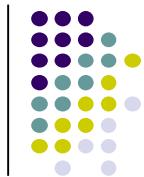
- Fractional binary numbers
- IEEE floating-point standard
- Floating-point operations and rounding
- Floating-point in C
- There are many more details that we won't cover
  - It's a 58-page standard...

# IEEE Floating Point



- IEEE 754
  - Established in 1985 as uniform standard for floating point arithmetic
  - Main idea: make numerically sensitive programs portable
  - Specifies two things: representation and result of floating operations
  - Now supported by all major CPUs
- Driven by numerical concerns
  - **Scientists**/numerical analysts want them to be as **real** as possible
  - **Engineers** want them to be **easy to implement** and **fast**
  - In the end:
    - Scientists mostly won out
    - Nice standards for rounding, overflow, underflow, but...
    - Hard to make fast in hardware
    - **Float operations can be an order of magnitude slower than integer ops**

# Floating Point Representation



- Numerical form:

$$V_{10} = (-1)^S * M * 2^E$$

- Sign bit  $S$  determines whether number is negative or positive
- Significand (mantissa)  $M$  normally a fractional value in range [1.0,2.0)
- Exponent  $E$  weights value by a (possibly negative) power of two

# Floating Point Representation



- Numerical form:

$$V_{10} = (-1)^S * M * 2^E$$

- Sign bit **s** determines whether number is negative or positive
  - Significand (mantissa) **M** normally a fractional value in range [1.0,2.0)
  - Exponent **E** weights value by a (possibly negative) power of two
- ❖ Representation in memory:
- MSB **s** is sign bit **s**
  - **exp** field encodes **E** (but is *not equal* to E)
  - **frac** field encodes **M** (but is *not equal* to M)





# Precisions

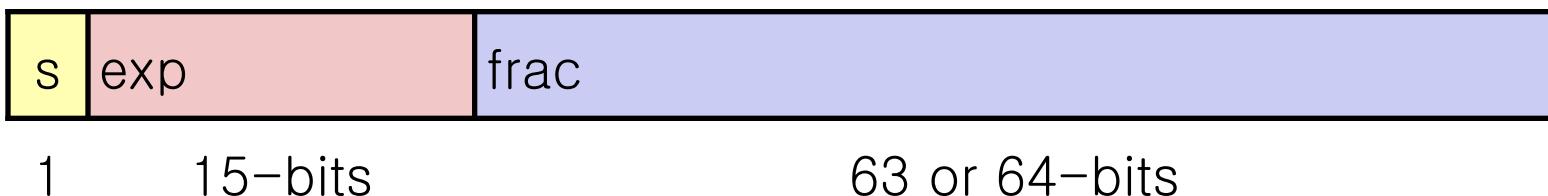
- Single precision: 32 bits



- Double precision: 64 bits



- Extended precision: 80 bits (Intel only)



Finite representation means not all values can be represented exactly. Some will be approximated.

# Normalization and Special Values



$$v = (-1)^s M 2^E$$



- “Normalized” = **M** has the form 1.xxxxx
  - As in scientific notation, but in binary
  - $0.011 \times 2^5$  and  $1.1 \times 2^3$  represent the same number, but the latter makes better use of the available bits
  - Since we know the mantissa starts with a 1, we don't bother to store it
- **How do we represent 0.0? Or special or undefined values like 1.0/0.0?**



# Normalized Values

$$v = (-1)^s M \cdot 2^e$$



- “Normalized” = **M** has the form 1.xxxxx
  - As in scientific notation, but in binary
  - $0.011 \times 2^5$  and  $1.1 \times 2^3$  represent the same number, but the latter makes better use of the available bits
  - Since we know the mantissa starts with a 1, we don't bother to store it
- Special values:

**zero:**    **s == 0**    **exp == 00...0**    **frac == 00...0**

**+∞, -∞ :**                **exp == 11...1**    **frac == 00...0**

$$1.0/0.0 = -1.0/-0.0 = +\infty, 1.0/-0.0 = -1.0/0.0 = -\infty$$

**NaN** (“Not a Number”):                **exp == 11...1**    **frac != 00...0**

Results from operations with undefined result:  $\sqrt{-1}$ ,  $\infty - \infty$ ,  $\infty * 0$ , etc.

Note: **exp=11...1** and **exp=00...0** are reserved, limiting exp range...



# Normalized Values

$$v = (-1)^s M 2^E$$



- When:  $\text{exp} \neq 000\dots0$  and  $\text{exp} \neq 111\dots1$
- Exponent coded as a ***biased*** value:  $E = \text{Exp} - \text{Bias}$ 
  - $\text{Exp}$ : unsigned value of exp field
  - $\text{Bias} = 2^{k-1} - 1$ , where  $k$  is number of exponent bits
    - Single precision: 127 (Exp: 1...254, E: -126...127)
    - Double precision: 1023 (Exp: 1...2046, E: -1022...1023)
- Significand coded with implied leading 1:  $M = 1.\text{xxx...x}_2$ 
  - $\text{xxx...x}$ : bits of frac field
  - Minimum when  $\text{frac}=000\dots0$  ( $M = 1.0$ )
  - Maximum when  $\text{frac}=111\dots1$  ( $M = 2.0 - \varepsilon$ )
  - Get extra leading bit for “free”

# Normalized Encoding Example

$$v = (-1)^s M \cdot 2^E$$
$$E = \text{Exp} - \text{Bias}$$

- Value: float F = 15213.0;
  - $15213_{10} = 11101101101101_2$   
 $= 1.1101101101101_2 \times 2^{13}$  (normalized form)

- Significand
  - $M = 1.\underline{1101101101101}_2$
  - $\text{frac} = \underline{1101101101101}0000000000_2$

- Exponent
  - $E = 13$
  - $\text{Bias} = 127$
  - $\text{Exp} = 140 = 10001100_2$

- Result:

0	10001100	1101101101101000000000000
s	exp	frac

# Question



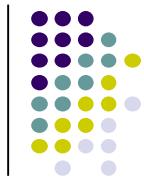
- What is the correct value encoded by the following floating point number?
    - 0b 0 10000000 11000000000000000000000000000000
- 
- A. + 0.75
- B. + 1.5
- C. + 2.75
- D. + 3.5

# Denormalized Values

$$v = (-1)^s M \cdot 2^E$$
$$E = 1 - Bias$$

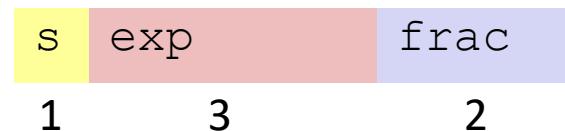
- Condition:  $exp = 000\dots0$
- Exponent value:  $E = 1 - Bias$  (instead of  $E = 0 - Bias$ )
- Significand coded with implied leading 0:  $M = 0.\text{xxx...x}_2$ 
  - $\text{xxx...x}$ : bits of  $\text{frac}$
- Cases
  - $\text{exp} = 000\dots0, \text{frac} = 000\dots0$ 
    - Represents zero value
    - Note distinct values: +0 and -0 (why?)
  - $\text{exp} = 000\dots0, \text{frac} \neq 000\dots0$ 
    - Numbers closest to 0.0
    - Equispaced

# Distribution of Values

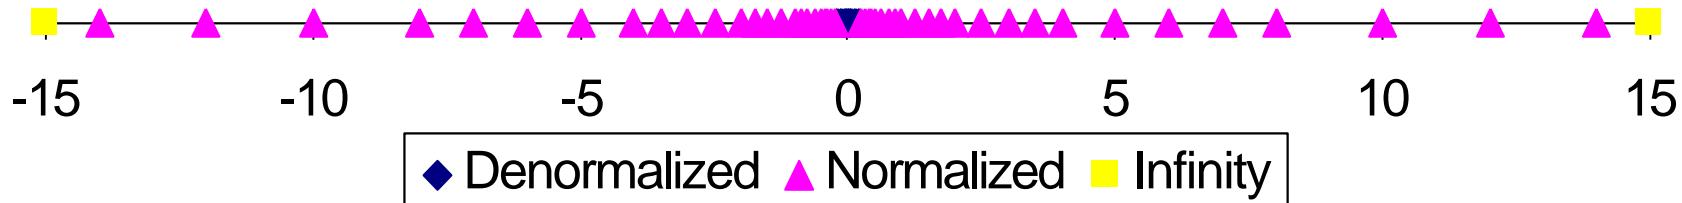


- 6-bit IEEE-like format

- $e = 3$  exponent bits
- $f = 2$  fraction bits
- Bias is  $2^{3-1}-1 = 3$



- Notice how the distribution gets denser toward zero.



# Floating point topics



- Fractional binary numbers
  - IEEE floating-point standard
  - **Floating-point operations** and rounding
  - Floating-point in C
- 
- There are many more details that we won't cover
    - It's a 58-page standard...

# Floating Point Operations



- Unlike the representation for integers, the representation for floating-point numbers is not exact

# Floating Point Operations: Basic Idea

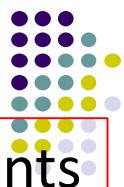


$$v = (-1)^s M 2^E$$



- $x +_f y = \text{Round}(x + y)$
- $x \times_f y = \text{Round}(x \times y)$
- Basic idea
  - First **compute exact result**
  - Then, **round** the result to make it fit into desired precision:
    - Possibly overflow if exponent too large
    - Possibly **drop least-significant bits** of significand to **fit into frac**

# Floating Point Addition



Line up the binary points

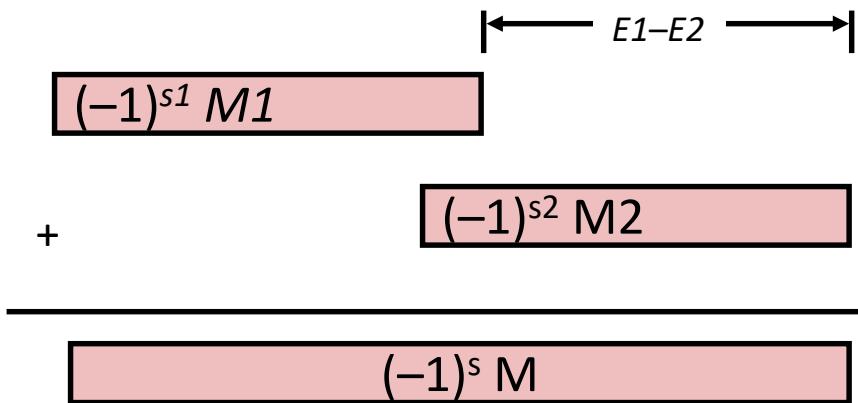
$$(-1)^{s1} M1 2^{E1} + (-1)^{s2} M2 2^{E2}$$

Assume  $E1 > E2$

- Exact Result:  $(-1)^s M 2^E$

- Sign  $s$ , significand  $M$ :
  - Result of signed align & add

- Exponent  $E$ :  $E1$



- Fixing
  - If  $M \geq 2$ , shift  $M$  right, increment  $E$
  - if  $M < 1$ , shift  $M$  left  $k$  positions, decrement  $E$  by  $k$
  - Overflow if  $E$  out of range
  - Round  $M$  to fit **frac** precision

# Floating Point Multiplication



$$(-1)^{s_1} M_1 2^{E_1} * (-1)^{s_2} M_2 2^{E_2}$$

- Exact Result:  $(-1)^s M 2^E$ 
  - Sign  $s$ :  $s_1 \wedge s_2$
  - Significand  $M$ :  $M_1 * M_2$
  - Exponent  $E$ :  $E_1 + E_2$
- Fixing
  - If  $M \geq 2$ , shift  $M$  right, increment  $E$
  - If  $E$  out of range, overflow
  - Round  $M$  to fit **frac** precision

# Rounding



- Rounding Modes (illustrate with \$ rounding)

	\$1.40	\$1.60	\$1.50	\$2.50	-\$1.50
Towards zero	\$1	\$1	\$1	\$2	-\$1
Round down ( $-\infty$ )	\$1	\$1	\$1	\$2	-\$2
Round up ( $+\infty$ )	\$2	\$2	\$2	\$3	-\$1
Nearest Even (default)	\$1	\$2	\$2	\$2	-\$2

- Round-to-even avoids statistical bias in repeated rounding.
  - Rounds up about half the time, down about half the time.
  - Default rounding mode for IEEE floating-point

# Mathematical Properties of FP Operations



- Exponent overflow yields  $+\infty$  or  $-\infty$
- Floats with value  $+\infty$ ,  $-\infty$ , and NaN can be used in operations
  - Result usually still  $+\infty$ ,  $-\infty$ , or NaN; but not always intuitive
- Floating point operations do not work like real math, due to **rounding!!**
  - Not associative:  $(3.14+1e100)-1e100 \neq 3.14 + (1e100-1e100)$   
$$0 \qquad \qquad \qquad \textcolor{red}{3.14}$$
  - Not distributive:  $100*(0.1+0.2) \neq 100*0.1+100*0.2$   
$$\textcolor{red}{30.0000000000003553} \qquad \qquad \qquad \textcolor{red}{30}$$
  - Not cumulative
    - Repeatedly adding a very small number to a large one may do nothing

# Floating Point Topics



- Fractional binary numbers
- IEEE floating-point standard
- Floating-point operations and rounding
- **Floating-point in C**
  
- There are many more details that we won't cover
  - It's a 58-page standard...

# Floating Point in C

!!!

- C offers two (well, 3) levels of precision

float

1.0f

single precision (32-bit)

double

1.0

double precision (64-bit)

long double

1.0L

(“*double double*” or *quadruple*)  
precision (64-128 bits)

- #include <math.h> to get INFINITY and NAN constants
- Equality (==) comparisons between floating point numbers are tricky, and often return unexpected results, so just avoid them!
- Instead use  $\text{abs}(f1 - f2) < 2^{-20}$  or some other threshold

# Floating Point Conversions in C

!!!

- Casting between `int`, `float`, and `double` changes the bit representation
  - `int` → `float`
    - May be rounded (not enough bits in mantissa: 23)
    - Overflow impossible
  - `int or float` → `double`
    - Exact conversion (all 32-bit `ints` representable; 52-bit frac)
  - `long` → `double`
    - Depends on word size (32-bit is exact, 64-bit may be rounded)
  - `double or float` → `int`
    - Truncates fractional part (rounded toward zero)
    - “Not defined” when out of range or NaN: generally sets to `Tmin` (even if the value is a very big positive)

# Number Representation Really Matters



- **1991:** Patriot missile targeting error
  - clock skew due to conversion from integer to floating point
- **1996:** Ariane 5 rocket exploded (\$1 billion)
  - overflow converting 64-bit floating point to 16-bit integer
- **2000:** Y2K problem
  - limited (decimal) representation: overflow, wrap-around
- **2038:** Unix epoch rollover
  - Unix epoch = seconds since 12am, January 1, 1970
  - signed 32-bit integer representation rolls over to TMin in 2038
- **Other related bugs:**
  - 1982: Vancouver Stock Exchange 10% error in less than 2 years
  - 1994: Intel Pentium FDIV (floating point division) HW bug (\$475 million)
  - 1997: USS Yorktown “smart” warship stranded: divide by zero
  - 1998: Mars Climate Orbiter crashed: unit mismatch (\$193 million)

# Floating Point and the Programmer



```
include <stdio.h>
```

```
int main(int argc, char* argv[]) {
    int a = 33554435;
    printf("a = %d\n(float) a = %f
           \n\n", a, (float) a);

    float f1 = 1.0;
    float f2 = 0.0;
    int i;
    for (i = 0; i < 10; i++)
        f2 += 1.0/10.0;

    printf("0x%08x  0x%08x\n", *(int*)&f1, *(int*)&f2);
    printf("f1 = %10.9f\n", f1);
    printf("f2 = %10.9f\n\n", f2);

    f1 = 1E30;
    f2 = 1E-30;
    float f3 = f1 + f2;
    printf("f1 == f3? %s\n", f1 == f3 ? "yes" : "no" );
    return 0;
}
```

```
$ ./a.out
a = 33554435
(float) a = 33554436.000000
0x3f800000 0x3f800001
f1 = 1.000000000
f2 = 1.000000119

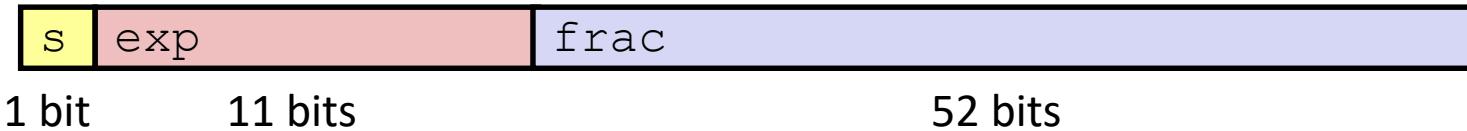
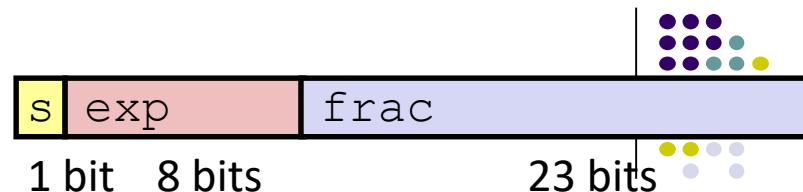
f1 == f3? yes
```

# Floating Point Summary



- As with integers, floats suffer from the fixed number of bits available to represent them
  - Can get overflow/underflow, just like ints
  - Some “simple fractions” have no exact representation (e.g., 0.2)
  - Can also lose precision, unlike ints
    - “Every operation gets a slightly wrong result”
- Floating point arithmetic not associative or distributive
  - Mathematically equivalent ways of writing an expression may compute different results
- **Never** test floating point values for equality!
- **Careful** when converting between ints and floats!

# Floating Point Puzzles



- For each of the following C expressions, either:
  - Argue that it is true for all argument values
  - Explain why not true

```
int x = ...;  
  
float f = ...;  
  
double d = ...;  
  
double d2 = ...;
```

Assume neither  
d nor f is NaN

- `x == (int)(float) x`
- `x == (int)(double) x`
- `f == (float)(double) f`
- `d == (double)(float) d`
- `f == -(-f);`
- `2/3 == 2/3.0`
- `(d+d2)-d == d2`

# Q&A

