

System Programming

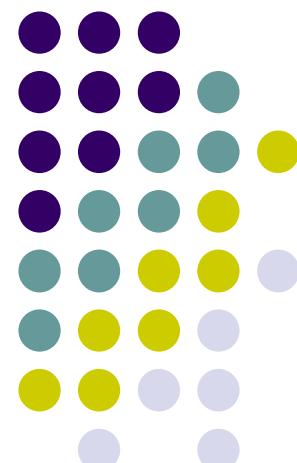
06. Machine-Level Programming I: Basics (ch 3.1-3.5)- part1

2019. Fall

Instructor: Joonho Kwon

jhwon@pusan.ac.kr

Data Science Lab @ PNU



Roadmap



C:

```
car *c = malloc(sizeof(car));  
c->miles = 100;  
c->gals = 17;  
float mpg = get_mpg(c);  
free(c);
```

Java:

```
Car c = new Car();  
c.setMiles(100);  
c.setGals(17);  
float mpg =  
    c.getMPG();
```

Assembly
language:

```
get_mpg:  
    pushq  %rbp  
    movq   %rsp, %rbp  
    ...  
    popq  %rbp  
    ret
```

Machine
code:

```
0111010000011000  
100011010000010000000010  
1000100111000010  
110000011111101000011111
```

Computer
system:



Memory & data
Integers & floats
Machine code & C
x86 assembly
Procedures & stacks
Arrays & structs
Memory & caches
Processes
Virtual memory
Memory allocation
Java vs. C

OS:

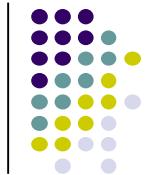


Machine Programming I: Basics

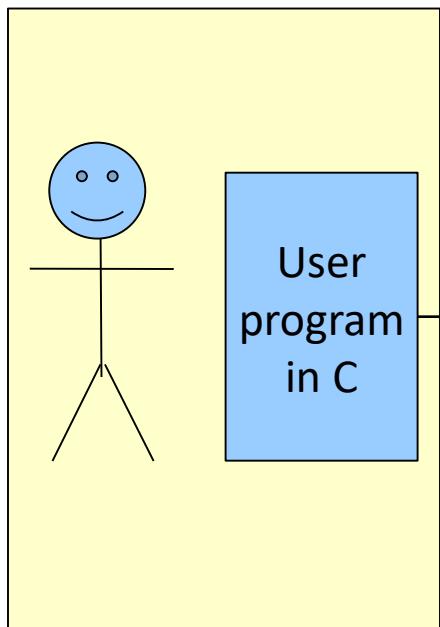


- **History of Intel processors and architectures**
- C, assembly, machine code
- Assembly Basics: Registers, operands, move
- Arithmetic & logical operations

Translation

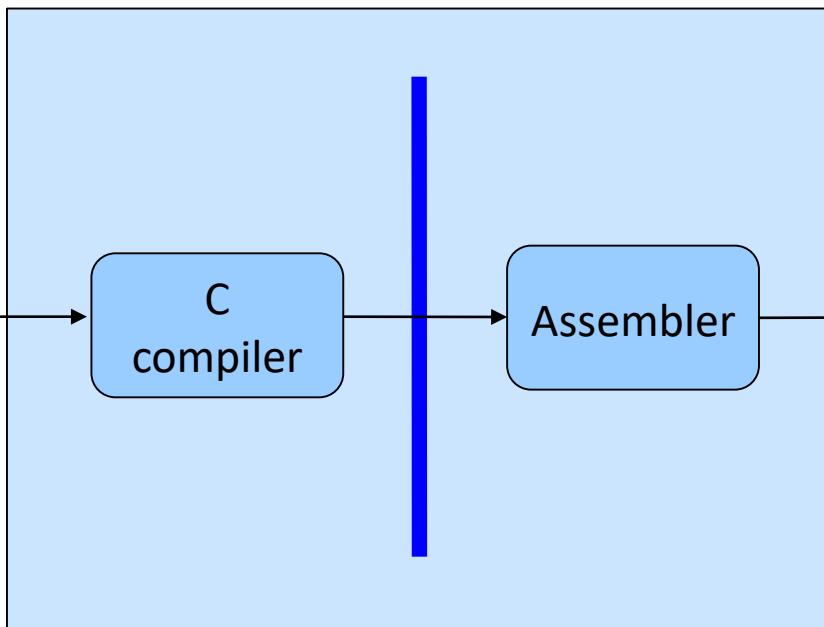


Code Time



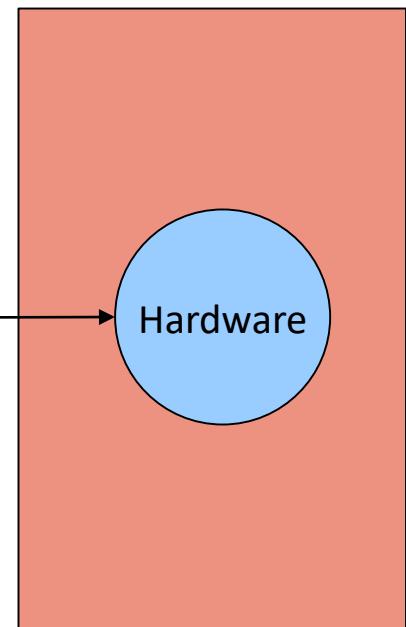
.c file

Compile Time



.exe file

Run Time



What makes programs run fast(er)?

Translation Impacts Performance



- The time required to execute a program depends on:
 - **The program** (as written in C, for instance)
 - **The compiler**: what set of assembler instructions it translates the C program into
 - **The instruction set architecture** (ISA): what set of instructions it makes available to the compiler
 - **The hardware implementation**: how much time it takes to execute an instruction

What should the HW/SW interface contain?

HW Interface Affects Performance



Source code

Different applications or algorithms

Compiler

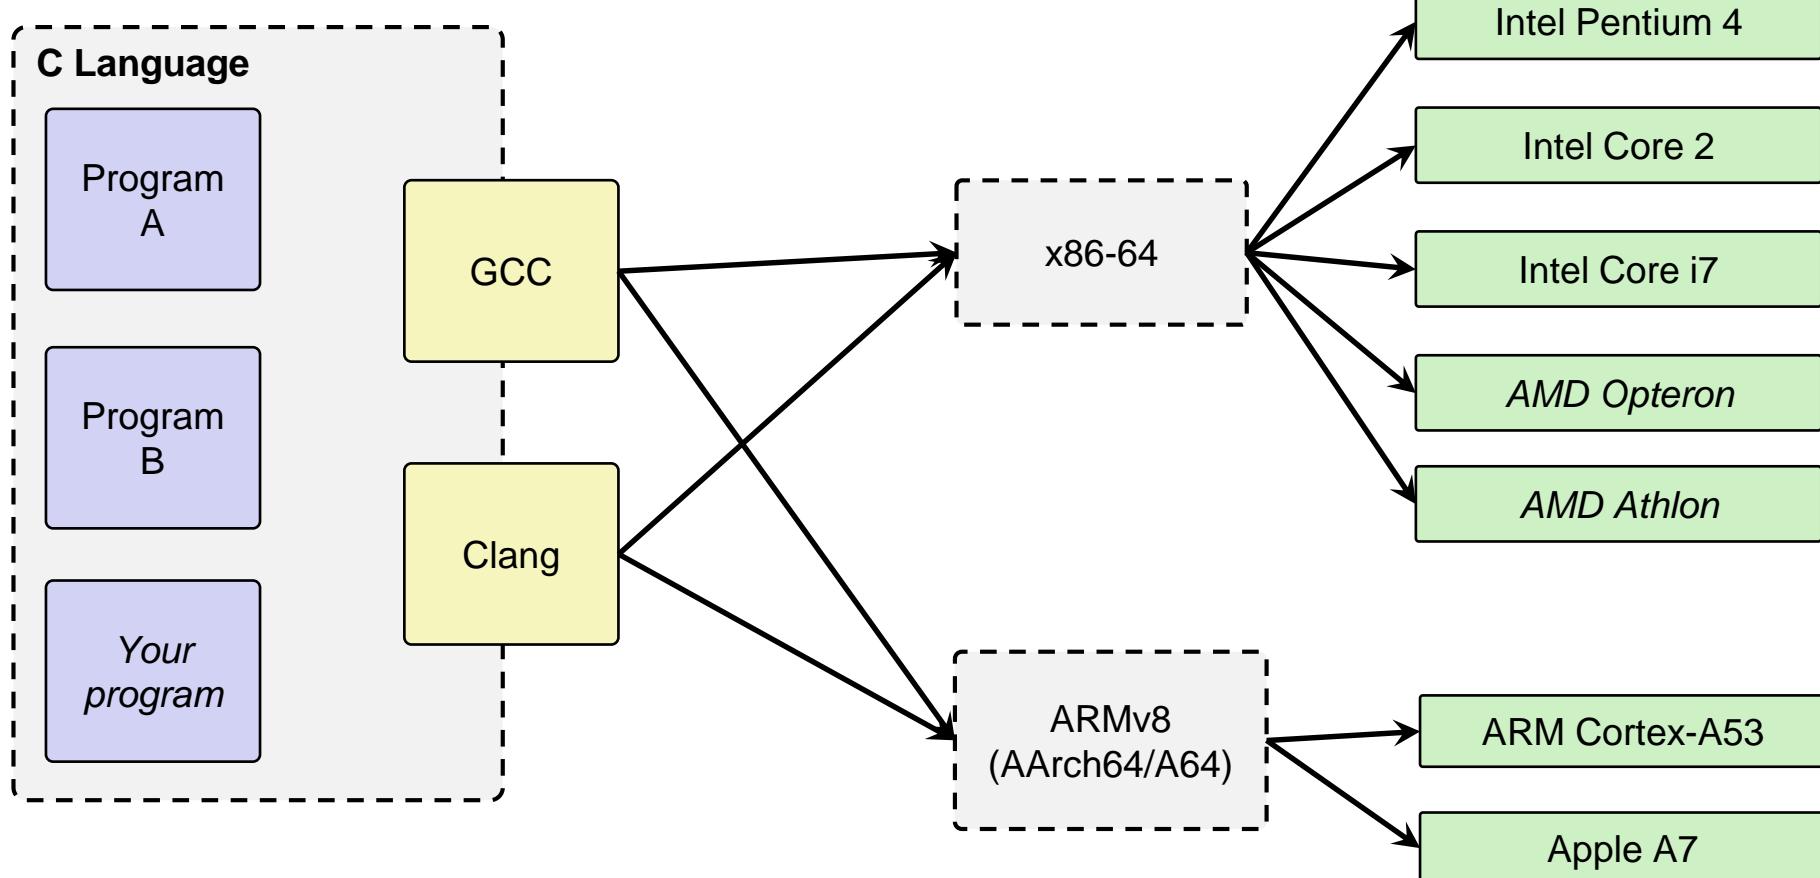
Perform optimizations, generate instructions

Architecture

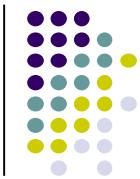
Instruction set

Hardware

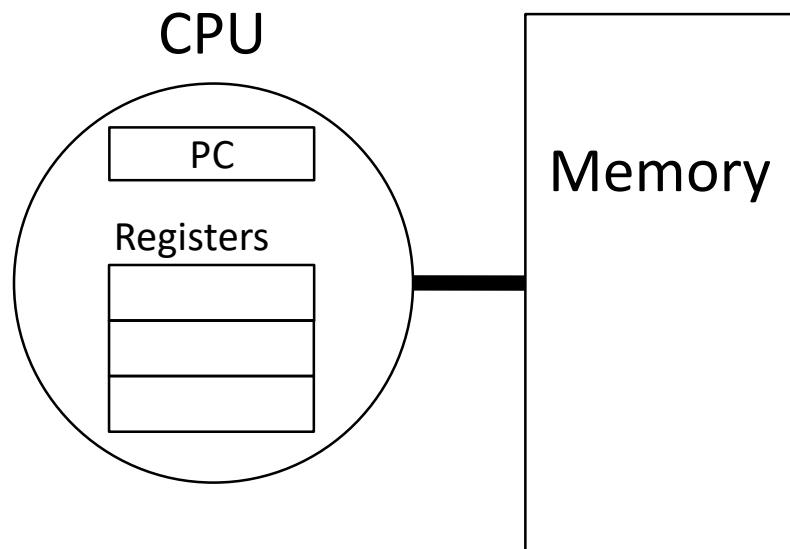
Different implementations



Instruction Set Architectures



- The ISA defines:
 - The system's state (e.g. registers, memory, program counter)
 - The instructions the CPU can execute
 - The effect that each of these instructions will have on the system state



Instruction Set Philosophies



- *Complex Instruction Set Computing (CISC):*
 - Add more and more elaborate and specialized instructions as needed
 - Lots of tools for programmers to use, but hardware must be able to handle all instructions
 - x86-64 is CISC, but only a small subset of instructions encountered with Linux programs
- *Reduced Instruction Set Computing (RISC):*
 - Keep instruction set small and regular
 - Easier to build fast hardware
 - Let software do the complicated operations by composing simpler ones

General ISA Design Decisions



- Instructions
 - What instructions are available? What do they do?
 - How are they encoded? **Instructions are data!**
- Registers
 - How many registers are there?
 - How wide are they? **Size of a word**
- Memory
 - How do you specify a memory location? **Different ways to build up an address**

Mainstream ISAs



x86

Designer	Intel, AMD
Bits	16-bit, 32-bit and 64-bit
Introduced	1978 (16-bit), 1985 (32-bit), 2003 (64-bit)
Design	CISC
Type	Register-memory
Encoding	Variable (1 to 15 bytes)
Endianness	Little

Macbooks & PCs
(Core i3, i5, i7, M)
[x86-64 Instruction Set](#)

ARM

ARM architectures

Designer	ARM Holdings
Bits	32-bit, 64-bit
Introduced	1985; 31 years ago
Design	RISC
Type	Register-Register
Encoding	AArch64/A64 and AArch32/A32 use 32-bit instructions, T32 (Thumb-2) uses mixed 16- and 32-bit instructions. ARMv7 user-space compatibility ^[1]
Endianness	Bi (little as default)

Smartphone-like devices
(iPhone, iPad, Raspberry Pi)
[ARM Instruction Set](#)

MIPS

Designer	MIPS Technologies, Inc.
Bits	64-bit (32→64)
Introduced	1981; 35 years ago
Design	RISC
Type	Register-Register
Encoding	Fixed
Endianness	Bi

Digital home & networking
equipment
(Blu-ray, PlayStation 2)
[MIPS Instruction Set](#)

X86 ISA



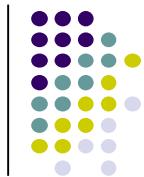
- Processors that implement the x86 ISA completely dominate the server, desktop and laptop markets
- **Evolutionary design**
 - Backwards compatible up until 8086, introduced in 1978
 - Added more features as time goes on
- **Complex instruction set computer (CISC)**
 - Many different instructions with many different formats
 - But, only small subset encountered with Linux programs
 - (as opposed to Reduced Instruction Set Computers (RISC), which use simpler instructions)

Intel x86 Evolution: Milestones



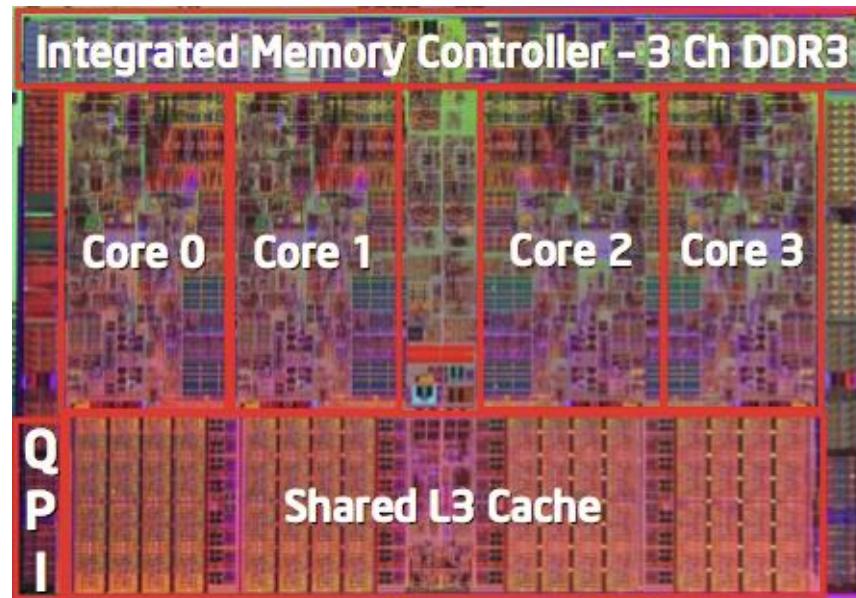
<i>Name</i>	<i>Date</i>	<i>Transistors</i>	<i>MHz</i>
● 8086	1978	29K	5-10
		● First 16-bit Intel processor. Basis for IBM PC & DOS	
		● 1MB address space	
● 386	1985	275K	16-33
		● First 32 bit Intel processor , referred to as IA32	
		● Added “flat addressing”, capable of running Unix	
● Pentium 4E	2004	125M	2800-3800
		● First 64-bit Intel x86 processor, referred to as x86-64	
● Core 2	2006	291M	1060-3500
		● First multi-core Intel processor	
● Core i7	2008	731M	1700-3900
		● Four cores (our shark machines)	

Intel x86 Processors, cont.



• Machine Evolution

• 386	1985	0.3M
• Pentium	1993	3.1M
• Pentium/MMX	1997	4.5M
• PentiumPro	1995	6.5M
• Pentium III	1999	8.2M
• Pentium 4	2001	42M
• Core 2 Duo	2006	291M
• Core i7	2008	731M



• Added Features

- Instructions to support multimedia operations
- Instructions to enable more efficient conditional operations
- Transition from 32 bits to 64 bits
- More cores

2015 State of the Art



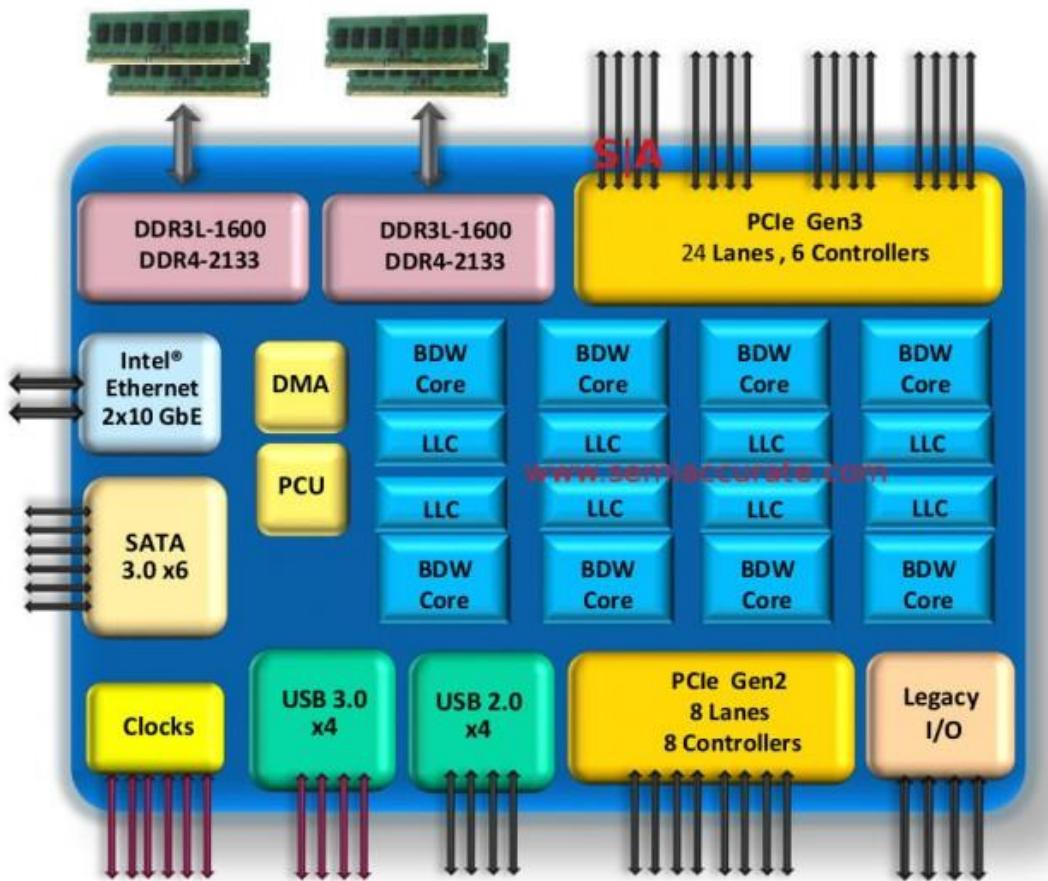
- Core i7 Broadwell 2015

- Desktop Model

- 4 cores
- Integrated graphics
- 3.3-3.8 GHz
- 65W

- Server Model

- 8 cores
- Integrated I/O
- 2-2.6 GHz
- 45W



x86 Clones: Advanced Micro Devices (AMD)

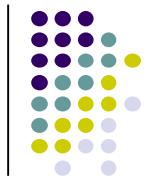


- Historically
 - AMD has followed just behind Intel
 - A little bit slower, a lot cheaper
- Then
 - Recruited top circuit designers from Digital Equipment Corp. and other downward trending companies
 - Built Opteron: tough competitor to Pentium 4
 - Developed x86-64, their own extension to 64 bits
- Recent Years
 - Intel got its act together
 - Leads the world in semiconductor technology
 - AMD has fallen behind
 - Relies on external semiconductor manufacturer

Intel's Transition to 64-Bit



- **2001: Intel Attempts Radical Shift from IA32 to IA64**
 - Totally different architecture (Itanium)
 - Executes IA32 code only as legacy
 - Performance disappointing
- **2003: AMD Steps in with Evolutionary Solution**
 - x86-64 (now called “AMD64”)
- **Intel Felt Obligated to Focus on IA64**
 - Hard to admit mistake or that AMD is better
- **2004: Intel Announces EM64T extension to IA32**
 - Extended Memory 64-bit Technology
 - Almost identical to x86-64!
- **All but low-end x86 processors support x86-64**
 - But, lots of code still runs in 32-bit mode



Our Coverage

- IA32
 - The traditional x86
- x86-64
 - The standard
 - \$ gcc hello.c
 - \$ gcc -m64 hello.c
- Presentation
 - Book covers x86-64
 - Web aside on IA32
 - We will only cover x86-64



Roadmap

C:

```
car *c = malloc(sizeof(car));
c->miles = 100;
c->gals = 17;
float mpg = get_mpg(c);
free(c);
```

Java:

```
Car c = new Car();
c.setMiles(100);
c.setGals(17);
float mpg =
    c.getMPG();
```

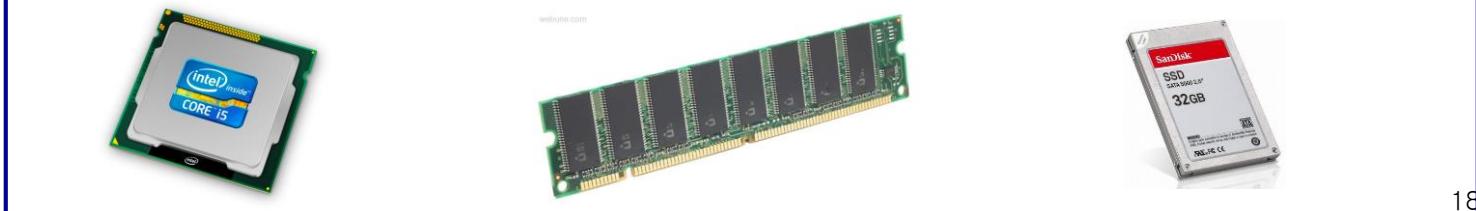
Assembly language:

```
get_mpg:
    pushq    %rbp
    movq    %rsp, %rbp
    ...
    popq    %rbp
    ret
```

Machine code:

```
0111010000011000
100011010000010000000010
1000100111000010
110000011111101000011111
```

Computer system:



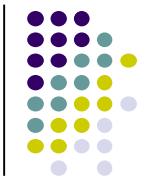
- Memory & data
- Integers & floats
- x86 assembly
- Procedures & stacks
- Executables
- Arrays & structs
- Memory & caches
- Processes
- Virtual memory
- Memory allocation
- Java vs. C

Machine Programming I: Basics



- History of Intel processors and architectures
- **C, assembly, machine code**
- Assembly Basics: Registers, operands, move
- Arithmetic & logical operations

Definitions



- **Architecture:** (also ISA: instruction set architecture) The parts of a processor design that one needs to understand or write assembly/machine code.
 - Examples: instruction set specification, registers.
- **Microarchitecture:** Implementation of the architecture.
 - Examples: cache sizes and core frequency.
- **Code Forms:**
 - **Machine Code:** The byte-level programs that a processor executes
 - **Assembly Code:** A text representation of machine code
- Is cache size “architecture”?
- How about CPU frequency?
- And number of registers?

Detour: Instruction Set Design Issues



- Where are operands stored?
 - registers, memory, stack, accumulator
- How many explicit operands are there?
 - 0, 1, 2, or 3
- How is the operand location specified?
 - register, immediate, indirect, . . .
- What type & size of operands are supported?
 - byte, int, float, double, string, vector. . .
- What operations are supported?
 - add, sub, mul, move, compare . . .

Detour: Evolution of Instruction Sets



Single Accumulator (EDSAC 1950, Maurice Wilkes)

Accumulator + Index Registers
(Manchester Mark I, IBM 700 series 1953)

Separation of Programming Model
from Implementation

High-level Language Based
(B5000 1963)

Concept of a Family
(IBM 360 1964)

General Purpose Register Machines

Complex Instruction Sets
(Vax, Intel 432 1977-80)

Load/Store Architecture
(CDC 6600, Cray 1 1963-76)

CISC
Intel x86, Pentium

RISC

(MIPS, Sparc, HP-PA, IBM RS6000, PowerPC . . . 1987)



Detour: Classifying ISAs

Accumulator (before 1960, e.g. 68HC11):

1-address	add A	$acc \leftarrow acc + mem[A]$
-----------	-------	-------------------------------

Stack (1960s to 1970s):

0-address	add	$tos \leftarrow tos + next$
-----------	-----	-----------------------------

Memory-Memory (1970s to 1980s):

2-address	add A, B	$mem[A] \leftarrow mem[A] + mem[B]$
3-address	add A, B, C	$mem[A] \leftarrow mem[B] + mem[C]$

Register-Memory (1970s to present, e.g. 80x86):

2-address	add R1, A	$R1 \leftarrow R1 + mem[A]$
	load R1, A	$R1 \leftarrow mem[A]$

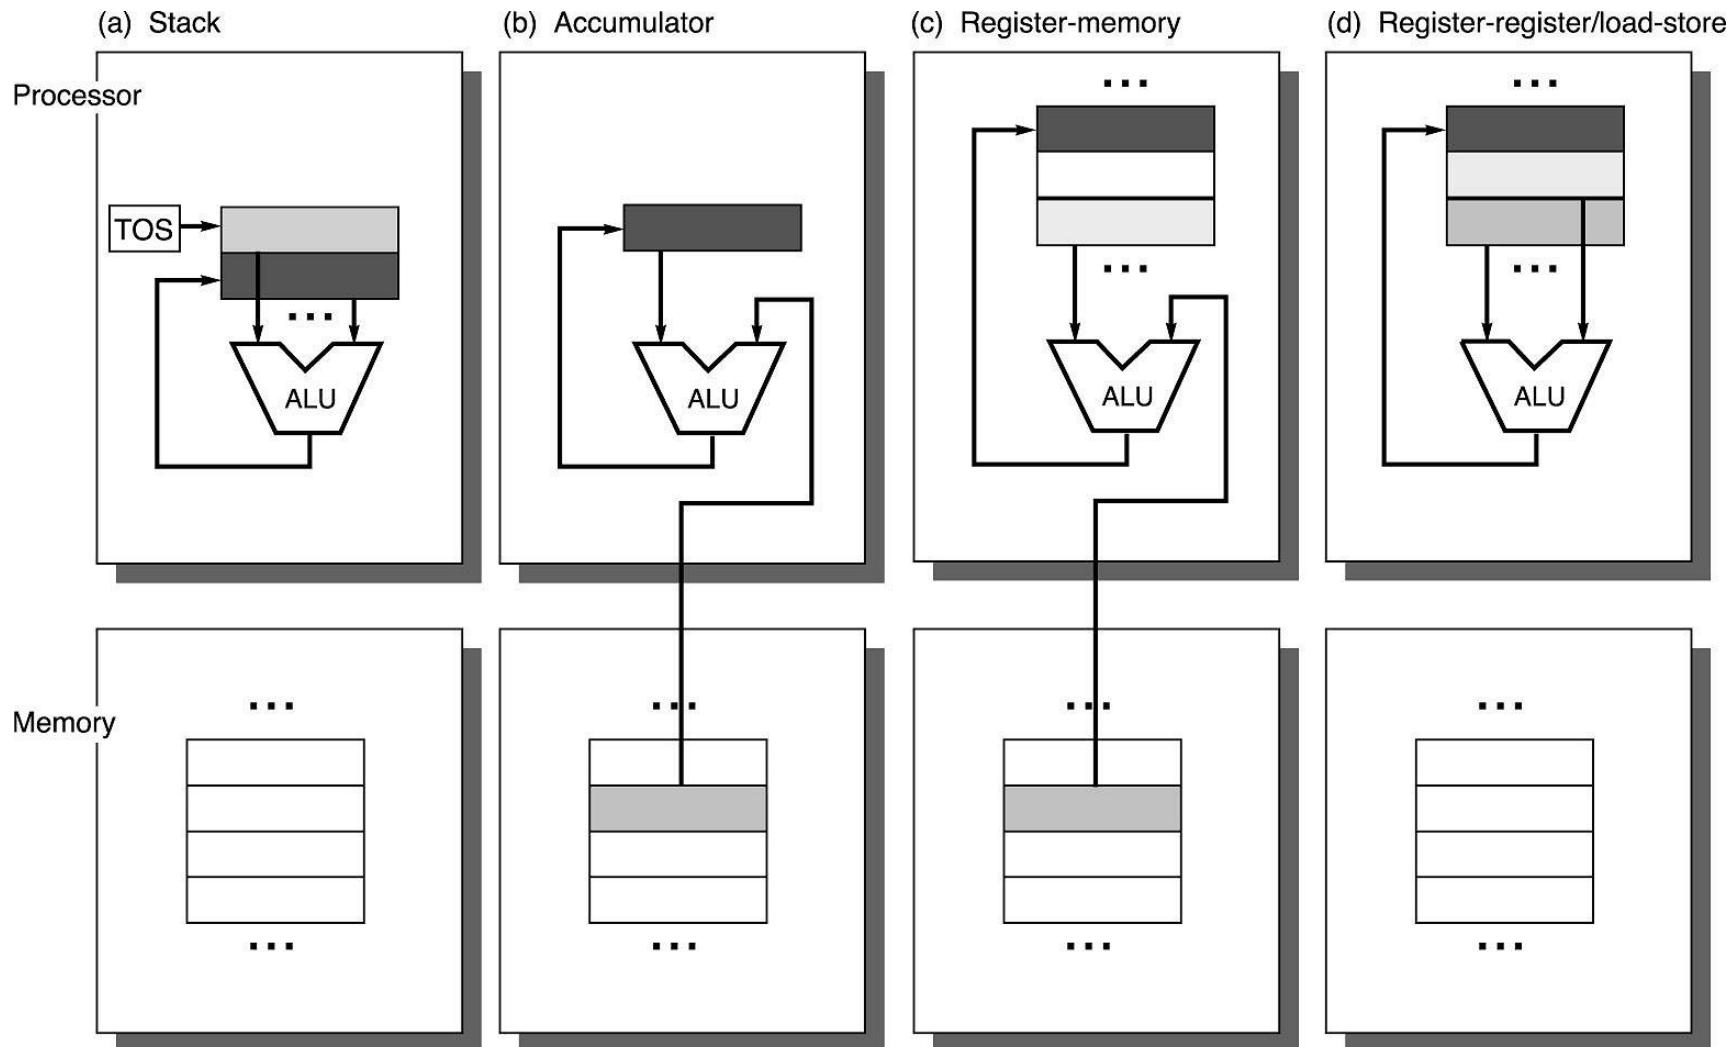
Register-Register (Load/Store) (1960s to present, e.g. MIPS):

3-address	add R1, R2, R3	$R1 \leftarrow R2 + R3$
	load R1, R2	$R1 \leftarrow mem[R2]$
	store R1, R2	$mem[R1] \leftarrow R2$

Detour: Operand Locations in Four ISA Class es



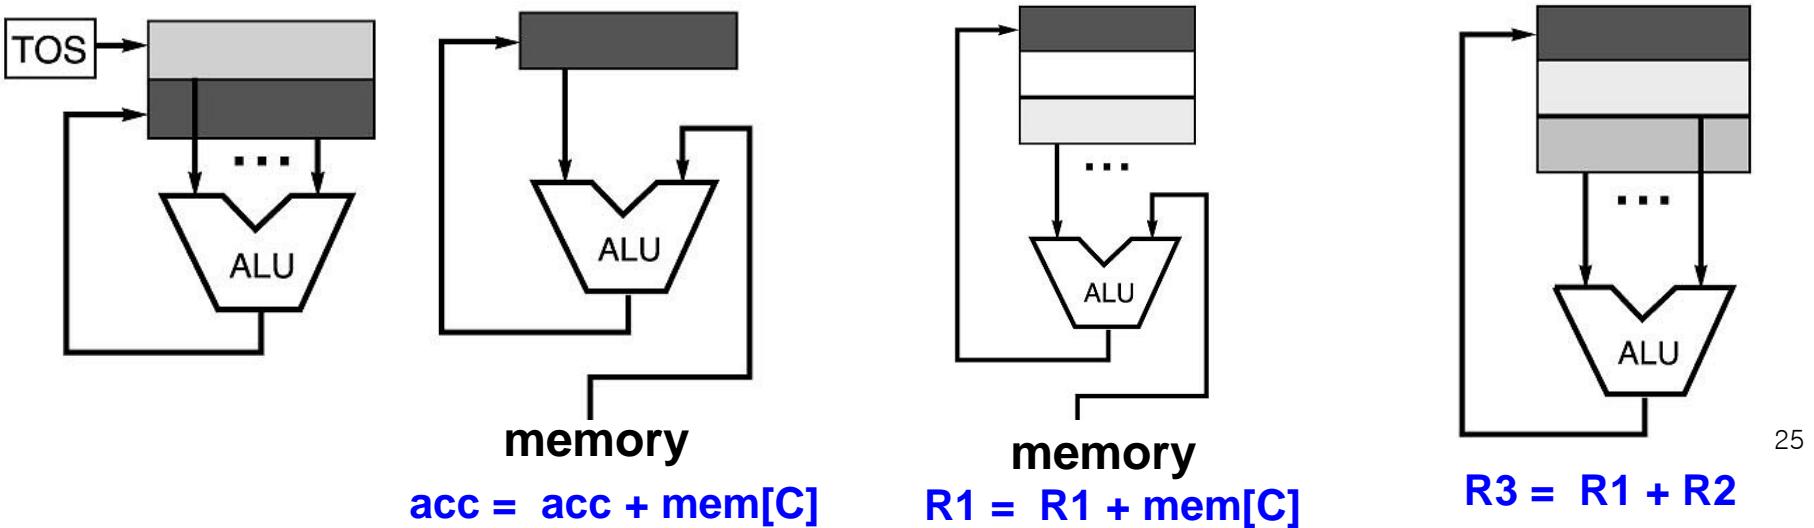
↔ GPR →



Detour: Code Sequence $C = A + B$ for Four Instruction Sets



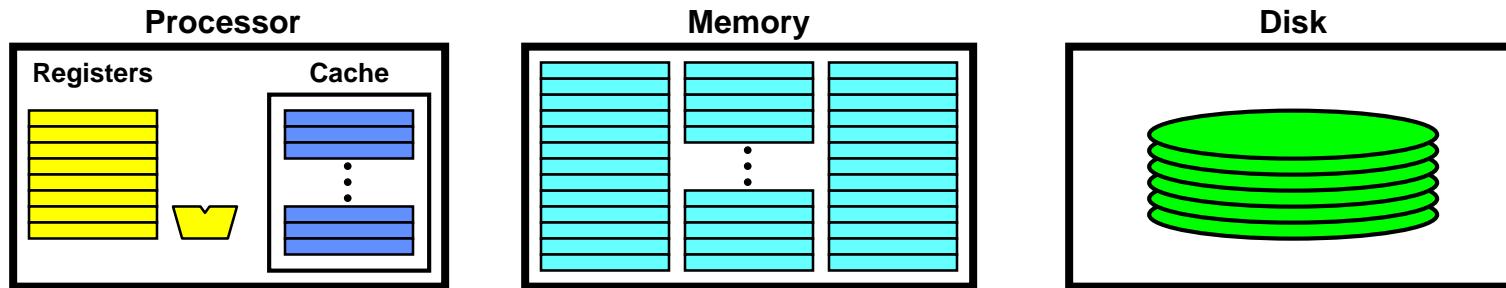
Stack	Accumulator	Register (register-memory)	Register (load-store)
Push A	Load A	Load R1, A	Load R1,A
Push B	Add B	Add R1, B	Load R2, B
Add	Store C	Store C, R1	Add R3, R1, R2
Pop C			Store C, R3



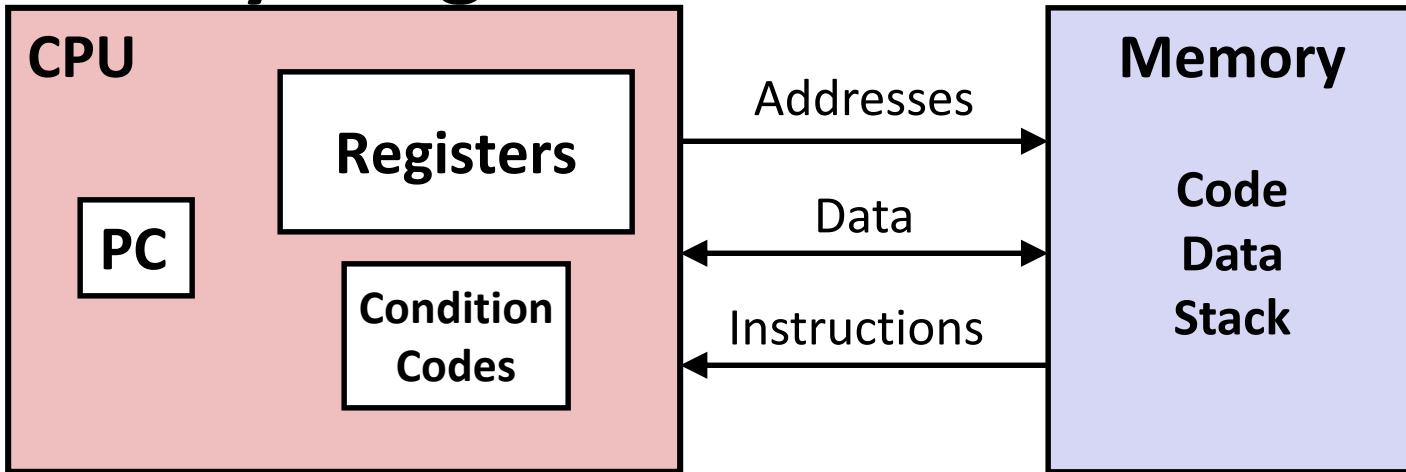
Detour: More About General Purpose Registers



- Why do almost all new architectures use GPRs?
 - Registers are much faster than memory (even cache)
 - Register values are available immediately
 - When memory isn't ready, processor must wait ("stall")
 - Registers are convenient for variable storage
 - Compiler assigns some variables just to registers
 - More compact code since small fields specify registers (compared to memory addresses)



Assembly Programmer's View



- **Programmer-Visible State**

- **PC: Program counter**
 - Address of next instruction
 - Called “RIP” (x86-64)

- **Register file**

- Heavily used program data

- **Condition codes**

- Store status information about most recent arithmetic or logical operation
 - Used for conditional branching

- **Memory**

- Byte addressable array
 - Code and user data
 - Stack to support procedures

Assembly Characteristics: Data Types



- Integral data of 1, 2, 4, or 8 bytes
 - Data values
 - Addresses (untyped pointers)
- Floating point data of 4, 8, 10 or 2x8 or 4x4 or 8x2
 - Different registers for those (*e.g.* %xmm1, %ymm2)
 - Come from *extensions to x86* (SSE, AVX, ...)
- No aggregate types such as arrays or structures
 - Just contiguously allocated bytes in memory
- Two common syntaxes
 - “AT&T”: used by our course, slides, textbook, gnu tools, ...
 - “Intel”: used by Intel documentation, Intel tools, ...
 - Must know which you’re reading



Roadmap

C:

```
car *c = malloc(sizeof(car));  
c->miles = 100;  
c->gals = 17;  
float mpg = get_mpg(c);  
free(c);
```

Java:

```
Car c = new Car();  
c.setMiles(100);  
c.setGals(17);  
float mpg =  
    c.getMPG();
```

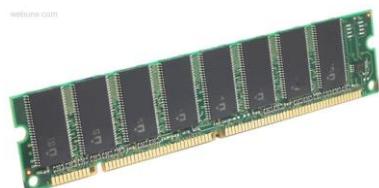
Assembly language:

```
get_mpg:  
    pushq  %rbp  
    movq   %rsp, %rbp  
    ...  
    popq  %rbp  
    ret
```

Machine code:

```
0111010000011000  
100011010000010000000010  
1000100111000010  
110000011111101000011111
```

Computer system:



Memory & data
Integers & floats
x86 assembly
Procedures & stacks
Executables
Arrays & structs
Memory & caches
Processes
Virtual memory
Memory allocation
Java vs. C

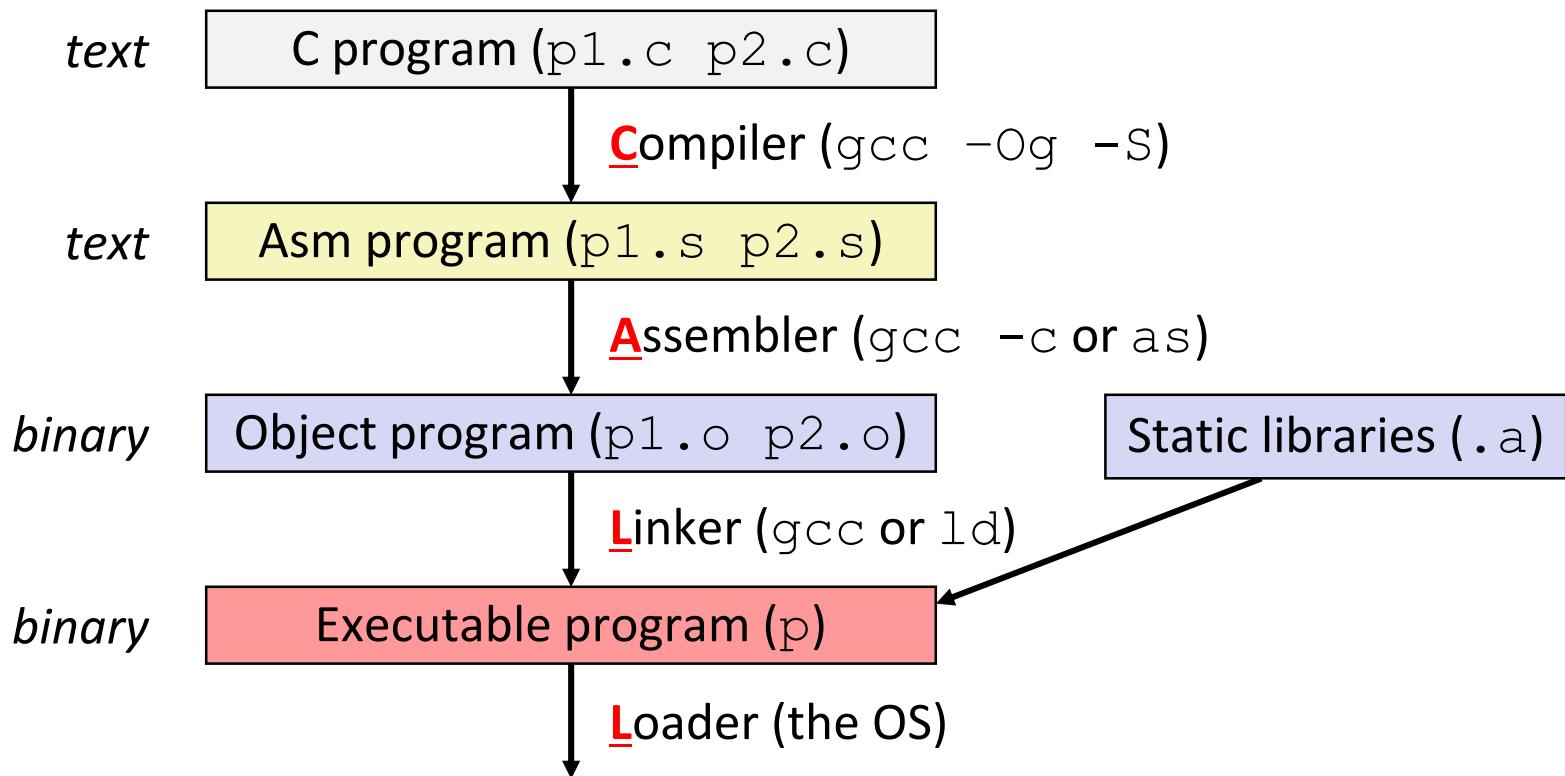
OS:



Building an Executable from a C File



- Code in files `p1.c p2.c`
- Compile with command: `gcc -Og p1.c p2.c -o p`
 - Put resulting machine code in file `p`
- Run with command: `./p`





Compiler

- **Input:** Higher-level language code (*e.g.* C, Java)
 - `foo.c`
- **Output:** Assembly language code (*e.g.* x86, ARM, MIPS)
 - `foo.s`
- First there's a preprocessor step to handle `#directives`
 - Macro substitution, plus other specialty directives
 - If curious/interested: <http://tigcc.ticalc.org/doc/cpp.html>
- Compiler is super complex, whole courses devoted to these!
- Compiler optimizations
 - “Level” of optimization specified by capital ‘O’ flag (*e.g.* `-Og`, `-O3`)
 - Options: <https://gcc.gnu.org/onlinedocs/gcc/Optimize-Options.html>

Compiling Into Assembly



- C Code (`sum.c`)

```
void sumstore(long x, long y, long *dest) {  
    long t = x + y;  
    *dest = t;  
}
```

- x86-64 assembly (`gcc -Og -S sum.c`)

- Generates file `sum.s` (see <https://godbolt.org/g/pQUhlZ>)

```
sumstore(long, long, long*):  
    addq    %rdi, %rsi  
    movq    %rsi, (%rdx)  
    ret
```

Warning: You may get different results with other versions of gcc and different compiler settings

Assembler



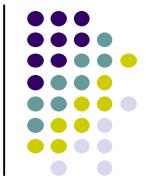
- **Input:** Assembly language code (*e.g.* x86, ARM, MIPS)
 - `foo.s`
- **Output:** Object files (*e.g.* ELF, COFF)
 - `foo.o`
 - Contains *object code* and *information tables*
- Reads and uses *assembly directives*
 - *e.g.* `.text`, `.data`, `.quad`
 - x86: https://docs.oracle.com/cd/E26502_01/html/E28388/eoiyg.html
- Produces “machine language”
 - Does its best, but object file is *not* a completed binary
- **Example:** `gcc -c foo.s`

Producing Machine Language



- **Simple cases:** arithmetic and logical operations, shifts, etc.
 - All necessary information is contained in the instruction itself
- What about the following?
 - Conditional jump
 - Accessing static data (*e.g.* global var or jump table)
 - call
- Addresses and labels are problematic because final executable hasn't been constructed yet!
 - So how do we deal with these in the meantime?

Object File Information Tables



- **Symbol Table** holds list of “items” that may be used by other files
 - *Non-local labels* – function names for `call`
 - *Static Data* – variables & literals that might be accessed across files
- **Relocation Table** holds list of “items” that this file needs the address of later (currently undetermined)
 - Any *label* or piece of *static data* referenced in an instruction in this file
 - Both internal and external
- Each file has its own symbol and relocation tables

Object File Format



- 1) object file header: size and position of the other pieces of the object file
 - 2) text segment: the machine code
 - 3) data segment: data in the source file (binary)
 - 4) relocation table: identifies lines of code that need to be “handled”
 - 5) symbol table: list of this file’s labels and data that can be referenced
 - 6) debugging information
-
- More info: ELF format
 - http://www.skyfree.org/linux/references/ELF_Format.pdf

Linker

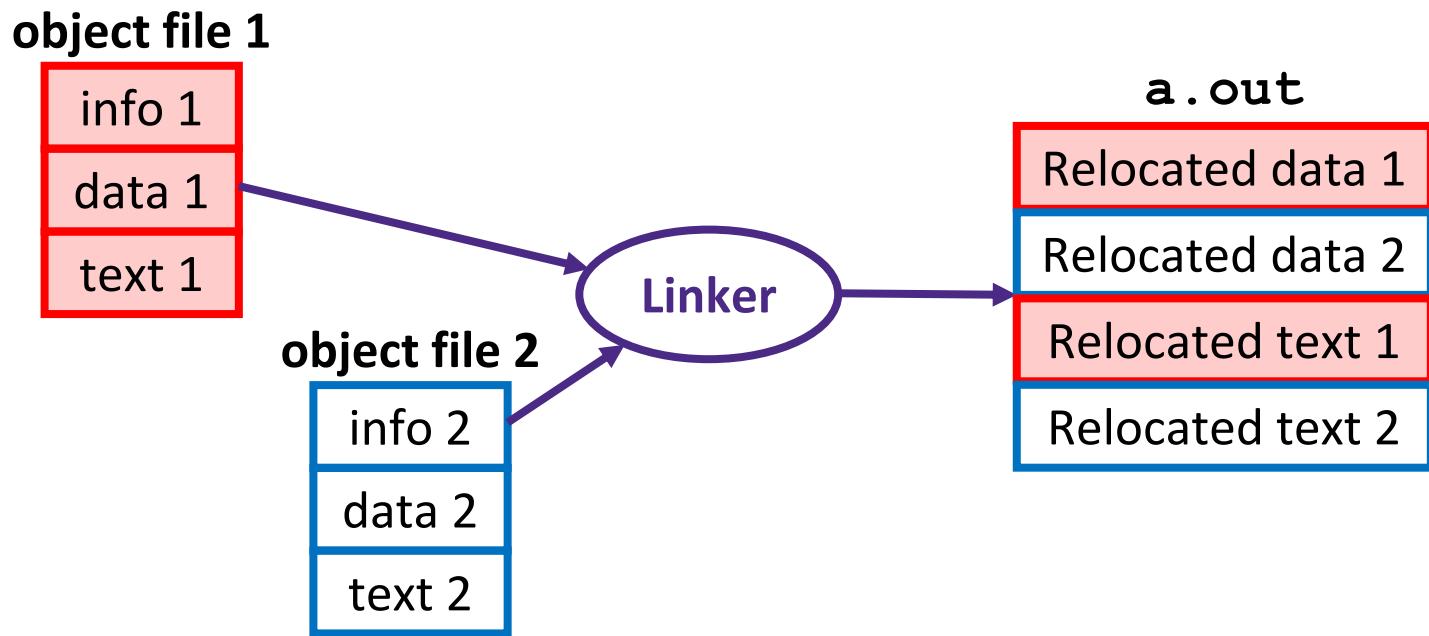


- **Input:** Object files (e.g. ELF, COFF)
 - foo.o
- **Output:** executable binary program
 - a.out
- Combines several object files into a single executable (*linking*)
- Enables separate compilation/assembling of files
 - Changes to one file do not require recompiling of whole program

Linking



- 1) Take text segment from each .o file and put them together
- 2) Take data segment from each .o file, put them together, and concatenate this onto end of text segments
- 3) Resolve References
 - Go through Relocation Table; handle each entry



Machine Instruction Example



```
*dest = t;
```

- C Code

- Store value **t** where designated by **dest**

```
movq %rax, (%rbx)
```

- Assembly

- Move 8-byte value to memory
 - Quad words in x86-64 parlance
- Operands:
 - t:** Register **%rax**
 - dest:** Register **%rbx**
 - *dest:** Memory **M[%rbx]**

```
0x40059e: 48 89 03
```

- Object Code

- 3-byte instruction
- Stored at address **0x40059e**

Object Code



Code for sumstore

```
0x0400536:  
 0x48  
 0x01  
 0xfe  
 0x48  
 0x89  
 0x32  
 0xc3
```

- Total of 7 bytes
- Each instruction
1, 3, or 5 bytes
(not shown)
- Starts at address
0x0400536

Disassembling Object Code



- Disassembled:

```
0000000000400536 <sumstore>:  
400536: 36 37 38 add    %rdi, %rsi  
400539: 39 3a 3b mov    %rsi, (%rdx)  
40053c: 3c                retq  
  
address of instruction   object code bytes (hex)   interpreted assembly instructions
```

- **Disassembler** (`objdump -d sum`)

- Useful tool for examining object code (`man 1 objdump`)
- Analyzes bit pattern of series of instructions
- Produces approximate rendition of assembly code
- Can run on either a `.out` (complete executable) or `.o` file



Alternate Disassembly in GDB

```
$ gdb sum
(gdb) disassemble sumstore
Dump of assembler code for function sumstore:
0x0000000000400536 <+0>:    add    %rdi,%rsi
0x0000000000400539 <+3>:    mov    %rsi,(%rdx)
0x000000000040053c <+6>:    retq
End of assembler dump.

(gdb) x/7bx sumstore
0x400536 <sumstore>: 0x48  0x01  0xfe  0x48  0x89  0x32  0xc3
```

- Within gdb debugger (gdb sum):
 - disassemble sumstore: disassemble procedure
 - x/7bx sumstore: show 7 bytes starting at sumstore⁴²

What Can be Disassembled?



```
% objdump -d WINWORD.EXE

WINWORD.EXE:      file format pei-i386

No symbols in "WINWORD.EXE".
Disassembly of section .text:

30001000 <.text>:
30001000:
30001001:
30001003:
30001005:
3000100a:          Reverse engineering forbidden by
                           Microsoft End User License Agreement
```

- Anything that can be interpreted as executable code
- Disassembler examines bytes and reconstructs assembly source

Loader



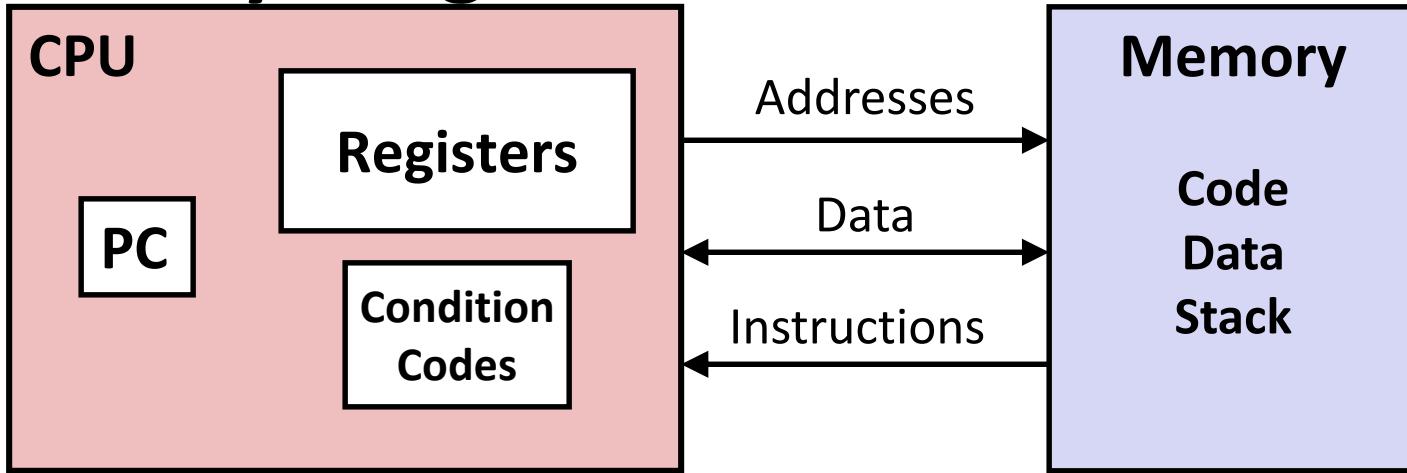
- **Input:** executable binary program, command-line arguments
 - ./.out arg1 arg2
- **Output:** <program is run>
- Loader duties primarily handled by OS/kernel
 - More about this when we learn about processes
- Memory sections (Instructions, Static Data, Stack) are set up
- Registers are initialized

Machine Programming I: Basics



- History of Intel processors and architectures
- C, assembly, machine code
- **Assembly Basics: Registers, operands, move**
- Arithmetic & logical operations

Assembly Programmer's View



- **Programmer-Visible State**

- **PC: Program counter**
 - Address of next instruction
 - Called “RIP” (x86-64)

- **Register file**

- Heavily used program data

- **Condition codes**

- Store status information about most recent arithmetic or logical operation
 - Used for conditional branching

- **Memory**

- Byte addressable array
 - Code and user data
 - Stack to support procedures

Assembly Characteristics: Data Types



- Integral data of 1, 2, 4, or 8 bytes
 - Data values
 - Addresses (untyped pointers)
- Floating point data of 4, 8, 10 or 2x8 or 4x4 or 8x2
 - Different registers for those (*e.g.* %xmm1, %ymm2)
 - Come from *extensions to x86* (SSE, AVX, ...)
- No aggregate types such as arrays or structures
 - Just contiguously allocated bytes in memory
- Two common syntaxes
 - “AT&T”: used by our course, slides, textbook, gnu tools, ...
 - “Intel”: used by Intel documentation, Intel tools, ...
 - Must know which you’re reading

Assembly Characteristics: Operations



- Perform arithmetic function on register or memory data
- Transfer data between memory and register
 - Load data from memory into register
 - Store register data into memory
- Transfer control
 - Unconditional jumps to/from procedures
 - Conditional branches

What is a Register?



- A location in the CPU that stores a small amount of data, which can be accessed very quickly (once every clock cycle)
- Registers have *names*, not *addresses*
 - In assembly, they start with % (e.g. %rsi)
- Registers are at the heart of assembly programming
 - They are a precious commodity in all architectures, but *especially* x86

x86-64 Integer Registers – 64 bits wide



%rax	%eax
%rbx	%ebx
%rcx	%ecx
%rdx	%edx
%rsi	%esi
%rdi	%edi
%rsp	%esp
%rbp	%ebp

%r8	%r8d
%r9	%r9d
%r10	%r10d
%r11	%r11d
%r12	%r12d
%r13	%r13d
%r14	%r14d
%r15	%r15d

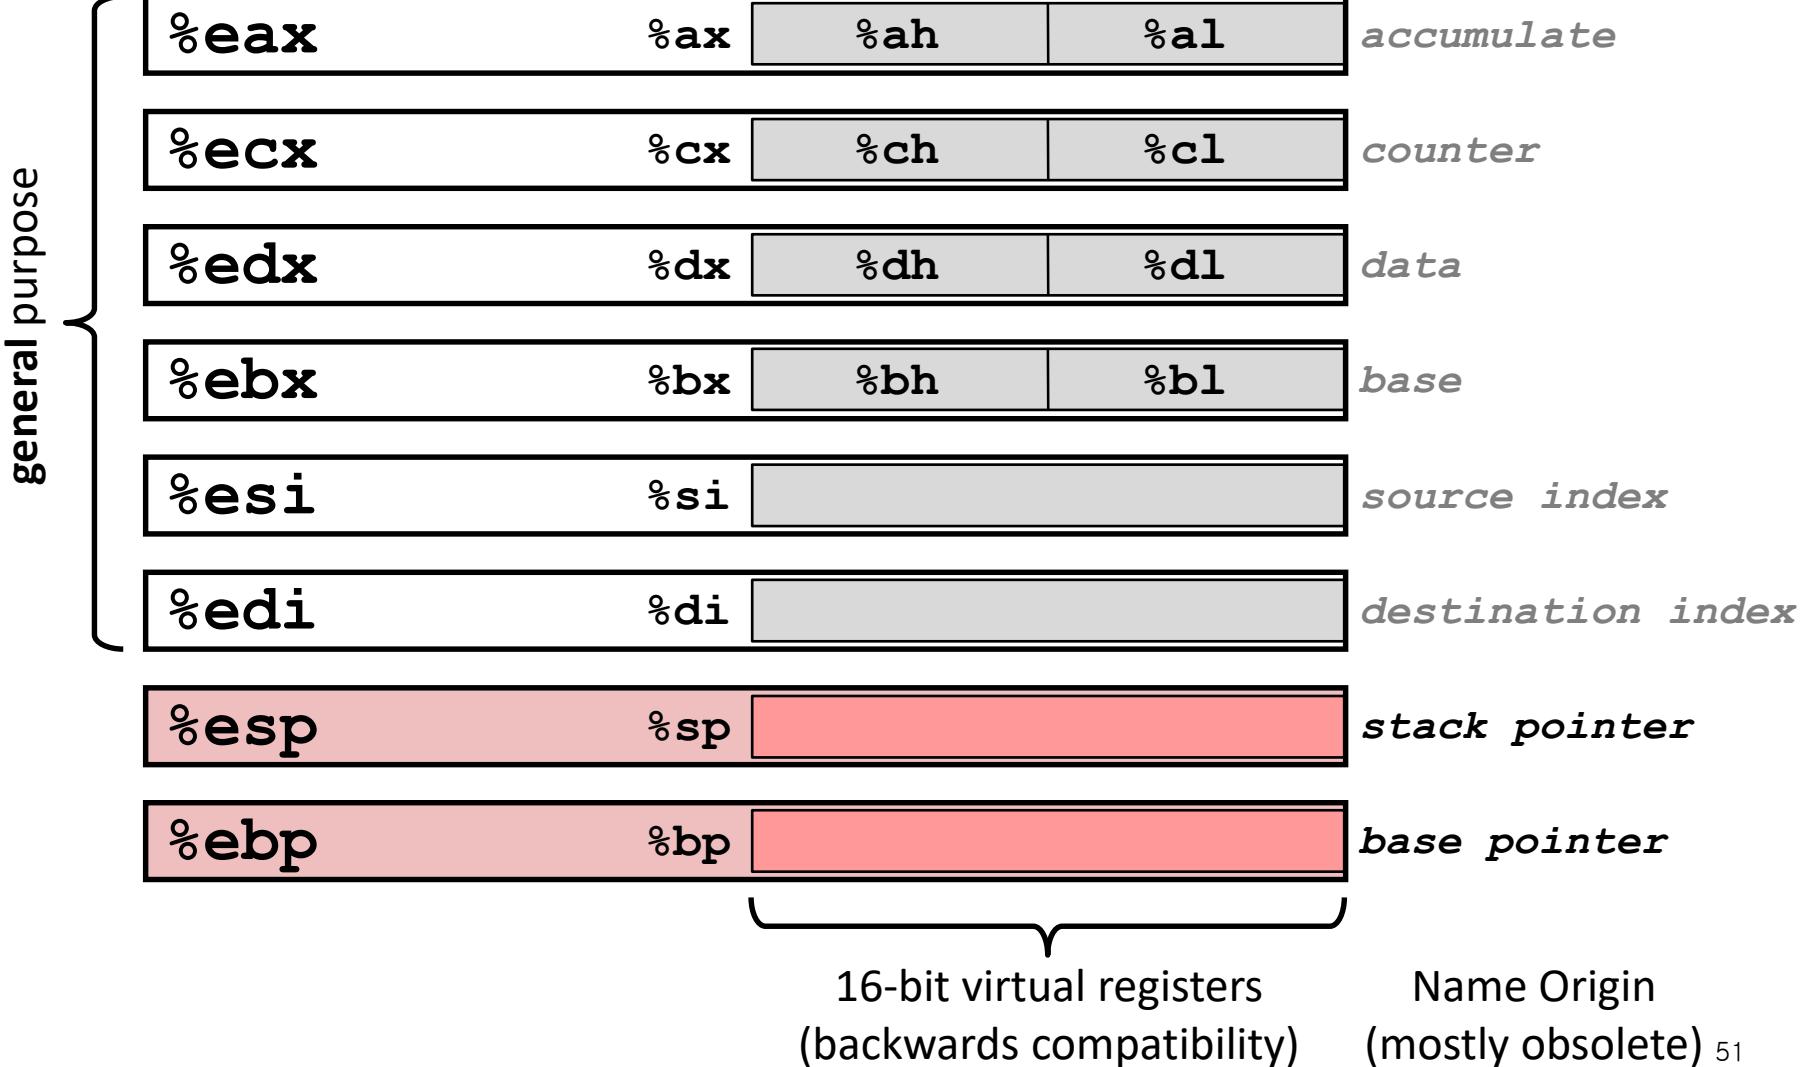
- Can reference low-order 4 bytes (also low-order 1 & 2 bytes)⁵⁰

Some History: IA32 Registers – 32 bits wide



Origin

(mostly obsolete)



Memory vs. Registers



- Addresses **vs.** Names
 - 0x7FFFD024C3DC %rdi
- Big **vs.** Small
 - ~ 8 GiB $(16 \times 8 \text{ B}) = 128 \text{ B}$
- Slow **vs.** Fast
 - ~50-100 ns sub-nanosecond timescale
- Dynamic **vs.** Static
 - Can “grow” as needed
while program runs fixed number in hardware

Three Basic Kinds of Instructions



- Transfer data between memory and register
 - **Load** data from memory into register
 - $\%reg = \text{Mem}[\text{address}]$
 - **Store** register data into memory
 - $\text{Mem}[\text{address}] = \%reg$
- Perform arithmetic function on register or memory data
 - $c = a + b;$ $z = x \ll y;$ $i = h \& g;$
- Control flow: what instruction to execute next
 - Unconditional jumps to/from procedures
 - Conditional branches

Remember: memory
is indexed just like
an array[] of bytes!



Operand types

- **Immediate:** Constant integer data
 - Examples: `$0x400`, `$-533`
 - Like C literal, but prefixed with ‘\$’
 - Encoded with 1, 2, 4, or 8 bytes
depending on the instruction
- **Register:** 1 of 16 integer registers
 - Examples: `%rax`, `%r13`
 - But `%rsp` reserved for special use
 - Others have special uses for particular instructions
- **Memory:** Consecutive bytes of memory at a computed address
 - Simplest example: (`%rax`)
 - Various other “address modes”

`%rax`

`%rcx`

`%rdx`

`%rbx`

`%rsi`

`%rdi`

`%rsp`

`%rbp`

`%rN`

Moving Data



- General form: `mov_ source, destination`
 - Missing letter (_) specifies size of operands
 - Note that due to backwards-compatible support for 8086 programs (16-bit machines!), “word” means 16 bits = 2 bytes in x86 instruction names
 - Lots of these in typical code

- ❖ `movb src, dst`
 - Move 1-byte “byte”
- ❖ `movw src, dst`
 - Move 2-byte “word”
- ❖ `movl src, dst`
 - Move 4-byte “long word”
- ❖ `movq src, dst`
 - Move 8-byte “quad word”⁵⁵

movq Operand Combinations



Source	Dest	Src, Dest	C Analog
movq	Imm	Reg movq \$0x4, %rax Mem movq \$-147, (%rax)	var_a = 0x4; *p_a = -147;
	Reg	Reg movq %rax, %rdx Mem movq %rax, (%rdx)	var_d = var_a; *p_d = var_a;
	Mem	Reg movq (%rax), %rdx	var_d = *p_a;

Cannot do memory-memory transfer with a single instruction

How would you do it?

Question



- Which of the following statements is TRUE?
 - A. For float f , $(f+2 == f+1+1)$ always returns TRUE
 - B. The width of a “word” is part of a system’s *architecture* (as opposed to *microarchitecture*)
 - C. Having more registers increases the performance of the hardware, but decreases the performance of the software
 - D. Mem to Mem (src to dst) is the only disallowed operand combination in x86-64

Summary



- x86-64 is a complex instruction set computing (CISC) architecture
- **Registers** are named locations in the CPU for holding and manipulating data
 - x86-64 uses 16 64-bit wide registers
- Assembly operands include immediates, registers, and data at specified memory locations
- History of Intel processors and architectures
 - Evolutionary design leads to many quirks and artifacts
- Assembly Basics: Registers, operands, move
 - The x86-64 move instructions cover wide range of data movement forms

Q&A

