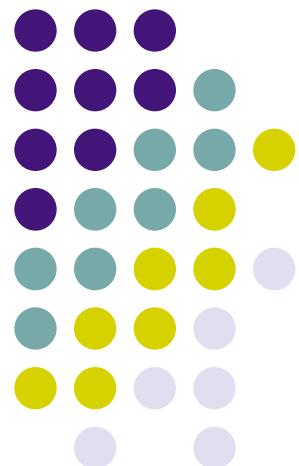


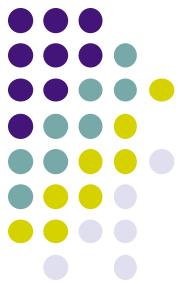
# Arrays and Pointer. Part 1

2019 Spring



# Lecture Outline

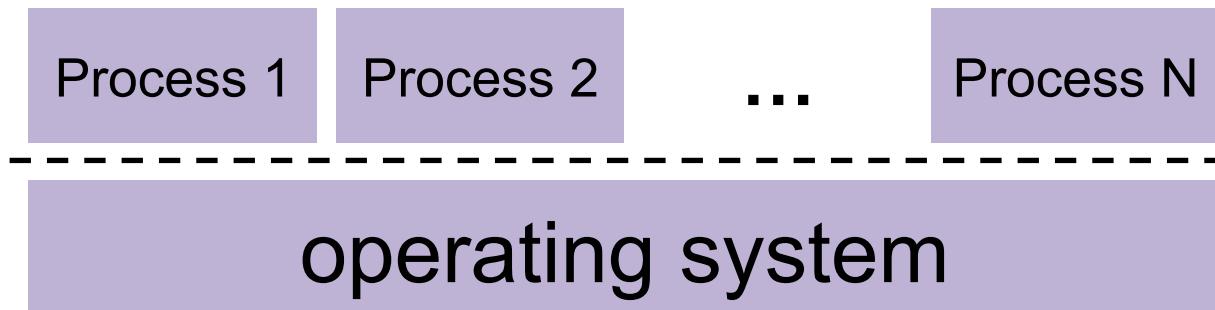
- C's Memory Model (refresher)
- Pointers (refresher)
- Arrays



# OS and Processes



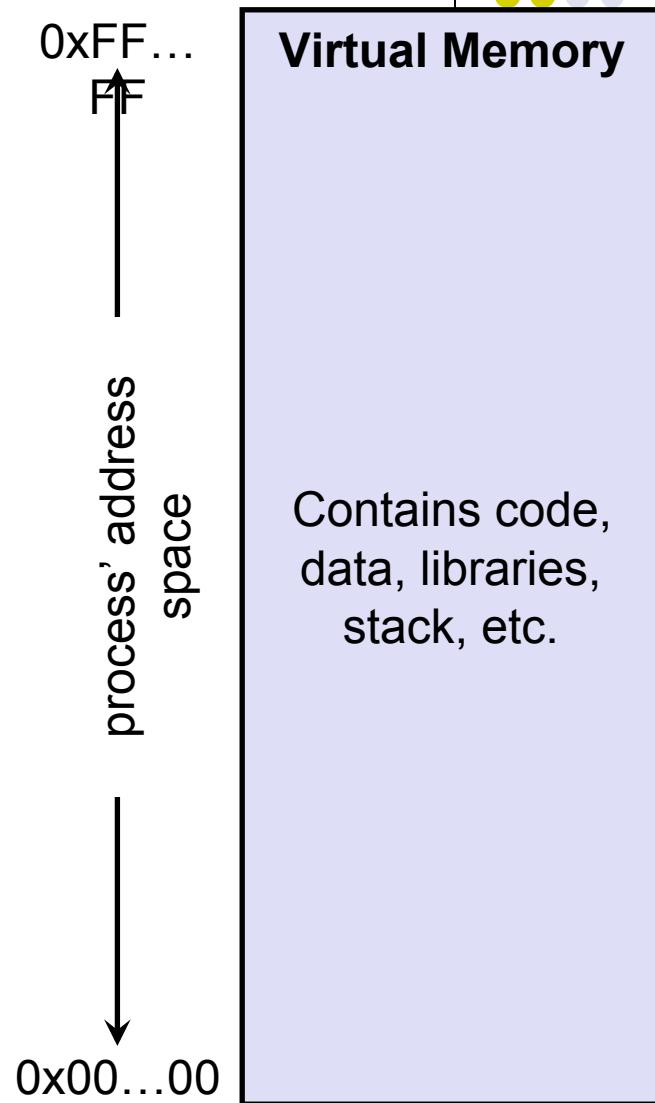
- The OS lets you run multiple applications at once
  - An application runs within an OS “process”
  - The OS timeslices each CPU between runnable processes
    - This happens *very quickly*: ~100 times per second



# Processes and Virtual Memory



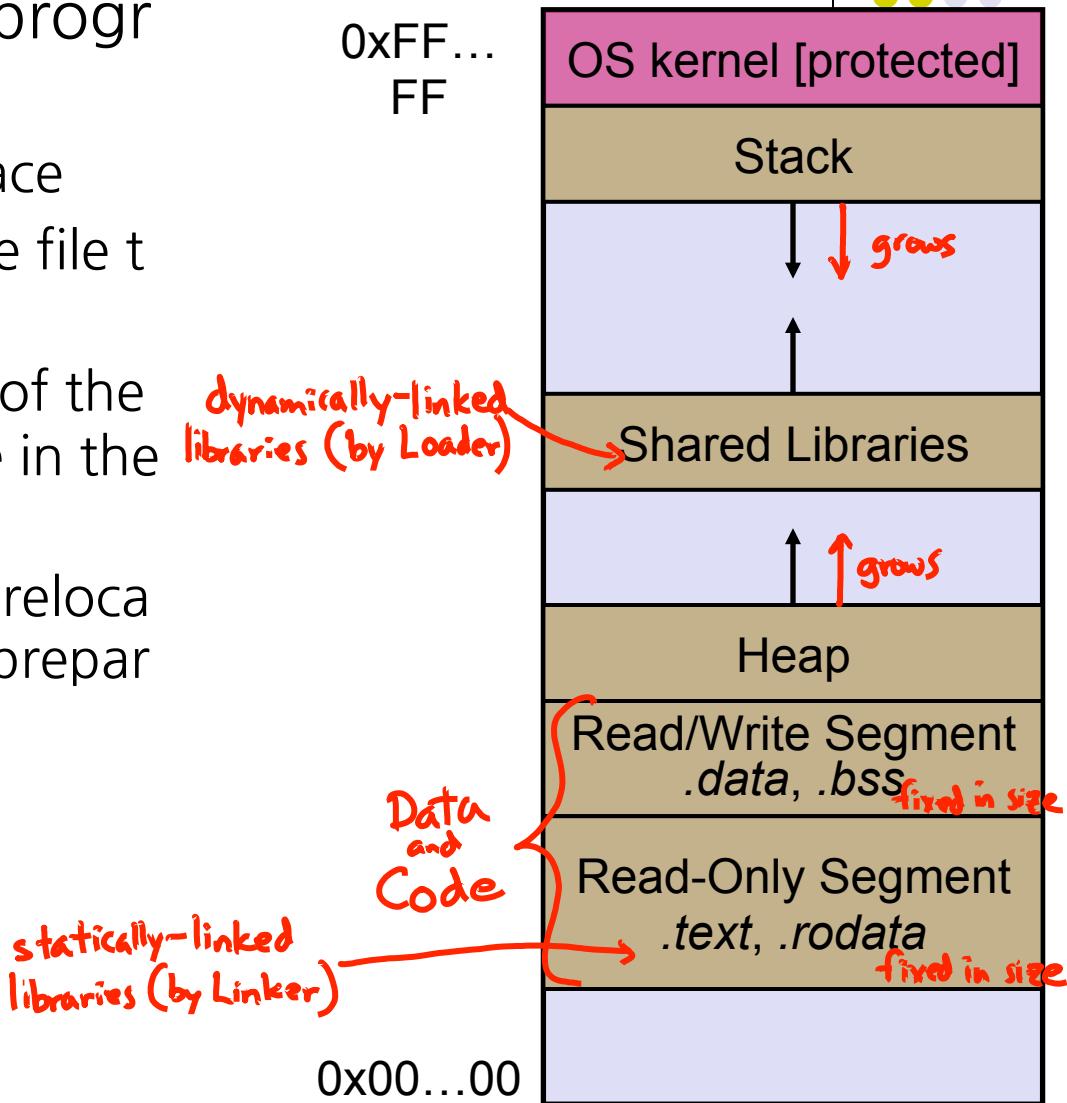
- The OS gives each process the illusion of its own private memory
  - Called the process' **address space**
  - Contains the process' virtual memory, visible only to it (via translation)
  - $2^{64}$  bytes on a 64-bit machine



# Loading

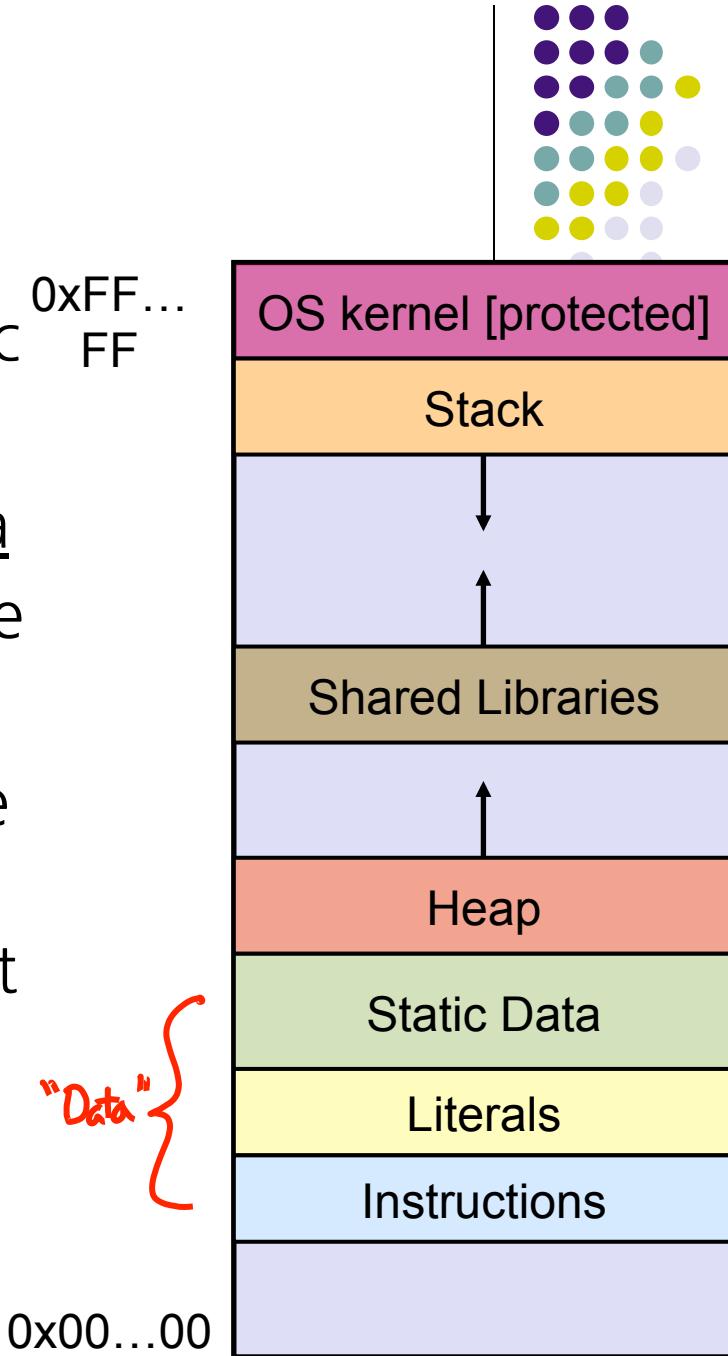


- When the OS loads a program it:
  - Creates an address space
  - Inspects the executable file to see what's in it
  - (Lazily) copies regions of the file into the right place in the address space
  - Does any final linking, relocation, or other needed preparation



# Memory Management

- *Local* variables on the Stack
  - Allocated and freed via calling conventions (push, pop, mov)
- *Global* and *static* variables in Data
  - Allocated/freed when the process starts/exits
- *Dynamically-allocated* data on the Heap
  - `malloc()` to request; `free()` to free, otherwise **memory leak**



# Stack in Action (1/3)

Note: arrow points to *next* instruction to be executed (like in `gdb`).



stack.c

```
#include <stdint.h>

int f(int, int);
int g(int);

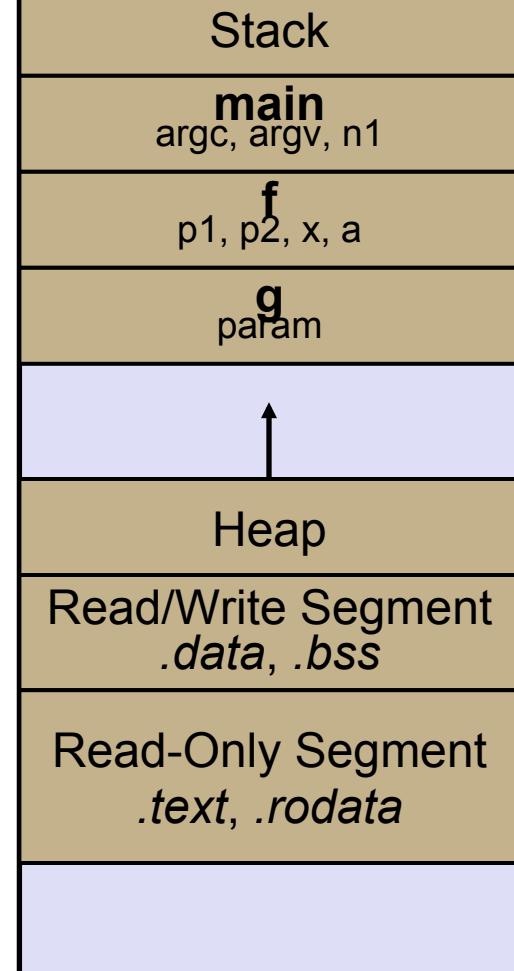
int main(int argc, char** argv) {
    int n1 = f(3, -5);
    n1 = g(n1);
}

int f(int p1, int p2) {
    int x;
    int a[3];
    ...
    x = g(a[2]);
    return x;
}

int g(int param) {
    return param * 2;
}
```



OS kernel [protected]



# Stack in Action (2/3)

Note: arrow points to *next* instruction to be executed (like in `gdb`).



## stack.c

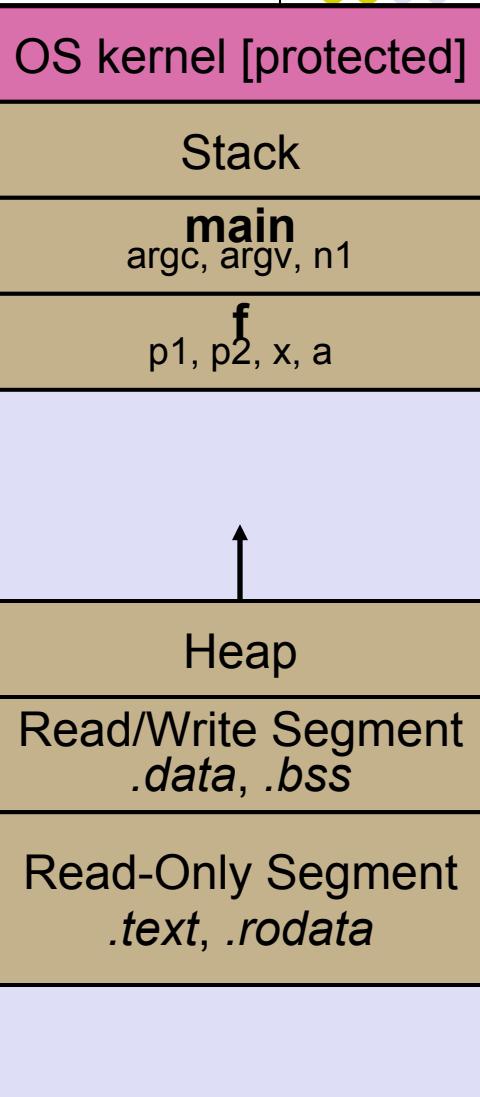
```
#include <stdint.h>

int f(int, int);
int g(int);

int main(int argc, char** argv) {
    int n1 = f(3, -5);
    n1 = g(n1);
}

int f(int p1, int p2) {
    int x;
    int a[3];
    ...
    x = g(a[2]);
    return x;
}

int g(int param) {
    return param * 2;
}
```



# Stack in Action (3/3)

Note: arrow points to *next* instruction to be executed (like in `gdb`).



```
#include <stdint.h>

int f(int, int);
int g(int);

int main(int argc, char** argv) {
    int n1 = f(3, -5);
    n1 = g(n1);
}

int f(int p1, int p2) {
    int x;
    int a[3];
    ...
    x = g(a[2]);
    return x;
}

int g(int param) {
    return param * 2;
}
```



OS kernel [protected]

Stack

main  
argc, argv, n1

g  
param



Heap

ReadWrite Segment  
.data, .bss

Read-Only Segment  
.text, .rodata



# Lecture Outline

- C's Memory Model (refresher)
- **Pointers** (refresher)
- Arrays

# Pointers



- Variables that store addresses
  - It points to somewhere in the process' virtual address space
  - `&foo` produces the virtual address of `foo`
- Generic definition:
  - Recommended to not define multiple pointers on same line:  
`int *p1, p2;` not the same as `int *p1, *p2;`
  - Instead, use:
- *Dereference* a pointer using the unary `*` operator
  - Access the memory referred to by a pointer

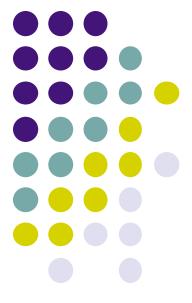
`type* name;`

`type *name;`

`int *p1, p2;`

`int *p1, *p2;`

`int *p1;`  
`int *p2;`



# Pointer Example

pointy.c

```
#include <stdio.h>
#include <stdint.h>

int main(int argc, char** argv) {
    int x = 351;
    int* p;          // p is a pointer to a int

    p = &x;          // p now contains the addr of x
    printf("&x is %p\n", &x);
    printf(" p is %p\n", p);
    printf(" x is %d\n", x);

    *p = 333;        // change value of x
    printf(" x is %d\n", x);

    return 0;
}
```

# Something Curious



- What happens if we run `pointy.c` several times?

```
$ gcc -Wall -std=c11 -o pointy pointy.c
```

Run 1: **bash\$ ./pointy**

```
&x is 0x7fffe9e28524
p is 0x7ffff9e28524
x is 351
x is 333
```

Run 2: **bash\$ ./pointy**

```
&x is 0x7fffe847be34
p is 0x7ffe847be34
x is 351
x is 333
```

Run 3: **bash\$ ./pointy**

```
&x is 0x7fffe7b14644
p is 0x7ffe7b14644
x is 351
x is 333
```

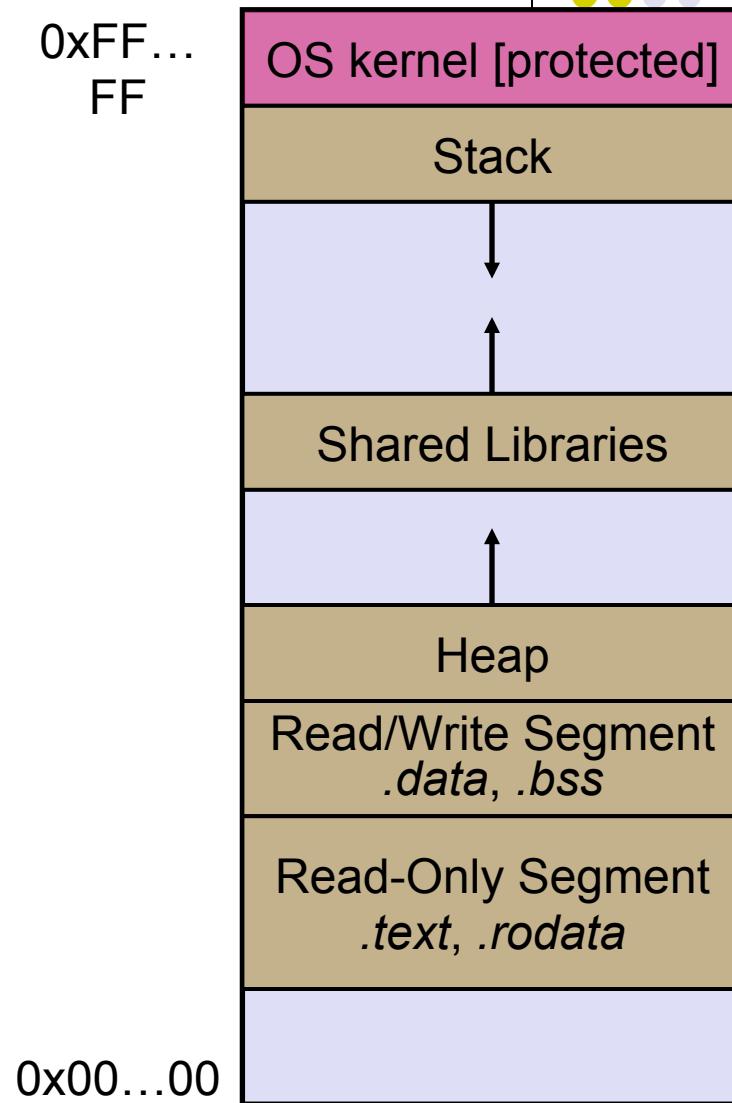
Run 4: **bash\$ ./pointy**

```
&x is 0x7fffef0dfe54
p is 0x7ffff0dfe54
x is 351
x is 333
```

# Address Space Layout Randomization



- Linux uses *address space layout randomization* (ASLR) for added security
  - Randomizes:
    - Base of stack
    - Shared library (`mmap`) location
  - Makes Stack-based buffer overflow attacks tougher
  - Makes debugging tougher
  - Can be disabled (`gdb` does this by default); Google if curious





# Lecture Outline

- C's Memory Model (refresher)
- Pointers (refresher)
- **Arrays**

# Arrays



- Definition: `type name [size]`
  - Allocates `size*sizeof(type)` bytes of *contiguous* memory
  - Normal usage is a **compile-time constant** for `size` (e.g. `int scores[175];`)
  - Initially, array values are “garbage”



# Size of an array

- Size of an array
  - Not stored anywhere - array does not know its own size!
    - `sizeof`(array) only works in variable scope of array definition
  - Recent versions of C allow for variable-length arrays
    - Uncommon and can be considered bad practice [*we won't use*]
      - <https://gcc.gnu.org/onlinedocs/gcc/Variable-Length.html>

```
int n = 175;
int scores[n]; // OK in C99
```

# Challenge Question

- The code snippets both use a ~~variable-length array~~. What will happen when we compile with C99?



should malloc  
instead of using  
vla's!

allocated in  
Static Data  
(can't change  
size)

```
int m = 175;
int scores[m];

void foo(int n) {
    ...
}
```

```
int m = 175;

void foo(int n) {
    int scores[n];
    ...
}
```

allocated on  
the Stack  
(can grow)

however, you  
don't want to  
put large arrays  
on the Stack

A. Compiler Error  
B. Compiler Error

**C. No Error**

D. No Error

E. We're lost...

Compiler Error  
No Error

**Compiler Error**

No Error

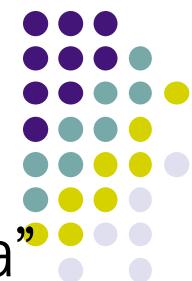


# Using Arrays (1/2)

- Initialization

```
type name[size] = {val0, ..., valN};
```

- {} initialization can *only* be used at time of definition
- If no **size** supplied, infers from length of array initializer



# Using Arrays (2/2)

- Array name used as identifier for “collection of data”
  - name [index] specifies an element of the array and can be used as an assignment target or as a value in an expression
  - Array name (by itself) produces the address of the start of the array
    - Cannot be assigned to / changed

```
int primes[6] = {2, 3, 5, 6, 11, 13};  
primes[3] = 7;  
primes[100] = 0; // memory smash!
```



# Multi-dimensional Arrays

- Generic 2D format:

```
type name[rows][cols] = { {values}, ..., {values} } ;
```

- Still allocates a single, contiguous chunk of memory
- C is *row-major*

```
// a 2-row, 3-column array of doubles
double grid[2][3];

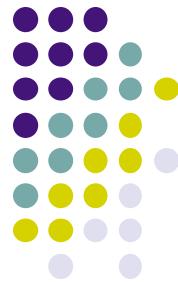
// a 3-row, 5-column array of ints
int matrix[3][5] = {
    {0, 1, 2, 3, 4},
    {0, 2, 4, 6, 8},
    {1, 3, 5, 7, 9}
};
```

# Parameters: reference vs. value (1/2)

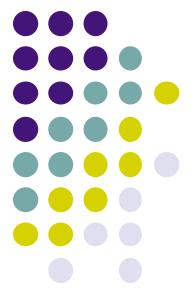


- There are two fundamental parameter-passing schemes in programming languages
- **Call-by-value**
  - Parameter is a local variable initialized when the function is called and gets a copy of the calling argument; manipulating the parameter only changes copy, *not* the calling argument
  - C, Java, C++ primitives

# Parameters: reference vs. value (2/2)



- Call-by-reference
  - Parameter is an alias for the supplied argument; manipulating the parameter manipulates the calling argument
  - C++ arrays and references (we'll see more later)



# Arrays as Parameters

- It's tricky to use arrays as parameters
  - What happens when you use an array name as an argument?
  - Arrays do not know their own size

```
int sumAll(int a[]); // prototype

int main(int argc, char** argv) {
    int numbers[] = {9, 8, 1, 9, 5};
    int sum = sumAll(numbers);
    return 0;
}

int sumAll(int a[]) {
    int i, sum = 0;
    for (i = 0; i < ...????
}
```



# Solution 1: Declare Array Size

```
int sumAll(int a[5]); // prototype

int main(int argc, char** argv) {
    int numbers[] = {9, 8, 1, 9, 5};
    int sum = sumAll(numbers);
    printf("sum is: %d\n", sum);
    return 0;
}

int sumAll(int a[5]) {
    int i, sum = 0;
    for (i = 0; i < 5; i++) {
        sum += a[i];
    }
    return sum;
}
```

- Problem: loss of generality/flexibility!

# Solution 2: Pass Size as Parameter



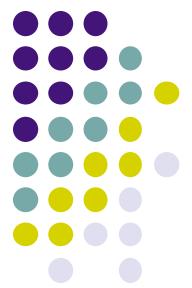
```
int sumAll(int a[], int size); //  
prototype  
  
int main(int argc, char** argv) {  
    int numbers[] = {9, 8, 1, 9, 5};  
    int sum = sumAll(numbers, 5);  
    printf("sum is: %d\n", sum);  
    return 0;  
}  
  
int sumAll(int a[], int size) {  
    int i, sum = 0;  
    for (i = 0; i < size; i++) {  
        sum += a[i];  
    }  
    return sum;  
}
```



# Returning an Array

- Local variables, including arrays, are allocated on the stack
  - They “disappear” when a function returns!
  - Can’t safely return local arrays from functions
    - Can’t return an array as a return value - why not?

```
int* copyArray(int src[], int size) {  
    int i, dst[size];    // OK in C99  
  
    for (i = 0; i < size; i++) {  
        dst[i] = src[i];  
    }  
  
    return dst; // no compiler error, but wrong!  
}
```



# Solution: Output Parameter

- Create the “returned” array in the caller
  - Pass it as an **output parameter** to `copyarray()`
    - A pointer parameter that allows the callee to leave values for the caller to use
  - Works because arrays are “passed” as pointers
    - “Feels” like call-by-reference, *but it’s not*

```
void copyArray(int src[], int dst[], int size) {  
    int i;  
  
    for (i = 0; i < size; i++) {  
        dst[i] = src[i];  
    }  
}
```



# Output Parameters

- Output parameters are common in library functions
  - `long int strtol(char* str, char** endptr, int base);`
  - `int sscanf(char* str, char* format, ...);`

```
int    num, i;
char* pEnd, str1 = "333 rocks";
char  str2[10];

// converts "333 rocks" into long -- pEnd is
// conversion end
num = (int) strtol(str1, &pEnd, 10);

// reads string into arguments based on format string
num = sscanf("3 blind mice", "%d %s", &i, str2);
```

outparam.c

# Questions?

