

# System Programming

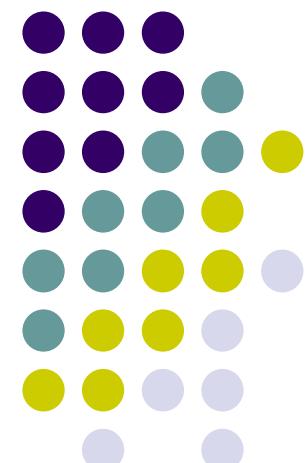
## 08. Machine-Level Programming III: Procedures

2019. Fall

Instructor: Joonho Kwon

[jhwon@pusan.ac.kr](mailto:jhwon@pusan.ac.kr)

Data Science Lab @ PNU



# Roadmap



Memory & data  
Integers & floats  
x86 assembly  
**Procedures & stacks**  
Executables  
Arrays & structs  
Memory & caches  
Processes  
Virtual memory  
Memory allocation  
Java vs. C

```
car *c = malloc(sizeof(car));  
c->miles = 100;  
c->gals = 17;  
float mpg = get_mpg(c);  
free(c);
```

Java:

```
Car c = new Car();  
c.setMiles(100);  
c.setGals(17);  
float mpg =  
    c.getMPG();
```

Assembly  
language:

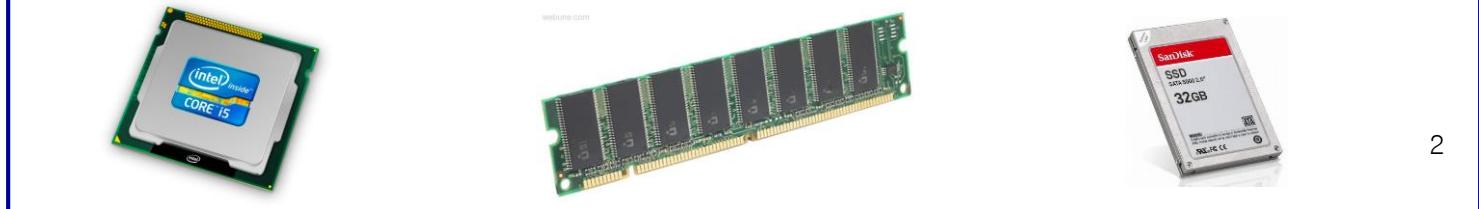
```
get_mpg:  
    pushq  %rbp  
    movq   %rsp, %rbp  
    ...  
    popq   %rbp  
    ret
```

Machine  
code:

```
0111010000011000  
100011010000010000000010  
1000100111000010  
1100000111110100001111
```

Computer  
system:

OS:



# Mechanisms required for *procedures*



## 1) Passing control

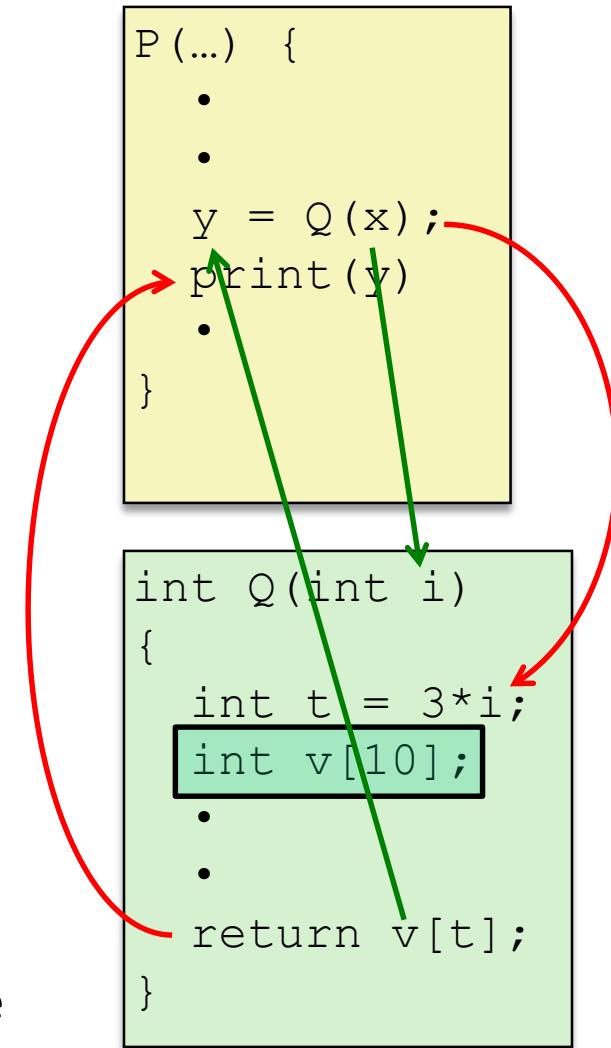
- To beginning of procedure code
- Back to return point

## 2) Passing data

- Procedure arguments
- Return value

## 3) Memory management

- Allocate during procedure execution
- Deallocate upon return
- All implemented with machine instructions!
- An x86-64 procedure uses only those mechanisms required for that procedure



# Questions to answer about procedures



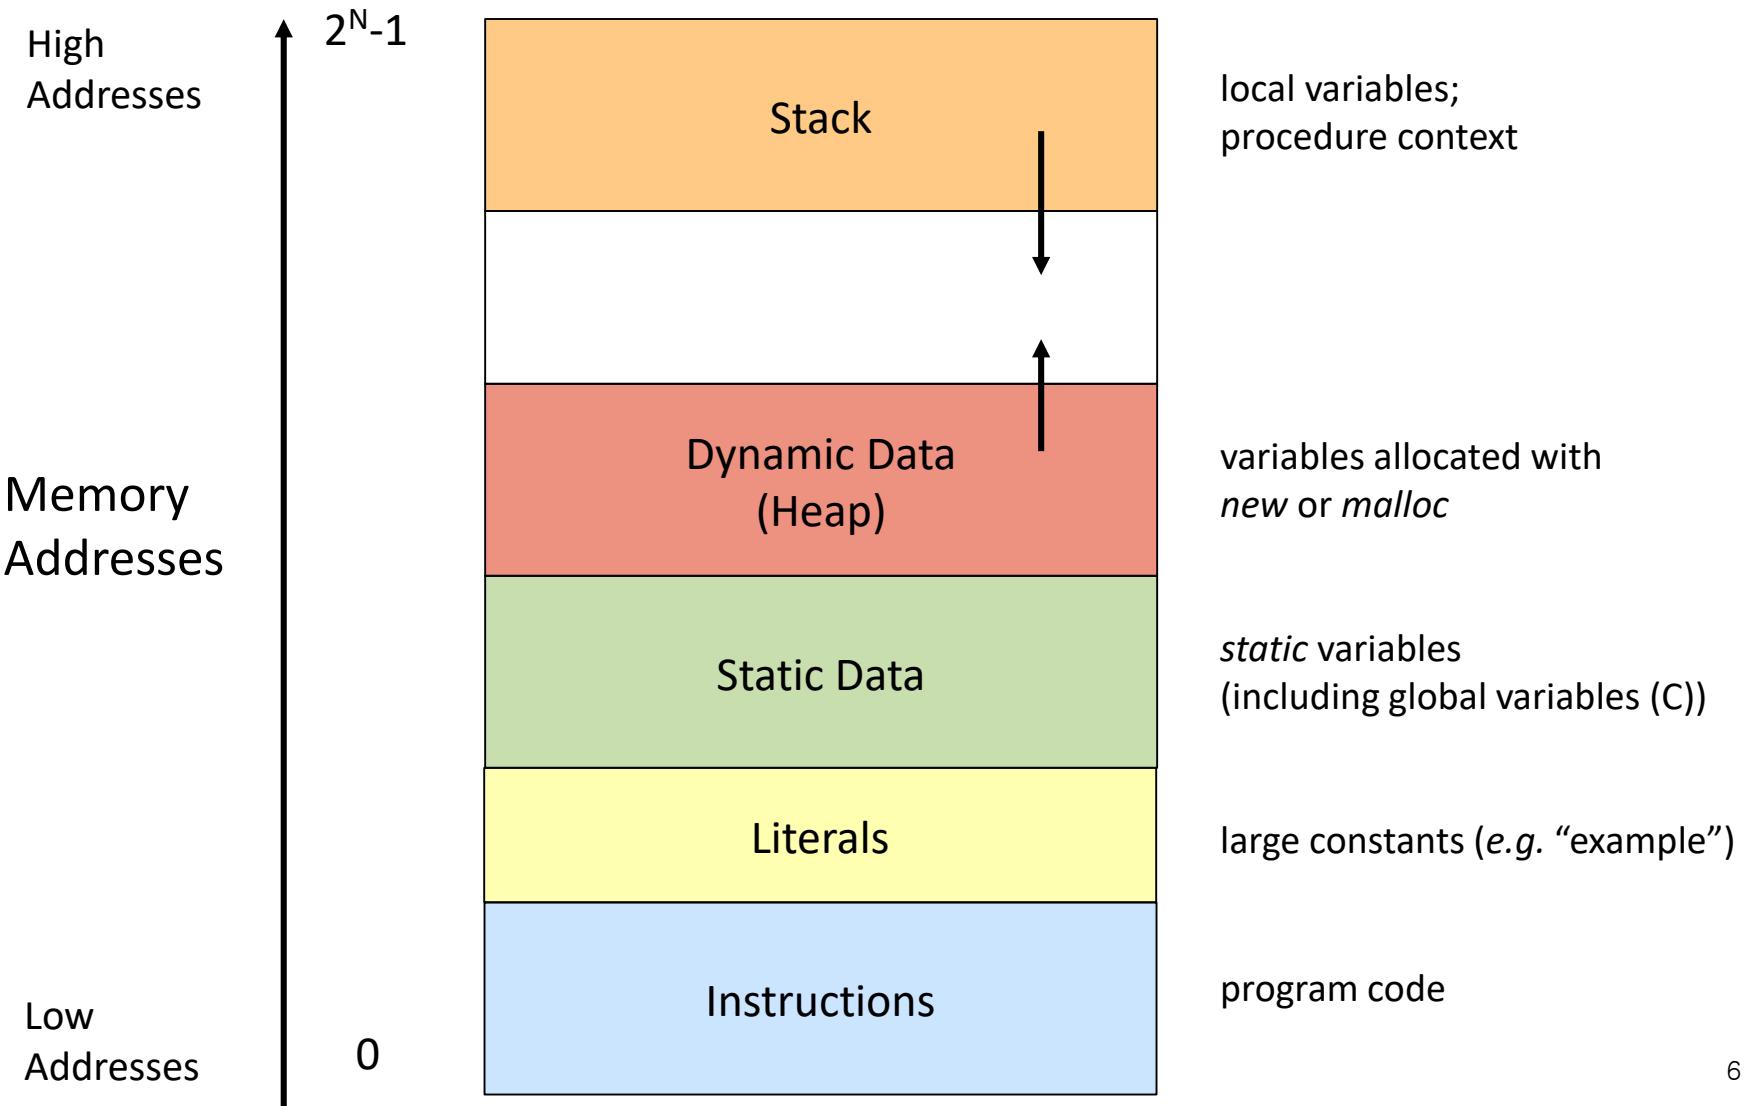
- How do I pass arguments to a procedure?
- How do I get a return value from a procedure?
- Where do I put local variables?
- When a function returns, how does it know where to return?
  - To answer some of these questions,  
we need a **call stack** ...

# Outline



- **Stack Structure**
- Calling Conventions
  - Passing control
  - Passing data
  - Managing local data
- Register Saving Conventions
- Illustration of Recursion

# Memory Layout



# Memory Permissions

segmentation faults?

writable; not executable

Stack

Managed “automatically”  
(by compiler)

writable; not executable

Dynamic Data  
(Heap)

Managed by programmer

writable; not executable

Static Data

Initialized when process starts

read-only; not executable

Literals

Initialized when process starts

read-only; executable

Instructions

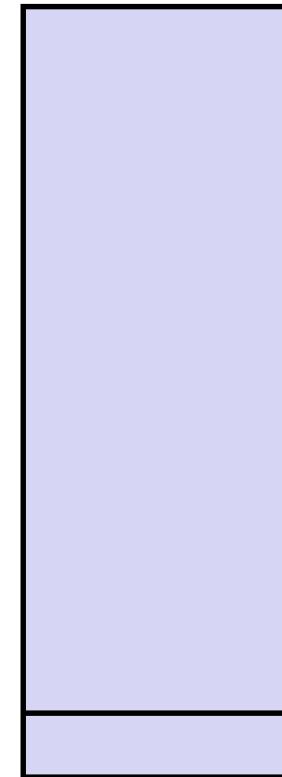
Initialized when process starts

# x86-64 Stack

- Region of memory managed with stack “discipline”
  - Grows toward lower addresses
  - Customarily shown “upside-down”
- Register `%rsp` contains *lowest* stack address
  - `%rsp` = address of *top* element, the most-recently-pushed item that is not-yet-popped

**Stack Pointer:** `%rsp` →

**Stack “Bottom”**

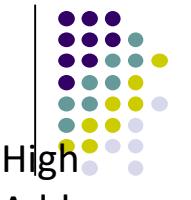


High Addresses

Increasing Addresses

Stack Grows Down

Low Addresses  
0x00...00



# x86-64 Stack: Push

- `pushq src`
  - Fetch operand at *src*
    - *Src* can be reg, memory, immediate
  - **Decrement**  $\%rsp$  by 8
  - Store value at address given by  $\%rsp$
- Example:
  - `pushq %rcx`
  - Adjust  $\%rsp$  and store contents of  $\%rcx$  on the stack

**Stack Pointer:**  $\%rsp$



**Stack “Bottom”**



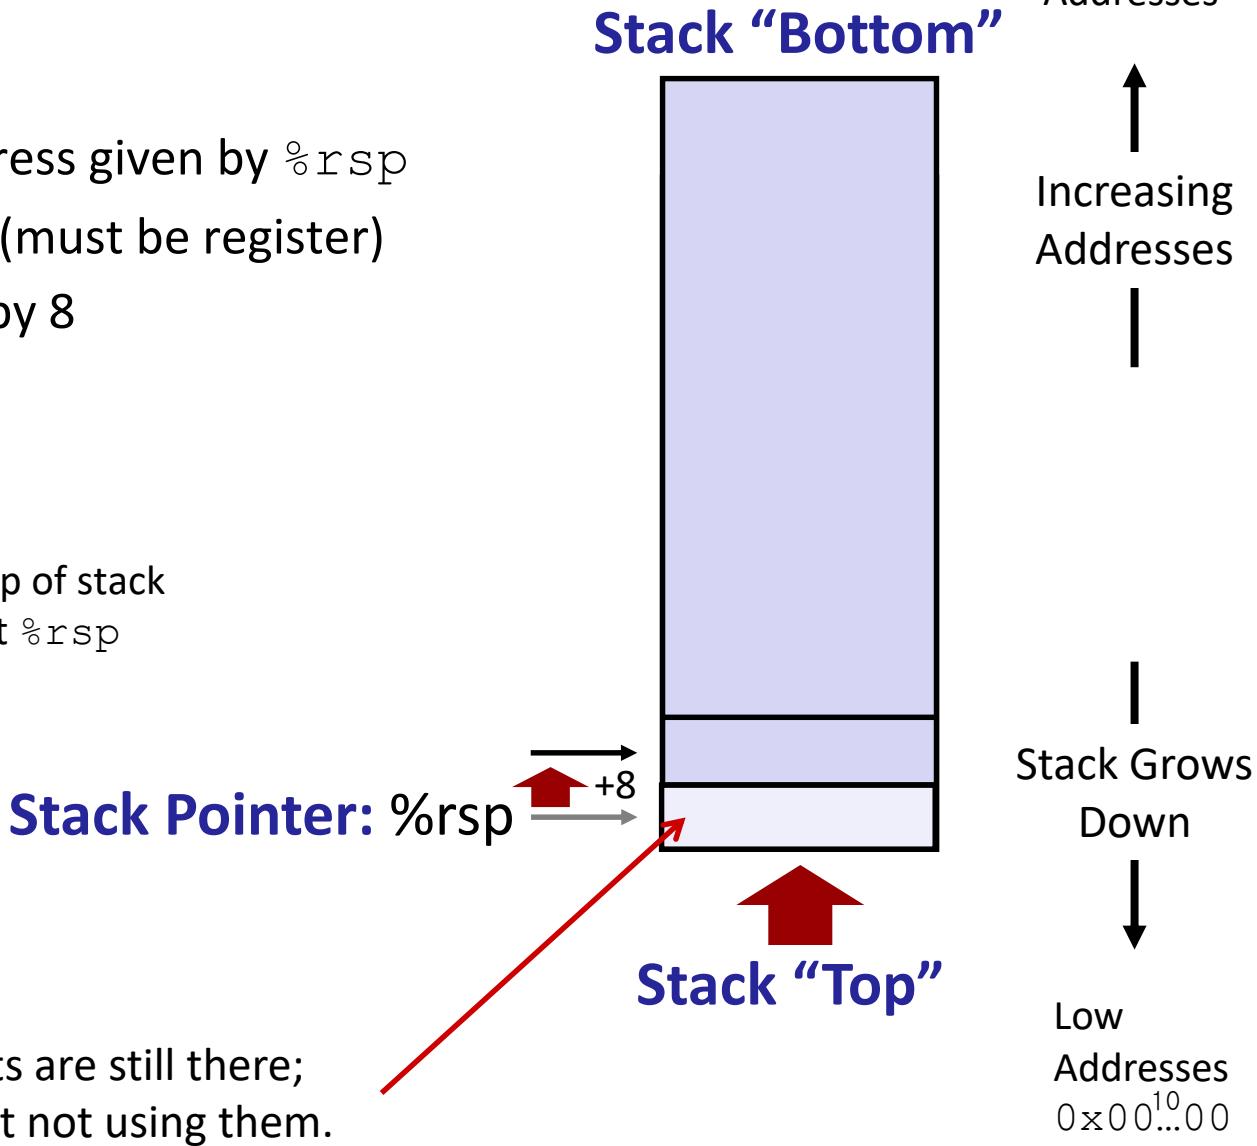
**Stack “Top”**

High Addresses  
Increasing Addresses  
Stack Grows Down  
Low Addresses  
0x00<sup>9</sup>..00



# x86-64 Stack: Pop

- `popq dst`
  - Load value at address given by `%rsp`
  - Store value at `dst` (must be register)
  - **Increment** `%rsp` by 8
- Example:
  - `popq %rcx`
  - Stores contents of top of stack into `%rcx` and adjust `%rsp`

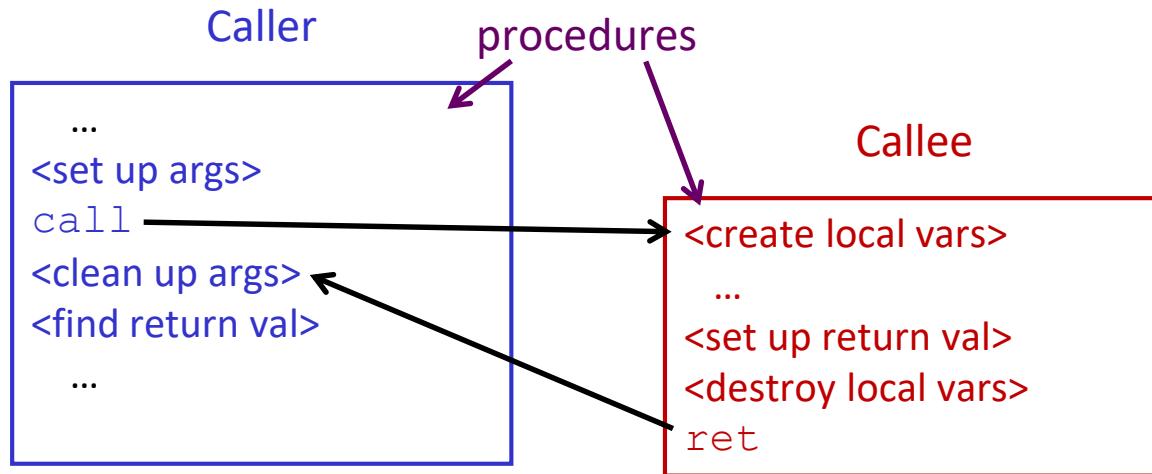


# Outline



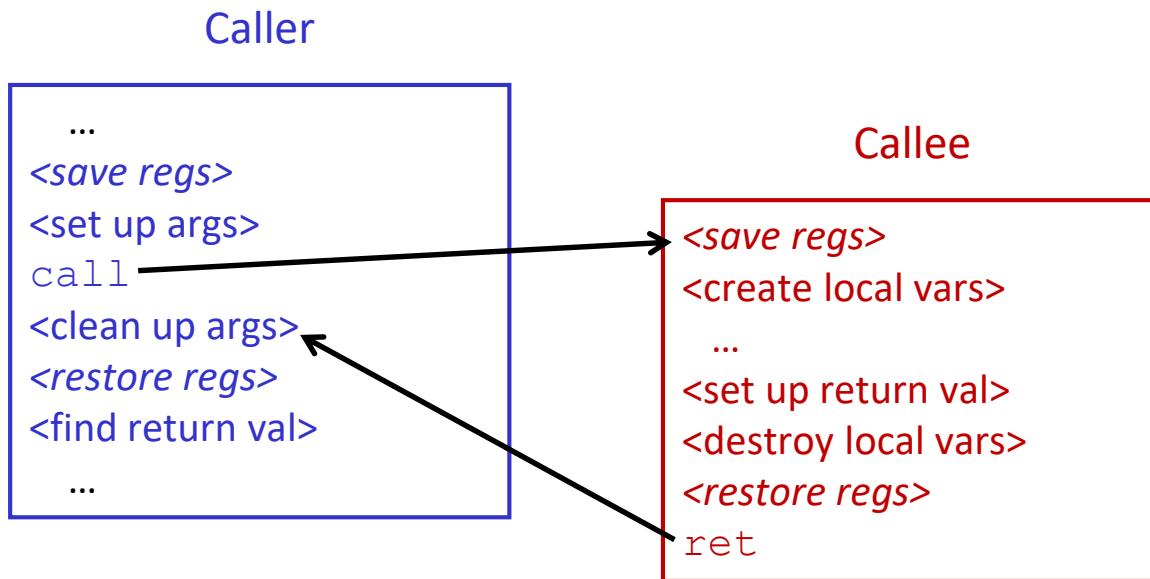
- Stack Structure
- Calling Conventions
  - **Passing control**
  - Passing data
  - Managing local data
- Illustration of Recursion

# Procedure Call Overview (1)



- **Callee** must know where to find args
- **Callee** must know where to find *return address*
- **Caller** must know where to find *return value*
- **Caller** and **Callee** run on same CPU, so use the same registers
  - How do we deal with register reuse?
- Unneeded steps can be skipped (*e.g.* no arguments)

# Procedure Call Overview (2)



- The *convention* of where to leave/find things is called the calling convention (or procedure call linkage)
  - Details vary between systems
  - We will see the convention for x86-64/Linux in detail
  - What could happen if our program didn't follow these conventions?

# Code Examples



```
void multstore  
    (long x, long y, long *dest)  
{  
    long t = mult2(x, y);  
    *dest = t;  
}
```

Compiler Explorer:

<https://godbolt.org/g/52Sqxj>

```
000000000400540 <multstore>:  
400540: push    %rbx          # Save %rbx  
400541: movq    %rdx,%rbx    # Save dest  
400544: call    400550 <mult2> # mult2(x,y)  
400549: movq    %rax,(%rbx)   # Save at dest  
40054c: pop     %rbx          # Restore %rbx  
40054d: ret             # Return
```

```
long mult2  
    (long a, long b)  
{  
    long s = a * b;  
    return s;  
}
```

```
000000000400550 <mult2>:  
400550: movq    %rdi,%rax    # a  
400553: imulq   %rsi,%rax    # a * b  
400557: ret             # Return
```

# Procedure Control Flow (1)



- Use stack to support procedure call and return
- **Procedure call:** `call label`
  - 1) Push return address on stack (*why? which address?*)
  - 2) Jump to `label`

# Procedure Control Flow (2)



- Use stack to support procedure call and return
- **Procedure call:** `call label`
  - 1) Push return address on stack (*why? which address?*)
  - 2) Jump to `label`
- Return address:
  - Address of instruction immediately after `call` instruction
  - Example from disassembly:

```
400544: call    400550 <mult2>
```

```
400549: movq    %rax, (%rbx)
```

Return address = **0x400549**

- **Procedure return:** `ret`
  - 1) Pop return address from stack
  - 2) Jump to address

next instruction  
happens to be a move,  
but could be anything

# Procedure Call Example (step 1)



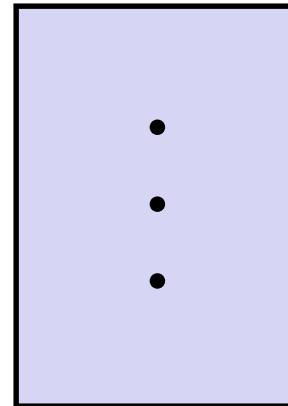
```
0000000000400540 <multstore>:
```

```
•  
•  
400544: call    400550 <mult2>  
400549: movq    %rax, (%rbx)  
•  
•
```

```
0000000000400550 <mult2>:
```

```
400550: movq    %rdi,%rax  
•  
•  
400557: ret
```

```
0x130  
0x128  
0x120
```



```
%rsp
```

```
0x120
```

```
%rip
```

```
0x400544
```

# Procedure Call Example (Step2)



```
0000000000400540 <multstore>:
```

```
•  
•  
400544: call    400550 <mult2>  
400549: movq    %rax, (%rbx) ←  
•  
•
```

```
0x130  
0x128  
0x120  
0x118 → 0x400549
```

```
0000000000400550 <mult2>:
```

```
400550: movq    %rdi,%rax ←  
•  
•  
400557: ret
```

```
%rsp 0x118  
%rip 0x400550 ←
```

# Procedure Return Example (Step1)



```
0000000000400540 <multstore>:
```

```
•  
•  
400544: call    400550 <mult2>  
400549: movq    %rax, (%rbx) ←
```

```
0000000000400550 <mult2>:
```

```
400550: movq    %rdi,%rax  
•  
•  
400557: ret ←
```

```
0x130  
0x128  
0x120  
0x118 0x400549
```

```
%rsp 0x118
```

```
%rip 0x400557
```

# Procedure Return Example (Step2)



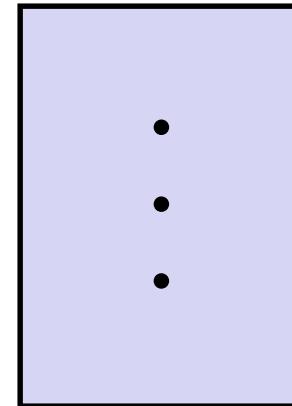
```
0000000000400540 <multstore>:
```

```
•  
•  
400544: call    400550 <mult2>  
400549: movq    %rax, (%rbx)  
•  
•
```

```
0000000000400550 <mult2>:
```

```
400550: movq    %rdi,%rax  
•  
•  
400557: ret
```

```
0x130  
0x128  
0x120
```



```
0x120
```

```
%rsp  
%rip  
0x400549
```

# Outline

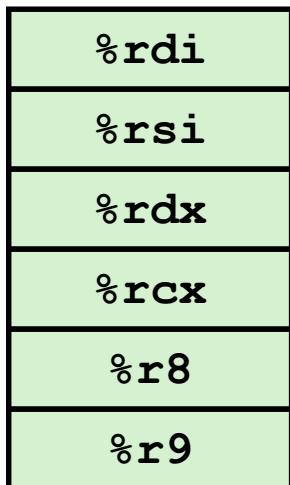


- Stack Structure
- Calling Conventions
  - Passing control
  - **Passing data**
  - Managing local data
- Illustration of Recursion

# Procedure Data Flow

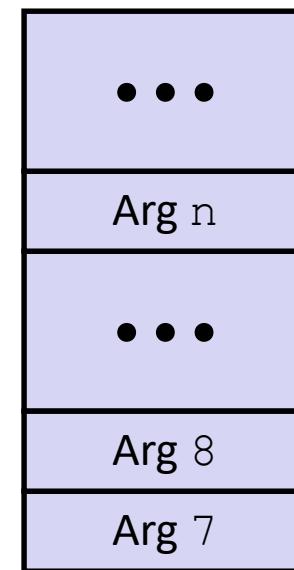


- Registers – **NOT in Memory!**
  - First 6 arguments



*Diane's  
Silk  
Dress  
Costs  
\$8 9*

- Stack – **in Memory!**



High Addresses  
↓  
Low Addresses  
0x00...00

- Return value



- Only allocate stack space when needed

# x86-64 Return Values



- By convention, values returned by procedures are placed in `%rax`
    - Choice of `%rax` is arbitrary
- 1) **Caller** must make sure to save the contents of `%rax` before calling a **callee** that returns a value
    - Part of register-saving convention
  - 2) **Callee** places return value into `%rax`
    - Any type that can fit in 8 bytes – integer, float, pointer, etc.
    - For return values greater than 8 bytes, best to return a *pointer* to them
  - 3) Upon return, **caller** finds the return value in `%rax`

# Data Flow Examples



```
void multstore  
    (long x, long y, long *dest)  
{  
    long t = mult2(x, y);  
    *dest = t;  
}
```

```
000000000400540 <multstore>:  
# x in %rdi, y in %rsi, dest in %rdx  
...  
400541: movq    %rdx,%rbx      # Save dest  
400544: call    400550 <mult2> # mult2(x,y)  
# t in %rax  
400549: movq    %rax,(%rbx)   # Save at dest  
...
```

```
long mult2  
    (long a, long b)  
{  
    long s = a * b;  
    return s;  
}
```

```
000000000400550 <mult2>:  
# a in %rdi, b in %rsi  
400550: movq    %rdi,%rax    # a  
400553: imulq   %rsi,%rax    # a * b  
# s in %rax  
400557: ret     # Return
```

# Outline



- Stack Structure
- Calling Conventions
  - Passing control
  - Passing data
  - **Managing local data**
- Illustration of Recursion

# Stack-Based Languages



- Languages that support recursion
  - e.g. C, Java, most modern languages
  - Code must be *re-entrant*
    - Multiple simultaneous instantiations of single procedure
  - Need some place to store *state* of each instantiation
    - Arguments, local variables, return pointer
- Stack allocated in *frames*
  - State for a single procedure instantiation
- Stack discipline
  - State for a given procedure needed for a limited time
    - Starting from when it is called to when it returns
  - Callee always returns before caller does

# Call Chain Example



```
yoo(...)
```

```
{  
    •  
    •  
    who();  
    •  
    •  
}
```

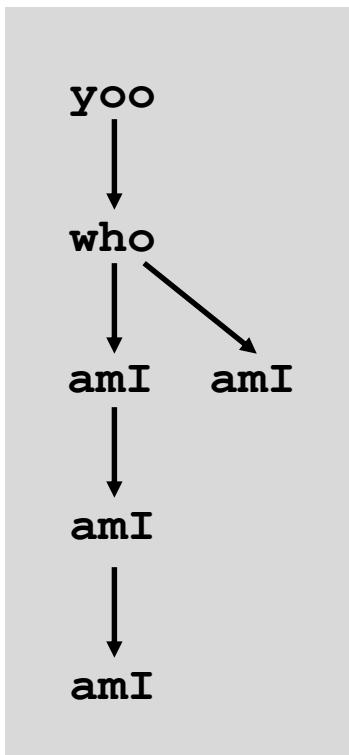
```
who(...)
```

```
{  
    •  
    amI();  
    •  
    amI();  
    •  
}
```

```
amI(...)
```

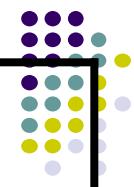
```
{  
    •  
    if (...) {  
        amI()  
    }  
    •  
}
```

Example  
Call Chain



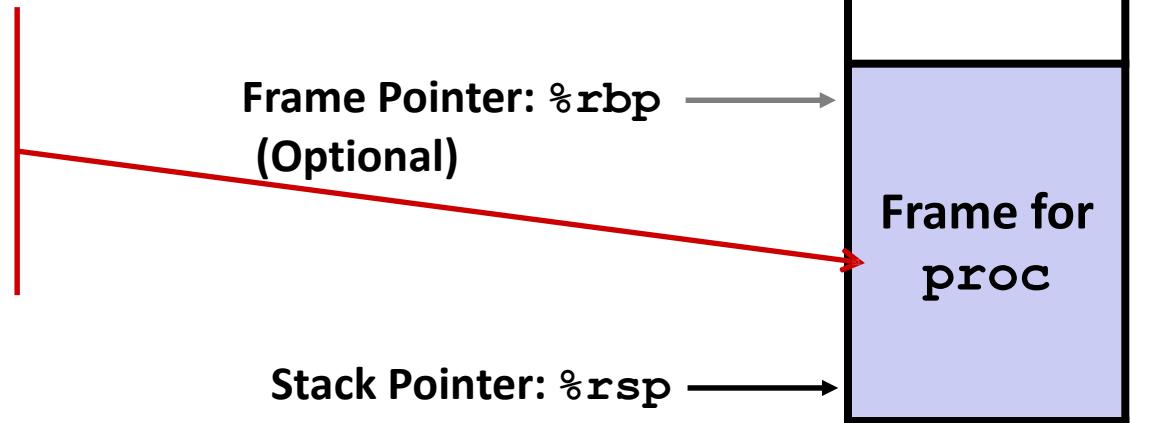
Procedure `amI` is recursive  
(calls itself)

# Stack Frames



- Contents

- Return information
- If needed
  - Local storage
  - Temporary space



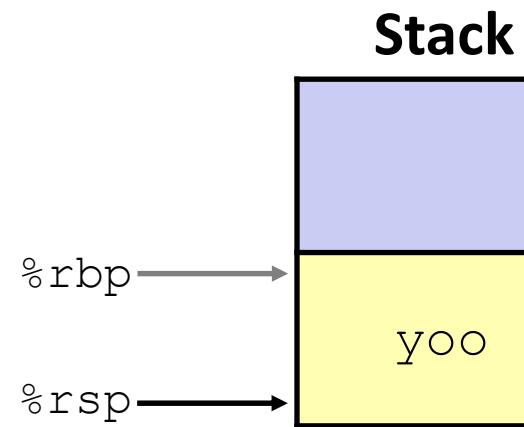
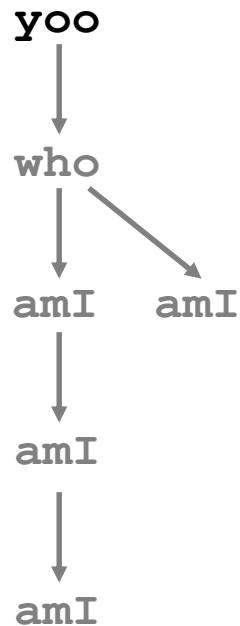
- Management

- Space allocated when enter procedure
  - “Set-up” code
  - Includes push by `call` instruction
- Deallocated when return
  - “Finish” code
  - Includes pop by `ret` instruction



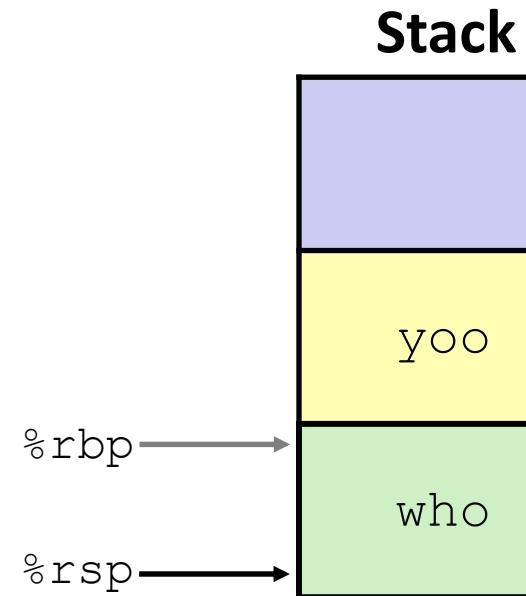
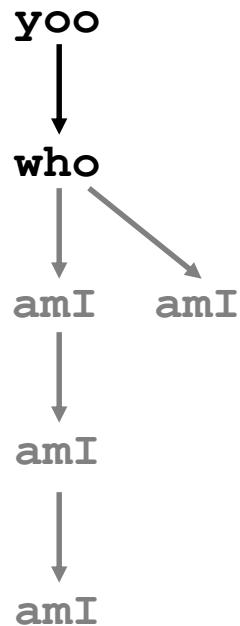
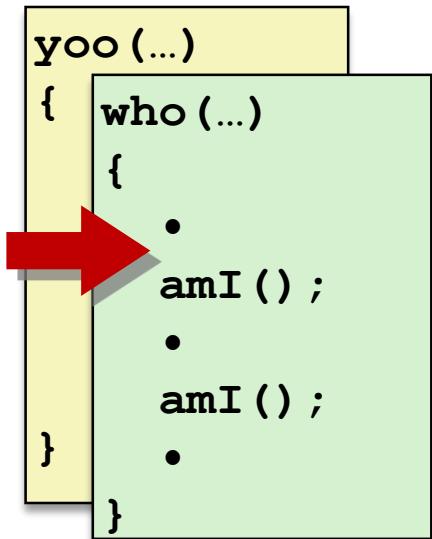
# 1) Call to yoo

```
yoo (...)  
{  
•  
•  
who () ;  
•  
}  
  
A red arrow points to the line "who () ;".
```



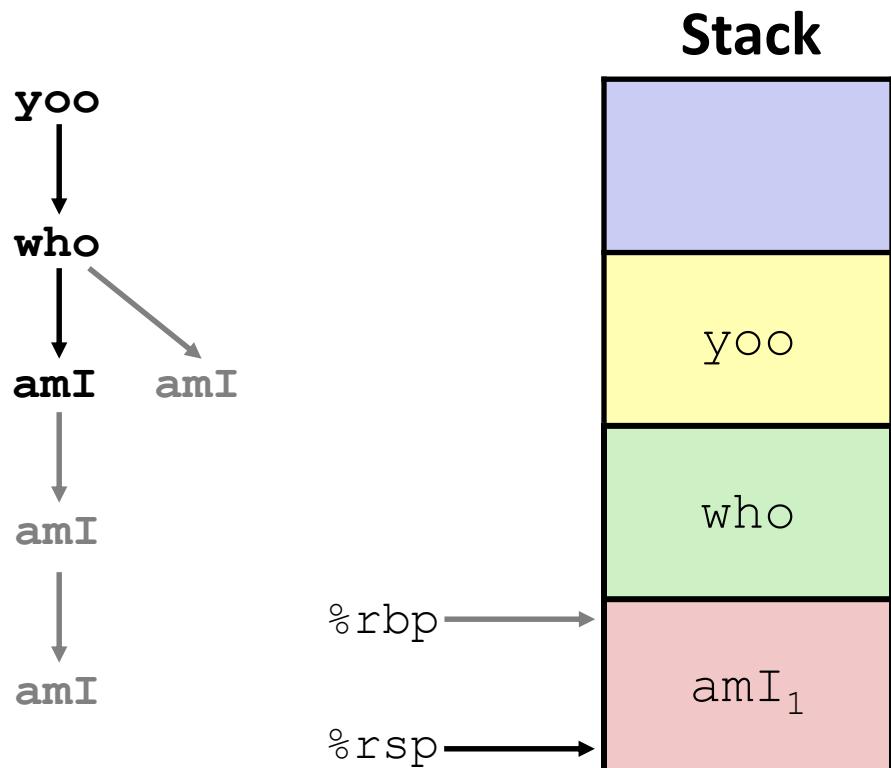
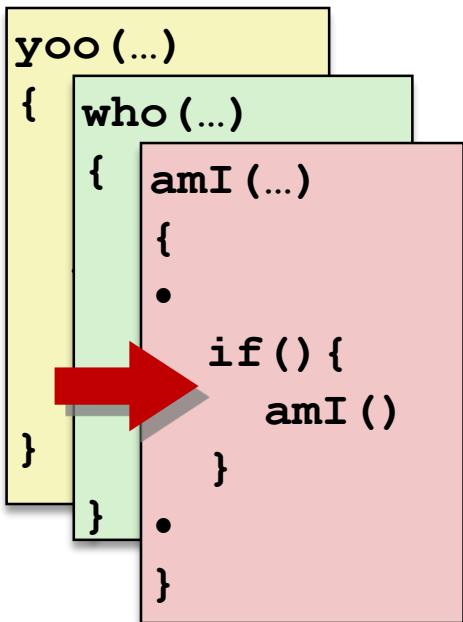


## 2) Call to who



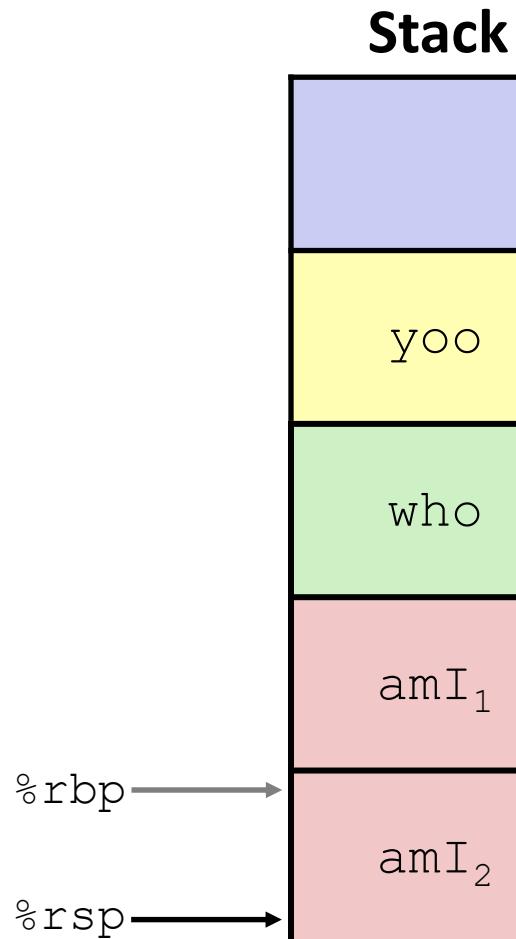
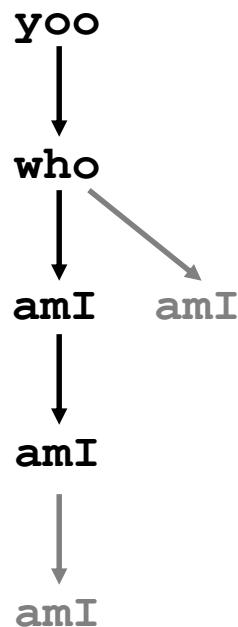
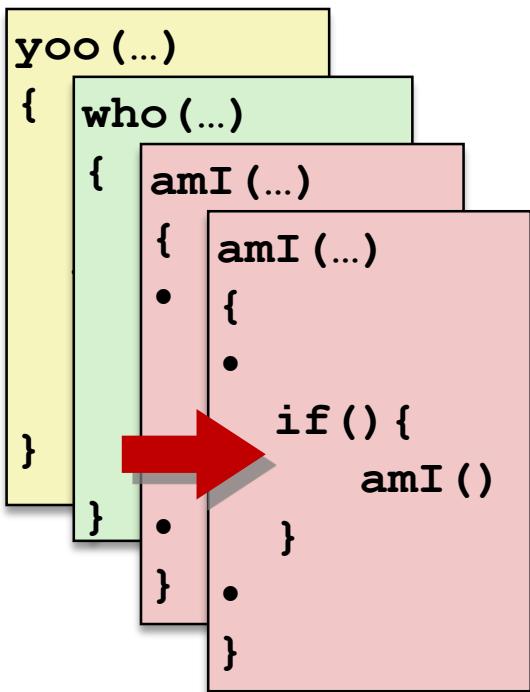


### 3) Call to amI (1)

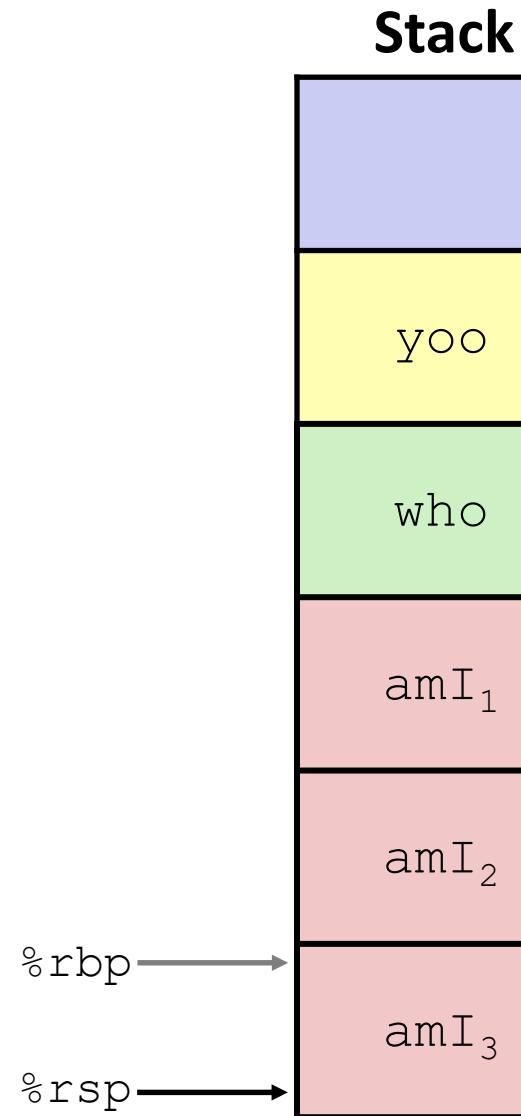
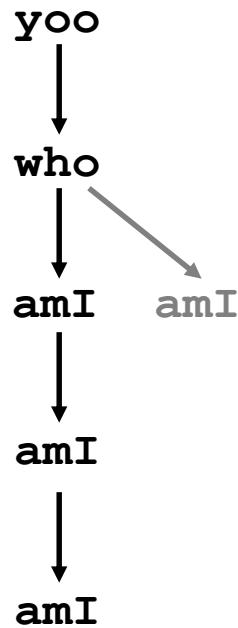
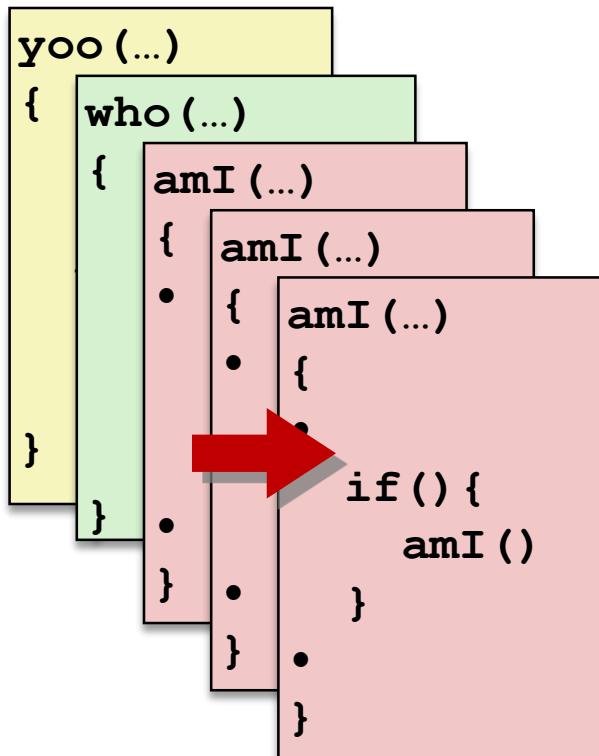




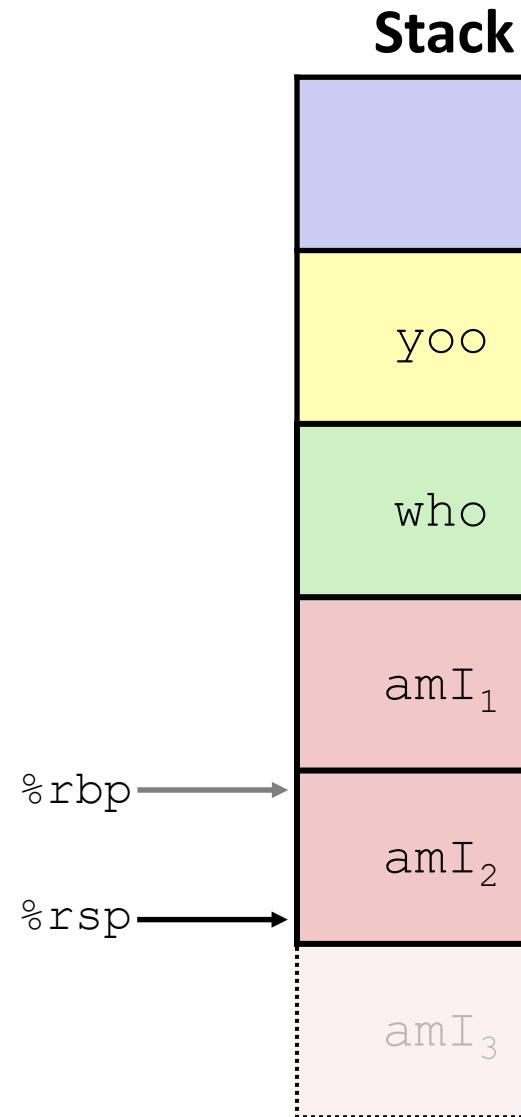
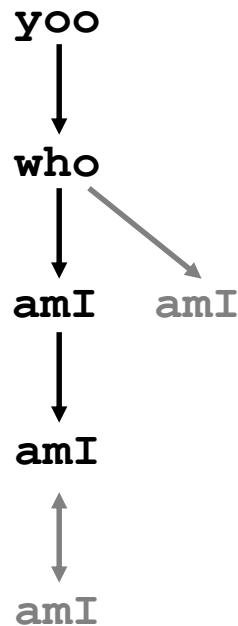
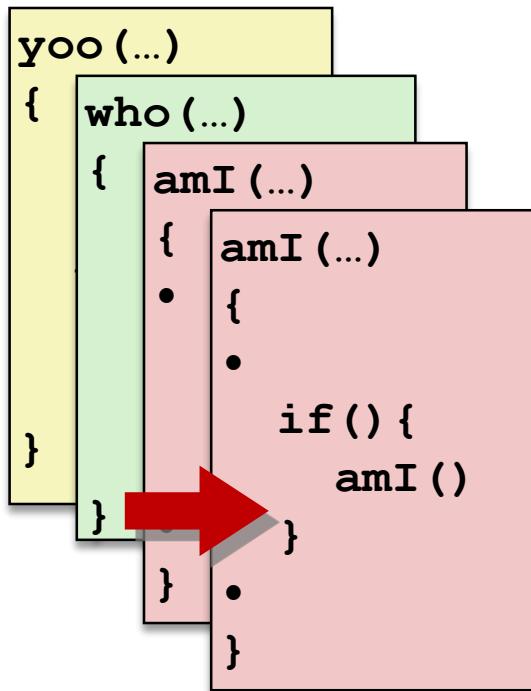
# 4) Recursive call to amI (2)



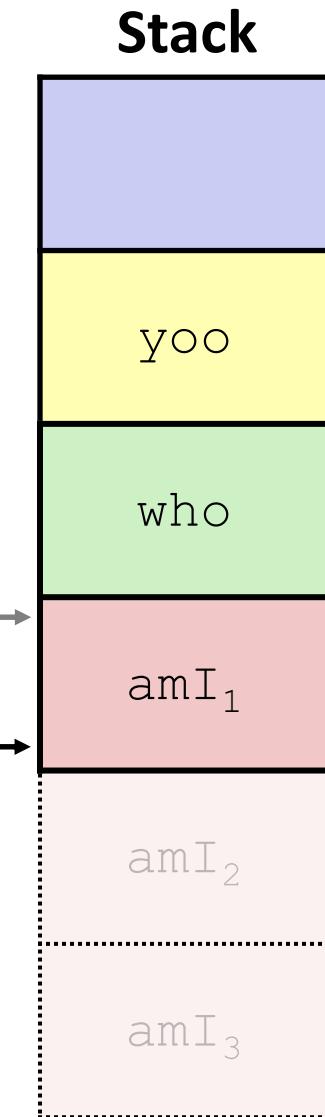
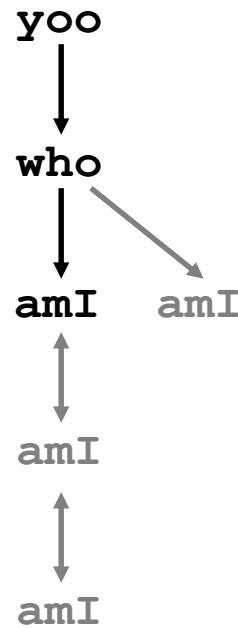
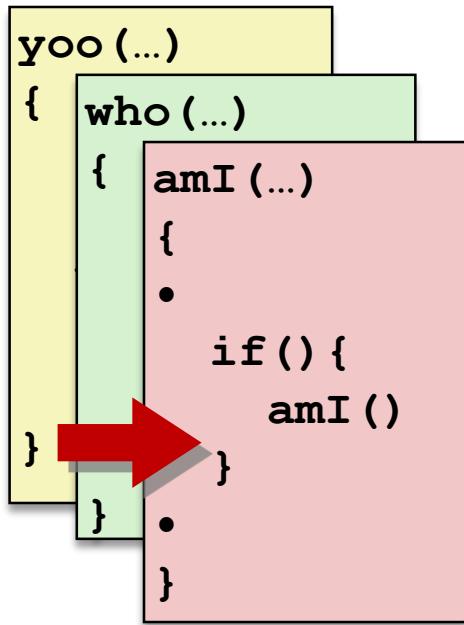
# 5) (another) Recursive call to amI (3)



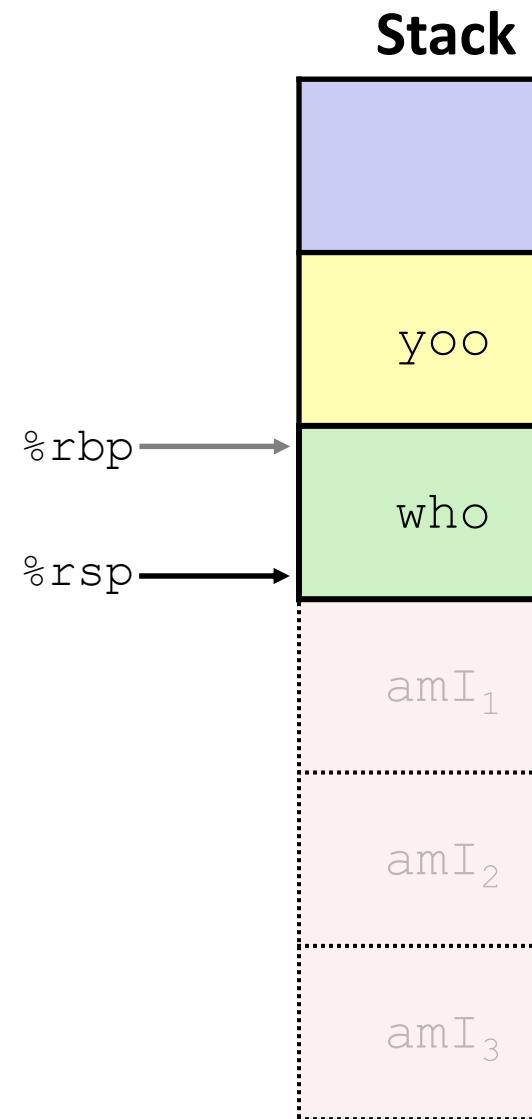
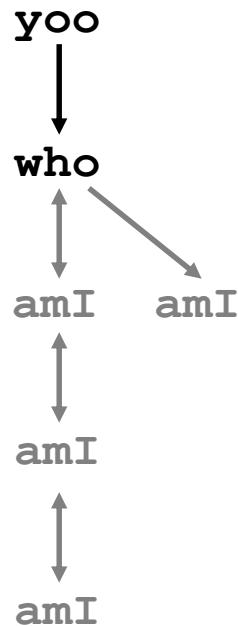
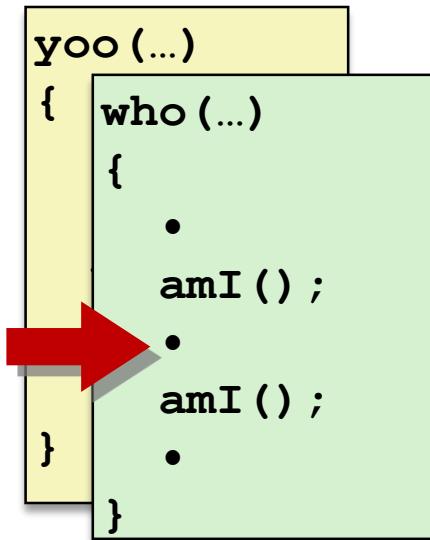
# 6) Return from (another) recursive call to amI



# 7) Return from recursive call to amI

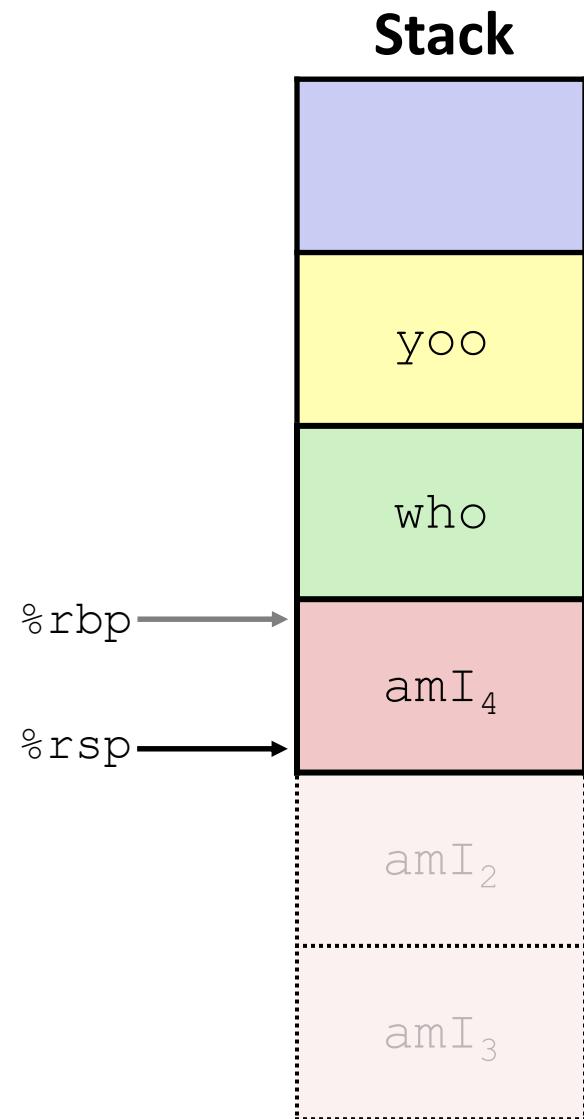
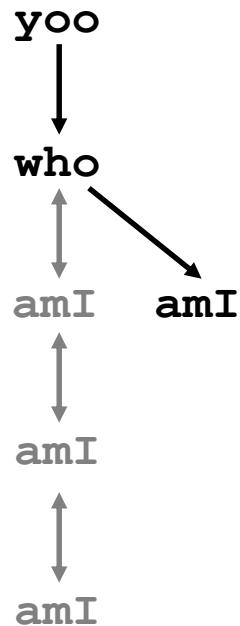
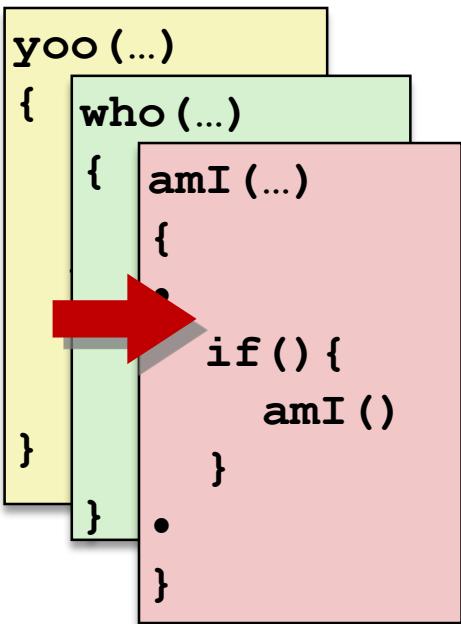


# 8) Return from call to amI





## 9) (second) Call to amI (4)

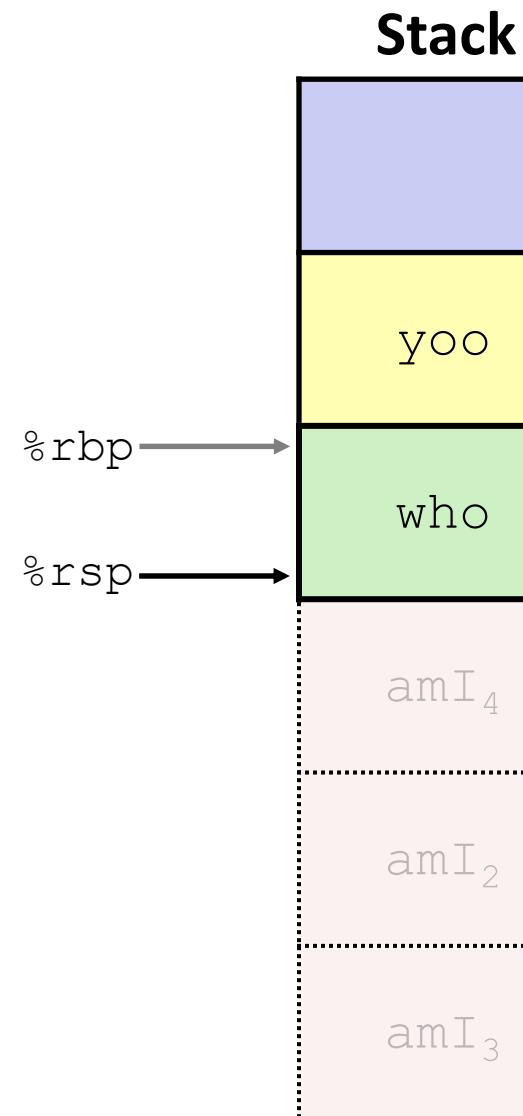
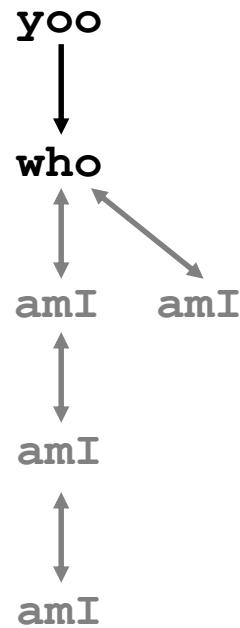


# 10) Return from (second) call to amI



```
yoo(...)  
{ who(...)  
{  
    •  
    amI();  
    •  
    amI();  
    •  
}
```

A red arrow points to the closing brace of the innermost block, indicating the return point.

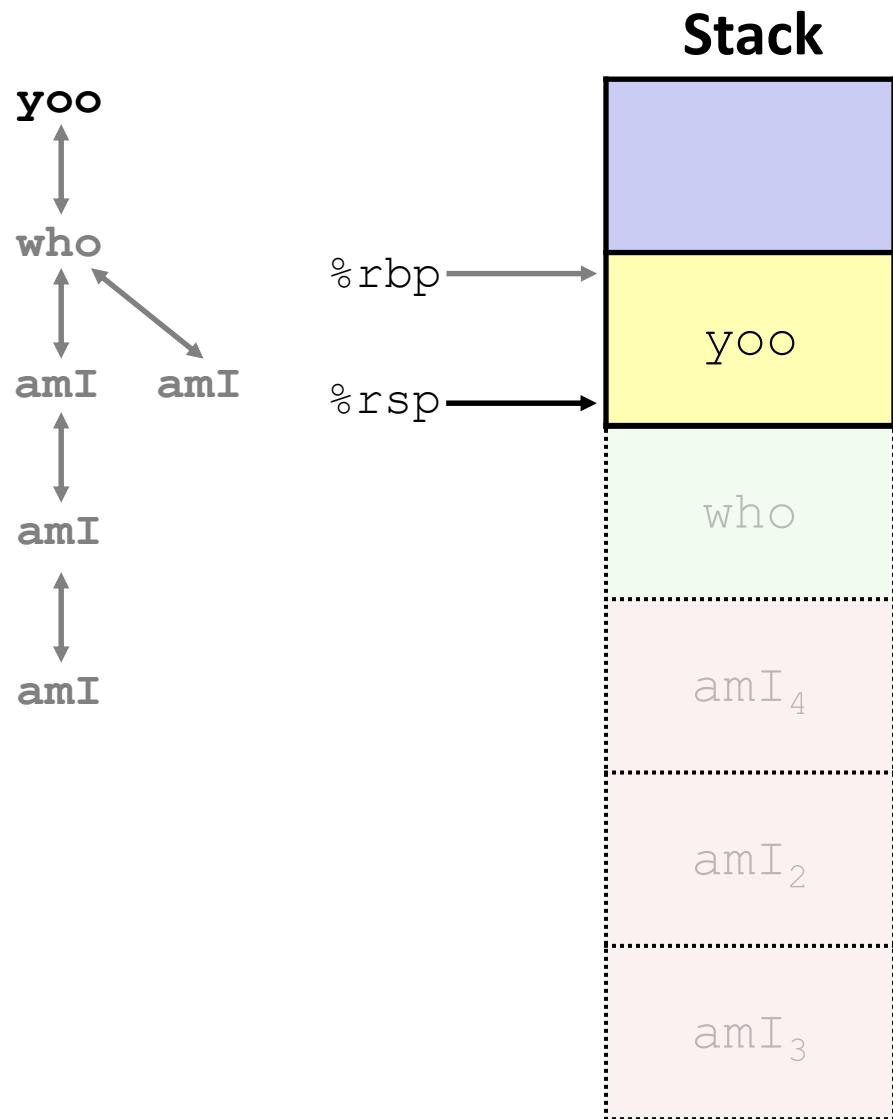


# 11) Return from call to who

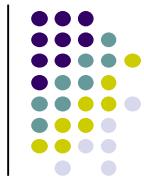


```
yoo(...)  
{  
    •  
    •  
    who();  
    •  
    •  
}
```

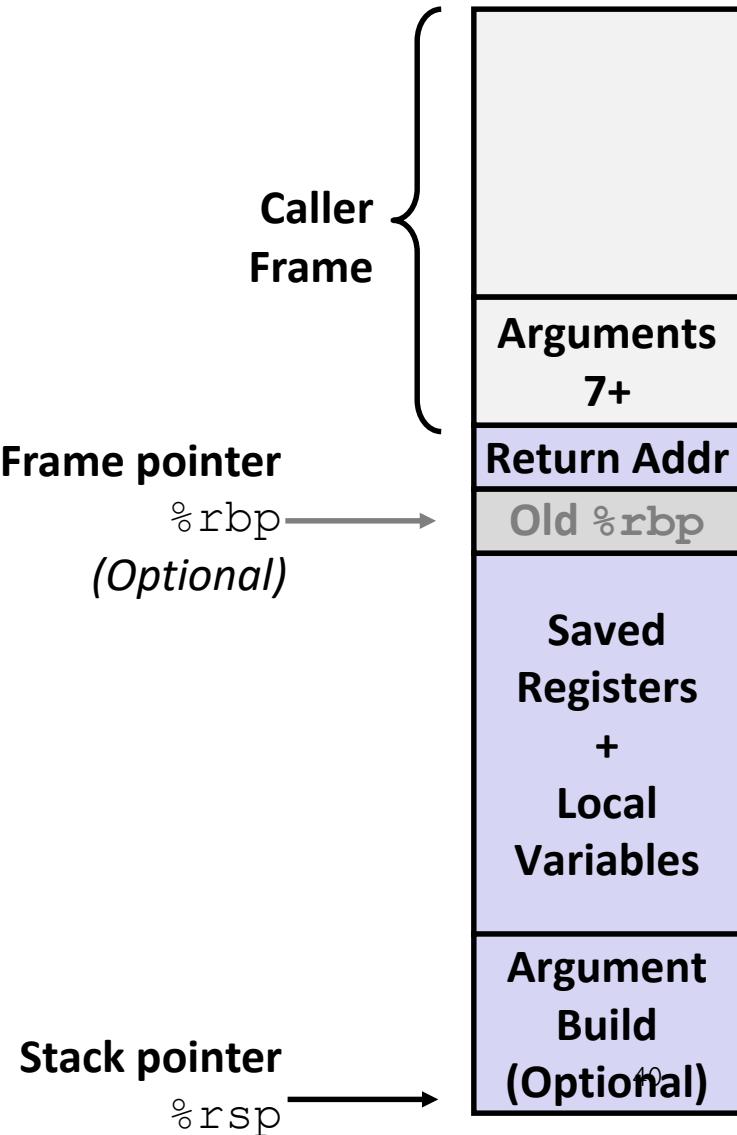
A red arrow points to the line "who();".



# x86-64/Linux Stack Frame



- Caller's Stack Frame
  - Extra arguments (if > 6 args) for this call
- Current/Callee Stack Frame
  - Return address
    - Pushed by `call` instruction
  - Old frame pointer (optional)
  - Saved register context (when reusing registers)
  - Local variables  
(If can't be kept in registers)
  - “Argument build” area  
(If callee needs to call another function -parameters for function about to call, if needed)



# Question



- Answer the following questions about when main() is run (assume x and y stored on the Stack):

```
int main() {  
    int i, x = 0;  
    for(i=0;i<3;i++)  
        x = randSum(x);  
    printf("x = %d\n", x);  
    return 0;  
}
```

```
int randSum(int n) {  
    int y = rand()%20;  
    return n+y;  
}
```

- Higher/larger address: x or y?
- How many total stack frames are *created*?
- What is the maximum *depth* (# of frames) of the Stack?

# Example: increment



```
long increment(long *p, long val) {  
    long x = *p;  
    long y = x + val;  
    *p = y;  
    return x;  
}
```

increment:

```
movq    (%rdi), %rax  
addq    %rax, %rsi  
movq    %rsi, (%rdi)  
ret
```

Register	Use(s)
%rdi	1 <sup>st</sup> arg (p)
%rsi	2 <sup>nd</sup> arg (val), y
%rax	x, return value

# Procedure Call Example (initial state)



```
long call_incr() {  
    long v1 = 351;  
    long v2 = increment(&v1, 100);  
    return v1+v2;  
}
```

```
call_incr:  
    subq    $16, %rsp  
    movq    $351, 8(%rsp)  
    movl    $100, %esi  
    leaq    8(%rsp), %rdi  
    call    increment  
    addq    8(%rsp), %rax  
    addq    $16, %rsp  
    ret
```

## Initial Stack Structure



- Return address on stack is the address of instruction immediately *following* the call to “call\_incr”
  - Shown here as main, but could be anything)

# Procedure Call Example (step 1)

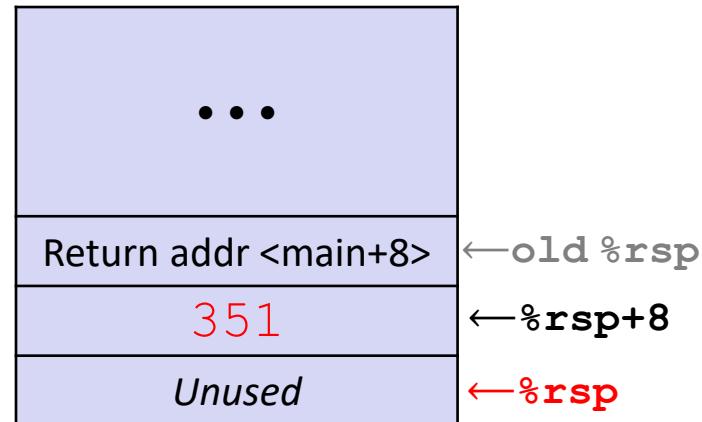


```
long call_incr() {  
    long v1 = 351;  
    long v2 = increment(&v1, 100);  
    return v1+v2;  
}
```

```
call_incr:  
    subq    $16, %rsp  
    movq    $351, 8(%rsp)  
    movl    $100, %esi  
    leaq    8(%rsp), %rdi  
    call    increment  
    addq    8(%rsp), %rax  
    addq    $16, %rsp  
    ret
```

} **Allocate space  
for local vars**

## Stack Structure



- Setup space for local variables
  - Only v1 needs space on the stack
- Compiler allocated extra space
  - Often does this for a variety of reasons, including alignment

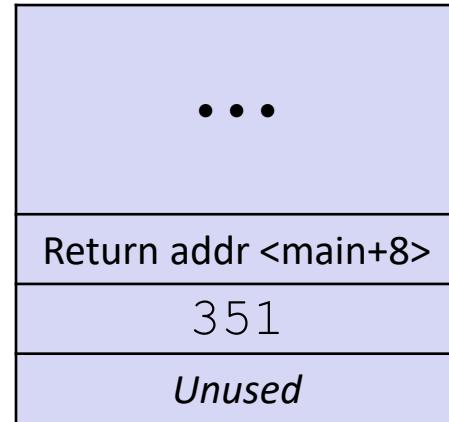


# Procedure Call Example (step 2)

```
long call_incr() {  
    long v1 = 351;  
    long v2 = increment(&v1, 100);  
    return v1+v2;  
}
```

```
call_incr:  
    subq    $16, %rsp  
    movq    $351, 8(%rsp)  
    movl    $100, %esi  
    leaq    8(%rsp), %rdi  
    call    increment  
    addq    8(%rsp), %rax  
    addq    $16, %rsp  
    ret
```

## Stack Structure



←%rsp+8  
←%rsp

} Set up parameters for call  
to increment

*Aside:* `movl` is used because 100 is a small positive value that fits in 32 bits. High order bits of `rsi` get set to zero automatically. It takes one less byte to encode a `movl` than a `movq`.

Register	Use(s)
%rdi	&v1
%rsi	100

# Procedure Call Example (step 3)

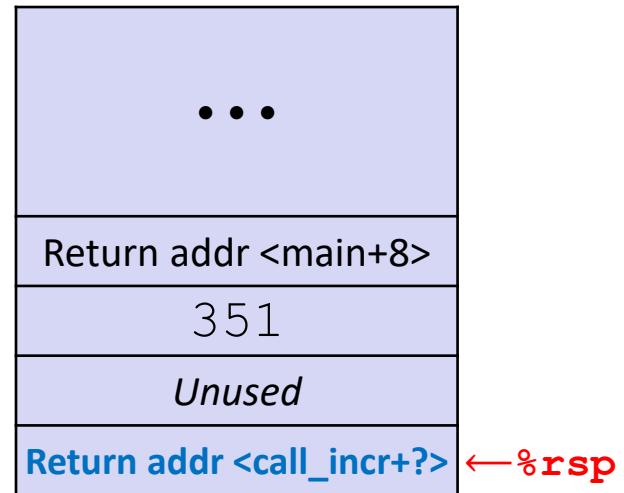


```
long call_incr() {  
    long v1 = 351;  
    long v2 = increment(&v1, 100);  
    return v1+v2;  
}
```

```
call_incr:  
    subq    $16, %rsp  
    movq    $351, 8(%rsp)  
    movl    $100, %esi  
    leaq    8(%rsp), %rdi  
    call    increment  
    addq    8(%rsp), %rax  
    addq    $16, %rsp  
    ret
```

```
increment:  
    movq    (%rdi), %rax  
    addq    %rax, %rsi  
    movq    %rsi, (%rdi)  
    ret
```

## Stack Structure



- State while inside increment
- Return address on top of stack is address of the addq instruction immediately following call to increment

Register	Use(s)
%rdi	&v1
%rsi	100
%rax	

# Procedure Call Example (step 4)

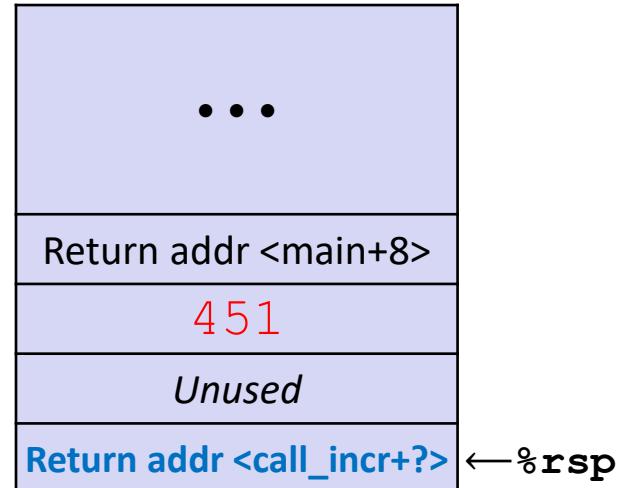


```
long call_incr() {  
    long v1 = 351;  
    long v2 = increment(&v1, 100);  
    return v1+v2;  
}
```

```
call_incr:  
    subq    $16, %rsp  
    movq    $351, 8(%rsp)  
    movl    $100, %esi  
    leaq    8(%rsp), %rdi  
    call    increment  
    addq    8(%rsp), %rax  
    addq    $16, %rsp  
    ret
```

```
increment:  
    movq    (%rdi), %rax # x = *p  
    addq    %rax, %rsi   # y = x+100  
    movq    %rsi, (%rdi) # *p = y  
    ret
```

## Stack Structure



- ❖ State while inside increment
  - After code in body has been executed

Register	Use(s)
%rdi	&v1
%rsi	451
%rax	351

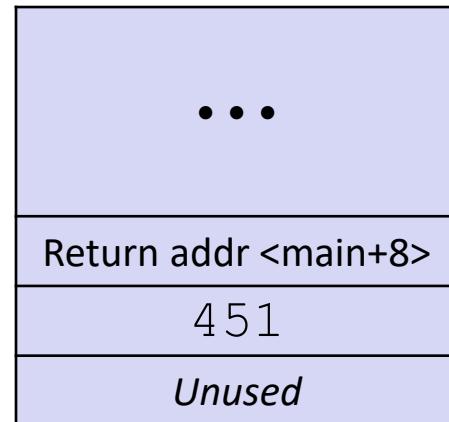
# Procedure Call Example (step 5)



```
long call_incr() {  
    long v1 = 351;  
    long v2 = increment(&v1, 100);  
    return v1+v2;  
}
```

```
call_incr:  
    subq    $16, %rsp  
    movq    $351, 8(%rsp)  
    movl    $100, %esi  
    leaq    8(%rsp), %rdi  
    call    increment  
    addq    8(%rsp), %rax  
    addq    $16, %rsp  
    ret
```

## Stack Structure



- After returning from call to `increment`
  - Registers and memory have been modified and return address has been popped off stack

Register	Use(s)
%rdi	&v1
%rsi	451
%rax	351

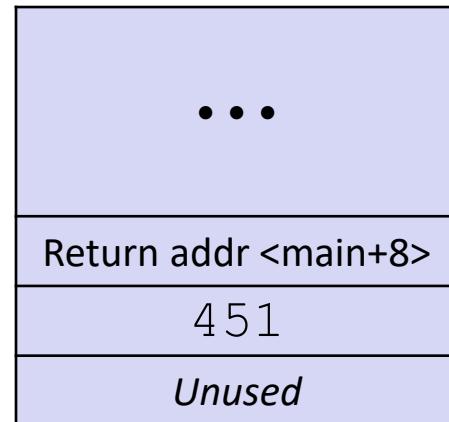
# Procedure Call Example (step 6)



```
long call_incr() {  
    long v1 = 351;  
    long v2 = increment(&v1, 100);  
    return v1+v2;  
}
```

```
call_incr:  
    subq    $16, %rsp  
    movq    $351, 8(%rsp)  
    movl    $100, %esi  
    leaq    8(%rsp), %rdi  
    call    increment  
    addq    8(%rsp), %rax  
    addq    $16, %rsp  
    ret
```

## Stack Structure



←%rsp+8  
←%rsp

← Update %rax to contain v1+v2

Register	Use(s)
%rdi	&v1
%rsi	451
%rax	451+351

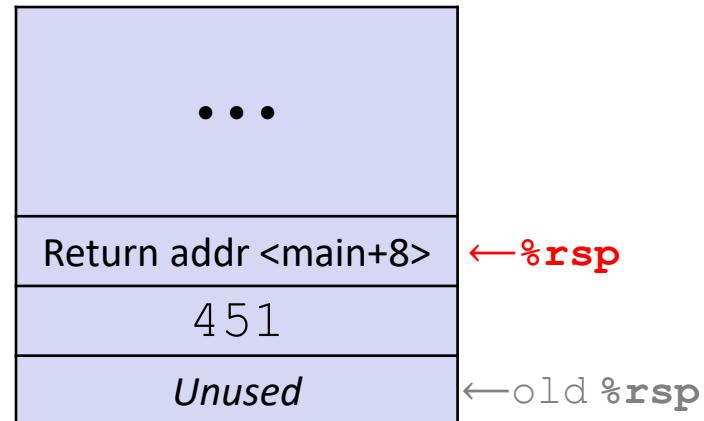
# Procedure Call Example (step 7)



```
long call_incr() {  
    long v1 = 351;  
    long v2 = increment(&v1, 100);  
    return v1+v2;  
}
```

```
call_incr:  
    subq    $16, %rsp  
    movq    $351, 8(%rsp)  
    movl    $100, %esi  
    leaq    8(%rsp), %rdi  
    call    increment  
    addq    8(%rsp), %rax  
    addq    $16, %rsp  
    ret
```

## Stack Structure



← De-allocate space for local vars

Register	Use(s)
%rdi	&v1
%rsi	451
%rax	802

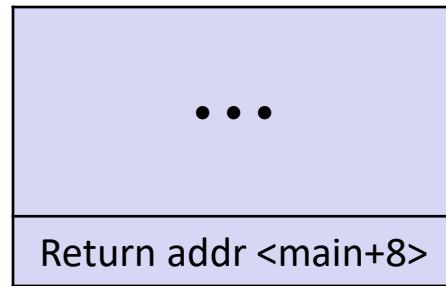
# Procedure Call Example (step 8)



```
long call_incr() {  
    long v1 = 351;  
    long v2 = increment(&v1, 100);  
    return v1+v2;  
}
```

```
call_incr:  
    subq    $16, %rsp  
    movq    $351, 8(%rsp)  
    movl    $100, %esi  
    leaq    8(%rsp), %rdi  
    call    increment  
    addq    8(%rsp), %rax  
    addq    $16, %rsp  
    ret
```

## Stack Structure



- State *just before* returning from call to call\_incr

Register	Use(s)
%rdi	&v1
%rsi	451
%rax	802

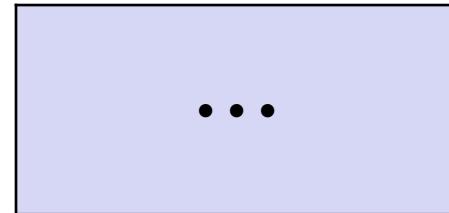
# Procedure Call Example (step 9)



```
long call_incr() {  
    long v1 = 351;  
    long v2 = increment(&v1, 100);  
    return v1+v2;  
}
```

```
call_incr:  
    subq    $16, %rsp  
    movq    $351, 8(%rsp)  
    movl    $100, %esi  
    leaq    8(%rsp), %rdi  
    call    increment  
    addq    8(%rsp), %rax  
    addq    $16, %rsp  
    ret
```

## Final Stack Structure



←%rsp

- State immediately *after* returning from call to `call_incr`
  - Return addr has been popped off stack
  - Control has returned to the instruction immediately following the call to `call_incr` (not shown here)

Register	Use(s)
%rdi	&v1
%rsi	451
%rax	802

# Outline



- Stack Structure
- Calling Conventions
  - Passing control
  - Passing data
  - Managing local data
- **Register Saving Conventions**
- Illustration of Recursion

# Register Saving Conventions



- When procedure **yoo** calls **who**:
  - yoo** is the *caller*
  - who** is the *callee*
- Can register be used for temporary storage?

```
yoo:
```

```
    . . .
    movq $15213, %rdx
    call who
    addq %rdx, %rax
    . . .
    ret
```

```
who:
```

```
    . . .
    subq $18213, %rdx
    . . .
    ret
```

- No! Contents of register %rdx overwritten by who
- This could be trouble → something should be done! Either:
  - caller* should save %rdx before the call (and restore it after the call)
  - callee* should save %rdx before using it (and restore it before returning)

# Register Saving Conventions



- “*Caller-saved*” registers

- It is the **caller**’s responsibility to save any important data in these registers before calling another procedure (*i.e.* the **callee** can freely change data in these registers)
- **Caller** saves values in its stack frame before calling **Callee**, then restores values after the call

- “*Callee-saved*” registers

- It is the callee’s responsibility to save any data in these registers before using the registers (*i.e.* the **caller** assumes the data will be the same across the **callee** procedure call)
- **Callee** saves values in its stack frame before using, then restores them before returning to **caller**

# Silly Register Convention Analogy



- 1) Parents (**caller**) leave for the weekend and give the keys to the house to their child (**callee**)
  - Being suspicious, they put away/hid the valuables (**caller**-saved) before leaving
  - Warn child to leave the bedrooms untouched: “These rooms better look the same when we return!”
- 2) Child decides to throw a wild party (*computation*), spanning the entire house
  - To avoid being disowned, child moves all of the stuff from the **bedrooms** to the backyard shed (**callee**-saved) before the guests trash the house
  - Child cleans up house after the party and moves stuff back to **bedrooms**
- 3) Parents return home and are satisfied with the state of the house
  - Move valuables back (**caller**-saved) and continue with their lives

# x86-64 Linux Register Usage, part1

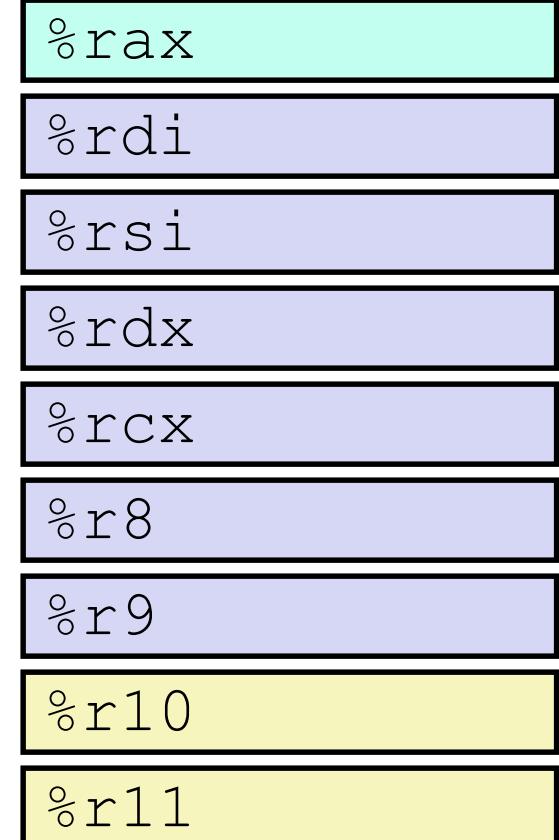


- **%rax**
  - Return value
  - Also **caller**-saved
  - Can be modified by procedure
- **%rdi, ..., %r9**
  - Arguments
  - Also **caller**-saved
  - Can be modified by procedure
- **%r10, %r11**
  - **Caller**-saved
  - Can be modified by procedure

Return value

Arguments

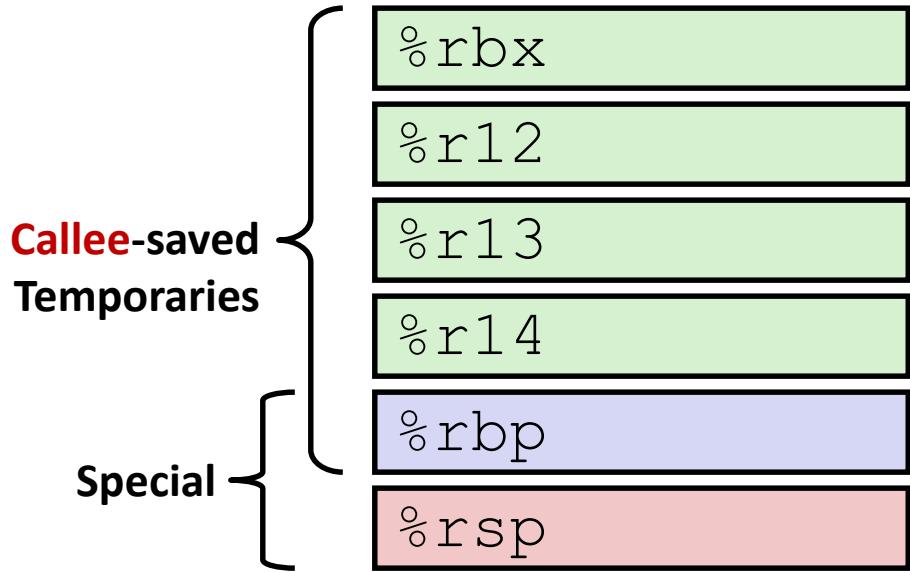
Caller-saved  
temporaries



# x86-64 Linux Register Usage, part2



- **%rbx, %r12, %r13, %r14**
  - **Callee**-saved
  - **Callee** must save & restore
- **%rbp**
  - **Callee**-saved
  - **Callee** must save & restore
  - May be used as frame pointer
  - Can mix & match
- **%rsp**
  - Special form of **callee** save
  - Restored to original value upon exit from procedure



# x86-64 64-bit Registers: Usage Conventions



%rax	Return value - <b>Caller saved</b>	%r8	Argument #5 - <b>Caller saved</b>
%rbx	<b>Callee saved</b>	%r9	Argument #6 - <b>Caller saved</b>
%rcx	Argument #4 - <b>Caller saved</b>	%r10	<b>Caller saved</b>
%rdx	Argument #3 - <b>Caller saved</b>	%r11	<b>Caller Saved</b>
%rsi	Argument #2- <b>Caller saved</b>	%r12	<b>Callee saved</b>
%rdi	Argument #1 - <b>Caller saved</b>	%r13	<b>Callee saved</b>
%rsp	Stack pointer	%r14	<b>Callee saved</b>
%rbp	<b>Callee saved</b>	%r15	<b>Callee saved</b>

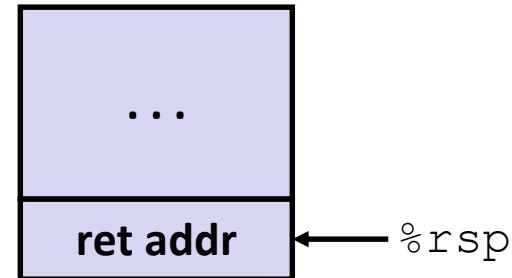
# Callee-Saved Example (step 1)



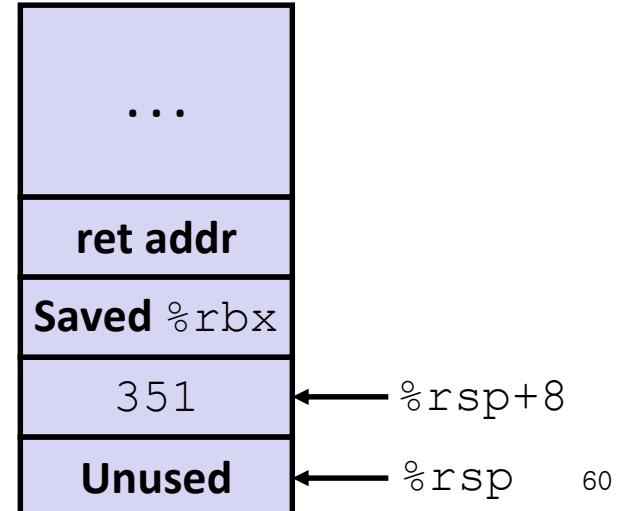
```
long call_incr2(long x) {  
    long v1 = 351;  
    long v2 = increment(&v1, 100);  
    return x+v2;  
}
```

```
call_incr2:  
    pushq  %rbx  
    subq    $16, %rsp  
    movq    %rdi, %rbx  
    movq    $351, 8(%rsp)  
    movl    $100, %esi  
    leaq    8(%rsp), %rdi  
    call    increment  
    addq    %rbx, %rax  
    addq    $16, %rsp  
    popq    %rbx  
    ret
```

Initial Stack Structure



Resulting Stack Structure



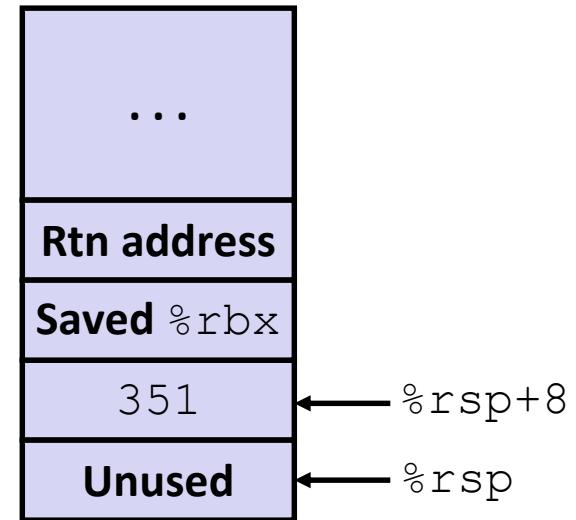
# Callee-Saved Example (step 2)



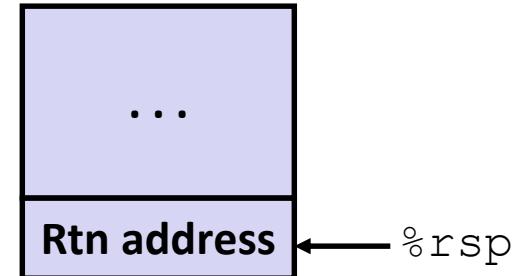
```
long call_incr2(long x) {  
    long v1 = 351;  
    long v2 = increment(&v1, 100);  
    return x+v2;  
}
```

```
call_incr2:  
    pushq    %rbx  
    subq    $16, %rsp  
    movq    %rdi, %rbx  
    movq    $351, 8(%rsp)  
    movl    $100, %esi  
    leaq    8(%rsp), %rdi  
    call    increment  
    addq    %rbx, %rax  
    addq    $16, %rsp  
    popq    %rbx  
    ret
```

## Stack Structure



## Pre-return Stack Structure



# Why Caller and Callee Saved?



- We want *one* calling convention to simply separate implementation details between caller and callee
- In general, neither caller-save nor callee-save is “best”:
  - If caller isn’t using a register, caller-save is better
  - If callee doesn’t need a register, callee-save is better
  - If “do need to save”, callee-save generally makes smaller programs
    - Functions are called from multiple places
- So... “some of each” and compiler tries to “pick registers” that minimize amount of saving/restoring

# Register Conventions Summary



- **Caller**-saved register values need to be pushed onto the stack before making a procedure call *only if the Caller needs that value later*
  - **Callee** may change those register values
- **Callee**-saved register values need to be pushed onto the stack *only if the Callee intends to use those registers*
  - **Caller** expects unchanged values in those registers
- Don't forget to restore/pop the values later!



# Outline

- Stack Structure
- Calling Conventions
  - Passing control
  - Passing data
  - Managing local data
- Register Saving Conventions
- **Illustration of Recursion**

# Recursive Function



```
/* Recursive popcount */
long pcount_r(unsigned long x) {
    if (x == 0)
        return 0;
    else
        return (x&1)+pcount_r(x >> 1);
}
```

## Compiler Explorer:

<https://godbolt.org/g/4ZJbz1>

- Compiled with `-O1` for brevity instead of `-Og`
- Try `-O2` instead!

```
pcount_r:
    movl    $0, %eax
    testq   %rdi, %rdi
    je     .L6
    pushq   %rbx
    movq    %rdi, %rbx
    shrq    %rdi
    call    pcount_r
    andl    $1, %ebx
    addq    %rbx, %rax
    popq    %rbx
.L6:
    rep ret
```

# Recursive Function: Base Case



```
/* Recursive popcount */
long pcount_r(unsigned long x) {
    if (x == 0)
        return 0;
    else
        return (x&1)+pcount_r(x >> 1);
}
```

Register	Use(s)	Type
%rdi	x	Argument
%rax	Return value	Return value

```
pcount_r:
    movl    $0, %eax
    testq   %rdi, %rdi
    je     .L6
    pushq   %rbx
    movq    %rdi, %rbx
    shrq    %rdi
    call    pcount_r
    andl    $1, %ebx
    addq    %rbx, %rax
    popq    %rbx
.L6:
    rep    ret
```

Trick because some AMD hardware doesn't like jumping to `ret`

# Recursive Function: **Callee** Register Save



```
/* Recursive popcount */
long pcount_r(unsigned long x) {
    if (x == 0)
        return 0;
    else
        return (x&1)+pcount_r(x >> 1);
}
```

Need original value  
of `x` *after* recursive  
call to `pcount_r`.

“Save” by putting in  
`%rbx` (**callee**  
saved), but need to  
save old value of  
`%rbx` before you  
change it.

`%rsp` →

## The Stack



Register	Use(s)	Type
<code>%rdi</code>	<code>x</code>	Argument

```
pcount_r:
    movl    $0, %eax
    testq   %rdi, %rdi
    je     .L6
    pushq   %rbx
    movq    %rdi, %rbx
    shrq    %rdi
    call    pcount_r
    andl    $1, %ebx
    addq    %rbx, %rax
    popq    %rbx
.L6:
    rep ret
```

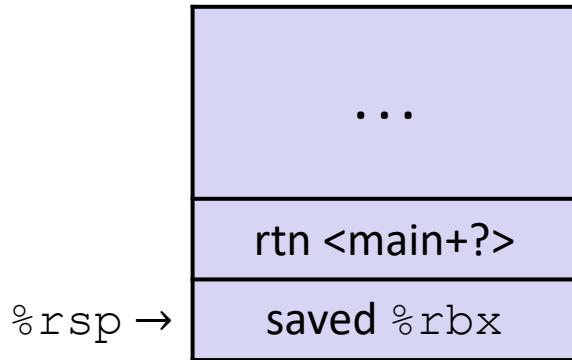
# Recursive Function: Call Setup



```
/* Recursive popcount */
long pcount_r(unsigned long x) {
    if (x == 0)
        return 0;
    else
        return (x&1)+pcount_r(x >> 1);
}
```

Register	Use(s)	Type
%rdi	x (new)	Argument
%rbx	x (old)	Callee saved

## The Stack



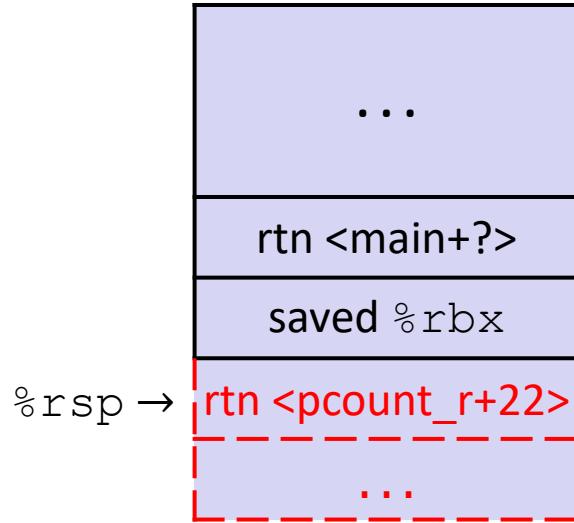
```
pcount_r:
    movl    $0, %eax
    testq   %rdi, %rdi
    je     .L6
    pushq   %rbx
    movq    %rdi, %rbx
    shrq    %rdi
    call    pcount_r
    andl    $1, %ebx
    addq    %rbx, %rax
    popq    %rbx
.L6:
    rep ret
```

# Recursive Function: Call (ink walkthrough)



```
/* Recursive popcount */
long pcount_r(unsigned long x) {
    if (x == 0)
        return 0;
    else
        return (x&1)+pcount_r(x >> 1);
}
```

## The Stack



Register	Contents
%rdi	
%rbx	
%rax	

## pcount\_r:

```
movl    $0, %eax
testq   %rdi, %rdi
je      .L6
pushq   %rbx
movq   %rdi, %rbx
shrq   %rdi
call   pcount_r
andl   $1, %ebx
addq   %rbx, %rax
popq   %rbx
.L6:
rep ret
```

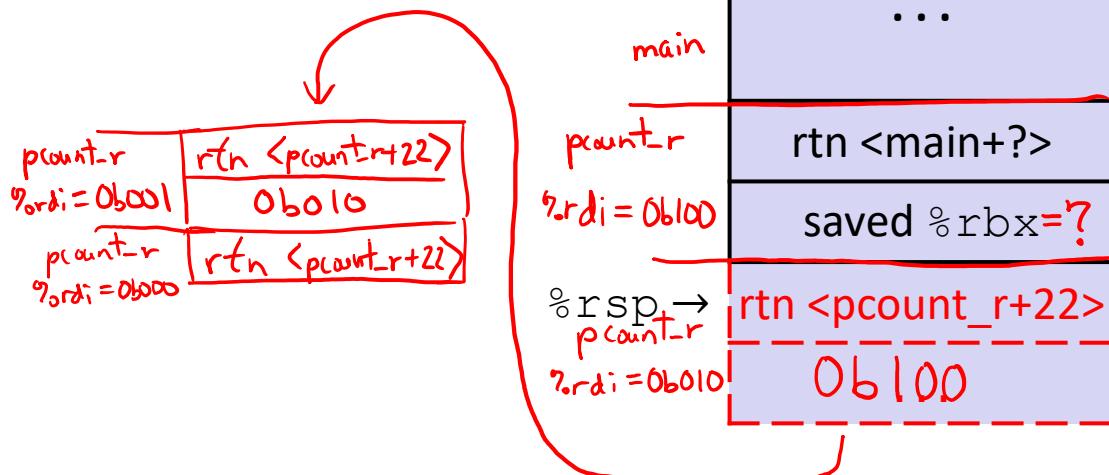
# Recursive Function: Call



```
/* Recursive popcorn */
long pcount_r(unsigned long x) {
    if (x == 0)
        return 0;
    else
        return (x&1)+pcount_r(x >> 1);
}
```

if original x=0b100:

## The Stack



Register	Use(s)	Type
%rax	Recursive call return value	Return value
%rbx	x (old)	Callee saved

pcount\_r:

<code>movl</code>	<code>\$0, %eax</code>
<code>testq</code>	<code>%rdi, %rdi</code>
<code>je</code>	<code>.L6</code>
<code>pushq</code>	<code>%rbx</code>
<code>movq</code>	<code>%rdi, %rbx</code>
<code>shrq</code>	<code>%rdi</code>
<code>call</code>	<code>pcount_r</code>
<code>andl</code>	<code>\$1, %ebx</code>
<code>addq</code>	<code>%rbx, %rax</code>
<code>popq</code>	<code>%rbx</code>
<code>.L6:</code>	
<code>rep ret</code>	

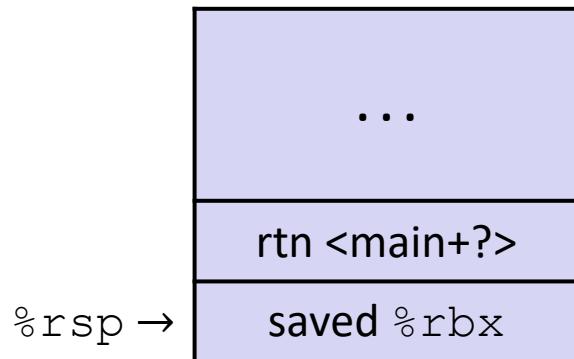


# Recursive Function: Result

```
/* Recursive popcount */
long pcount_r(unsigned long x) {
    if (x == 0)
        return 0;
    else
        return (x&1)+pcount_r(x >> 1);
}
```

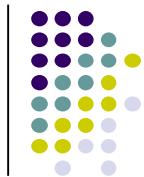
Register	Use(s)	Type
%rax	Return value	Return value
%rbx	x&1	Callee saved

## The Stack

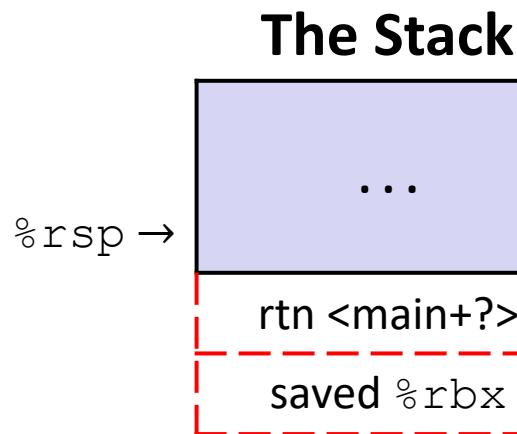


```
pcount_r:
    movl    $0, %eax
    testq   %rdi, %rdi
    je     .L6
    pushq   %rbx
    movq    %rdi, %rbx
    shrq    %rdi
    call    pcount_r
    andl    $1, %ebx
    addq    %rbx, %rax
    popq    %rbx
.L6:
    rep ret
```

# Recursive Function: Completion



```
/* Recursive popcount */
long pcount_r(unsigned long x) {
    if (x == 0)
        return 0;
    else
        return (x&1)+pcount_r(x >> 1);
}
```



Register	Use(s)	Type
%rax	Return value	Return value
%rbx	Previous %rb x value	Callee restored

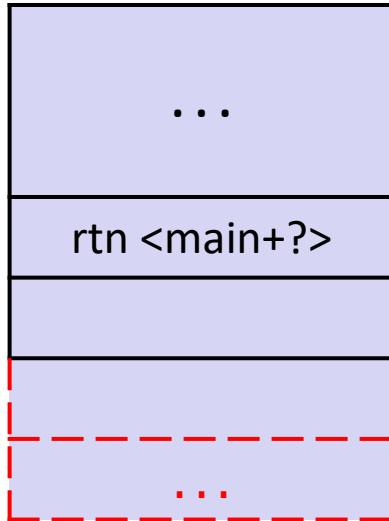
```
pcount_r:
    movl    $0, %eax
    testq   %rdi, %rdi
    je     .L6
    pushq   %rbx
    movq    %rdi, %rbx
    shrq    %rdi
    call    pcount_r
    andl    $1, %ebx
    addq    %rbx, %rax
    popq    %rbx
.L6:
    rep ret
```



# Recursive Function: Handout

```
/* Recursive popcorn */
long pcount_r(unsigned long x) {
    if (x == 0)
        return 0;
    else
        return (x&1)+pcount_r(x >> 1);
}
```

## The Stack



Register	Contents
%rdi	
%rbx	
%rax	

pcount\_r:

```
movl    $0, %eax
testq   %rdi, %rdi
je      .L6
pushq   %rbx
movq   %rdi, %rbx
shrq   %rdi
call   pcount_r
andl   $1, %ebx
addq   %rbx, %rax
popq   %rbx
.L6:
rep ret
```

# Observations About Recursion



- Handled Without Special Consideration
  - Stack frames mean that each function call has private storage
    - Saved registers & local variables
    - Saved return address
  - Register saving conventions prevent one function call from corrupting another's data
    - Unless the C code explicitly does so (e.g., buffer overflow in Lecture 9)
  - Stack discipline follows call / return pattern
    - If P calls Q, then Q returns before P
    - Last-In, First-Out
- Also works for mutual recursion
  - P calls Q; Q calls P

# x86-64 Stack Frames

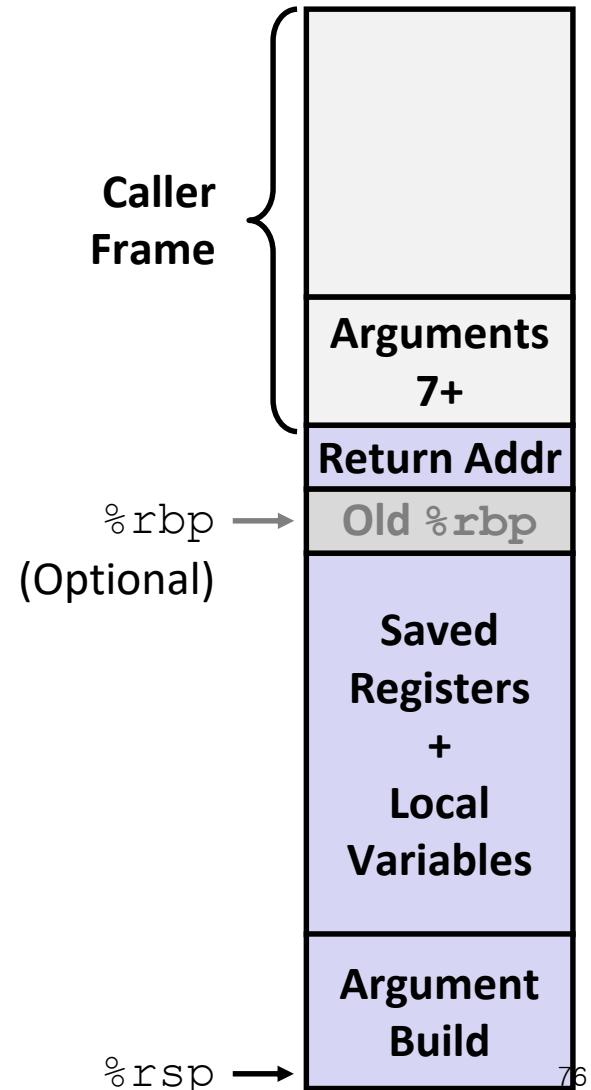


- Many x86-64 procedures have a minimal stack frame
  - Only return address is pushed onto the stack when calling procedure is called
- A procedure *needs* to grow its stack frame when it:
  - Has too many local variables to hold in **caller**-saved registers
  - Has local variables that are arrays or structs
  - Uses & to compute the address of a local variable
  - Calls another function that takes more than six arguments
  - Is using **caller**-saved registers and then calls a procedure
  - Modifies/uses **callee**-saved registers

# x86-64 Procedure Summary



- Important Points
  - Procedures are a **combination of *instructions and conventions***
    - Conventions prevent functions from disrupting each other
  - Stack is the right data structure for procedure call/return
    - If P calls Q, then Q returns before P
  - Recursion handled by normal calling conventions
- Heavy use of registers
  - Faster than using memory
  - Use limited by data size and conventions
- Minimize use of the Stack



# One more x86-64 example



- Example of passing more than 6 parameters and passing addresses of local variables
- The following example, along with a brief re-cap of x86-64 calling conventions is in this video:

# x86-64 Example (1)



```
long int call_proc()
{
    long  x1 = 1;
    int   x2 = 2;
    short x3 = 3;
    char  x4 = 4;
    proc(x1, &x1, x2, &x2,
          x3, &x3, x4, &x4);
    return (x1+x2) * (x3-x4);
}
```

```
call_proc:
    subq $32,%rsp
    movq $1,16(%rsp) # x1
    movl $2,24(%rsp) # x2
    movw $3,28(%rsp) # x3
    movb $4,31(%rsp) # x4
    • • •
```

Return address to caller of call\_proc

←%rsp

**Note:** Details may vary  
depending on compiler.

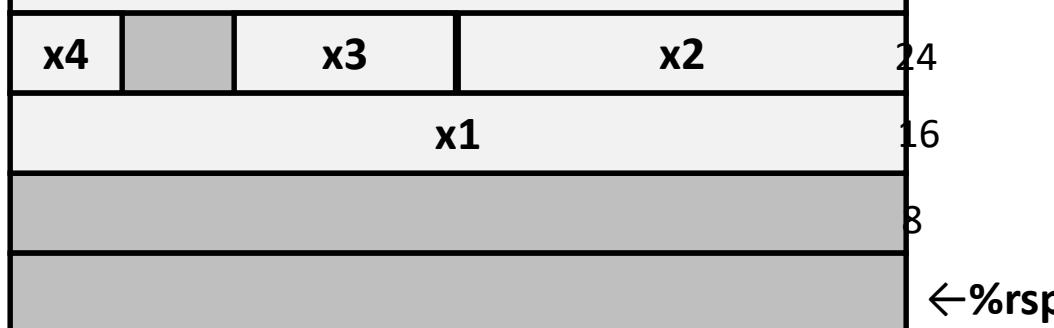
# x86-64 Example (2) – Allocate local vars



```
long int call_proc()
{
    long x1 = 1;
    int x2 = 2;
    short x3 = 3;
    char x4 = 4;
    proc(x1, &x1, x2, &x2,
          x3, &x3, x4, &x4);
    return (x1+x2) * (x3-x4);
}
```

```
call_proc:
    subq $32,%rsp
    movq $1,16(%rsp) # x1
    movl $2,24(%rsp) # x2
    movw $3,28(%rsp) # x3
    movb $4,31(%rsp) # x4
    . . .
```

Return address to caller of call\_proc



# x86-64 Example (3) – setup params to proc



```
long int call_proc()
{
    long  x1 = 1;
    int   x2 = 2;
    short x3 = 3;
    char  x4 = 4;
    proc(x1, &x1, x2, &x2,
          x3, &x3, x4, &x4);
    return (x1+x2) * (x3-x4);
}
```

Return address to caller of call\_proc

x4		x3	x2	24
		x1		16
		Arg 8		8
		Arg 7		

call\_proc:

```
• • •  
leaq  24(%rsp),%rcx # %rcx=&x2  
leaq  16(%rsp),%rsi # %rsi=&x1  
leaq  31(%rsp),%rax # %rax=&x4  
movq  %rax,8(%rsp)  # arg8=&4  
movl  $4,%rsp        # arg7=4  
leaq  28(%rsp),%r9 # %r9=&x3  
movl  $3,%r8d        # %r8 = 3  
movl  $2,%edx         # %rdx = 2  
movq  $1,%rdi        # %rdi = 1  
call  proc  
• • •
```

Arguments passed in (in order):  
rdi, rsi, rdx, rcx, r8, r9

←%rsp

Same  
instructions as in  
video, just a  
different order.<sup>80</sup>

# x86-64 Example (4) – setup params to proc



```
long int call_proc()
{
    long  x1 = 1;
    int   x2 = 2;
    short x3 = 3;
    char  x4 = 4;
    proc(x1, &x1, x2, &x2,
          x3, &x3, x4, &x4);
    return (x1+x2) * (x3-x4);
}
```

```
call_proc:
    . . .
    leaq   24(%rsp),%rcx
    leaq   16(%rsp),%rsi
    leaq   31(%rsp),%rax
    movq   %rax,8(%rsp)
    movl   $4,%rsp
    leaq   28(%rsp),%r9
    movl   $3,%r8d
    movl   $2,%edx
    movq   $1,%rdi
    call   proc
    . . .
```

Note  
sizes

Return address to caller of call\_proc

x4		x3	x2	24
		x1		16
		Arg 8		8
		Arg 7		←%rsp

Arguments passed in (in order):  
rdi, rsi, rdx, rcx, r8, r9

# x86-64 Example (5) – call proc



```
long int call_proc()
{
    long  x1 = 1;
    int   x2 = 2;
    short x3 = 3;
    char  x4 = 4;
    proc(x1, &x1, x2, &x2,
          x3, &x3, x4, &x4);
    return (x1+x2) * (x3-x4);
}
```

```
call_proc:
    . . .
    leaq   24(%rsp),%rcx
    leaq   16(%rsp),%rsi
    leaq   31(%rsp),%rax
    movq   %rax,8(%rsp)
    movl   $4,%rsp
    leaq   28(%rsp),%r9
    movl   $3,%r8d
    movl   $2,%edx
    movq   $1,%rdi
    call   proc
    . . .
```

Return address to caller of call\_proc

x4		x3	x2
----	--	----	----

x1

Arg 8

Arg 7

Return address to line after call to proc

←%rsp

# x86-64 Example (6) – after call to proc



```
long int call_proc()
{
    long x1 = 1;
    int x2 = 2;
    short x3 = 3;
    char x4 = 4;
    proc(x1, &x1, x2, &x2,
          x3, &x3, x4, &x4);
    return (x1+x2)*(x3-x4);
}
```

```
call_proc:
    • • •
    movswl 28(%rsp), %eax # %eax=x3
    movsbl 31(%rsp), %edx # %edx=x4
    subl    %edx, %eax # %eax=x3-x4
    cltq
    movslq 24(%rsp), %rdx # %rdx=x2
    addq    16(%rsp), %rdx # %rdx=x1+x2
    imulq   %rdx, %rax # %rax=rax*rdx
    addq    $32, %rsp
    ret
```

Return address to caller of call_proc				
x4		x3	x2	24
		x1		16
		Arg 8		8
		Arg 7		←%rsp

**movs\_\_:**  
move and sign extend

**cltq:**  
sign extend %eax into %rax  
(special-case to save space)

# x86-64 Example (7) – de-allocate local vars



```
long int call_proc()
{
    long  x1 = 1;
    int   x2 = 2;
    short x3 = 3;
    char  x4 = 4;
    proc(x1, &x1, x2, &x2,
          x3, &x3, x4, &x4);
    return (x1+x2)*(x3-x4);
}
```

```
call_proc:
    • • •
    movswl 28(%rsp),%eax
    movsbl 31(%rsp),%edx
    subl    %edx,%eax
    cltq
    movslq 24(%rsp),%rdx
    addq    16(%rsp),%rdx
    imulq  %rdx,%rax
    addq    $32,%rsp
    ret
```

Return address to caller of call\_proc

←%rsp

# Q&A

