

System Programming

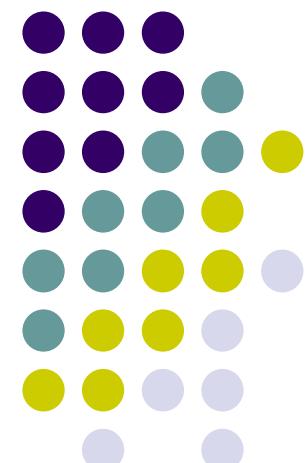
08. Machine-Level Programming IV: Data (ch 3.8)

2019. Fall

Instructor: Joonho Kwon

jhwon@pusan.ac.kr

Data Science Lab @ PNU





Roadmap

C:

```
car *c = malloc(sizeof(car));
c->miles = 100;
c->gals = 17;
float mpg = get_mpg(c);
free(c);
```

Java:

```
Car c = new Car();
c.setMiles(100);
c.setGals(17);
float mpg =
    c.getMPG();
```

Assembly
language:

```
get_mpg:
    pushq    %rbp
    movq    %rsp, %rbp
    ...
    popq    %rbp
    ret
```

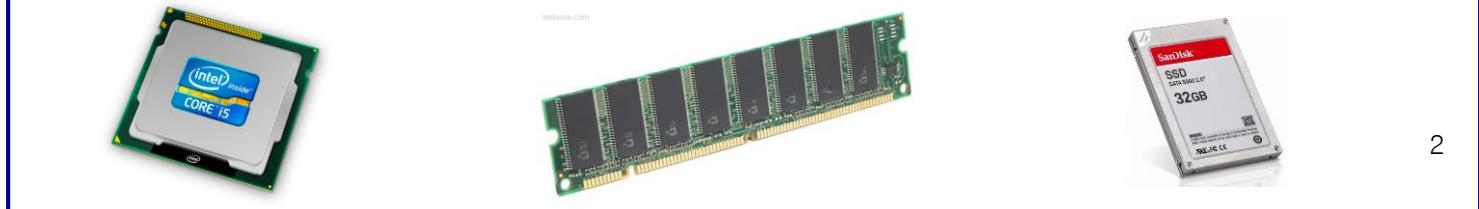
Machine
code:

```
0111010000011000
1000110100000100000000010
1000100111000010
110000011111101000011111
```

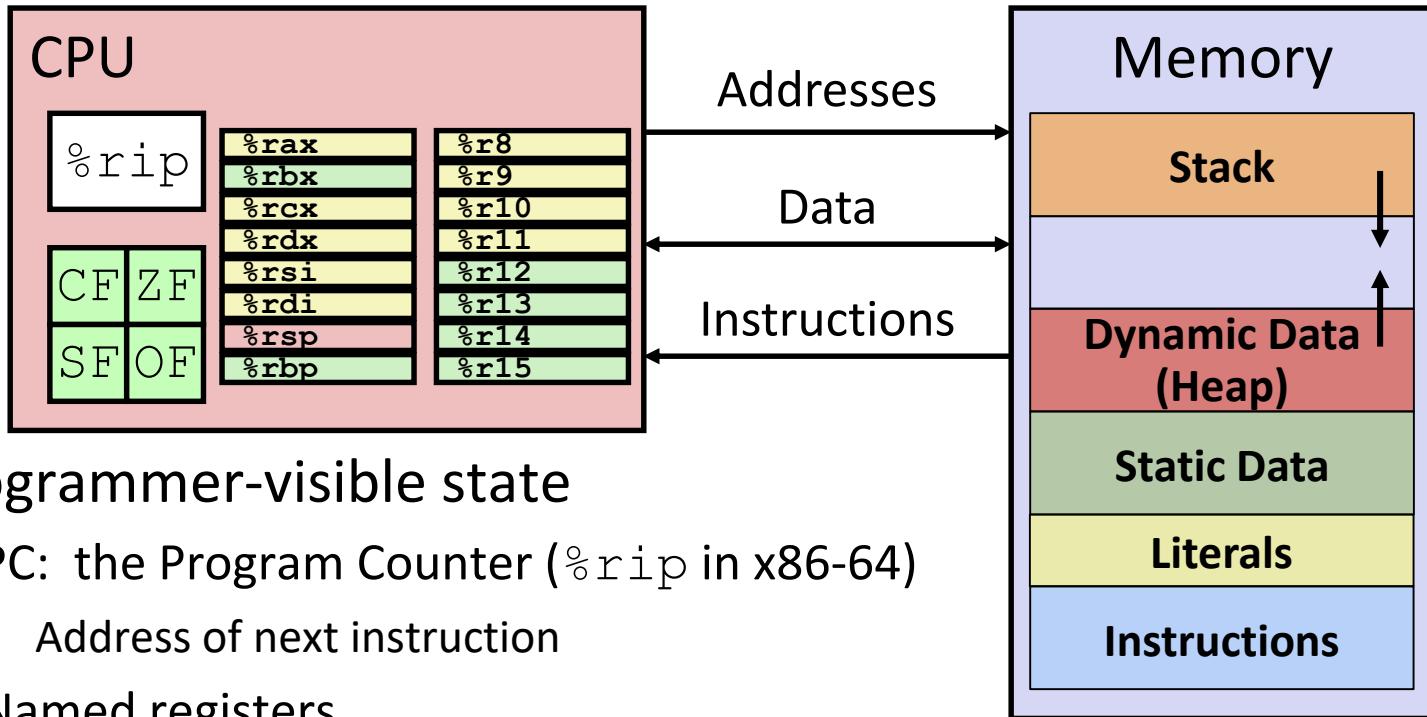
Computer
system:

- Memory & data
- Integers & floats
- x86 assembly
- Procedures & stacks
- Executables
- Arrays & **structs**
- Memory & caches
- Processes
- Virtual memory
- Memory allocation
- Java vs. C

OS:



Assembly Programmer's View



❖ Programmer-visible state

- PC: the Program Counter (`%rip` in x86-64)
 - Address of next instruction
- Named registers
 - Together in “register file”
 - Heavily used program data
- Condition codes
 - Store status information about most recent arithmetic operation
 - Used for conditional branching

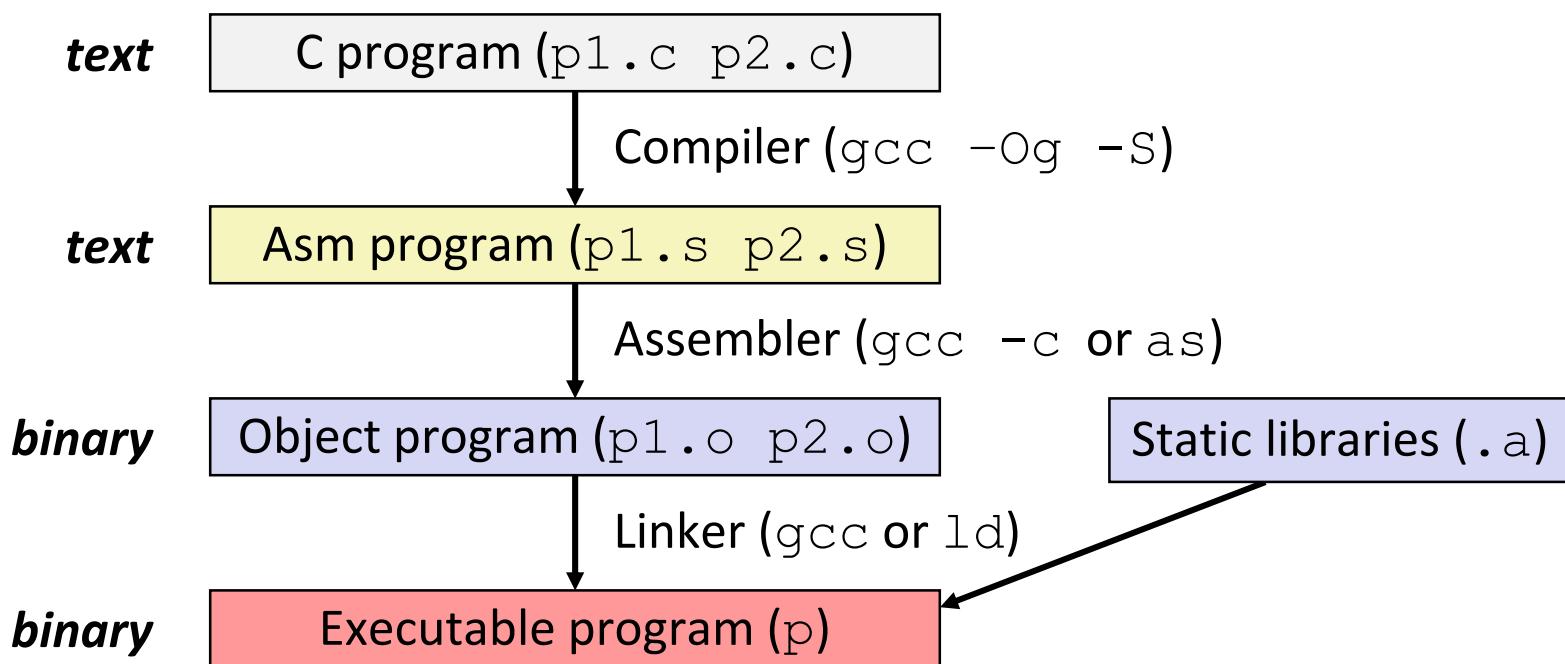
❖ Memory

- Byte-addressable array
- Code and user data
- Includes *the Stack* (for supporting procedures)

Turning C into Object Code



- Code in files `p1.c p2.c`
- Compile with command: `gcc -Og p1.c p2.c -o p`
 - Use basic optimizations (`-Og`) [New to recent versions of GCC]
 - Put resulting machine code in file `p`



Assembling



- Executable has **addresses**

assembler

```
00000000004004f6 <pcount_r>:  
 4004f6: b8 00 00 00 00    mov    $0x0,%eax  
 4004fb: 48 85 ff          test   %rdi,%rdi  
 4004fe: 74 13              je     400513 <pcount_r+0x1d>  
 400500: 53                 push   %rbx  
 400501: 48 89 fb          mov    %rdi,%rbx  
 400504: 48 d1 ef          shr    %rdi  
 400507: e8 ea ff ff ff    callq  4004f6 <pcount_r>  
 40050c: 83 e3 01          and    $0x1,%ebx  
 40050f: 48 01 d8          add    %rbx,%rax  
 400512: 5b                 pop    %rbx  
 400513: f3 c3              rep    ret
```

- gcc -g pcount.c -o pcount
- objdump -d pcount



Picture of Memory (64-bit view)

```
00000000004004f6 <pcount_r>:  
4004f6: b8 00 00 00 00    mov    $0x0,%eax  
4004fb: 48 85 ff        test   %rdi,%rdi  
4004fe: 74 13            je     400513 <pcount_r+0x1d>  
400500: 53                push   %rbx  
400501: 48 89 fb        mov    %rdi,%rbx  
400504: 48 d1 ef        shr    %rdi  
400507: e8 ea ff ff ff  callq  4004f6 <pcount_r>  
40050c: 83 e3 01        and    $0x1,%ebx  
40050f: 48 01 d8        add    %rbx,%rax  
400512: 5b                pop    %rbx  
400513: f3 c3            rep    ret
```

Roadmap



C:

```
car *c = malloc(sizeof(car));
c->miles = 100;
c->gals = 17;
float mpg = get_mpg(c);
free(c);
```

Java:

```
Car c = new Car();
c.setMiles(100);
c.setGals(17);
float mpg =
    c.getMPG();
```

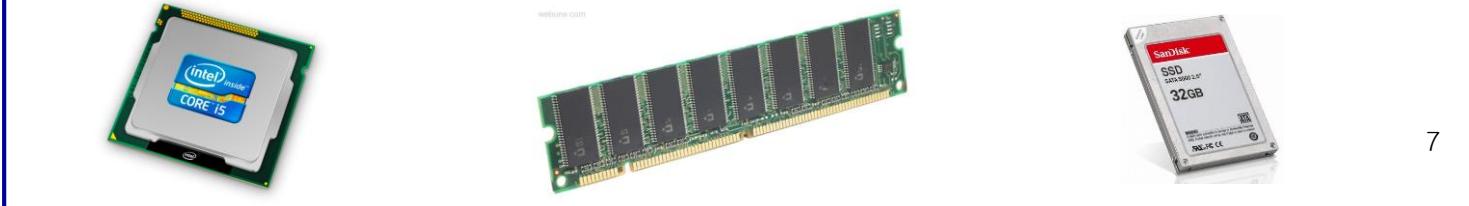
Assembly language:

```
get_mpg:
pushq %rbp
movq %rsp, %rbp
...
popq %rbp
ret
```

Machine code:

```
0111010000011000
1000110100000100000000010
1000100111000010
110000011111101000011111
```

Computer system:



Memory & data
Integers & floats
x86 assembly
Procedures & stacks
Executables
Arrays & structs
Memory & caches
Processes
Virtual memory
Operating Systems

OS:



Data Structures in Assembly



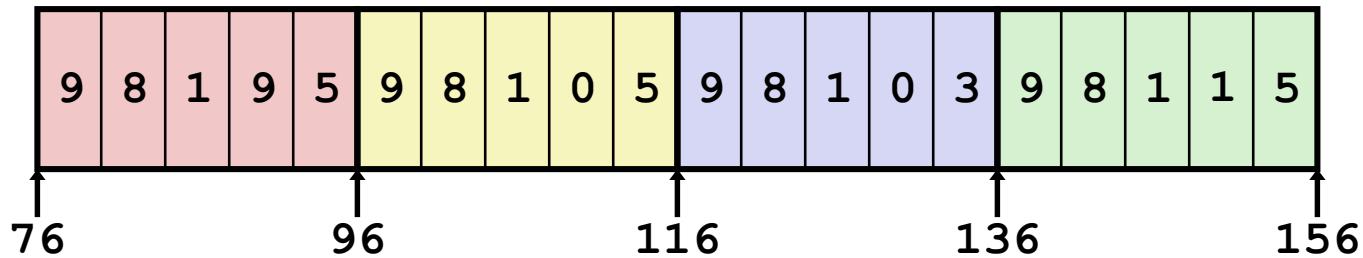
- Arrays
 - One-dimensional
 - Multi-dimensional (nested)
 - Multi-level
- Structs
 - Alignment
- Unions

Question



- Which of the following statements is FALSE?

```
int sea[4][5];
```



- A. **sea[4][-2]** is a *valid* array reference
- B. **sea[1][1]** makes *two* memory accesses
- C. **sea[2][1]** will *always* be a higher address than **sea[1][2]**
- D. **sea[2]** is calculated using *only* 1ea

Structs in C



- Way of defining compound data types
- A structured group of variables, possibly including other structs

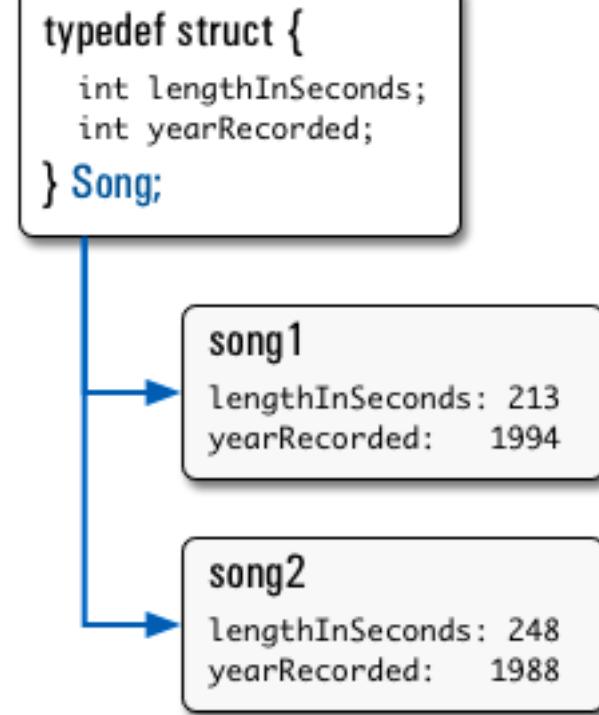
```
typedef struct {
    int lengthInSeconds;
    int yearRecorded;
} Song;

Song song1;

song1.lengthInSeconds = 213;
song1.yearRecorded     = 1994;

Song song2;

song2.lengthInSeconds = 248;
song2.yearRecorded     = 1988;
```



Review: Structs



```
// Use typedef to create a type: FourInts
typedef struct {
    int a, b, c, d;
} FourInts; // Name of type is "FourInts"

int main(int argc, char* argv[]) {
    FourInts f1; // Allocates memory to hold a FourInts
                    // (16 bytes) on stack (local variable)
    f1.a = 0;      // Assign first field in f1 to be zero

    FourInts* f2; // Declare f2 as a pointer to FourInts

    // Allocate space for a FourInts on the heap,
    // f2 is a "pointer to"/"address of" this space.
    f2 = (FourInts*) malloc(sizeof(FourInts));
    f2->b = 17;   // Assign the second field to be 17
    ...
}
```

Aside: Syntax for structs without `typedef`



```
struct rec {           // Declares the type "struct rec"  
    int a[4];         // Total size = _____ bytes  
    long i;  
    struct rec *next;  
};  
struct rec r1;        // Allocates memory to hold a struct rec  
                      // named r1, on stack or globally,  
                      // depending on where this code appears  
  
struct rec *r;         // Allocates memory for a pointer  
r = &r1;              // Initializes r to "point to" r1
```

More Structs Syntax



Declaring a struct **struct rec**, then declaring a variable **r1**:

```
struct rec {           // Declares the type "struct rec"  
    int a[4];  
    long i;  
    struct rec *next;  
};  
struct rec r1;         // Declares r1 as a struct rec
```

Equivalent to:

```
struct rec {           // Declares the type "struct rec"  
    int a[4];  
    long i;  
    struct rec *next;  
} r1;                 // Declares r1 as a struct rec
```

Declare type **struct rec** and variable **r1** at the same time! 13

Another Syntax Example



Declaring a struct **struct rec**, then declaring a variable **r**:

```
struct rec {           // Declares the type "struct rec"
    int a[4];
    long i;
    struct rec *next;
};
struct rec *r;         // Declares r as pointer to a struct rec
```

Equivalent to:

```
struct rec {           // Declares the type "struct rec"
    int a[4];
    long i;
    struct rec *next;
} *r;                 // Declares r as pointer to a struct rec
```

Declare type **struct rec** and variable **r** at the same time!

Struct Definitions



- Structure definition:

- Does NOT declare a variable

- Variable definitions:

- Variable type is “struct name”

```
struct name {  
    /* fields */  
};
```

Easy to forget
semicolon!

```
struct name name1, *pn, name_ar[3];
```

pointer

array

- Joint struct definition and typedef

```
struct nm {  
    /* fields */  
};  
typedef struct nm name;  
name n1;
```



```
typedef struct {  
    /* fields */  
} name;  
name n1;
```

Scope of Struct Definition



- Why is placement of struct definition important?
 - What actually happens when you declare a variable?
 - Creating space for it somewhere!
 - Without definition, program doesn't know how much space

4B x 4
8B

```
struct data {  
    int ar[4];  
    long d;  
};
```

← Size = 24 bytes

Size = 32 bytes →

```
struct rec {  
    int a[4];  
    long i;  
    8B struct rec* next;  
};
```

- Almost always define structs in global scope near the top of your C file
 - Struct definitions follow normal rules of scope

Accessing Structure Members



- Given a struct instance, access member using the . operator:

```
struct rec r1;  
r1.i = val;
```

```
struct rec {  
    int a[4];  
    long i;  
    struct rec *next;  
};
```

- Given a *pointer* to a struct:

```
struct rec *r;  
r = &r1; // or malloc space for r to point to
```

We have two options:

- Use * and . operators: (*r).i = val;
- Use -> operator for short: r->i = val;

- In assembly: register holds address of the first byte
 - Access members with offsets

Java side-note

Java:

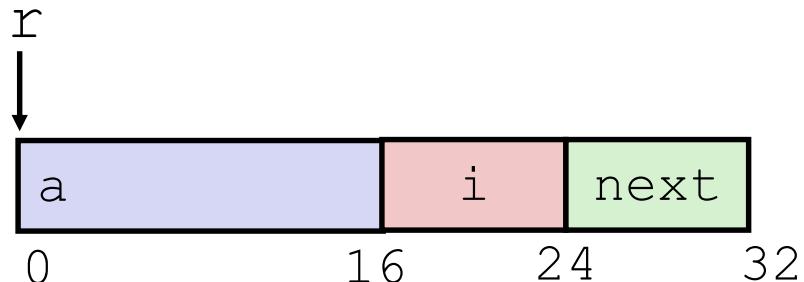
```
class Record { ... }  
Record x = new Record();
```

- An instance of a class is like a *pointer to* a struct containing the fields
 - (Ignoring methods and subclassing for now)
 - So Java's `x.f` is like C's `x->f` or `(*x).f`
- In Java, almost everything is a pointer ("reference") to an object
 - Cannot declare variables or fields that are structs or arrays
 - Always a *pointer* to a struct or array
 - So every Java variable or field is ≤ 8 bytes (but can point to lots of data)

Structure Representation



```
struct rec {  
    int a[4];  
    long i;  
    struct rec *next;  
} *r;
```

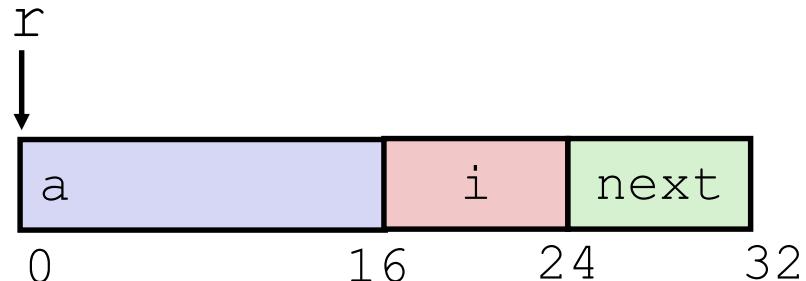


- Characteristics
 - Contiguously-allocated region of memory
 - Refer to members within structure by names
 - Members may be of different types

Structure Representation



```
struct rec {  
    int a[4];  
    long i;  
    struct rec *next;  
} *r;
```

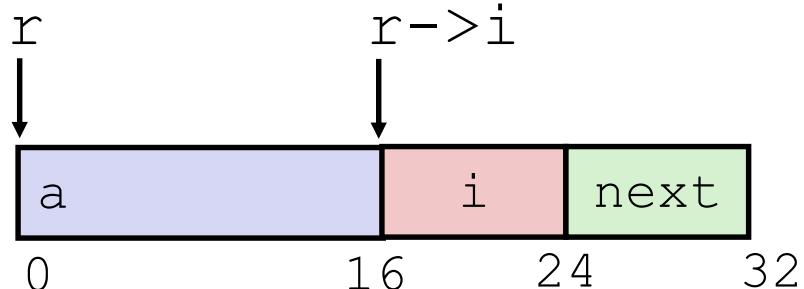


- Structure represented as block of memory
 - Big enough to hold all of the fields
- Fields ordered according to declaration order
 - Even if another ordering could yield a more compact representation
- Compiler determines overall size + positions of fields
 - Machine-level program has no understanding of the structures in the source code

Accessing a Structure Member



```
struct rec {
    int a[4];
    long i;
    struct rec *next;
} *r;
```



- Compiler knows the *offset* of each member within a struct.
 - Compute as $\ast(r + \text{offset})$
 - Referring to absolute offset, so no pointer arithmetic

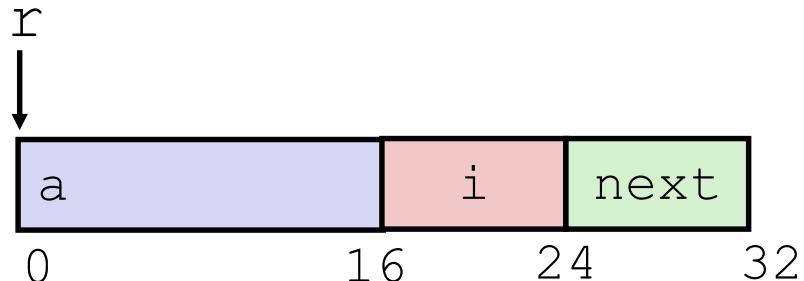
```
long get_i(struct rec *r)
{
    return r->i;
}
```

```
# r in %rdi
movq 16(%rdi), %rax
ret
```

Exercise: Pointer to Structure Member



```
struct rec {
    int a[4];
    long i;
    struct rec *next;
} *r;
```



```
long* addr_of_i(struct rec *r)
{
    return &(r->i);
}
```

```
# r in %rdi
_____,_____,%rax
ret
```

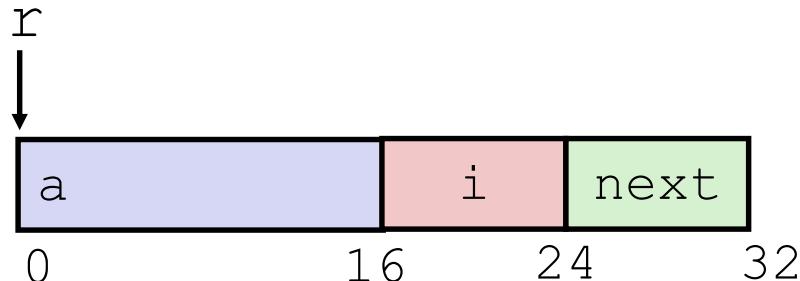
```
struct rec** addr_of_next(struct rec *r)
{
    return &(r->next);
}
```

```
# r in %rdi
_____,_____,%rax
ret
```

Exercise: Pointer to Structure Member



```
struct rec {  
    int a[4];  
    long i;  
    struct rec *next;  
} *r;
```



```
long* addr_of_i(struct rec *r)  
{  
    return &(r->i);  
}
```

want address

```
struct rec** addr_of_next(struct rec *r)  
{  
    return &(r->next);  
}
```

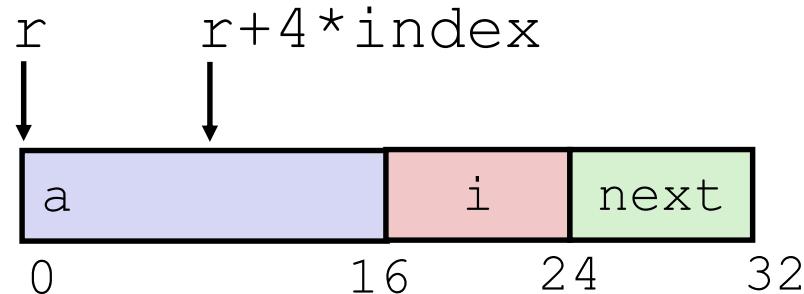
r in %rdi
leaq 16(%rdi),%rax
ret

r in %rdi
leaq 24(%rdi),%rax
ret

Generating Pointer to Array Element



```
struct rec {  
    int a[4];  
    long i;  
    struct rec *next;  
} *r;
```



- Generating Pointer to Array Element
 - Offset of each structure member determined at compile time
 - Compute as: $r + 4 * \text{index}$

```
int* find_addr_of_array_elem  
(struct rec *r, long index)  
{  
    return &r->a[index];  
}
```

A purple bracket underlines the expression `&r->a[index]`, and a red bracket highlights the same expression in red. A purple arrow points from the purple bracket to the red bracket.

```
# r in %rdi, index in %rsi  
leaq (%rdi,%rsi,4), %rax  
ret
```

Review: Memory Alignment in x86-64



- For good memory system performance, Intel recommends data be aligned
 - However the x86-64 hardware will work correctly regardless of alignment of data.
- *Aligned* means:
 - Any primitive object of K bytes must have an address that is a multiple of K.
- This means we could expect these types to have starting addresses that are the following multiples:

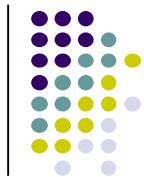
K	Type	Addresses
1	char	No restrictions
2	short	Lowest bit must be zero: ...0 ₂
4	int, float	Lowest 2 bits zero: ...00 ₂
8	long, double, * (pointers)	Lowest 3 bits zero: ...000 ₂
16	long double	Lowest 4 bits zero: ...0000 ₂

Alignment Principles

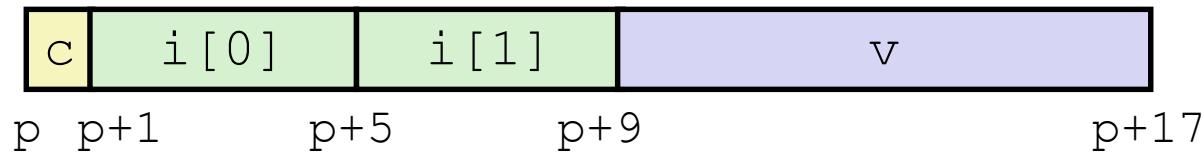


- Aligned Data
 - Primitive data type requires K bytes
 - Address must be multiple of K
 - Required on some machines; advised on x86-64
- Motivation for Aligning Data
 - Memory accessed by (aligned) chunks of 4 or 8 bytes (system dependent)
 - Inefficient to load or store value that spans quad word boundaries
 - Virtual memory trickier when value spans 2 pages (more on this later)

Structures & Alignment



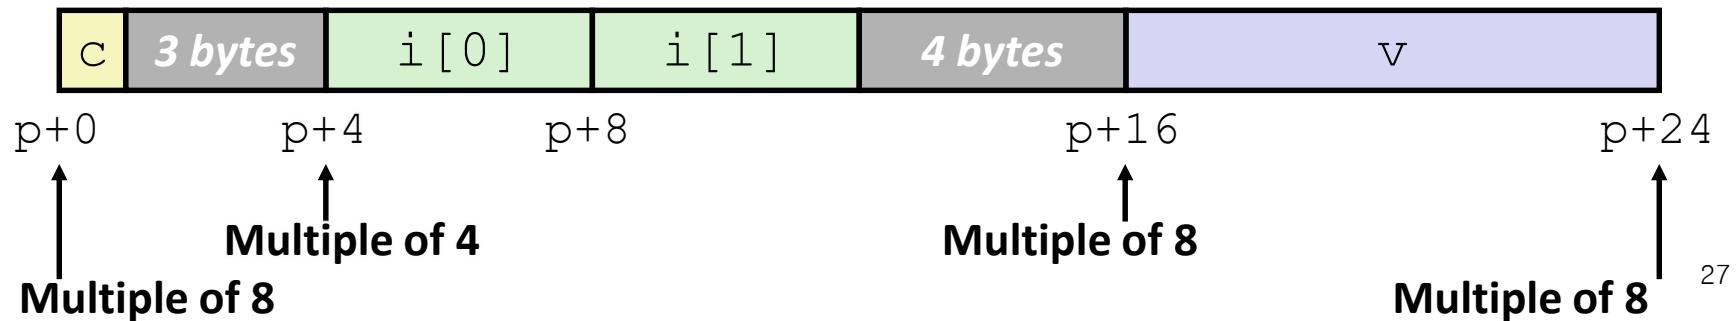
- Unaligned Data



```
struct S1 {  
    char c;  
    int i[2];  
    double v;  
} *p;
```

- Aligned Data

- Primitive data type requires K bytes
- Address must be multiple of K

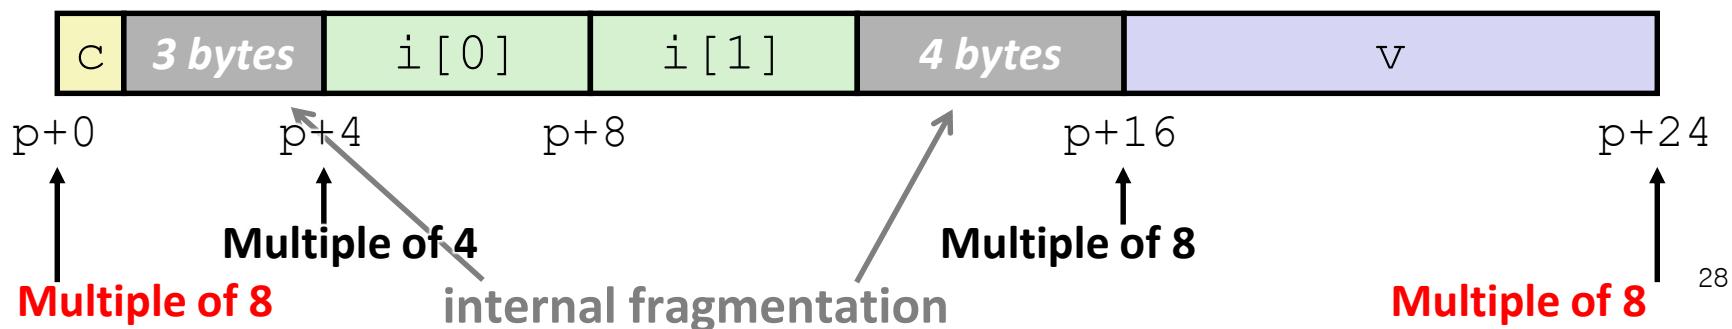


Satisfying Alignment with Structures (1)



- Within structure:
 - Must satisfy each element's alignment requirement
- Overall structure placement
 - Each structure has alignment requirement K_{\max}
 - K_{\max} = Largest alignment of any element
 - Counts individual items in the array as elements (entire array is not an “element”)
 - **Initial address of structure & structure length must be multiples of K**
- Example:
 - $K_{\max} = 8$, due to double element

```
struct S1 {  
    char c;  
    int i[2];  
    double v;  
} *p;
```



Satisfying Alignment Requirements:

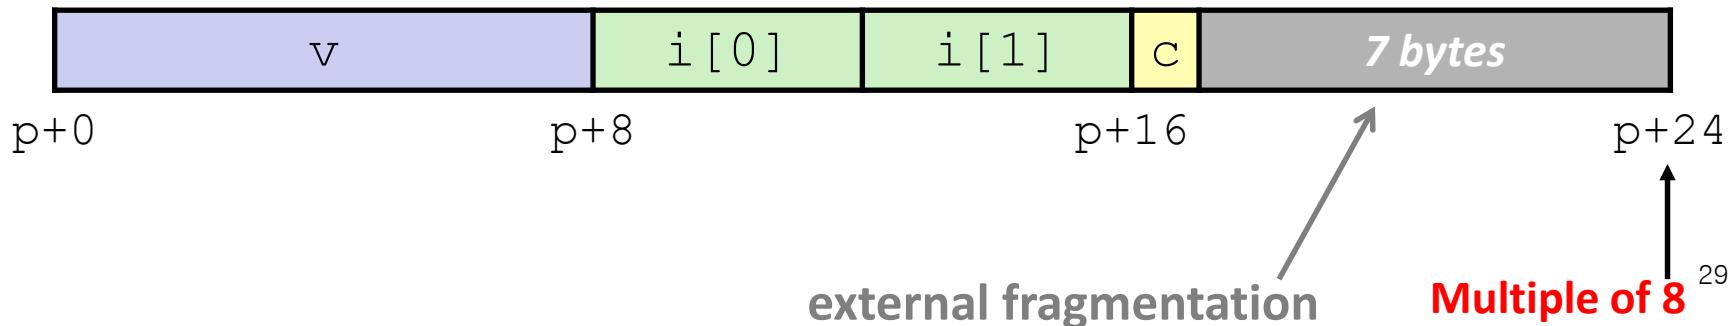


Another Example

- Can find offset of individual fields using `offsetof()`
 - Need to `#include <stddef.h>`
 - Example: `offsetof(struct S2, c)` returns 16

```
struct S2 {  
    double v;  
    int i[2];  
    char c;  
} *p;
```

- For largest alignment requirement K_{\max} ,
overall structure size must be multiple of K_{\max}
 - Compiler will add padding **at end** of structure to meet overall structure alignment requirement



Alignment of Structs



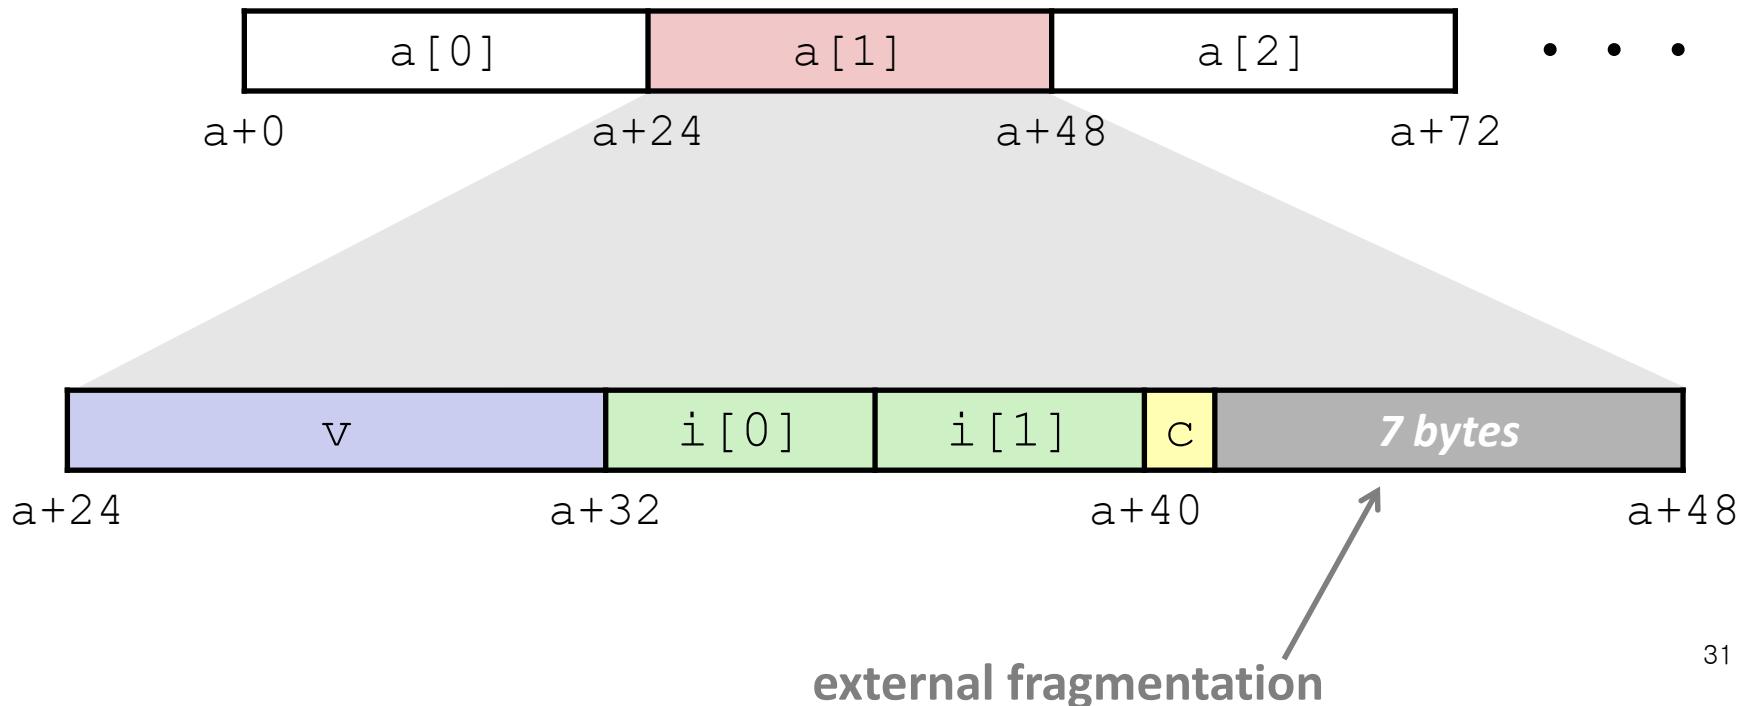
- Compiler will do the following:
 - Maintains declared *ordering* of fields in struct
 - Each **field** must be aligned *within* the struct
(*may insert padding*)
 - offsetof can be used to get actual field offset
 - Overall struct must be **aligned** according to largest field
 - Total struct **size** must be multiple of its alignment
(*may insert padding*)
 - sizeof should be used to get true size of structs

Arrays of Structures

Create an array of ten S2 structs called "a"

- Overall structure length multiple of K_{max}
- Satisfy alignment requirement for every element in array

```
struct S2 {  
    double v;  
    int i[2];  
    char c;  
} a[10];
```

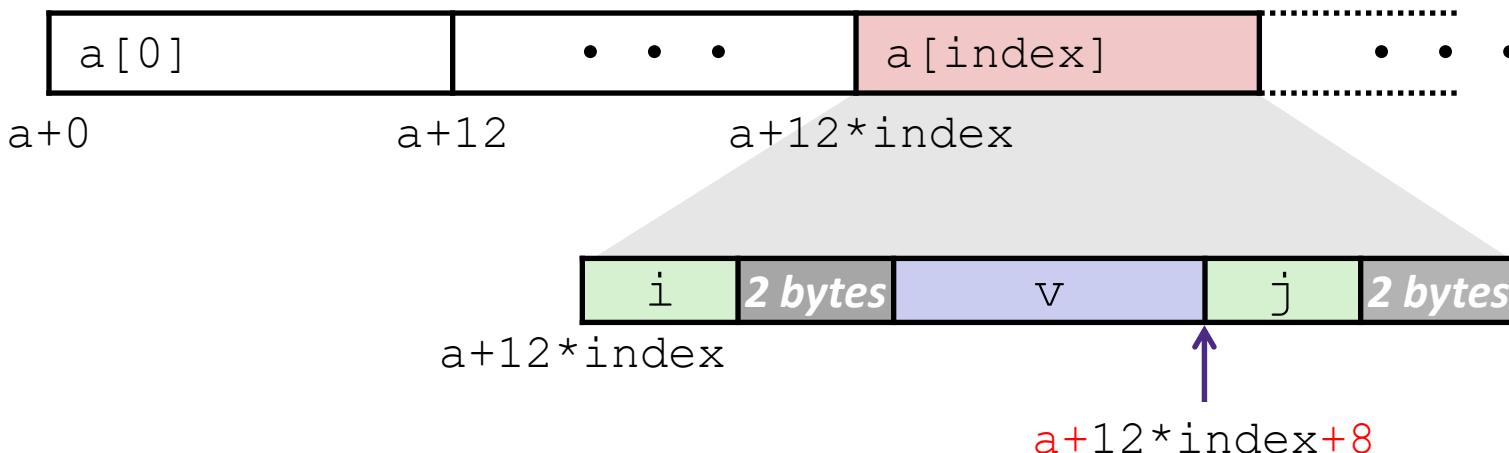


Accessing Array Elements

Create an array of
ten S3 structs
called "a"

- Compute start of array element as: $12 * \text{index}$
 - `sizeof(S3) = 12`, including alignment padding
- Element j is at offset 8 within structure
- Assembler gives offset $a+8$

```
struct S3 {  
    short i;  
    float v;  
    short j;  
} a[10];
```



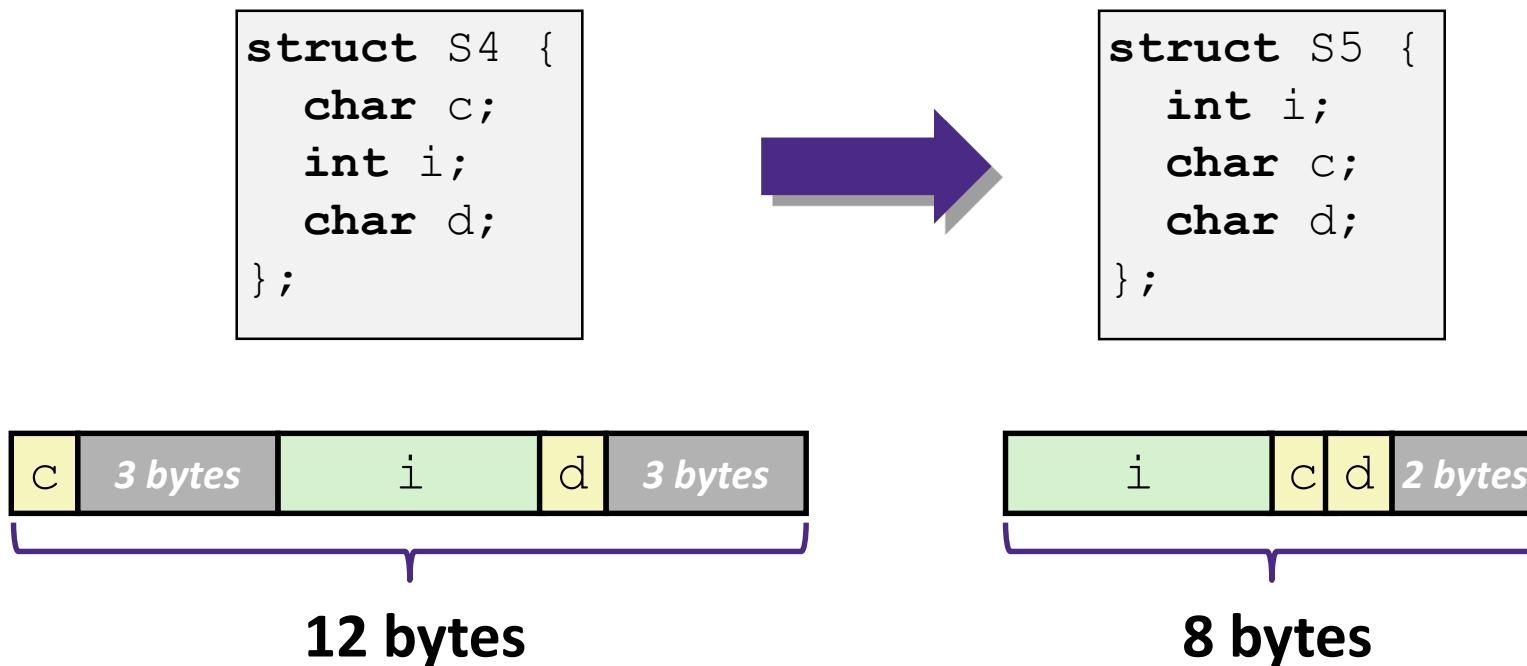
```
short get_j(int index)  
{  
    return a[index].j;  
}
```

```
# %rdi = index  
leaq (%rdi,%rdi,2),%rax # 3*index  
movzwl a+8(%rax,4),%eax
```

How the Programmer Can Save Space



- Compiler must respect order elements are declared in
 - Sometimes the programmer can save space by declaring large data types first

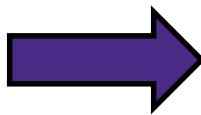


Question



- Minimize the size of the struct by re-ordering the vars

```
struct old {  
    int i;  
  
    short s[3];  
  
    char *c;  
  
    float f;  
};
```



```
struct new {  
    int i;  
  
    _____ _____;  
    _____ _____;  
    _____ _____;  
};
```

- What are the old and new sizes of the struct?

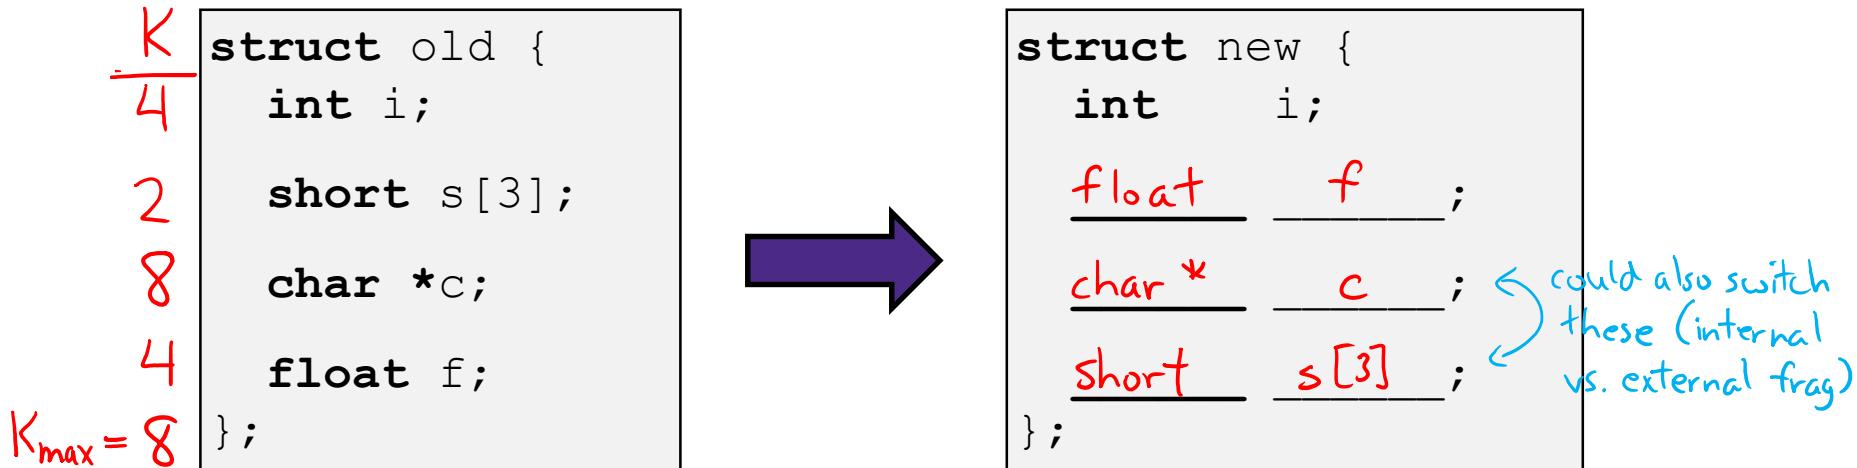
`sizeof(struct old) = _____`

`sizeof(struct new) = _____`

Answers



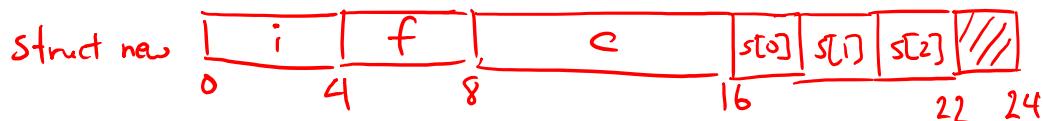
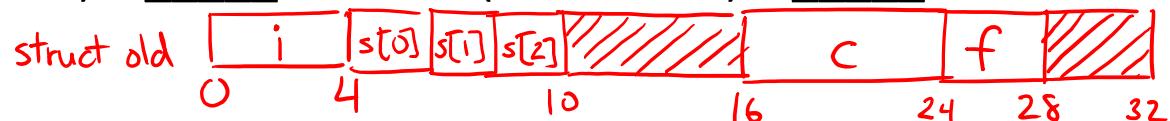
- Minimize the size of the struct by re-ordering the vars



- What are the old and new sizes of the struct?

32 B 24 B

- `sizeof(struct old) = _____` `sizeof(struct new) = _____`

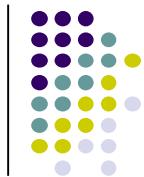


Data Structures in Assembly

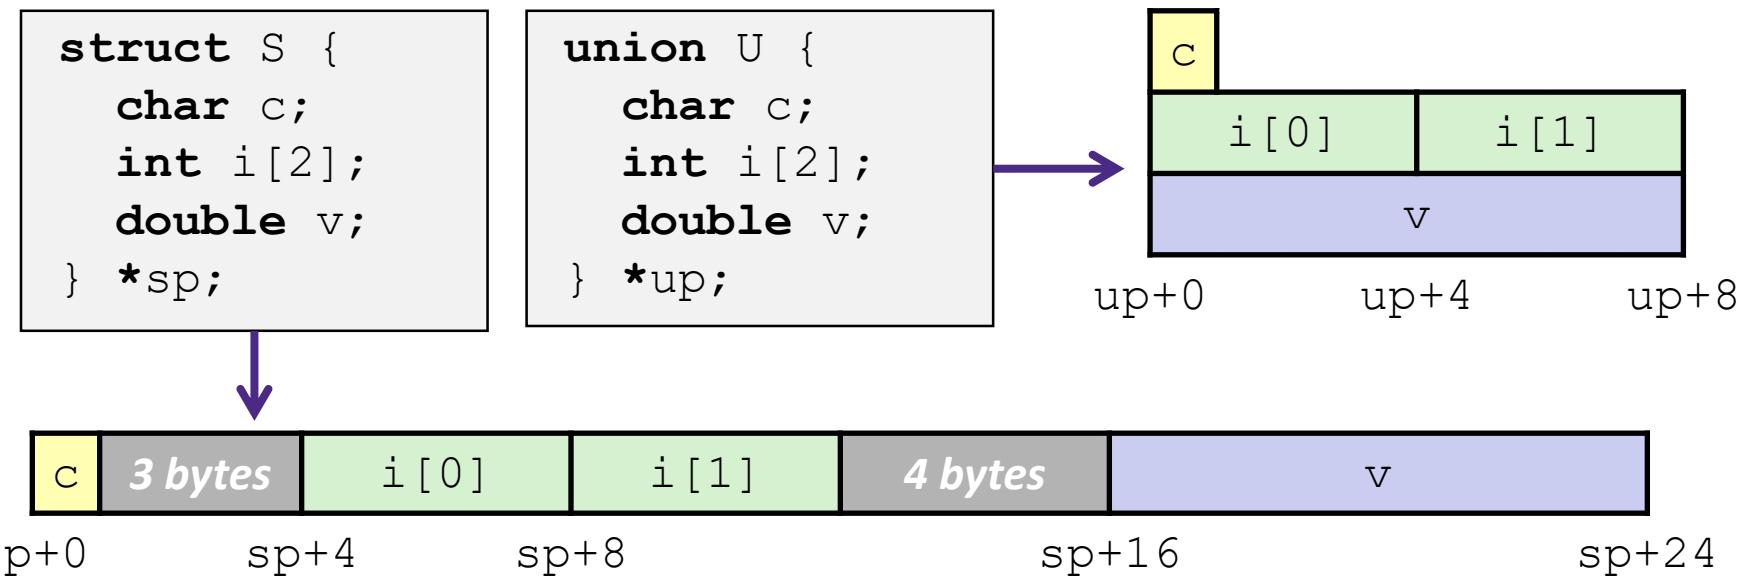


- Arrays
 - One-dimensional
 - Multi-dimensional (nested)
 - Multi-level
- Structs
 - Alignment
- **Unions**

Unions



- Only allocates enough space for the **largest element** in union
- Can only use one member at a time



What Are Unions Good For?



- Unions allow the same region of memory to be referenced as different types
 - Different “views” of the same memory location
 - Can be used to circumvent C’s type system (bad idea and technically not guaranteed to work)
- Better idea: use a struct inside a union to access some memory location either as a whole or by its parts
- But watch out for endianness at a small scale...
- Layout details are implementation/machine-specific

```
union int_or_bytes {
    int i;
    struct bytes {
        char b0, b1, b2, b3;
    }
}
```

Using Union to Access Bit Patterns



```
typedef union {
    float f;
    unsigned u;
} bit_float_t;
```



```
float bit2float(unsigned u)
{
    bit_float_t arg;
    arg.u = u;
    return arg.f;
}
```

```
unsigned float2bit(float f)
{
    bit_float_t arg;
    arg.f = f;
    return arg.u;
}
```

Same as `(float) u` ?

Same as `(unsigned) f` ?

Byte Ordering Revisited



- Idea
 - Short/long/quad words stored in memory as 2/4/8 consecutive bytes
 - Which byte is most (least) significant?
 - Can cause problems when exchanging binary data between machines
- Big Endian
 - Most significant byte has lowest address
 - Sparc
- Little Endian
 - Least significant byte has lowest address
 - Intel x86, ARM Android and IOS
- Bi Endian
 - Can be configured either way
 - ARM

Byte Ordering Example



```
union {
    unsigned char c[8];
    unsigned short s[4];
    unsigned int i[2];
    unsigned long l[1];
} dw;
```

32-bit

c[0]	c[1]	c[2]	c[3]	c[4]	c[5]	c[6]	c[7]				
s[0]		s[1]		s[2]		s[3]					
i[0]				i[1]							
l[0]											

64-bit

c[0]	c[1]	c[2]	c[3]	c[4]	c[5]	c[6]	c[7]				
s[0]		s[1]		s[2]		s[3]					
i[0]				i[1]							
l[0]											

Byte Ordering Example (Cont).



```
int j;
for (j = 0; j < 8; j++)
    dw.c[j] = 0xf0 + j;

printf("Characters 0-7 ==
[0x%x,0x%x,0x%x,0x%x,0x%x,0x%x,0x%x,0x%x]\n",
       dw.c[0], dw.c[1], dw.c[2], dw.c[3],
       dw.c[4], dw.c[5], dw.c[6], dw.c[7]);

printf("Shorts 0-3 == [0x%x,0x%x,0x%x,0x%x]\n",
       dw.s[0], dw.s[1], dw.s[2], dw.s[3]);

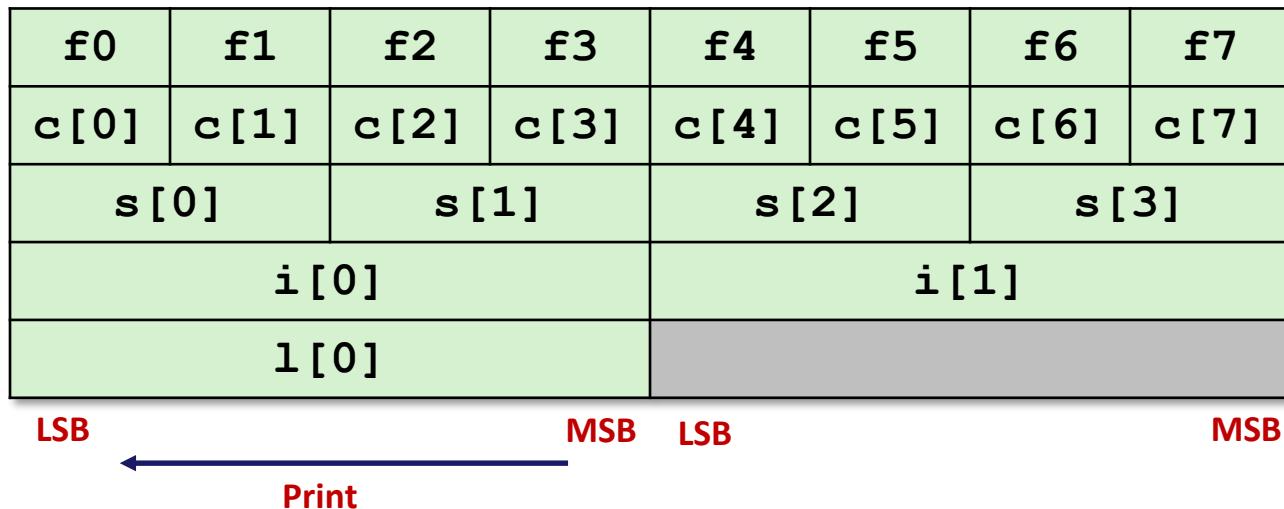
printf("Ints 0-1 == [0x%x,0x%x]\n",
       dw.i[0], dw.i[1]);

printf("Long 0 == [0x%lx]\n",
       dw.l[0]);
```



Byte Ordering on IA32

Little Endian



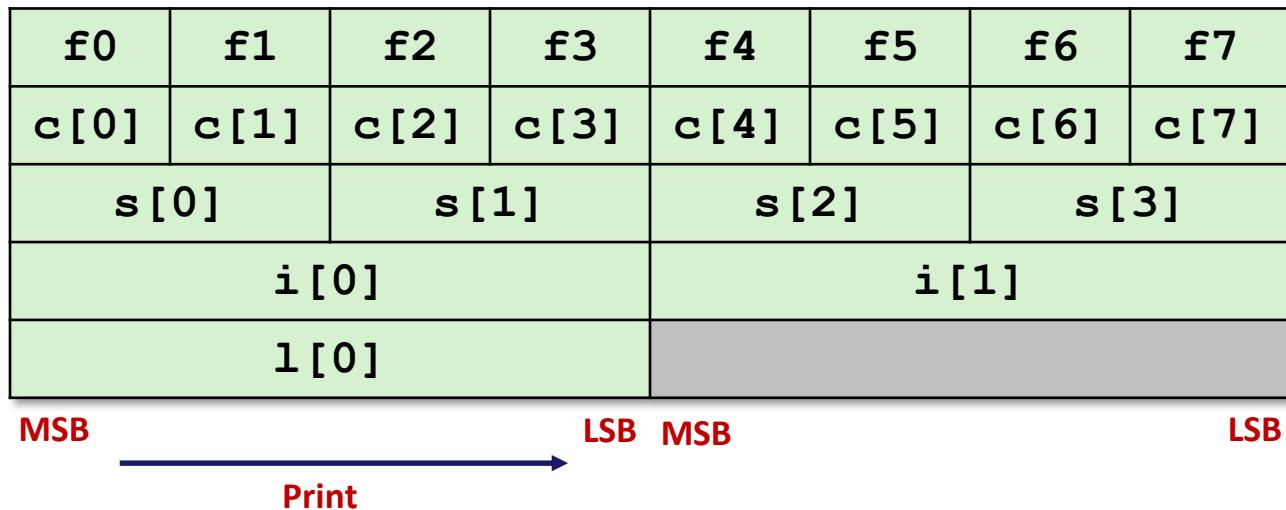
Output:

```
Characters 0-7 == [0xf0,0xf1,0xf2,0xf3,0xf4,0xf5,0xf6,0xf7]
Shorts      0-3 == [0xf1f0,0xf3f2,0xf5f4,0xf7f6]
Ints        0-1 == [0xf3f2f1f0,0xf7f6f5f4]
Long         0    == [0xf3f2f1f0]
```

Byte Ordering on Sun



Big Endian



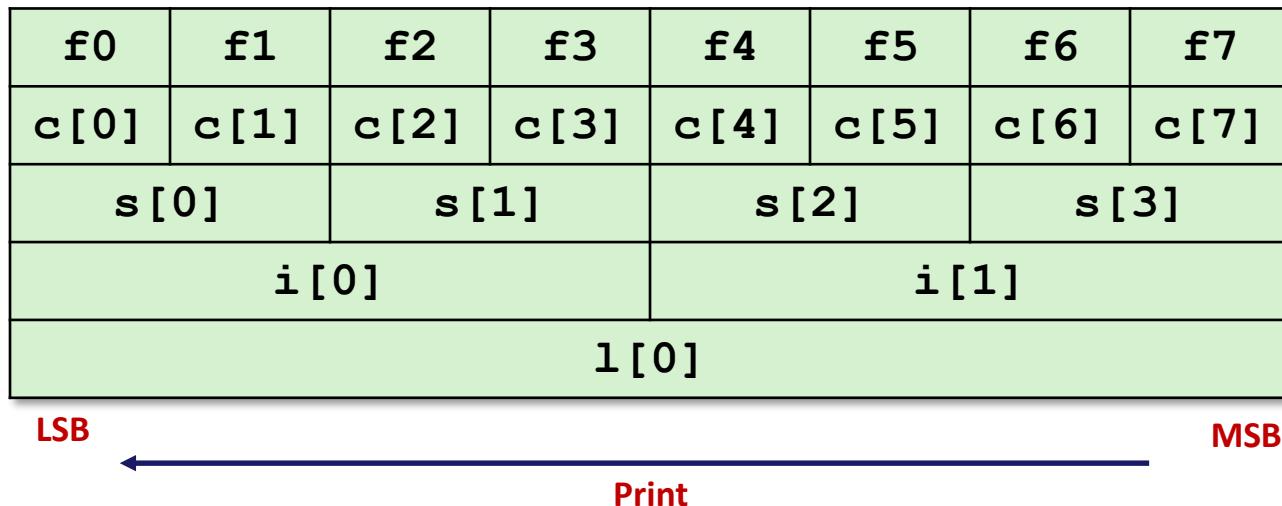
Output on Sun:

Characters 0-7 == [0xf0,0xf1,0xf2,0xf3,0xf4,0xf5,0xf6,0xf7]
Shorts 0-3 == [0xf0f1,0xf2f3,0xf4f5,0xf6f7]
Ints 0-1 == [0xf0f1f2f3,0xf4f5f6f7]
Long 0 == [0xf0f1f2f3]

Byte Ordering on x86-64



Little Endian



Output on x86-64:

```
Characters 0-7 == [0xf0,0xf1,0xf2,0xf3,0xf4,0xf5,0xf6,0xf7]
Shorts    0-3 == [0xf1f0,0xf3f2,0xf5f4,0xf7f6]
Ints      0-1 == [0xf3f2f1f0,0xf7f6f5f4]
Long      0    == [0xf7f6f5f4f3f2f1f0]
```

Unions For Embedded Programming



```
typedef union
{
    unsigned char byte;
    struct {
        unsigned char reserved:4;
        unsigned char b3:1;
        unsigned char b2:1;
        unsigned char b1:1;
        unsigned char b0:1;
    } bits;
} hw_register;

hw_register reg;
reg.byte = 0x3F;           // 001111112
reg.bits.b2 = 0;           // 001110112
reg.bits.b3 = 0;           // 001100112
unsigned short a = reg.byte;
printf("0x%X\n", a);      // output: 0x33
```

(Note: the placement of these fields and other parts of this example are implementation-dependent)

Summary



- Arrays in C
 - Aligned to satisfy every element's alignment requirement
- Structures
 - Allocate bytes in order declared
 - Pad in middle and at end to satisfy alignment
- Unions
 - Provide different views of the same memory location

Q&A

