

# System Programming

Ch00. Introduction to HW/SW

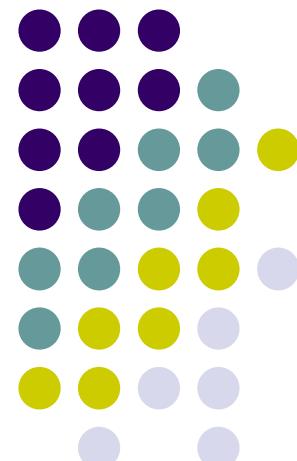
Ch01. A Tour of Computer Systems

2019. Fall

Instructor: Joonho Kwon

[jhwon@pusan.ac.kr](mailto:jhwon@pusan.ac.kr)

Data Science Lab @ PNU



**data**lab

data science laboratory

# The Hardware/Software Interface



- What do we mean by hardware? software?
  - What is an interface?
  - Why do we need a hardware/software interface?
  - Why do we need to understand both sides of this interface?

1000001101111100001001000001110000000000  
0111010000011000  
10001011010001000010010000010100  
100010110100011000100100100010100



1111011101111000010010000011100  
10001001010001000010010000011000

**HW/SW Interface**



# C/Java, assembly, and machine code



```
if (x != 0) y = (y+z)/x;
```



High Level Language  
(e.g. C, Java)

```
cmpl    $0, -4(%ebp)
je      .L2
movl    -12(%ebp), %eax
movl    -8(%ebp), %edx
leal    (%edx, %eax), %eax
movl    %eax, %edx
sarl    $31, %edx
idivl   -4(%ebp)
movl    %eax, -8(%ebp)
```

.L2:

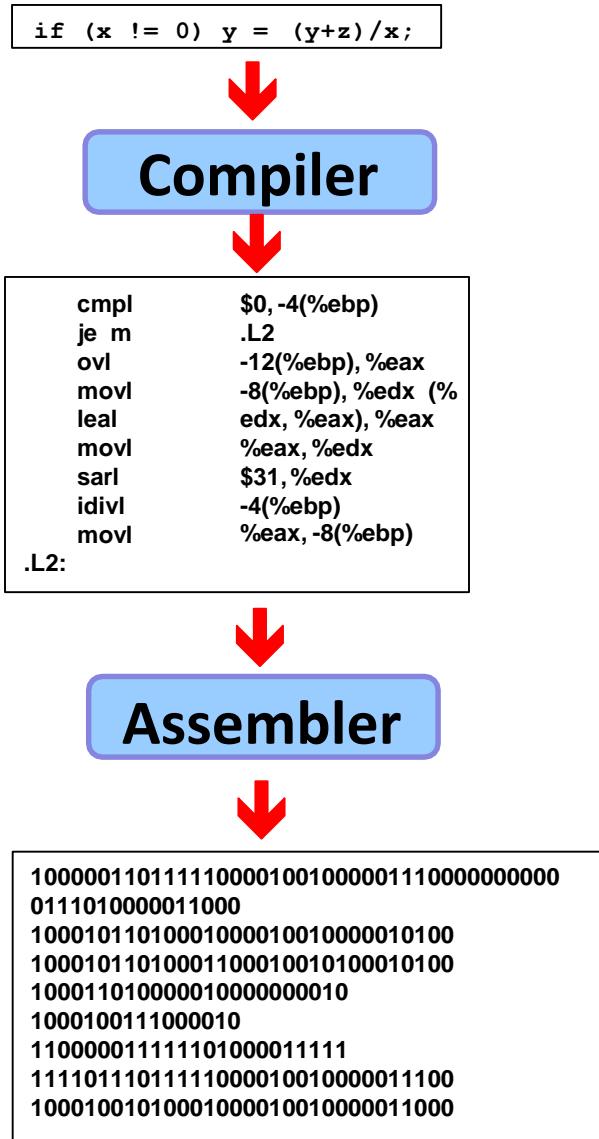


Assembly Language

```
100000110111100001001000001110000000000
0111010000011000
10001011010001000010010000010100
10001011010001100010010100010100
1000110100000100000000010
1000100111000010
11000001111101000011111
1111011101111000010010000011100
10001001010001000010010000011000
```

Machine Code

# C/Java, assembly, and machine code



High Level Language  
(e.g. C, Java)

Assembly Language

Machine Code

# C/Java, assembly, and machine code



```
if (x != 0) y = (y+z)/x;
```



```
cmpl $0, -4(%ebp)
je .L2
movl -12(%ebp), %eax
movl -8(%ebp), %edx
leal (%edx, %eax), %eax
movl %eax, %edx
sarl $31, %edx
idivl -4(%ebp)
movl %eax, -8(%ebp)
```

.L2:



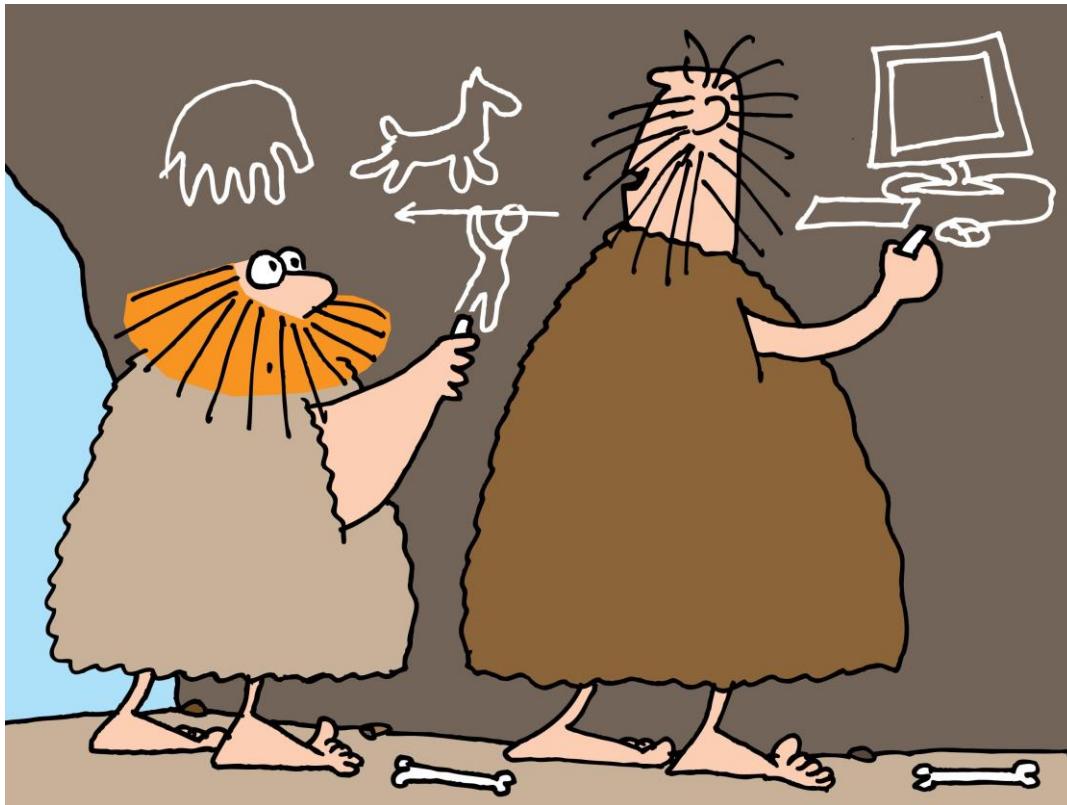
```
100000110111100001001000001110000000000
0111010000011000
10001011010001000010010000010100
10001011010001100010010100010100
1000110100000100000000010
1000100111000010
11000001111101000011111
11110111011111000010010000011100
10001001010001000010010000011000
```

- The three program fragments are equivalent
- You'd rather write C! - a more human-friendly language
- The hardware likes bit strings! - everything is voltages
  - Everything is voltages
  - The machine instructions are actually much shorter than the number of bits we would need to represent the characters in the assembly language

# HW/SW Interface: Historical Perspective



- Hardware started out quite primitive



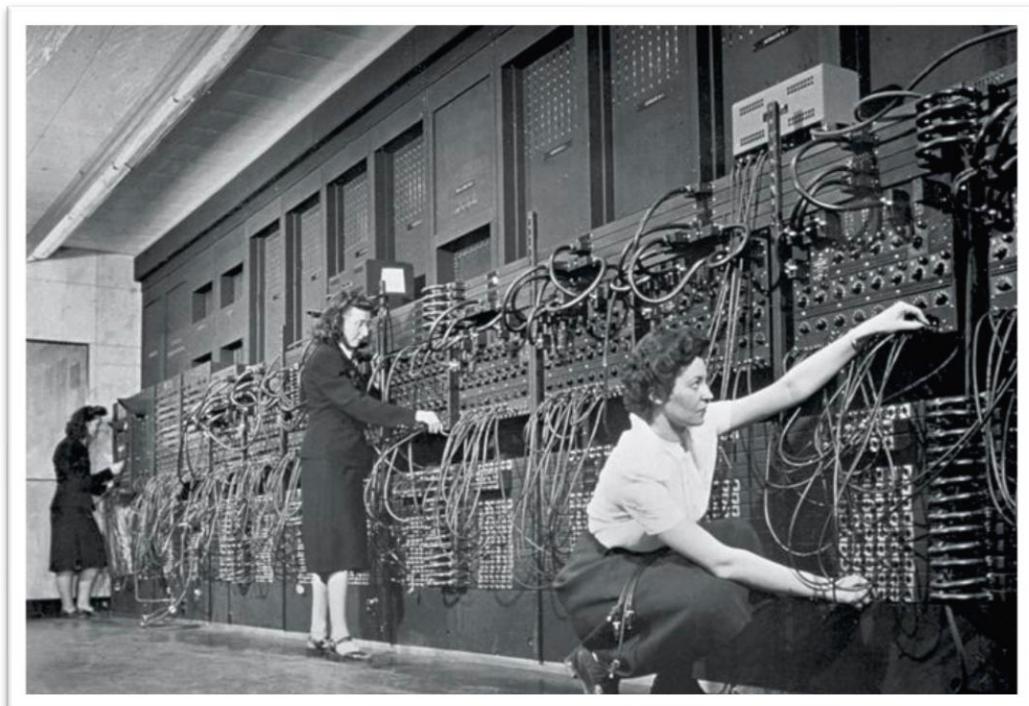
# HW/SW Interface: Historical Perspective



- Hardware started out quite primitive



<http://hightechhistory.com/2011/10/24/the-late-dennis-macalystair-ritchie-innovator-of-the-%E2%80%9Cc%E2%80%9D-programming-language-and-the-unix-operating-system-an-appreciation/>

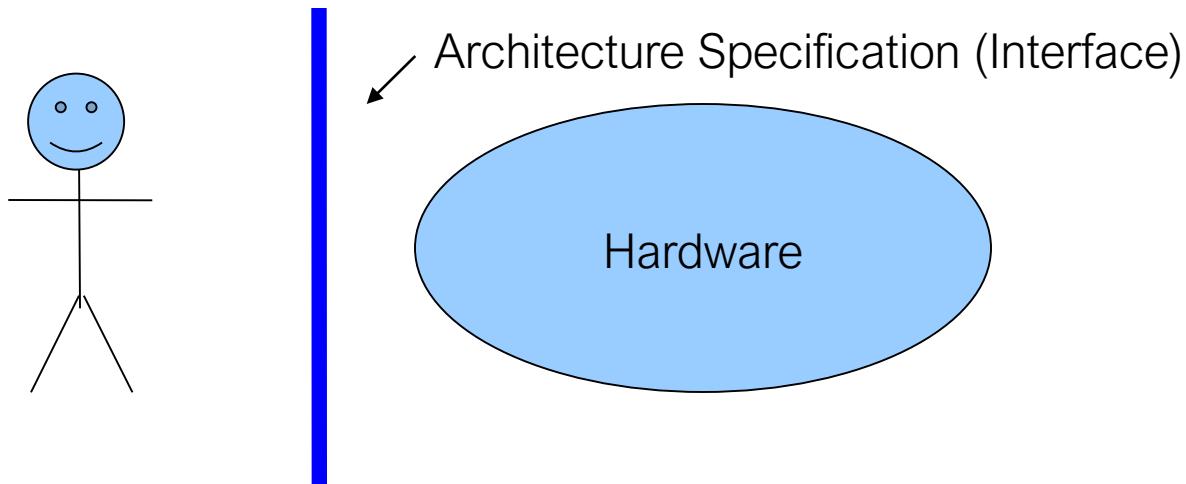


Jean Jennings (left), Marlyn Wescoff (center), and Ruth Lichterman program ENIAC at the University of Pennsylvania, circa 1946. Photo: Corbis

# HW/SW Interface: The Historical Perspective



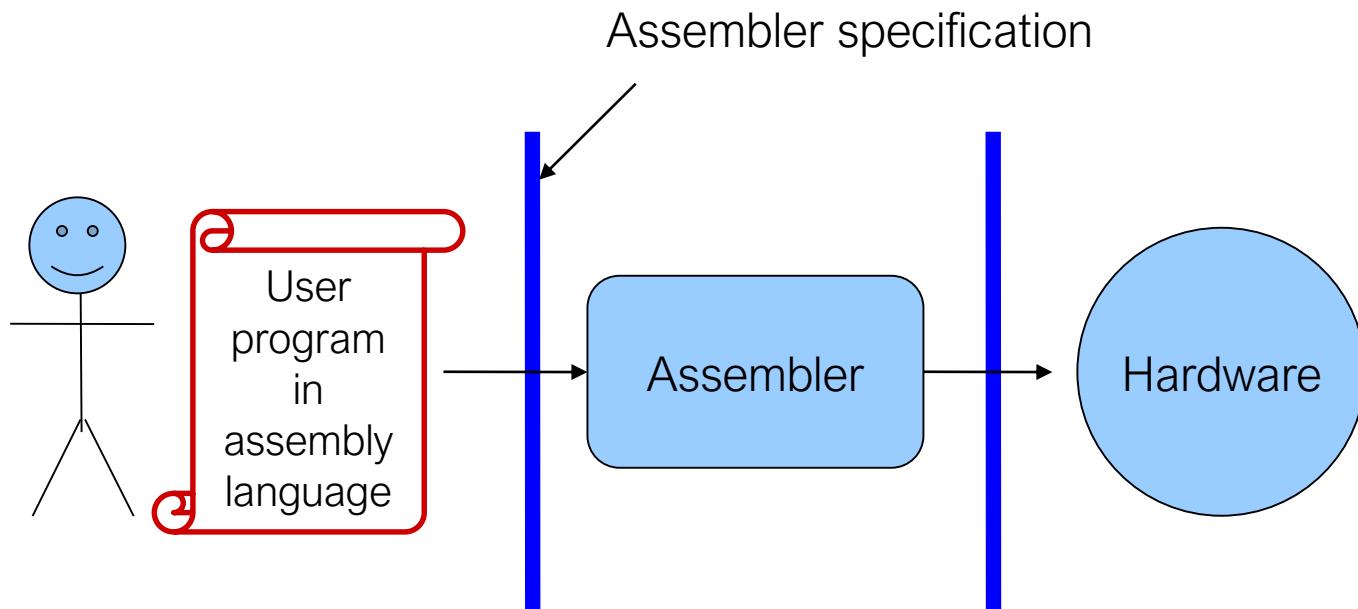
- Hardware started out quite primitive
  - Programmed with very basic instructions.
  - E.g. a single instruction for adding two integers
- Software was also very basic
  - Closely reflected the actual hardware it was running on
  - Specify each step manually



# HW/SW Interface: Assemblers



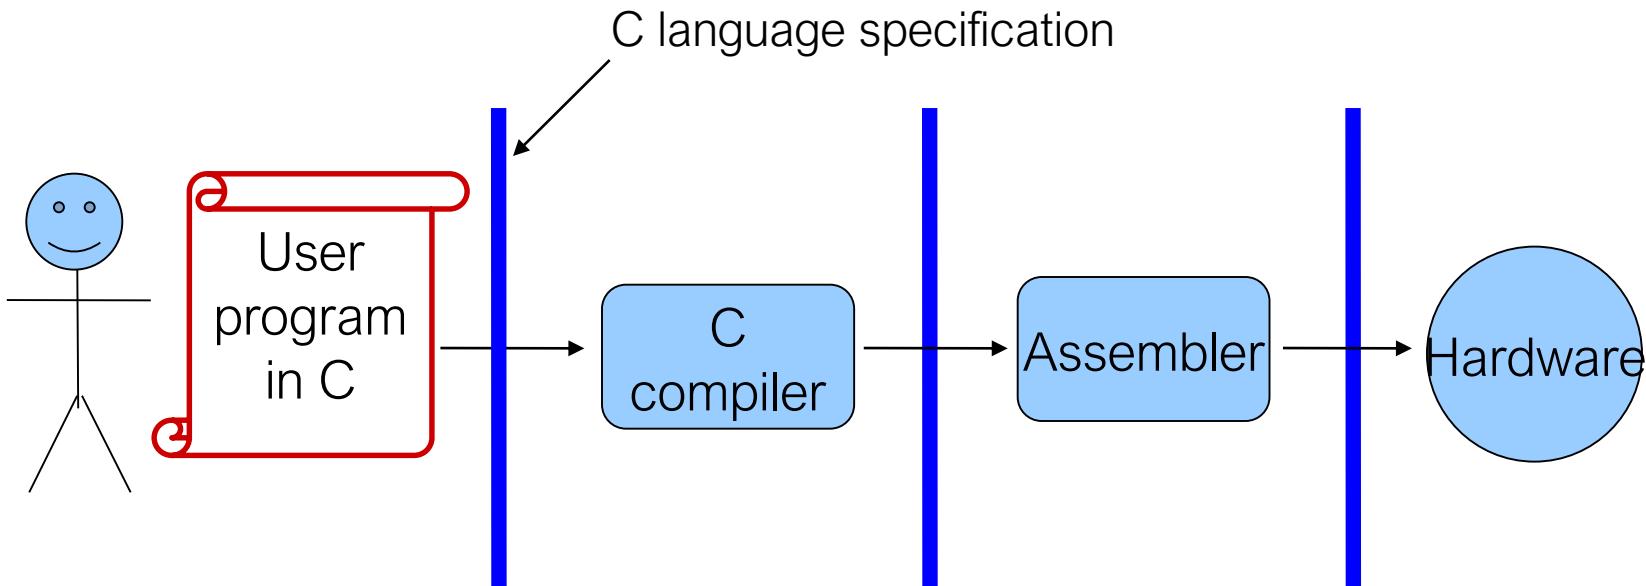
- Life was made a lot better by assemblers
  - 1 assembly instruction = 1 machine instruction, but...
  - different syntax: assembly instructions are character strings, not bit strings, a lot easier to read/write by humans
  - can use symbolic names



# HW/SW Interface: Higher-Level Languages



- Higher level of abstraction:
  - 1 line of a high-level language is compiled into many (sometimes very many) lines of assembly language



# HW/SW Interface: Code / Compile / Run

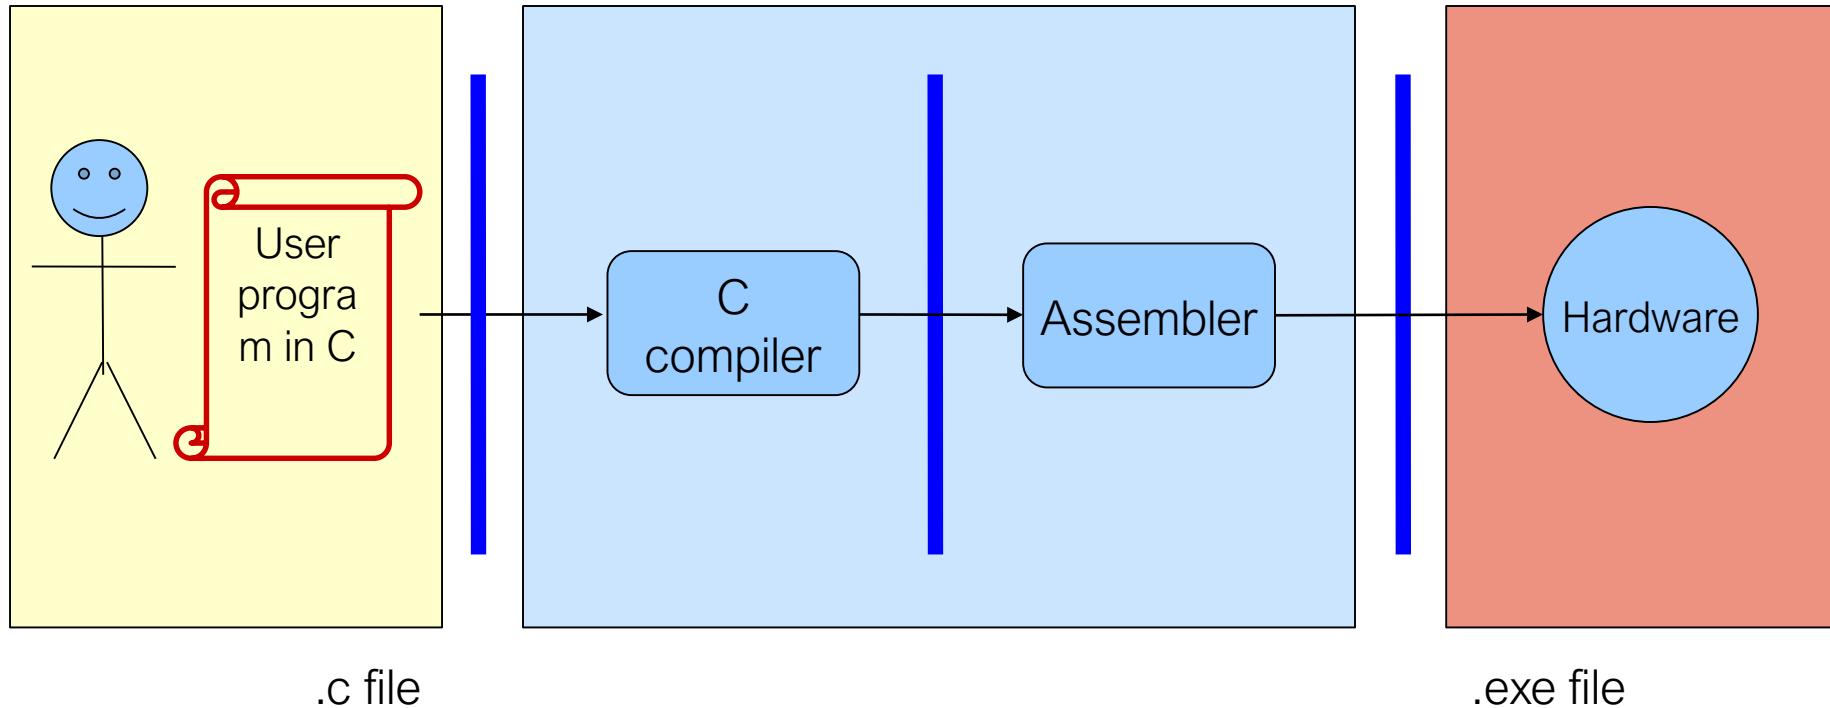
## Times



Code Time

Compile Time

Run Time



Note: The compiler and assembler are just programs, developed using this same process.

# The Big Theme: Abstractions and Interfaces



- Computing is about abstractions
  - (but we can't forget reality)
- What are the abstractions that we use?
- What do YOU need to know about them?
  - When do they break down and you have to peek under the hood?
  - What bugs can they cause and how do you find them?
- How does the hardware (0s and 1s, processor executing instructions) relate to the software (C/Java programs)?
  - Become a better programmer and begin to understand the important concepts that have evolved in building ever more complex computer systems

# Roadmap



C:

```
car *c = malloc(sizeof(car));  
c->miles = 100;  
c->gals = 17;  
float mpg = get_mpg(c);  
free(c);
```

Java:

```
Car c = new Car();  
c.setMiles(100);  
c.setGals(17);  
float mpg =  
    c.getMPG();
```

Assembly  
language:

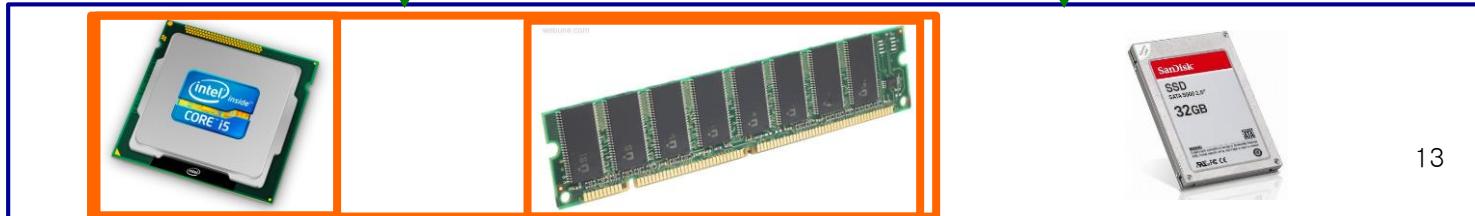
```
get mpg:  
pushq %rbp  
movq %rsp, %rbp  
...  
popq %rbp  
ret
```

Machine  
code:

```
0111010000011000  
10001101000010000000010  
1000100111000010  
11000001111101000011111
```

Computer  
system:

OS:



# Little Theme 1: Representation



- All digital systems represent everything as 0s and 1s
  - The 0 and 1 are really two different voltage ranges in the wires
  - Or magnetic positions on a disc, or hole depths on a dvd, or...
- “Everything” includes:
  - Numbers – integers and floating point
  - Characters – the building blocks of strings
  - Instructions – the directives to the CPU that make up a program
  - Pointers – addresses of data objects stored away in memory
- These encodings are stored throughout a computer system
  - In registers, caches, memories, disks, etc.
- They all need addresses
  - A way to find them
  - Find a new place to put a new item
  - Reclaim the place in memory when data no longer needed

# Little Theme 2: Translation



- There is a **big gap** between how we think about programs and data and the 0s and 1s of computers
- Need **languages** to describe what we mean
- Languages need to be **translated** one step at a time
- We know C/Java as a programming language
  - Have to work our way down to the 0s and 1s of computers
  - Try not to lose anything in translation!
  - We'll encounter Java byte-codes, C language, assembly language, and machine code (for the X86 family of CPU architectures)
    - Not in that order, but will all connect by the last lecture!!!

# Little Theme 3: Control Flow



- How do computers orchestrate the many things they are doing?
- In one program:
  - How do we implement if/else, loops, switches?
  - What do we have to keep track of when we call a procedure, and then another, and then another, and so on?
  - How do we know what to do upon “return”?
- Across programs and operating systems:
  - Multiple user programs
  - Operating system has to orchestrate them all
    - Each gets a share of computing cycles
    - They may need to share system resources (memory, I/O, disks)
  - Yielding and taking control of the processor
    - Voluntary or “by force”?

# Writing Assembly Code??? In 2019??? (1/2)



- Chances are, you'll never write a program in assembly code
  - Compilers are much better and more patient than you are

# Writing Assembly Code??? In 2019??? (2/2)



- But: understanding assembly is the key to the machine-level execution model
  - Behavior of programs in presence of bugs
    - High-level language model breaks down
  - Tuning program performance
    - Understand optimizations done/not done by the compiler
    - Understanding sources of program inefficiency
  - Implementing system software
    - Operating systems must manage process state
  - Fighting malicious software
  - Using special units (timers, I/O co-processors, etc.) inside processor!

# Course Perspective



- This course will make you a better programmer
  - Purpose is to show how software really works
  - Understanding the underlying system makes you more effective
  - Better debugging
  - Better basis for evaluating performance
  - How multiple activities work in concert (e.g., OS and user programs)
- Not just a course for hardware enthusiasts!
  - What every CSE major needs to know (plus many more details)
  - Job interviewers love to ask questions from this course!
- Like other courses, “stuff everybody learns and uses and forgets not knowing”



[HTTP://XKCD.COM/676/](http://xkcd.com/676/)

AN x64 PROCESSOR IS SCREAMING ALONG AT BILLIONS OF CYCLES PER SECOND TO RUN THE XNU KERNEL, WHICH IS FRANTICALLY WORKING THROUGH ALL THE POSIX-SPECIFIED ABSTRACTION TO CREATE THE DARWIN SYSTEM UNDERLYING OS X, WHICH IN TURN IS STRAINING ITSELF TO RUN FIREFOX AND ITS GECKO RENDERER, WHICH CREATES A FLASH OBJECT WHICH RENDERS DOZENS OF VIDEO FRAMES EVERY SECOND

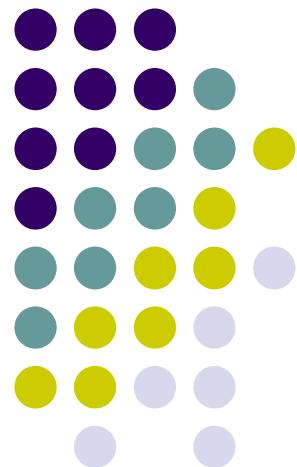
BECAUSE I WANTED TO SEE A CAT JUMP INTO A BOX AND FALL OVER.



I AM A GOD.

# Ch01. A Tour of Computer Systems

---



# The hello Program



```
#include <stdio.h>

int main() {
    printf("hello, world\n");
}
```

# Information is Bits



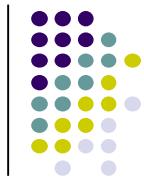
- Source File
  - The hello program begins life as a source program
  - Most programs consist of multiple source files
- Representation
  - A sequence of bits with a value of 0 or 1
  - Organized into bytes of 8 bits each
  - Each byte represents a character in the program

# Character Interpretation



- Bytes and Characters
  - Characters are encoded in bytes
  - Bytes are 8 bits
  - A bit is a 1 or a 0
- Bytes and Numbers
  - A byte can represent a number in base-2:  
 $00000110_2 = 6_{10}$ ,  $00100000_2 = 64_{10}$ ,  $10000111_2 = 135_{10}$ ,  $11111111_2 = 255_{10}$
  - Numbers can be used to represent characters

# Representing Characters



- ASCII Standard (ASCII = American Standard Code for Information Interchange)
  - Represents each character with a unique byte--sized integer value

The ASCII text representation of **hello** source file

#	i	n	c	l	u	d	e	<sp>	<	s	t	d	i	o	.
35	105	110	99	108	117	100	101	32	60	115	116	100	105	111	46
h	>	\n	\n	i	n	t	<sp>	m	a	i	n	(	)	\n	{
104	62	10	10	105	110	116	32	109	97	105	110	40	41	10	123
\n	<sp>	<sp>	<sp>	<sp>	p	r	i	n	t	f	(	"	h	e	l
10	32	32	32	32	112	114	105	110	116	102	40	34	104	101	108
l	o	,	<sp>	w	o	r	l	d	\	n	"	)	;	\n	}
108	111	44	32	119	111	114	108	100	92	110	34	41	59	10	125

# Unicode



- ASCII has limitations (only 128 characters).
- Unicode is an extension of ASCII.
- Unicode characters can be stored in 32 bits, but there are representations of them that use fewer bits.
- Java uses Unicode, though Linux does not.

# Program Translation



- Program Source Files
  - Beginning of life for a C program
  - Represented as ASCII character text
  - “Easy” for humans to understand
  - Not understood by machines
- Executable Object File
  - Low--level primitive machine operations
  - Understood by the machine
- Program Translation
  - Translates source file into object (machine code) file!
  - Also known as **compilation**

# Compilation System



```
unix> gcc -o hello hello.c
```

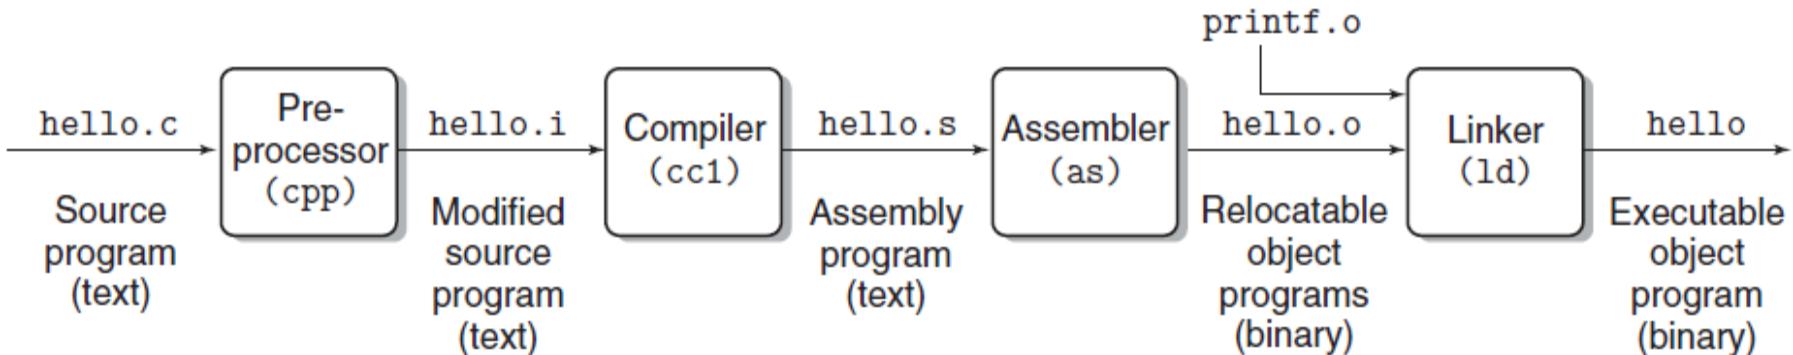


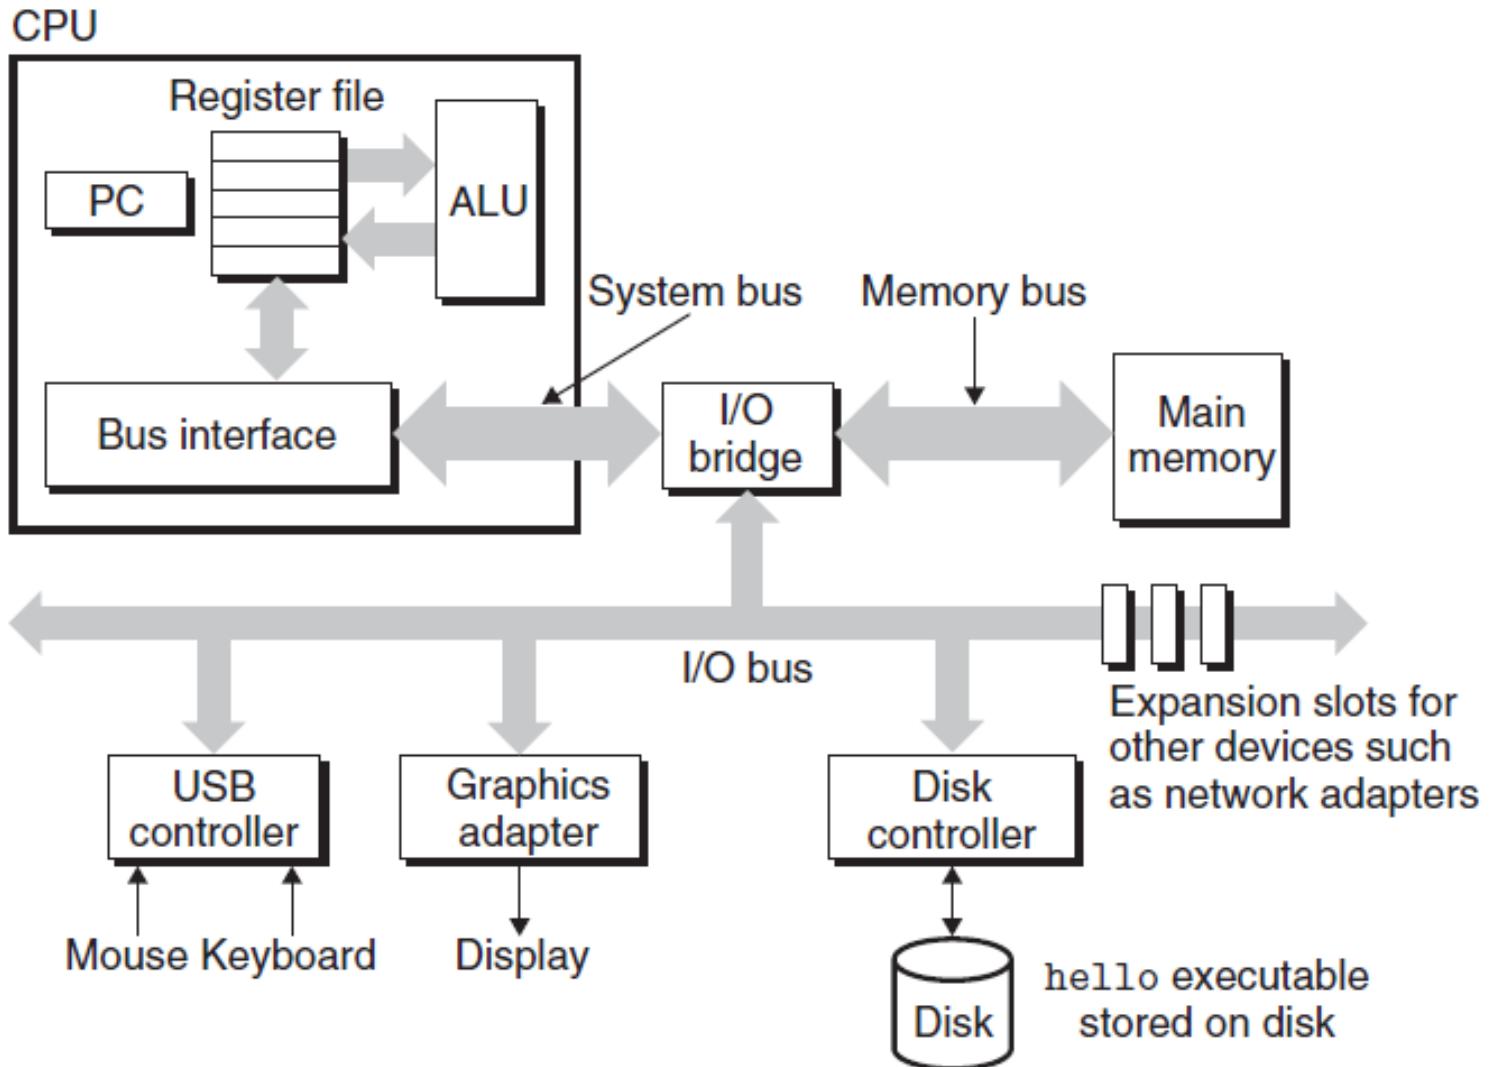
Figure 1.3 The compilation system.

# Processors

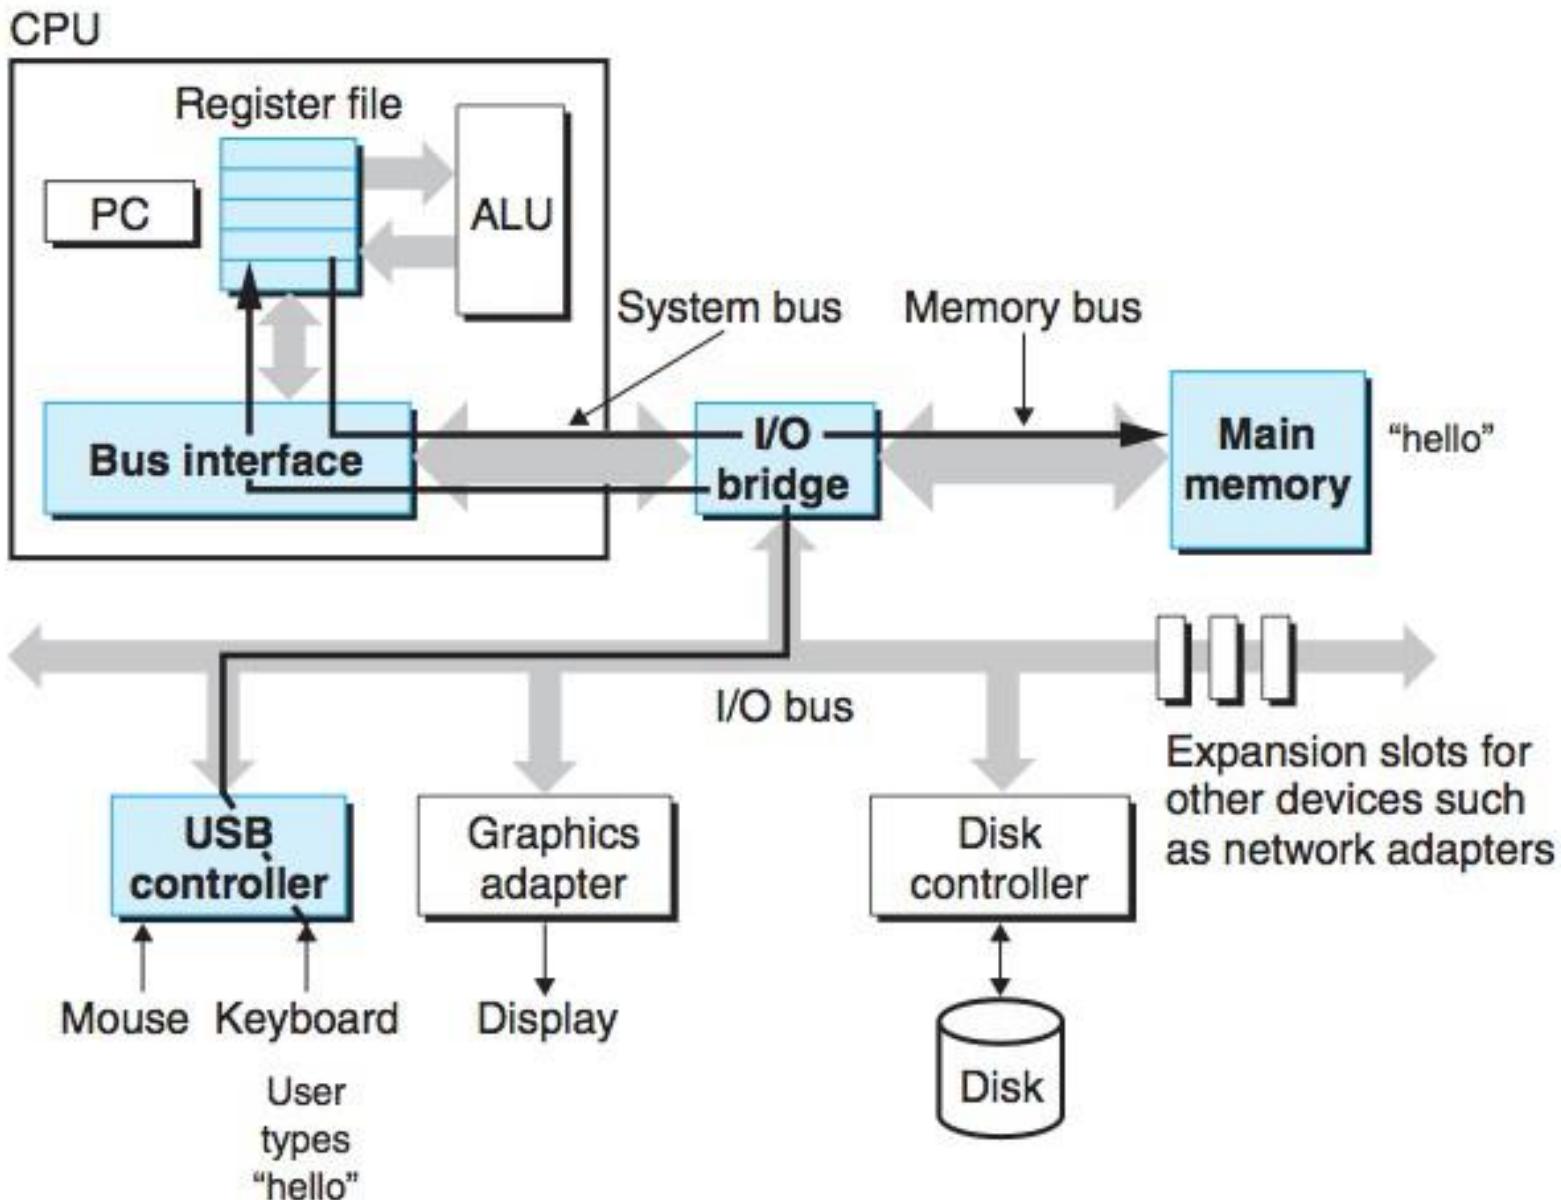


- Machine Code Instructions
  - Programs at the only level a machine can understand
  - Stored in memory
- Processors
  - Read instructions from memory
  - Interpret those instructions (do what the instructions say to do)
  - Implemented in hardware

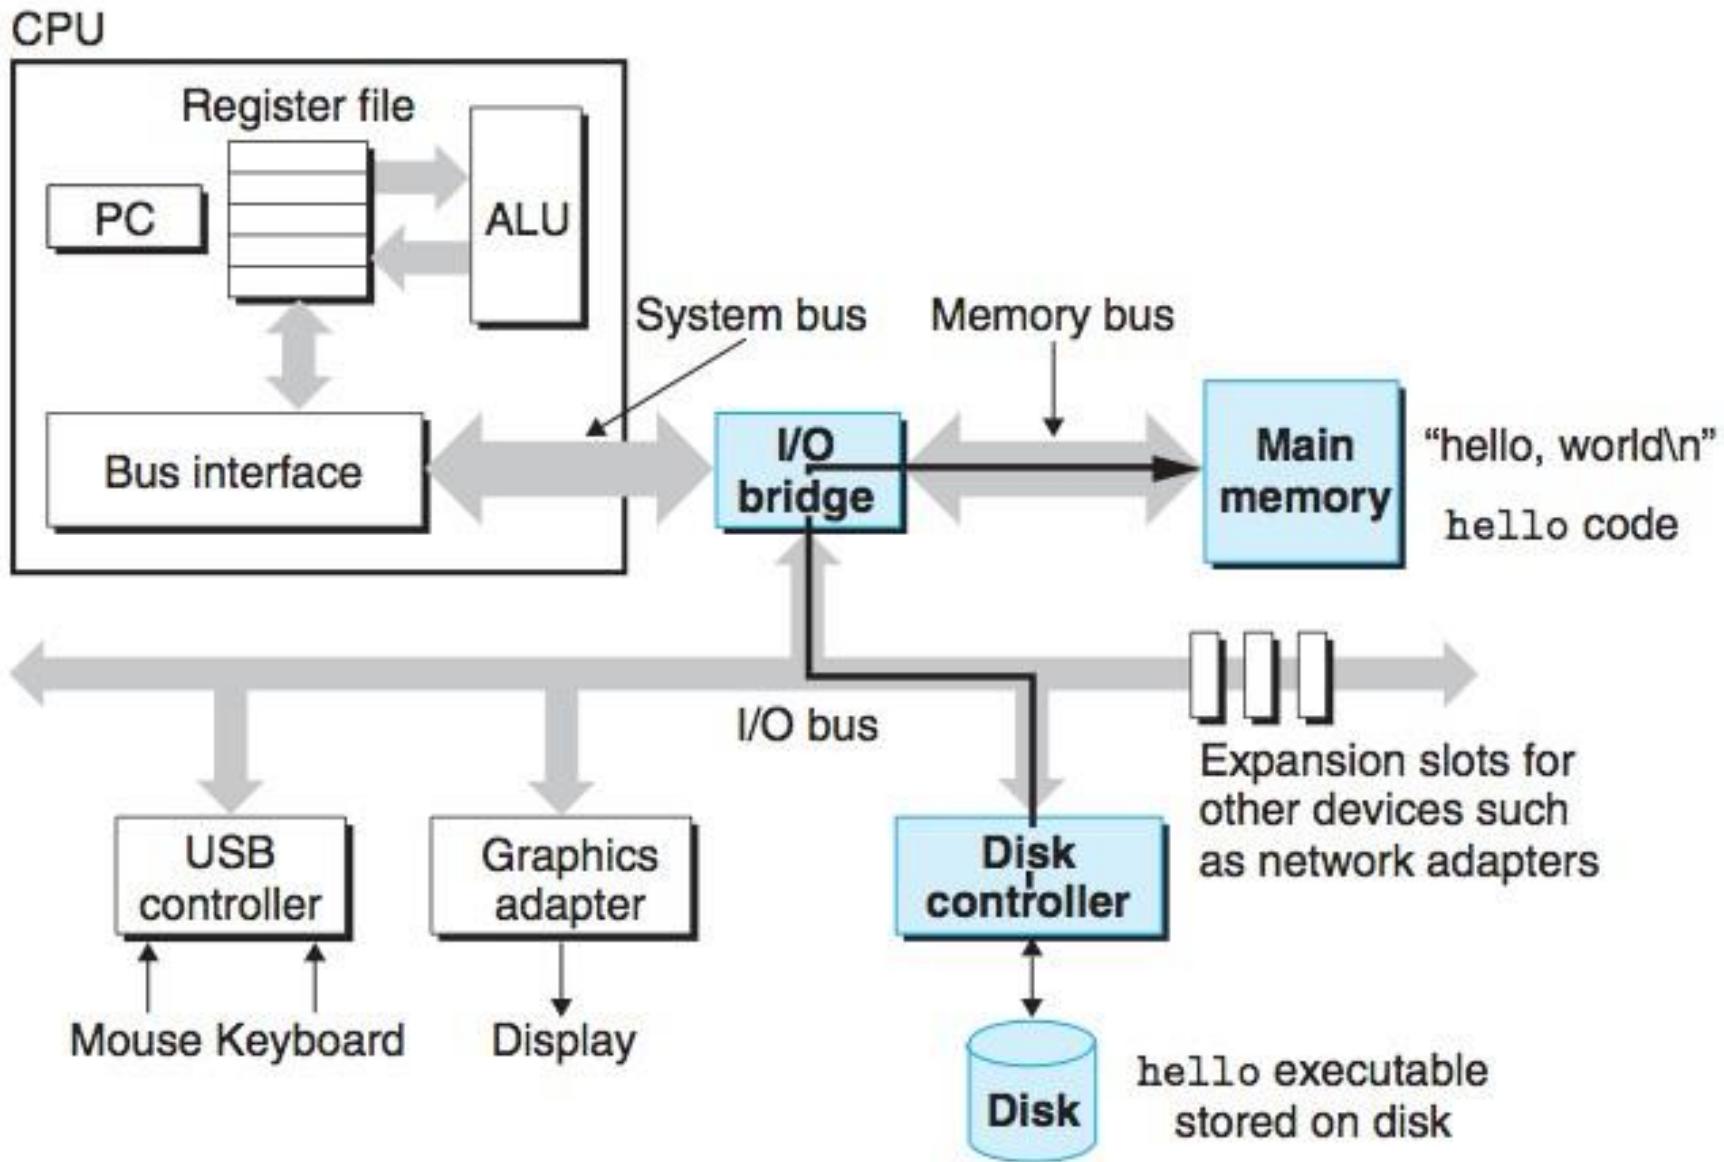
# Hardware Organization



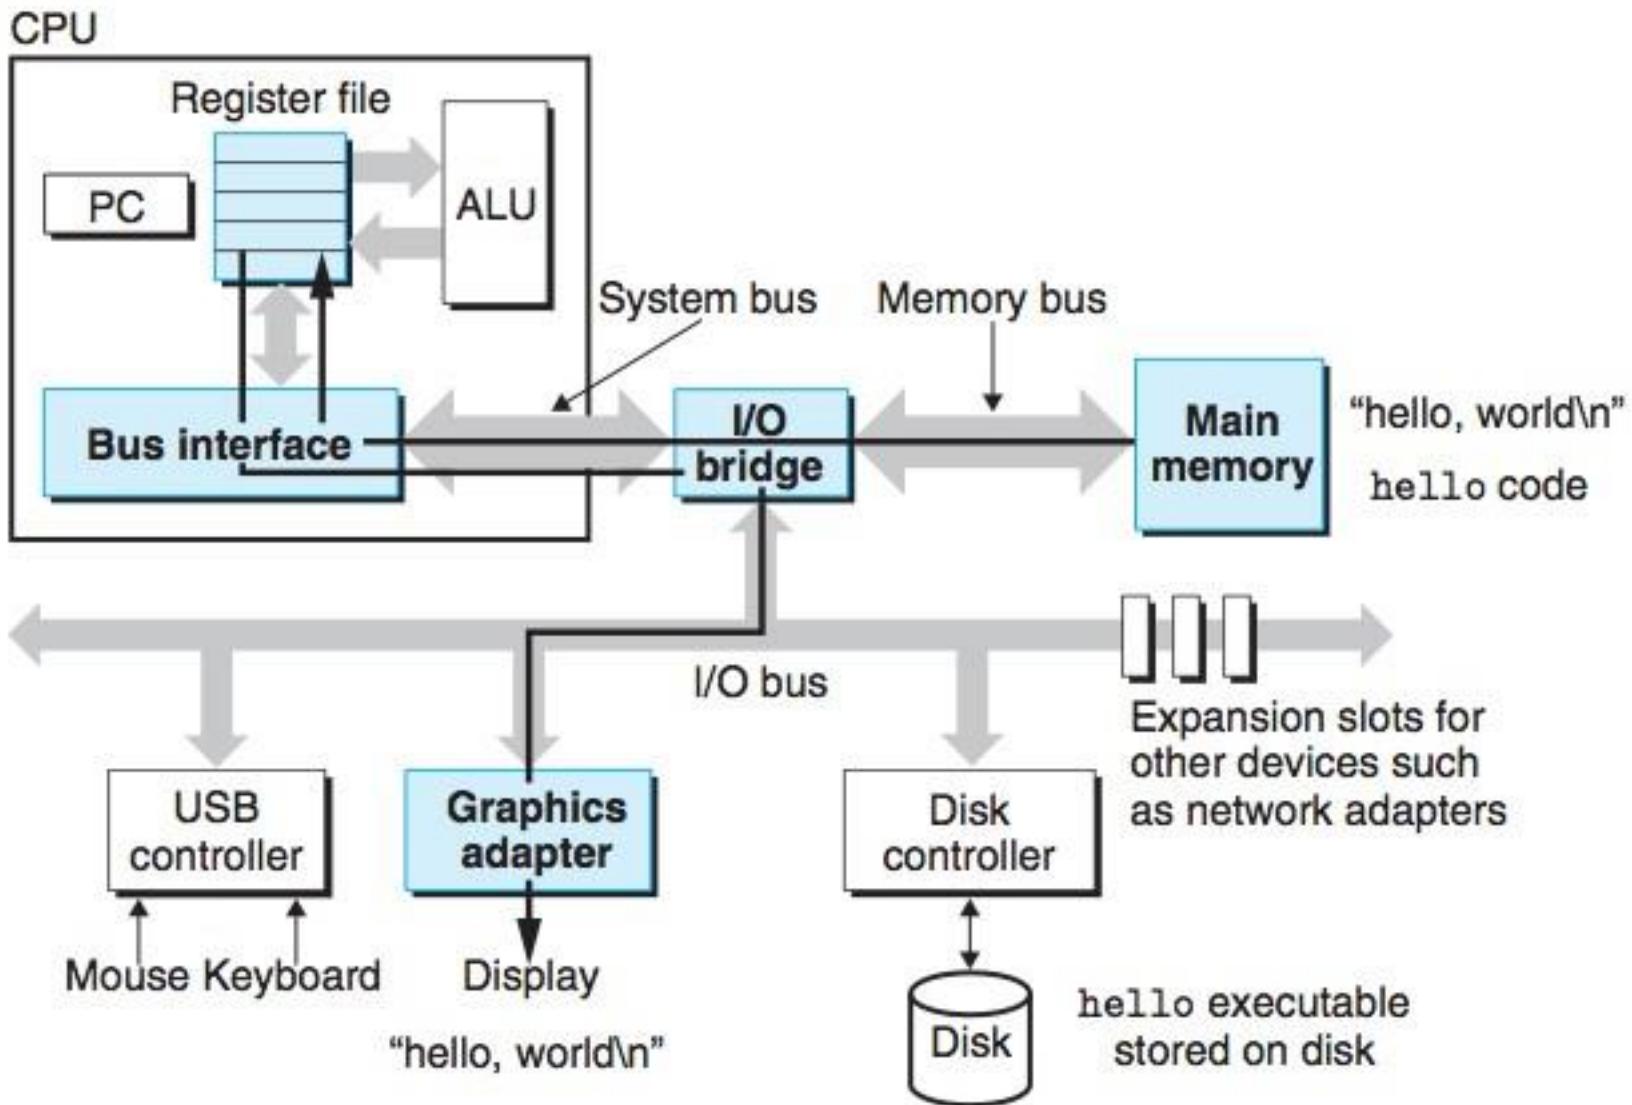
# Reading hello command from keyboard

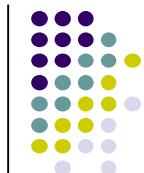


# Loading the executable from disk into main memory



# Writing the output string from memory to the display





# Memory this, Memory that

- Memory is Important
  - Stores program code
  - Stores program data
  - Accesses required for execution
- Memory is Slow
  - Yep, it takes a long time to access memory
  - Need a mechanism to reduce memory latency

# Cache

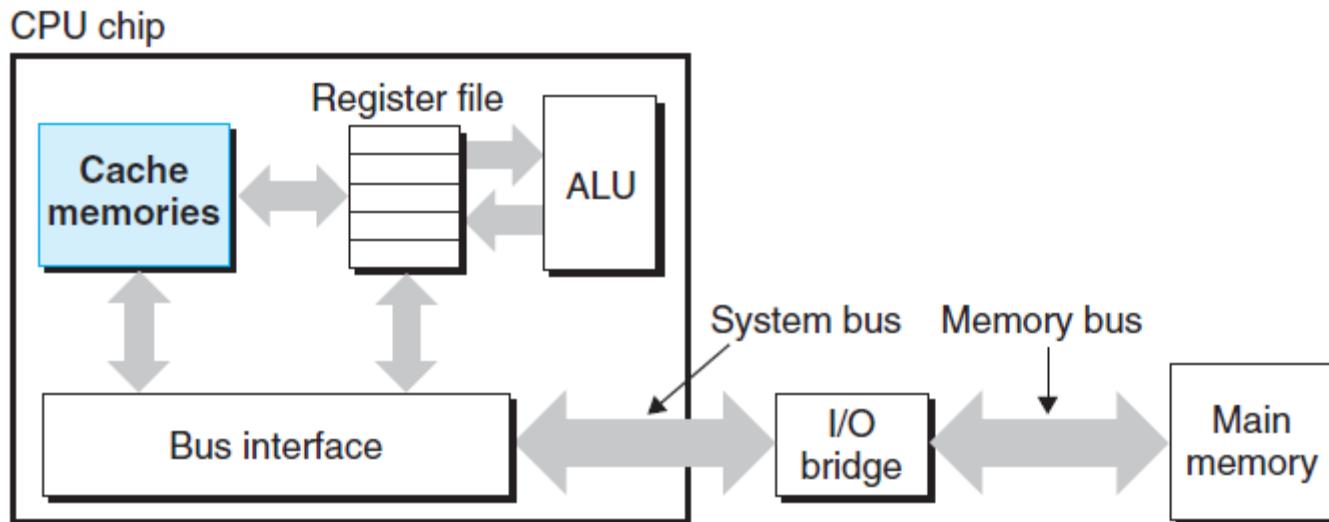


- Smaller Memories

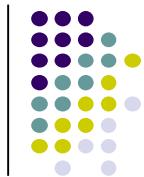
- Resides on CPU chip
- Larger than register file
- Smaller than RAM

- Locality

- Access to program code and data tends to exhibit a high degree of locality, on both space and time
- Caches exploit this!



# Memory Hierarchy



Smaller,  
faster,  
and  
costlier  
(per byte)  
storage  
devices

Larger,  
slower,  
and  
cheaper  
(per byte)  
storage  
devices

L6:

Remote secondary storage  
(e.g., Web servers)

L5:

Local secondary storage  
(local disks)

L4:

Main memory  
(DRAM)

L3:  
L2:  
L1:

Regs

L1 cache  
(SRAM)

L2 cache  
(SRAM)

L3 cache  
(SRAM)

CPU registers hold words  
retrieved from the L1 cache.

L1 cache holds cache lines  
retrieved from the L2 cache.

L2 cache holds cache lines  
retrieved from L3 cache

L3 cache holds cache lines  
retrieved from main memory.

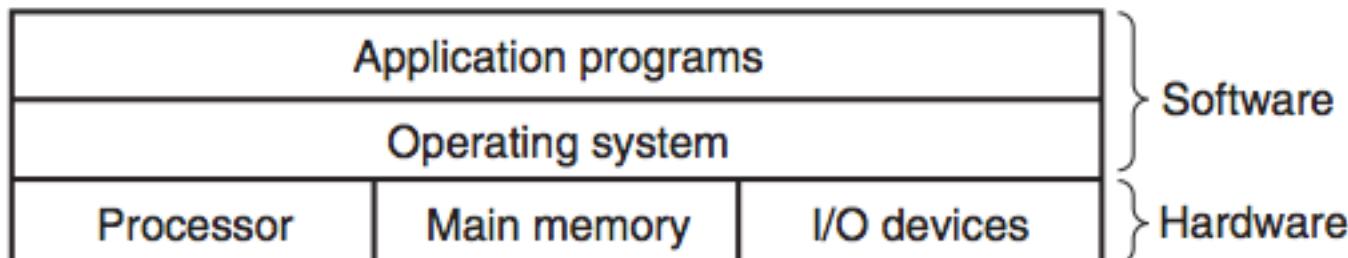
Main memory holds  
disk blocks retrieved  
from local disks.

Local disks hold files  
retrieved from disks  
on remote servers

# Operating System



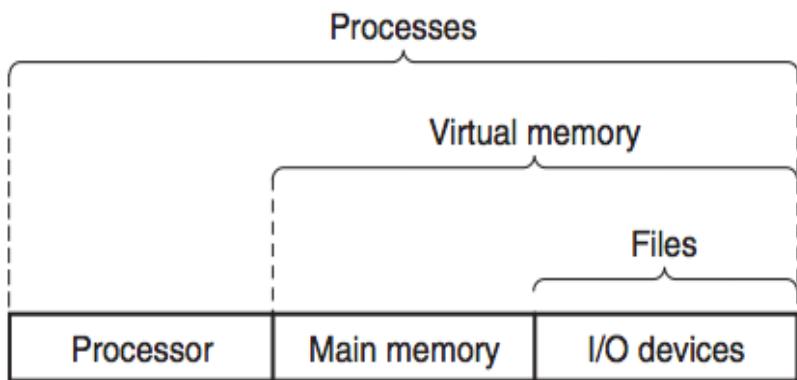
- Two Primary Purposes
  - to protect the hardware (and other programs and files) from misuse by runaway applications
  - provide applications with simple and uniform mechanisms for manipulating complicated and often wildly different low--level hardware devices



# OS Abstractions



- How does the OS do this?
  - Three fundamental abstractions
- 1: Files
  - Abstraction for I/O devices
- 2: Virtual Memory
  - Abstraction for main memory
  - Abstraction for I/O devices
- 3: Processes
  - Abstraction for the processor
  - Abstraction for main memory
  - Abstraction for I/O devices



# Processes

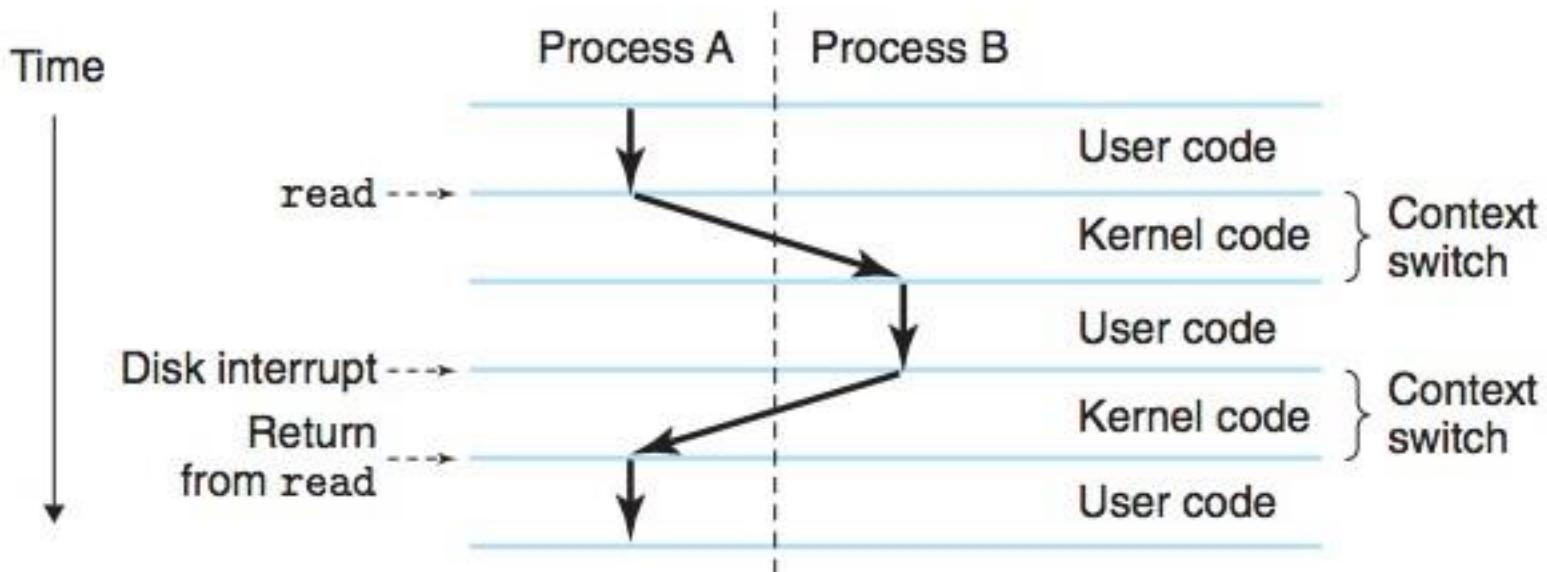


- What are they?
  - An abstraction for a running program
- How many?
  - Lots of them
  - Multiple processes can run concurrently
- What do they give us?
  - Illusion that each program has exclusive access to the processor and memory

# What does “concurrently” mean?



- The machine code instructions of one process are interleaved with the machine code instructions of another process.



# Aside: What about multiple “cores”?



- Each “core” in a multi--core CPU
  - is effectively a separate CPU
  - can context--switch independently
- If enough processes are ready to run, two or more cores can be running programs at the same time
- Can be thought of as multiple computers on the same chip, but managed by the same OS and sharing the same memory and I/O devices

# Processes and Threads



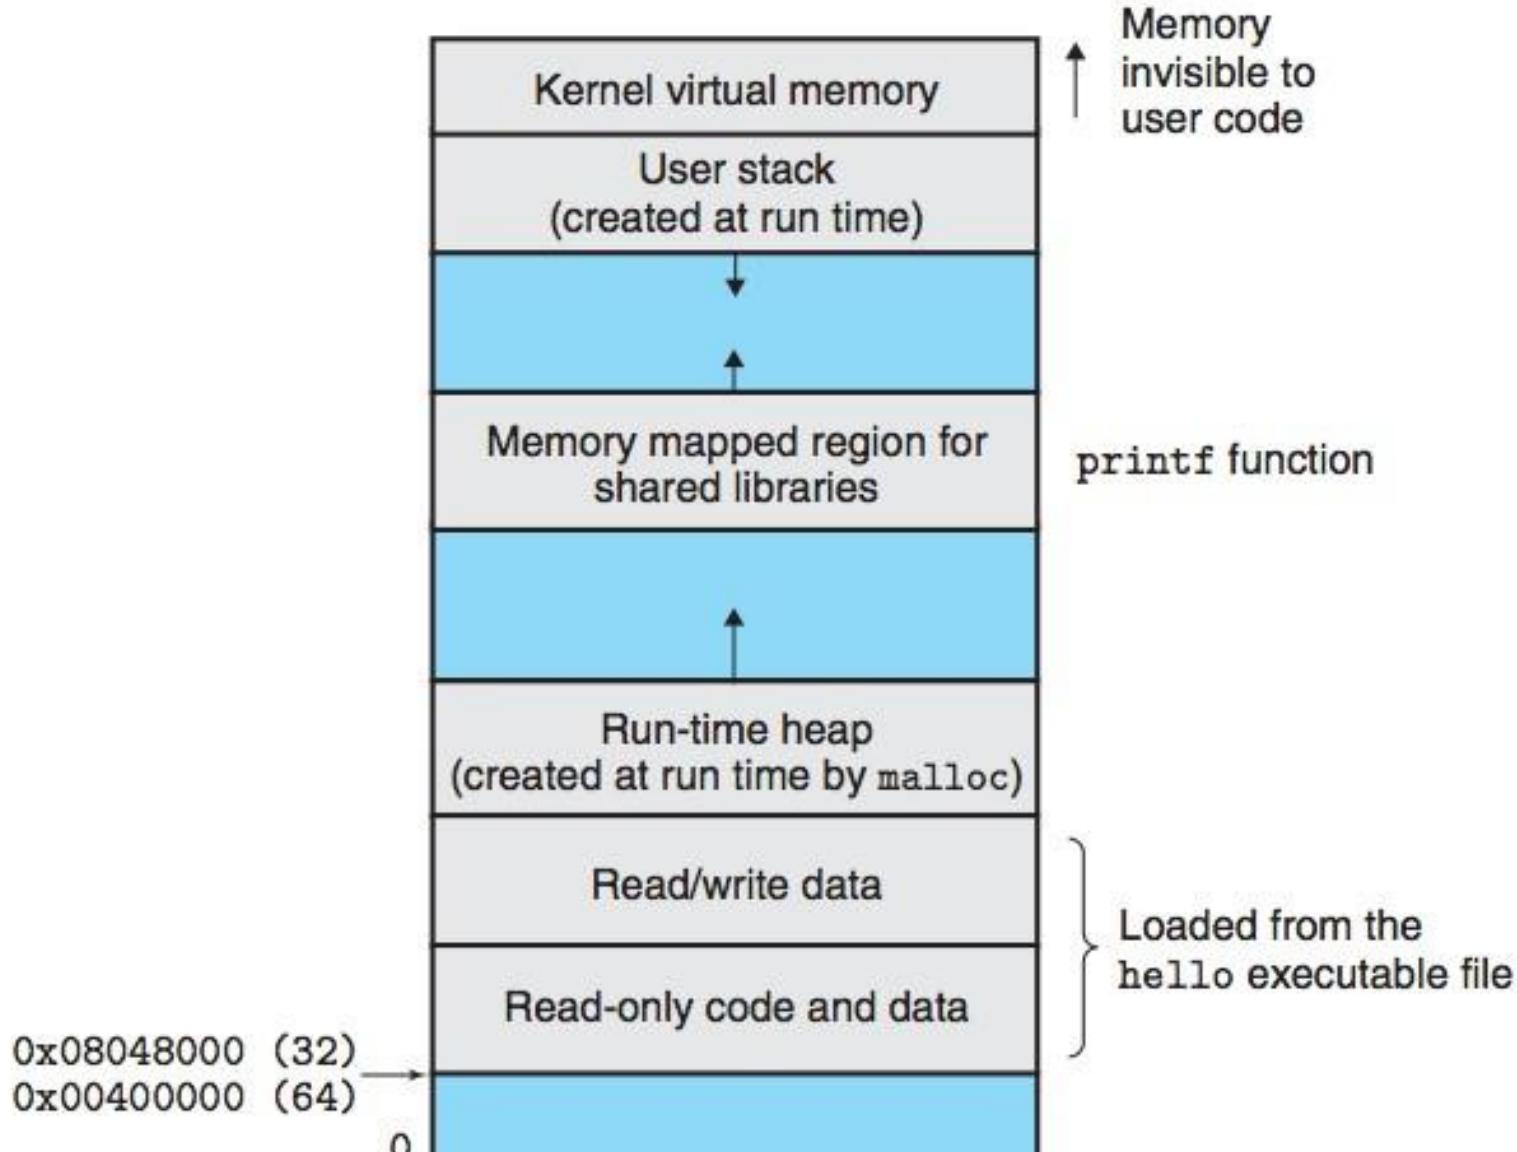
- Processes

- The illusion is great, but what if I want to share my memory with another process?
- You can't!

- Threads

- Associated with each process
- Can be lots of them
- Can share memory between them

# Virtual Memory



# Files



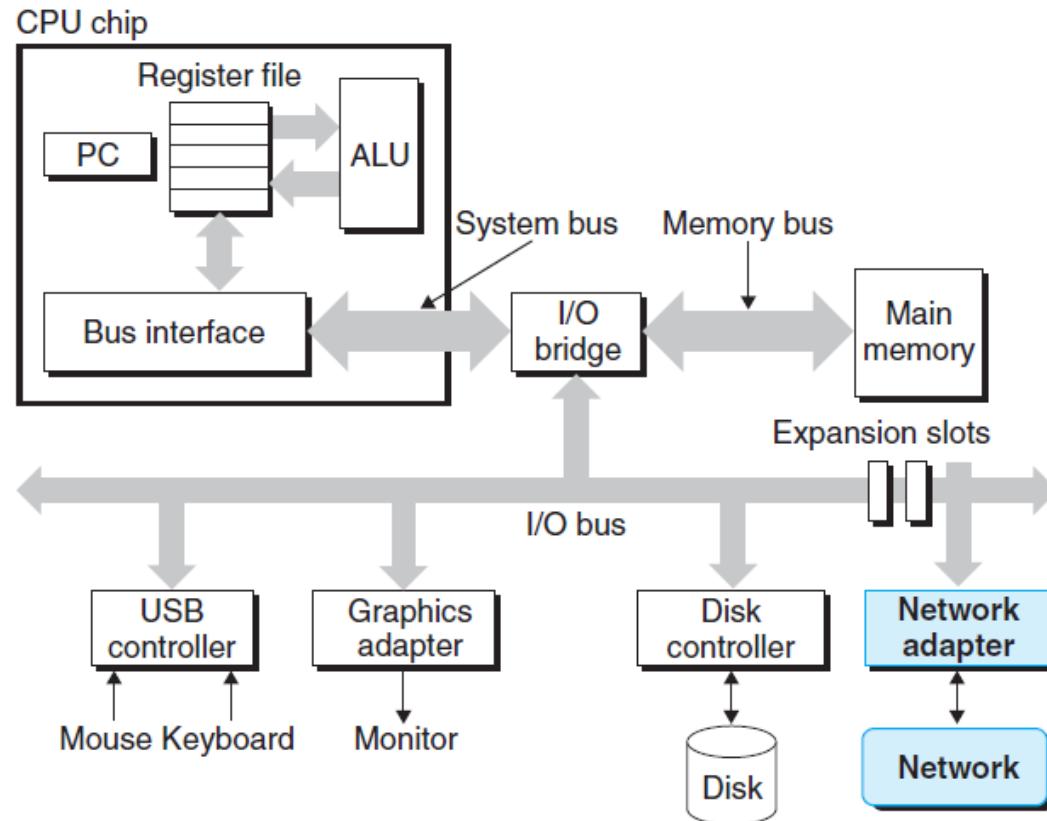
- Sequence of bytes... nothing more, nothing less
- **File System** resides on secondary storage (disks)



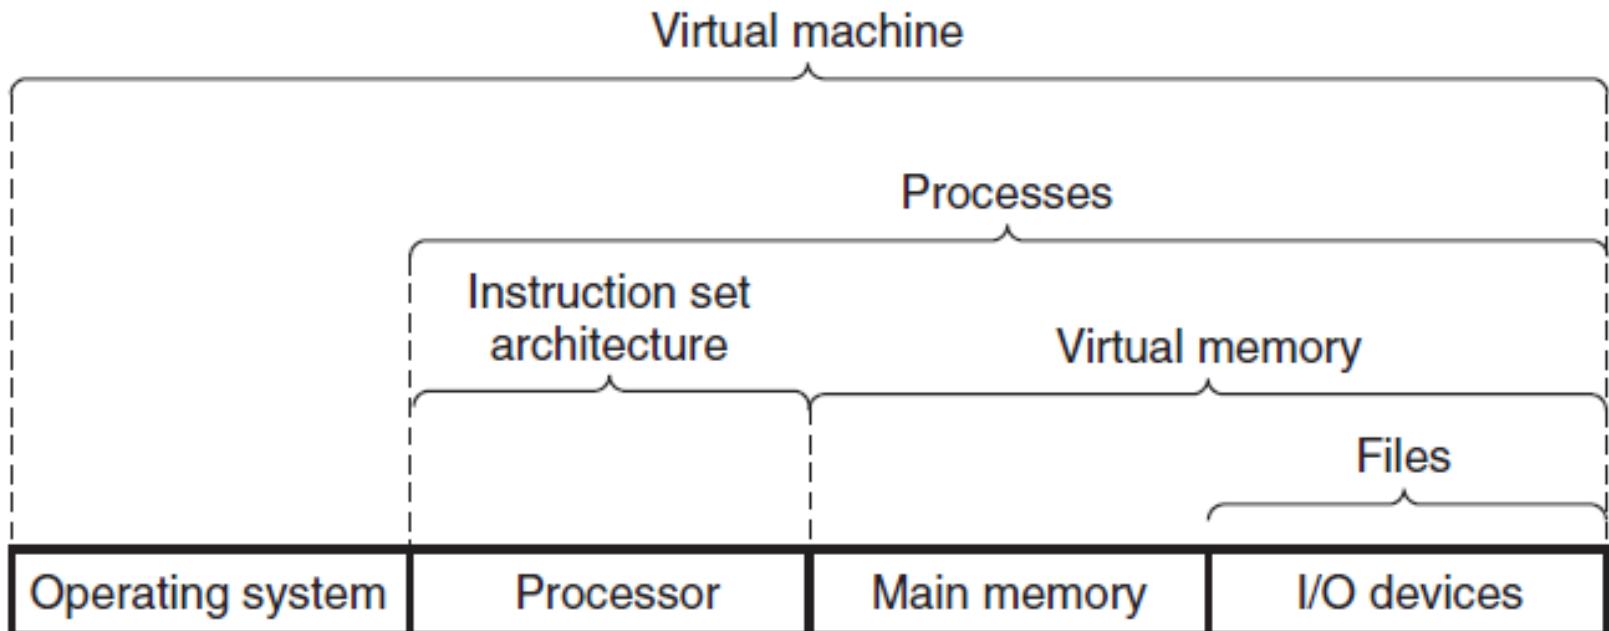
# Network Communication



- Processes like to talk to other processes



# Abstractions in Computer Systems



# Q&A

