

# System Programming

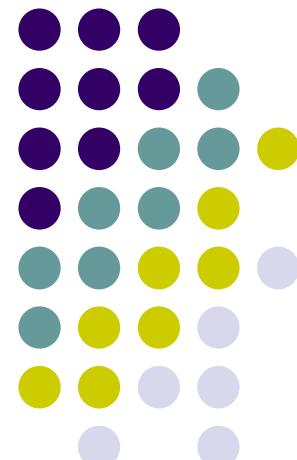
## 03. B. Memory (ch 2.1)

2019. Fall

Instructor: Joonho Kwon

[jhwon@pusan.ac.kr](mailto:jhwon@pusan.ac.kr)

Data Science Lab @ PNU



# Pointers



- <http://xkcd.com/138/>





# Roadmap

C:

```
car *c = malloc(sizeof(car));  
c->miles = 100;  
c->gals = 17;  
float mpg = get_mpg(c);  
free(c);
```

Java:

```
Car c = new Car();  
c.setMiles(100);  
c.setGals(17);  
mpg =  
    c.getMPG();
```

Assembly language:

```
get_mpg:  
    pushq  %rbp  
    movq   %rsp, %rbp  
    ...  
    popq  %rbp  
    ret
```

Machine code:

```
0111010000011000  
10001101000010000000010  
1000100111000010  
11000001111101000011111
```

Computer system:



OS:



Memory & data

Integers & floats

x86 assembly

Procedures & stacks

Executables

Arrays & structs

Memory & caches


Processes

Virtual memory


Memory allocation

Java vs. C


# Hardware: Logical View



# Hardware: Physical View




# Hardware: View (version 0)




- CPU executes instructions; memory stores data
- To execute an instruction, the CPU must:
  - fetch an instruction;
  - fetch the data used by the instruction; and, finally,
  - execute the instruction on the data...
  - which may result in writing data back to memory.

# Hardware: View (version 1)



- The CPU holds instructions temporarily in the **instruction cache**
- The CPU holds data temporarily in a fixed number of **registers**
- **Instruction and operand fetching** is HW-controlled
- **Data movement** is (assembly language) programmer-controlled
- We'll learn about the instructions the CPU executes
  - take computer architecture course to find out how it actually executes them

# Hardware: View (version 1)



instructions

data

this week...  
Memory

How are data and instructions represented?

- **Data movement** is (assembly language)
- We'll learn about the instructions the CPU
  - take computer architecture course to find out how it actually executes them

temporarily in the

temporarily in a fixe

atching is HW

tion cache


How does a program find its data in memory?

# Memory, Data, and Addressing




- Representing information as bits and bytes
- Organizing and addressing data in memory
- Manipulating data in memory using C

# Question 1:




**How are data and  
instructions  
represented?**


# Binary Representations



- Everything is bits
- Base 2 number representation
  - A base 2 digit (0 or 1) is called a bit.
  - Represent  $351_{10}$  as  $0000000101011111_2$  or  $10101111_2$
- Electronic implementation
  - Easy to store with bi-stable elements
  - Reliably transmitted on noisy and inaccurate wires




# Review: Number Bases



- **Key terminology:** digit ( $d$ ) and base ( $B$ )
  - In base  $B$ , each digit is one of  $B$  possible symbols
- Value of  $i$ -th digit is  $d \times B^i$  where  $i$  starts at 0 and increases from right to left
  - $n$  digit number  $d_{n-1}d_{n-2} \dots d_1d_0$
  - value =  $d_{n-1} \times B^{n-1} + d_{n-2} \times B^{n-2} + \dots + d_1 \times B^1 + d_0 \times B^0$
  - In a *fixed-width* representation, left-most digit is called the **most-significant** and the right-most digit is called the **least-significant**
- **Notation:** Base is indicated using either a prefix or a subscript

# Describing Byte Values



- Binary (00000000, – 11111111,)

- Byte = 8 bits (binary digits)

MSB




- Decimal ( $0_{10}$  –  $255_{10}$ )

- Hexadecimal ( $00_{16}$  –  $FF_{16}$ )

- Byte = 2 hexadecimal (or “hex” or base 16) digits
  - Base 16 number representation
  - Use characters ‘0’ to ‘9’ and ‘A’ to ‘F’
  - Write FA1D37B<sub>16</sub> in the C language
    - as 0xFA1D37B or 0xfa1d37b


Hex	Decimal	Binary
0	0	0000
1	1	0001
2	2	0010
3	3	0011
4	4	0100
5	5	0101
6	6	0110
7	7	0111
8	8	1000
9	9	1001
A	10	1010
B	11	1011
C	12	1100
D	13	1101
E	14	1110
F	15	1111

## Question 2:



How does a  
program find its  
data in memory?

# Byte-Oriented Memory Organization



- Conceptually, memory is a single, **large array of bytes**, each with an unique **address** (index)
- The value of each byte in memory can be read and written
- Programs refer to bytes in memory by their **addresses**
  - Domain of possible addresses = address space
- But not all values (e.g., 351) fit in a single byte...
  - Store addresses to “remember” where other data is in memory
  - How much memory can we address with 1-byte (8-bit) addresses?
- Many operations actually use multi-byte values

# Machine Words



- Word size = address size = register size
- Word size bounds the size of the address space and memory
  - word size = w bits =>  **$2^w$**  addresses
  - Until recently, most machines used **32-bit (4-byte) words**
    - Potential address space:  $2^{32}$  addresses  
 $2^{32}$  bytes  $\approx 4 \times 10^9$  bytes = 4 billion bytes = **4GB**
    - Became too small for memory-intensive applications
  - Current x86 systems use **64-bit (8-byte) words**
    - Potential address space:  $2^{64}$  addresses  
 $2^{64}$  bytes  $\approx 1.8 \times 10^{19}$  bytes = 18 billion billion bytes = **18 EB** (exabytes)
    - Actual physical address space: **48 bits**

# Aside: Units and Prefixes



- Here focusing on large numbers (exponents > 0)
- Note that  $10^3 \approx 2^{10}$
- SI prefixes are ambiguous if base 10 or 2
- IEC prefixes are unambiguously base 2


**SIZE PREFIXES (10<sup>x</sup> for Disk, Communication; 2<sup>x</sup> for Memory)**

SI Size	Prefix	Symbol	IEC Size	Prefix	Symbol
$10^3$	Kilo-	K	$2^{10}$	Kibi-	Ki
$10^6$	Mega-	M	$2^{20}$	Mebi-	Mi
$10^9$	Giga-	G	$2^{30}$	Gibi-	Gi
$10^{12}$	Tera-	T	$2^{40}$	Tebi-	Ti
$10^{15}$	Peta-	P	$2^{50}$	Pebi-	Pi
$10^{18}$	Exa-	E	$2^{60}$	Exbi-	Ei
$10^{21}$	Zetta-	Z	$2^{70}$	Zebi-	Zi
$10^{24}$	Yotta-	Y	$2^{80}$	Yobi-	Yi



# Word-Oriented Memory Organization


- Addresses specify locations of bytes in memory
  - Address of word = **address of first byte in word**
  - Addresses of successive words differ by word size (in bytes): e.g., 4 (32-bit) or 8 (64-bit)
  - Address of word 0, 1, .. 10?





# Word-Oriented Memory Organization


- Addresses specify locations of bytes in memory
  - Address of word = **address of first byte in word**
  - Addresses of successive words differ by word size (in bytes): e.g., 4 (32-bit) or 8 (64-bit)
  - Address of word 0, 1, .. 10?



# A Picture of Memory (32-bit view)




- A “32-bit (4-byte) word-aligned” view of memory:



# A Picture of Memory (64-bit view)




- A “64-bit (8-byte) word-aligned” view of memory:
  - In this type of picture, each row is composed of 8 bytes
  - Each cell is a byte
  - A 64-bit pointer will fit on one row



# Picture of Memory (64-bit view)



- A “64-bit (8-byte) word-aligned” view of memory:
  - In this type of picture, each row is composed of 8 bytes
  - Each cell is a byte
  - A 64-bit pointer will fit on one row



# Addresses and Pointers (1)

32-bit example  
(pointers are 32-bits wide)



- An **address** is a location in memory
- A **pointer** is a data object that holds an address
- The value 351 is stored at address 0x04
  - $351_{10} = 15F_{16} = 0x\ 00\ 00\ 01\ 5F$

Address
0x00
0x04
0x08
0x0C
0x10
0x14
0x18
0x1C
0x20
0x24


The diagram illustrates a memory dump with 16 bytes. The first four bytes are highlighted in red and labeled '0x04'. The value '15F' is shown in the fourth column of the dump. The dump consists of four columns separated by vertical dashed lines, representing bytes 0 through 3. The addresses on the left range from 0x00 to 0x24.

# Addresses and Pointers (2)

## 32-bit example



- An **address** is a location in memory
  - A **pointer** is a data object that holds an address
  - The value 351 is stored at address 0x04
    - $351_{10} = 15F_{16} = 0x\ 00\ 00\ 01\ 5F$
  - A **pointer stored at address 0x1C points to address 0x04**




# Addresses and Pointers (3)

32-bit example  
(pointers are 32-bits wide)



- An **address** is a location in memory
- A **pointer** is a data object that holds an address
- The value 351 is stored at address 0x04
  - $351_{10} = 15F_{16} = 0x00\ 00\ 01\ 5F$
- A pointer stored at address 0x1C points to address 0x04
- **A pointer to a pointer is stored at address 0x24**




# Addresses and Pointers (4)

32-bit example  
(pointers are 32-bits wide)



- An **address** is a location in memory
- A **pointer** is a data object that holds an address
- The value 351 is stored at address 0x04
  - $351_{10} = 15F_{16} = 0x00\ 00\ 01\ 5F$
- A pointer stored at address 0x1C points to address 0x04
- A pointer to a pointer is stored at address 0x24
- **The value 12 is stored at address 0x14**
  - Is it a pointer?
    - Could be, depending on how you use it




# Addresses and Pointers (5)

64-bit example  
(pointers are 64-bits wide)



- A **64-bit (8-byte) word-aligned view of memory**
- Value 351 stored at address **0x08**
  - $351_{10} = 15F_{16} = 0x\ 00\ 00\ 01\ 5F$  **(note hex addresses)**
- Pointer stored at **0x38** points to address **0x08**




# Addresses and Pointers (6)

64-bit example  
(pointers are 64-bits wide)



- A 64-bit (8-byte) word-aligned view of memory
- Value 351 stored at address 0x08
  - $351_{10} = 15F_{16} = 0x\ 00\ 00\ 01\ 5F$  (note hex addresses)
- Pointer stored at 0x48 points to address 0x38
  - Pointer to a pointer!
- Is the data stored at 0x08 a pointer?
  - Could be, depending on how you use it



# Data Representations



- Sizes of data types (in bytes)

Java Data Type	C Data Type	32-bit (old)	x86-64
boolean	bool	1	1
byte	char	1	1
char		2	2
short	short int	2	2
int	int	4	4
float	float	4	4
	long int	4	8
double	double	8	8
long	long	8	8
	long double	8	16
(reference)	pointer *	4	8

To use “bool” in C, you must #include <stdbool.h>

address size = word size

# More on Memory Alignment in x86-64



- For good memory system performance, Intel recommends data be aligned
  - However the x86-64 hardware will work correctly regardless of alignment of data
  - Design choice: x86-64 instructions are *variable* bytes long
- **Aligned:** Primitive object of  $K$  bytes must have an address that is a multiple of  $K$

$K$	Type
1	char
2	short
4	int, float
8	long, double, pointers

More about alignment later in the course

# Byte Ordering





- How should bytes within a word be ordered in memory?
- Example:
  - **Store the 4-byte (32-bit) word: 0xa1 b2 c3 d4**
    - In what order will the bytes be stored?
- By convention, ordering of bytes called *endianness*
  - The two options are big-endian and little-endian
  - Based on *Gulliver's Travels*: tribes cut eggs on different sides (big, little)

# Byte Ordering



- **Big-Endian** (SPARC, z/Architecture, The Internet)
  - Least significant byte has highest address
- **Little-Endian** (x86, x86-64)
  - Least significant byte has lowest address
- Bi-endian (ARM, PowerPC)
  - Endianness can be specified as big or little
- **Example:** 4-byte data 0xa1b2c3d4 at address 0x100




# Byte Ordering Examples


Decimal:	12345
Binary:	0011 0000 0011 1001
Hex:	3 0 3 9

```
int x = 12345;  
// or x = 0x3039;
```

IA32, x86-64  
(little endian)




SPARC  
(big endian)




```
long int y = 12345;  
// or y = 0x3039;
```

(A long int is  
the size of a word)

IA32      x86-64




32-bit      64-bit  
SPARC      SPARC



# Endianness



- Often programmer can ignore endianness because it is handled for you
  - Bytes wired into correct place when reading or storing from memory (hardware)
  - Compiler and assembler generate correct behavior (software)
- Endianness still shows up:
  - Logical issues: accessing different amount of data than how you stored it (e.g. store int, access byte as a char)
  - When running down memory errors, need to know exact values
  - Manual translation to and from machine code



# Reading Byte-Reversed Listings

32-bit example

- Disassembly
  - Take binary machine code and generate an assembly code version
  - Does the reverse of the assembler
- Example instruction in memory
  - add value 0x12ab to register ‘ebx’ (a special location in the CPU)

Address	Instruction Code	Assembly Rendition
8048366:	81 c3 ab 12 00 00	add \$0x12ab,%ebx

## Deciphering numbers

# Reading Byte-Reversed Listings



- Disassembly
  - Take binary machine code and generate an assembly code version
  - Does the reverse of the assembler
- Example instruction in memory
  - add value 0x12ab to register ‘ebx’ (*a special location in the CPU*)

Address	Instruction Code	Assembly Rendition
8048366:	81 c3 ab 12 00 00	add \$0x12ab,%ebx

## Deciphering numbers


- Value:
- Pad to 32 bits:
- Split into bytes:
- Reverse (little-endian):

0x12ab  
0x000012ab  
00 00 12 ab  
ab 12 00 00

# Question:



- We store the value 0x 00 01 02 03 as a long int (size of long int = size of pointer) at address 0x100 and then get back 0x00 when we read a **byte** at address 0x102
- What machine setup are we using?



- (A) 32-bit, big-endian
- (B) 32-bit, little-endian
- (C) 64-bit, big-endian
- (D) 64-bit, little-endian

# Summary



- Memory is a long, *byte-addressed* array
  - Word size bounds the size of the *address space* and memory
  - Different data types use different number of bytes
  - Address of chunk of memory given by address of lowest byte in chunk
  - Object of  $K$  bytes is *aligned* if it has an address that is a multiple of  $K$
- IEC prefixes refer to powers of  $2^{10}$
- Pointers are data objects that holds addresses
- Endianness determines storage order for multi-byte objects

# Q&A

