

System Programming

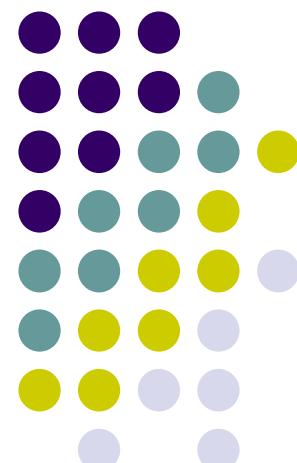
04. Integers (ch2.2-2.3)

2019. Fall

Instructor: Joonho Kwon

jhwon@pusan.ac.kr

Data Science Lab @ PNU





Roadmap

C:

```
car *c = malloc(sizeof(car));  
c->miles = 100;  
c->gals = 17;  
float mpg = get_mpg(c);  
free(c);
```

Java:

```
Car c = new Car();  
c.setMiles(100);  
c.setGals(17);  
float mpg =  
    c.getMPG();
```

Assembly
language:

```
get_mpg:  
    pushq  %rbp  
    movq   %rsp, %rbp  
    ...  
    popq   %rbp  
    ret
```

Machine
code:

```
0111010000011000  
100011010000010000000010  
1000100111000010  
11000001111101000011111
```

Computer
system:




- Memory & data
- Integers & floats
- x86 assembly
- Procedures & stacks
- Executables
- Arrays & structs
- Memory & caches
- Processes
- Virtual memory
- Memory allocation
- Java vs. C

But before we get to integers....



- Encode a standard deck of playing cards
- 52 cards in 4 suits, 13 cards per suit
 - How do we encode suits, face cards?
- What operations do we want to make easy to implement?
 - Which is the higher value card?
 - Are they the same suit?



Two possible representations



- 1) 1 bit per card (52): bit corresponding to card set to 1



- “One-hot” encoding (similar to set notation)
- Drawbacks:
 - Hard to compare values and suits
 - Large number of bits required

- 2) 1 bit per suit (4), 1 bit per number (13): 2 bits set




- Pair of one-hot encoded values
- Easier to compare suits and values, but still lots of bits used
- Can we do better?

Two better representations (1)



- 3) Binary encoding of all 52 cards – only 6 bits needed



low-order 6 bits of a byte

- $2^6 = 64 \geq 52$
- Fits in one byte (smaller than one-hot encodings)
- How can we make value and suit comparisons easier?

Two better representations (2)



- 4) Separate binary encodings of suit (2 bits) and value (4 bits)



- Also fits in one byte, and easy to do comparisons

K	Q	J	...	3	2	A
1101	1100	1011	...	0011	0010	0001

♣	00
♦	01
♥	10
♠	11

Compare Card Suits

mask: a bit vector designed to achieve a desired behavior when used with a bitwise operator on another bit vector v .

Here we turn all *but* the bits of interest in v to 0.


```
char hand[5];           // represents a 5-card hand
char card1, card2;     // two cards to compare
card1 = hand[0];
card2 = hand[1];
...
if ( sameSuitP(card1, card2) ) { ... }
```

```
#define SUIT_MASK 0x30
```

```
int sameSuitP(char card1, char card2) {
    return (!(card1 & SUIT_MASK) ^ (card2 & SUIT_MASK));
    //return (card1 & SUIT_MASK) == (card2 & SUIT_MASK);
}
```

returns int

SUIT_MASK = 0x30 =



equivalent


Compare Card Suits

mask: a bit vector designed to achieve a desired behavior when used with a bitwise operator on another bit vector v .

Here we turn all *but* the bits of interest in v to 0.

```
#define SUIT_MASK 0x30

int sameSuitP(char card1, char card2) {
    return (!( (card1 & SUIT_MASK) ^ (card2 & SUIT_MASK) ) );
    //return (card1 & SUIT_MASK) == (card2 & SUIT_MASK);
}
```



Compare Card Values

mask: a bit vector designed to achieve a desired behavior when used with a bitwise operator on another bit vector v .

```
char hand[5];           // represents a 5-card hand
char card1, card2;    // two cards to compare
card1 = hand[0];
card2 = hand[1];
...
if ( greaterValue(card1, card2) ) { ... }
```

```
#define VALUE_MASK 0x0F

int greaterValue(char card1, char card2) {
    return ((unsigned int)(card1 & VALUE_MASK) >
            (unsigned int)(card2 & VALUE_MASK));
}
```

VALUE_MASK = 0x0F = 

Compare Card Values

mask: a bit vector designed to achieve a desired behavior when used with a bitwise operator on another bit vector v .

```
#define VALUE_MASK 0x0F

int greaterValue(char card1, char card2) {
    return ((unsigned int)(card1 & VALUE_MASK) >
            (unsigned int)(card2 & VALUE_MASK));
}
```


0	0	1	0	0	0	1	0
---	---	---	---	---	---	---	---

&

0	0	0	0	1	1	1	1
---	---	---	---	---	---	---	---

=

0	0	0	0	0	0	1	0
---	---	---	---	---	---	---	---



VALUE_MASK

0	0	1	0	1	1	0	1
---	---	---	---	---	---	---	---

&

0	0	0	0	1	1	1	1
---	---	---	---	---	---	---	---

=

0	0	0	0	1	1	0	1
---	---	---	---	---	---	---	---

$$2_{10} > 13_{10}$$


0 (false)

Integers




- **Binary representation of integers**
 - Unsigned and signed
 - Casting in C
- Consequences of finite width representation
 - Overflow, sign extension
- Shifting and arithmetic operations

Encoding Integers



- The hardware (and C) supports two flavors of integers
 - *unsigned* – only the non-negatives
 - *signed* – both negatives and non-negatives
- Cannot represent all integers with w bits
 - Only 2^w distinct bit patterns
 - Unsigned values: $0 \dots 2^w - 1$
 - Signed values: $-2^{w-1} \dots 2^{w-1} - 1$
- **Example:** 8-bit integers (e.g. `char`)



Unsigned Integers



- Unsigned values follow the standard base 2 system
 - $b_7b_6b_5b_4b_3b_2b_1b_0 = b_72^7 + b_62^6 + \dots + b_12^1 + b_02^0$
- Add and subtract using the normal “carry” and “borrow” rules, just in binary

63	00111111
+ 8	+00001000
<hr/>	<hr/>
71	01000111

- Useful formula: $2^{N-1} + 2^{N-2} + \dots + 2 + 1 = 2^N - 1$
 - i.e. N ones in a row = $2^N - 1$
- How would you make *signed* integers?



Sign and Magnitude (1)

Most Significant Bit



- Designate the high-order bit (MSB) as the “sign bit”
 - sign=0: positive numbers; sign=1: negative numbers
- Benefits:
 - Using MSB as sign bit matches positive numbers with unsigned
 - All zeros encoding is still = 0
- Examples (8 bits):
 - 0x00 = 00000000_2 is non-negative, because the sign bit is 0
 - 0x7F = 01111111_2 is non-negative ($+127_{10}$)
 - 0x85 = 10000101_2 is negative (-5_{10})
 - 0x80 = 10000000_2 is negative... zero???




Sign and Magnitude (2)

Most Significant Bit


- MSB is the sign bit, rest of the bits are magnitude
- Drawbacks?




Sign and Magnitude (3)



- MSB is the sign bit, rest of the bits are magnitude
- Drawbacks:
 - Two representations of 0 (bad for checking equality)



Sign and Magnitude (4)




- MSB is the sign bit, rest of the bits are magnitude
- Drawbacks:
 - Two representations of 0 (bad for checking equality)
 - Arithmetic is cumbersome
 - Example: $4 - 3 \neq 4 + (-3)$

$$\begin{array}{r} 4 \\ - 3 \\ \hline 1 \end{array}$$

$$\begin{array}{r} 4 \\ + -3 \\ \hline -7 \end{array}$$


- Negatives “increment” in wrong direction



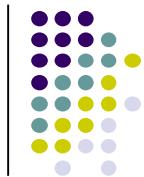
Two's Complement (1)



- Let's fix these problems:
 - “Flip” negative encodings so incrementing works




Two's Complement (2)



- Let's fix these problems:
 - “Flip” negative encodings so incrementing works
 - “Shift” negative numbers to eliminate -0


- MSB *still* indicates sign!
 - This is why we represent one more negative than positive number (-2^{N-1} to $2^{N-1} - 1$)



Two's Complement Negatives (1)



- How should we represent -1 in binary?



Two's Complement Negatives (2)



- Accomplished with one neat mathematical trick!

b_{w-1} has weight -2^{w-1} , other bits have usual weights $+2^i$



- 4-bit Examples:

- 1010_2 unsigned:

$$1*2^3 + 0*2^2 + 1*2^1 + 0*2^0 = 10$$


- 1010_2 two's complement:

$$-1*2^3 + 0*2^2 + 1*2^1 + 0*2^0 = -6$$

- 1 represented as:

$$1111_2 = -2^3 + (2^3 - 1)$$


- MSB makes it super negative, add up all the other bits to get back up to -1



Why Two's Complement is So Great



- Roughly same number of (+) and (-) numbers
- Positive number encodings match unsigned
- Single zero
- All zeros encoding = 0
- Simple negation procedure:
 - Get negative representation of any integer by taking bitwise complement and then adding one!
 $(\sim x + 1 == -x)$




Question



- Take the 4-bit number encoding $x = 0b1011$
 - Which of the following numbers is NOT a valid interpretation of x using any of the number representation schemes discussed today?
 - Unsigned, Sign and Magnitude, Two's Complement
- A. -4
- B. -5
- C. 11
- D. -3



4-bit Unsigned vs. Two's Complement (1)



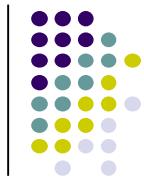
1 0 1 1

$$2^3 \times 1 + 2^2 \times 0 + 2^1 \times 1 + 2^0 \times 1$$

$$-2^3 \times 1 + 2^2 \times 0 + 2^1 \times 1 + 2^0 \times 1$$



4-bit Unsigned vs. Two's Complement (2)





1 0 1 1


$$2^3 \times 1 + 2^2 \times 0 + 2^1 \times 1 + 2^0 \times 1$$

$$-2^3 \times 1 + 2^2 \times 0 + 2^1 \times 1 + 2^0 \times 1$$

11 ← (math) difference = 16 = 2^4 → -5



4-bit Unsigned vs. Two's Complement (3)



1 0 1 1



$$2^3 \times 1 + 2^2 \times 0 + 2^1 \times 1 + 2^0 \times 1$$

$$-2^3 \times 1 + 2^2 \times 0 + 2^1 \times 1 + 2^0 \times 1$$

11

-5

(math) difference = 16 = 2⁴



Summary



- Choice of *encoding scheme* is important
 - Tradeoffs based on size requirements and desired operations
- Integers represented using unsigned and two's complement representations
 - Limited by fixed bit width
 - We'll examine arithmetic operations next lecture

Q&A

