

System Programming

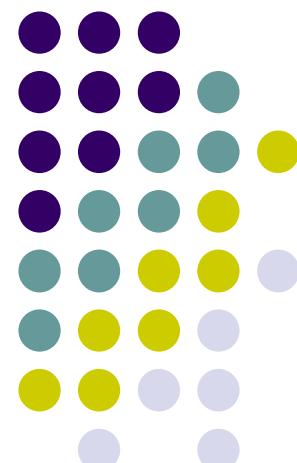
04. Integers II (ch2.2-2.3)

2019. Fall

Instructor: Joonho Kwon

jhwon@pusan.ac.kr

Data Science Lab @ PNU





Roadmap

C:

```
car *c = malloc(sizeof(car));  
c->miles = 100;  
c->gals = 17;  
float mpg = get_mpg(c);  
free(c);
```

Java:

```
Car c = new Car();  
c.setMiles(100);  
c.setGals(17);  
float mpg =  
    c.getMPG();
```


Assembly
language:

```
get_mpg:  
    pushq  %rbp  
    movq   %rsp, %rbp  
    ...  
    popq   %rbp  
    ret
```

Machine
code:

```
0111010000011000  
100011010000010000000010  
1000100111000010  
11000001111101000011111
```

Computer
system:



OS:



Memory & data

Integers & floats

x86 assembly

Procedures & stacks

Executables

Arrays & structs

Memory & caches

Processes

Virtual memory

Memory allocation

Java vs. C

Integers



- Binary representation of integers
 - Unsigned and signed
 - **Casting in C**
- Consequences of finite width representations
 - Overflow, sign extension
- Shifting and arithmetic operations

Two's Complement Arithmetic



- The same addition procedure works for both unsigned and two's complement integers
 - **Simplifies hardware:** only one algorithm for addition
 - **Algorithm:** simple addition, **discard the highest carry bit**
 - Called modular addition: result is sum *modulo* 2^w
- **4-bit Examples:**

4 +3 <hr/> $=7$	-4 +3 <hr/> $=-1$	4 -3 <hr/> $=1$
-----------------------	-------------------------	-----------------------

Why Does Two's Complement Work? (1)



- For all representable positive integers x , we want:

$$\begin{array}{r} \textit{bit representation of } x \\ + \textit{bit representation of } -x \\ \hline 0 \end{array} \quad (\text{ignoring the carry-out bit})$$

- What are the 8-bit negative encodings for the following?

$$\begin{array}{r} 00000001 \\ + \text{????????} \\ \hline 00000000 \end{array}$$

$$\begin{array}{r} 00000010 \\ + \text{????????} \\ \hline 00000000 \end{array}$$

$$\begin{array}{r} 11000011 \\ + \text{????????} \\ \hline 00000000 \end{array}$$

Why Does Two's Complement Work? (2)



- For all representable positive integers x , we want:

$$\begin{array}{r} \text{bit representation of } x \\ + \text{bit representation of } -x \\ \hline 0 \end{array} \quad (\text{ignoring the carry-out bit})$$

- What are the 8-bit negative encodings for the following?

$$\begin{array}{r} 00000001 \\ + 11111111 \\ \hline 100000000 \end{array} \quad \begin{array}{r} 00000010 \\ + 11111110 \\ \hline 100000000 \end{array} \quad \begin{array}{r} 11000011 \\ + 00111101 \\ \hline 100000000 \end{array}$$


These are the bitwise complement plus 1!

$$-x == \sim x + 1$$

Mapping Signed ↔ Unsigned (1)



Bits	Signed	Unsigned
0000	0	0
0001	1	1
0010	2	2
0011	3	3
0100	4	4
0101	5	5
0110	6	6
0111	7	7
1000	-8	8
1001	-7	9
1010	-6	10
1011	-5	11
1100	-4	12
1101	-3	13
1110	-2	14
1111	-1	15




Mapping Signed \leftrightarrow Unsigned (2)



Bits
0000
0001
0010
0011
0100
0101
0110
0111
1000
1001
1010
1011
1100
1101
1110
1111

Signed
0
1
2
3
4
5
6
7
-8
-7
-6
-5
-4
-3
-2
-1




Unsigned
0
1
2
3
4
5
6
7
8
9
10
11
12
13
14
15

Signed/Unsigned Conversion Visualized



- Two's Complement \rightarrow Unsigned

- Ordering Inversion
- Negative \rightarrow Big Positive





Values To Remember

- Unsigned Values

- $U_{Min} = 0b00\dots0$
= 0

- $U_{Max} = 0b11\dots1$
= $2^w - 1$

- Two's Complement Values

- $T_{Min} = 0b10\dots0$
= -2^{w-1}

- $T_{Max} = 0b01\dots1$
= $2^{w-1} - 1$

- -1 = 0b11...1

Example: Values for $w = 64$

	Decimal	Hex
UMax	18,446,744,073,709,551,615	FF FF FF FF FF FF FF FF
TMax	9,223,372,036,854,775,807	7F FF FF FF FF FF FF FF
TMin	-9,223,372,036,854,775,808	80 00 00 00 00 00 00 00
-1	-1	FF FF FF FF FF FF FF FF
0	0	00 00 00 00 00 00 00 00

In C: Signed vs. Unsigned



- Casting
 - Bits are unchanged, just interpreted differently!
 - `int tx, ty;`
 - `unsigned int ux, uy;`
 - *Explicit* casting
 - `tx = (int) ux;`
 - `uy = (unsigned int) ty;`
 - *Implicit* casting can occur during assignments or function calls
 - `tx = ux;`
 - `uy = ty;`

Casting Surprises

!!!

- Integer literals (constants)
 - By default, integer constants are considered *signed* integers
 - Hex constants already have an explicit binary representation
 - Use “U” (or “u”) suffix to explicitly force *unsigned*
 - Examples: 0U, 4294967259u
- Expression Evaluation
 - When you mixed unsigned and signed in a single expression, then **signed values are implicitly cast to unsigned**
 - Including comparison operators <, >, ==, <=, >=

Casting Surprises

REMINDER: If you mix unsigned and signed in a single expression, then **signed values are implicitly cast to unsigned**.
(The bit pattern does not change, bits are just interpreted differently.).

!!!

● 32-bit examples:

- TMin = -2,147,483,648, TMax = 2,147,483,647

Constant ₁	Constant ₂	Interpretation of bits	Relation
0 0000 0000 0000 0000 0000 0000 0000 0000	0U 0000 0000 0000 0000 0000 0000 0000 0000	Unsigned	==
-1 1111 1111 1111 1111 1111 1111 1111 1111	0 0000 0000 0000 0000 0000 0000 0000 0000	Signed	<
-1 1111 1111 1111 1111 1111 1111 1111 1111	0U 0000 0000 0000 0000 0000 0000 0000 0000	Unsigned	>
2147483647 0111 1111 1111 1111 1111 1111 1111 1111	-2147483648 1000 0000 0000 0000 0000 0000 0000 0000	Signed	>
2147483647U 0111 1111 1111 1111 1111 1111 1111 1111	-2147483648 1000 0000 0000 0000 0000 0000 0000 0000	Unsigned	<
-1 1111 1111 1111 1111 1111 1111 1111 1111	-2 1111 1111 1111 1111 1111 1111 1111 1110	Signed	>
(unsigned) -1 1111 1111 1111 1111 1111 1111 1111 1111	-2 1111 1111 1111 1111 1111 1111 1111 1110	Unsigned	>
2147483647 0111 1111 1111 1111 1111 1111 1111 1111	2147483648U 1000 0000 0000 0000 0000 0000 0000 0000	Unsigned	<
2147483647 0111 1111 1111 1111 1111 1111 1111 1111	(int) 2147483648U 1000 0000 0000 0000 0000 0000 0000 0000	Signed	>

Integers



- Binary representation of integers
 - Unsigned and signed
 - Casting in C
- **Consequences of finite width representations**
 - **Overflow, sign extension**
- Shifting and arithmetic operations




Arithmetic Overflow

Bits	Unsigned	Signed
0000	0	0
0001	1	1
0010	2	2
0011	3	3
0100	4	4
0101	5	5
0110	6	6
0111	7	7
1000	8	-8
1001	9	-7
1010	10	-6
1011	11	-5
1100	12	-4
1101	13	-3
1110	14	-2
1111	15	-1

- **Arithmetic Overflow** - when a calculation produces a result that can't be represented in the current encoding scheme
 - Integer range limited by fixed width
 - Can occur in both the positive and negative directions
- C and Java ignore overflow exceptions
 - You end up with a bad value in your program and no warning/indication... oops!

Overflow: Unsigned



- **Addition:** drop carry bit (-2^N)


$$\begin{array}{r} 15 \\ + 2 \\ \hline \cancel{17} \\ 1 \end{array}$$

$$\begin{array}{r} 1111 \\ + 0010 \\ \hline \cancel{10001} \end{array}$$

- **Subtraction:** borrow ($+2^N$)


$$\begin{array}{r} 1 \\ - 2 \\ \hline -1 \\ 15 \end{array}$$

$$\begin{array}{r} 10001 \\ - 0010 \\ \hline 1111 \end{array}$$



$\pm 2^N$ because of
modular arithmetic

Overflow: Two's Complement




- **Addition:** $(+) + (+) = (-)$ result?

$$\begin{array}{r} 6 \\ + 3 \\ \hline \cancel{9} \\ - 7 \end{array} \qquad \begin{array}{r} 0110 \\ + 0011 \\ \hline 1001 \end{array}$$

- **Subtraction:** $(-) + (-) = (+)?$

$$\begin{array}{r} -7 \\ - 3 \\ \hline -10 \end{array} \qquad \begin{array}{r} 1001 \\ - 0011 \\ \hline 0110 \end{array}$$

6



For signed: overflow if operands have same sign and result's sign is different



Sign Extension (1)

- What happens if you convert a *signed* integral data type to a larger one?
 - e.g. `char` → `short` → `int` → `long`
- **4-bit → 8-bit Example:**
 - Positive Case
 - Add 0's?

4-bit:	0010	=	+2
8-bit:	0000 0010	=	+2

Question




- Which of the following 8-bit numbers has the same *signed* value as the (signed) 4-bit number **0b1100**?
 - Underlined digit = MSB
- A. 0b 0000 1100
- B. 0b 1000 1100
- C. 0b 1111 1100
- D. 0b 1100 1100

Sign Extension (2)



- **Task:** Given a w -bit signed integer X , convert it to $w+k$ -bit signed integer X' *with the same value*
- **Rule:** Add k copies of sign bit

- Let x_i be the i -th digit of X in binary
- $X' = x_{w-1}, \dots, x_{w-1}, x_{w-1}, x_{w-2}, \dots, x_1, x_0$





Sign Extension (3)

0 0 1 0 4-bit 2

0 0 0 0 0 1 0 8-bit 2

1 1 0 0 4-bit -4

? ? ? ? 1 1 0 0 8-bit -4



Sign Extension (4)

0 0 1 0 4-bit 2

0 0 0 0 0 1 0 8-bit 2

1 1 0 0 4-bit -4

0 0 0 0 1 1 0 0 8-bit 12

Just adding zeroes to the front does not work



Sign Extension (5)

0 0 1 0 4-bit 2




0 0 0 0 0 1 0 8-bit 2

1 1 0 0 4-bit -4

1 0 0 0 1 1 0 0 8-bit -116

Just making the first bit=1 also does not work

Sign Extension (6)



Need to extend the sign bit to all “new” locations

Sign Extension Example



- Convert from smaller to larger integral data types
- C automatically performs sign extension
 - Java too

```
short int x = 12345;
int ix = (int) x;
short int y = -12345;
int iy = (int) y;
```

Var	Decimal	Hex	Binary			
x	12345	30 39		00110000	00111001	
ix	12345	00 00 30 39	00000000	00000000	00110000	00111001
y	-12345	CF C7			11001111	11000111
iy	-12345	FF FF CF C7	11111111	11111111	11001111	11000111

Integers



- Binary representation of integers
 - Unsigned and signed
 - Casting in C
- Consequences of finite width representations
 - Overflow, sign extension
- **Shifting and arithmetic operations**

Shift Operations



- Left shift ($x \ll n$) bit vector x by n positions
 - Throw away (drop) extra bits on left
 - Fill with 0s on right
- Right shift ($x \gg n$) bit-vector x by n positions
 - Throw away (drop) extra bits on right
 - Logical shift (for **unsigned** values)
 - Fill with 0s on left
 - Arithmetic shift (for **signed** values)
 - Replicate most significant bit on left
 - Maintains sign of x

Shift Operations



- Left shift ($x \ll n$)
 - Fill with 0s on right
- Right shift ($x \gg n$)
 - Logical shift (for **unsigned** values)
 - Fill with 0s on left
 - Arithmetic shift (for **signed** values)
 - Replicate most significant bit on left
- Notes:
 - Shifts by $n < 0$ or $n \geq w$ (bit width of x) are *undefined*
 - **In C:** behavior of \gg is determined by compiler
 - In gcc / C lang, depends on data type of x (signed/unsigned)
 - **In Java:** logical shift is \ggg and arithmetic shift is \gg

x	0010	0010
$x \ll 3$	0001	0 000
logical:	00 00	1000
arithmetic:	00 00	1000


x	1010	0010
$x \ll 3$	0001	0 000
logical:	00 10	1000
arithmetic:	11 10	1000

Shifting Arithmetic?



- What are the following computing?
 - $x \gg n$
 - 0b 0100 $\gg 1 = 0b 0010$
 - 0b 0100 $\gg 2 = 0b 0001$
 - Divide by 2^n
 - $x \ll n$
 - 0b 0001 $\ll 1 = 0b 0010$
 - 0b 0001 $\ll 2 = 0b 0100$
 - Multiply by 2^n
- Shifting is faster than general multiply and divide operations

Left Shifting Arithmetic 8-bit Example



- No difference in left shift operation for unsigned and signed numbers (just manipulates bits)
 - Difference comes during interpretation: $x * 2^n$?

		Signed	Unsigned
$x = 25;$	00011001	=	25 25
$L1=x<<2;$	 0001100100	=	100 100
$L2=x<<3;$	 00011001000	=	-56 200
$L3=x<<4;$	 000110010000	=	-112 144

signed overflow

unsigned overflow

Right Shifting 8-bit Examples



- **Reminder:** C operator `>>` does *logical shift* on **unsigned** values and *arithmetic shift* on **signed** values
 - Logical Shift: $x/2^n?$

$xu = 240u; \quad 11110000 = 240$

$R1u=xu>>3; \quad 00011110000 = 30$

$R2u=xu>>5; \quad 0000011110000 = 7$

rounding
(down)

Right Shifting 8-bit Examples



- **Reminder:** C operator `>>` does *logical shift* on **unsigned** values and *arithmetic shift* on **signed** values
 - Arithmetic Shift: $x/2^n$?

`xs = -16;` `11110000` $= -16$

`R1s=xu>>3;` `11111110000` $= -2$

`R2s=xu>>5;` `111111110000` $= -1$

rounding
(down)

Question



For the following expressions, find a value of **signed char** x , if there exists one, that makes the expression TRUE. Compare with your neighbor(s)!

- Assume we are using 8-bit arithmetic:

- $x == (\text{unsigned char}) x$
- $x \geq 128U$
- $x != (x >> 2) << 2$
- $x == -x$
 - Hint: there are two solutions
- $(x < 128U) \&\& (x > 0x3F)$

Question: A



- Assume we are using 8-bit arithmetic:

• $x \overset{\text{unsigned}}{==} (\text{unsigned char}) x$

• $x \overset{\text{unsigned}}{\geq} 128U$
 $0b1000 0000$

• $x != (x >> 2) << 2$ } will fill lowest two bits with zero

• $x == -x$ "flip bits and add one"

• Hint: there are two solutions

• $(x \overset{\text{unsigned}}{<} 128U) \&\& (x > 0x3F)$
 $0b0011 1111$

X cannot have a one in first bit in order to be $< 128U$

works for all x Both sides treated as unsigned

any $x < 0$ (need 1st bit set to one)

any x where lowest two bits are not 0b00

① $x = 0b0.....0 = 0$

② $x = 0b10....0 = -128$

any x where upper two bits are exactly 0b01

X Needs a one in the second bit position in order to be $> 0x3F$

Using Shifts and Masks (1)



- Extract the 2nd most significant *byte* of an `int`:

- First shift, then mask: $(x \gg 16) \& 0xFF$

x	00000001	00000010	00000011	00000100
x>>16	00000000	00000000	00000001	00000010
0xFF	00000000	00000000	00000000	11111111
(x>>16) & 0xFF	00000000	00000000	00000000	00000010

- Or first mask, then shift: $(x \& 0xFF0000) \gg 16$

x	00000001	00000010	00000011	00000100
0xFF0000	00000000	11111111	00000000	00000000
x & 0xFF0000	00000000	00000010	00000000	00000000
(x&0xFF0000)>>16	00000000	00000000	00000000	00000010

Using Shifts and Masks (2)



- Extract the *sign bit* of a signed int:
 - First shift, then mask: $(x \gg 31) \& 0x1$
 - Assuming arithmetic shift here, but this works in either case
 - Need mask to clear 1s possibly shifted in

x	00000001 00000010 00000011 00000100
x>>31	00000000 00000000 00000000 00000000 →
0x1	00000000 00000000 00000000 00000001
(x>>31) & 0x1	00000000 00000000 00000000 00000000

x	10000001 00000010 00000011 00000100
x>>31	11111111 11111111 11111111 11111111 →
0x1	00000000 00000000 00000000 00000001
(x>>31) & 0x1	00000000 00000000 00000000 00000001

Using Shifts and Masks (3)




- Conditionals as Boolean expressions (**assuming x is 0 or 1**)

- For **int** x , what does $(x \ll 31) \gg 31$ do?
 - $!123 = 0x\ 00, \,!123 = 0x01$


$x = \,!123$	00000000 00000000 00000000 00000000 1
$x \ll 31$	1 0000000 00000000 00000000 00000000
$(x \ll 31) \gg 31$	11111111 11111111 11111111 11111111
$\!x$	00000000 00000000 00000000 00000000 0
$\!x \ll 31$	00000000 00000000 00000000 00000000
$(\!x \ll 31) \gg 31$	00000000 00000000 00000000 00000000

- Can use in place of conditional:
 - In C: `if (x) {a=y;} else {a=z;}` equivalent to `a=x?y:z;`
 - `a=(((x<<31)>>31)&y) | (((!x<<31)>>31)&z);`

Unsigned Multiplication in C



Operands:
w bits



True Product:
2w bits



Discard w bits:
w bits




- Standard Multiplication Function
 - Ignores high order w bits
- Implements Modular Arithmetic
 - $\text{UMult}_w(u, v) = u \cdot v \bmod 2^w$

Multiplication with shift and add



- Operation $u \ll k$ gives $u * 2^k$
 - Both signed and unsigned


Operands: w bits



True Product: $w + k$ bits



Discard k bits: w bits



- Examples:

- $u \ll 3 == u * 8$

- $u \ll 5 - u \ll 3 == u * 24$

- Most machines shift and add faster than multiply

- *Compiler generates this code automatically*

Why Should I Use Unsigned?



- *Don't* use without understanding implications
 - Easy to make mistakes

```
unsigned i;  
for (i = cnt-2; i >= 0; i--)  
    a[i] += a[i+1];
```

- Can be very subtle

```
#define DELTA sizeof(int)  
int i;  
for (i = CNT; i-DELTA >= 0; i-= DELTA)  
    . . .
```

Counting Down with Unsigned



- Proper way to use unsigned as loop index

```
unsigned i;  
for (i = cnt-2; i < cnt; i--)  
    a[i] += a[i+1];
```

- See Robert Seacord, *Secure Coding in C and C++*
 - C Standard guarantees that unsigned addition will behave like modular arithmetic
 - $0 - 1 \rightarrow UMax$
- Even better

```
size_t i;  
for (i = cnt-2; i < cnt; i--)  
    a[i] += a[i+1];
```

- Data type `size_t` defined as unsigned value with length = word size
- Code will work even if `cnt = UMax`
- What if `cnt` is signed and < 0?

Why Should I Use Unsigned? (cont.)



- *Do Use When Performing Modular Arithmetic*
 - Multiprecision arithmetic
- *Do Use When Using Bits to Represent Sets*
 - Logical right shift, no sign extension

Code Security Example



```
/* Kernel memory region holding user-accessible data */
#define KSIZE 1024
char kbuf[KSIZE];

/* Copy at most maxlen bytes from kernel region to user buffer */
int copy_from_kernel(void *user_dest, int maxlen) {
    /* Byte count len is minimum of buffer size and maxlen */
    int len = KSIZE < maxlen ? KSIZE : maxlen;
    memcpy(user_dest, kbuf, len);
    return len;
}
```

- Similar to code found in FreeBSD's implementation of `getpeername`
- There are legions of smart people trying to find vulnerabilities in programs

Typical Usage



```
/* Kernel memory region holding user-accessible data */
#define KSIZE 1024
char kbuf[KSIZE];

/* Copy at most maxlen bytes from kernel region to user buffer */
int copy_from_kernel(void *user_dest, int maxlen) {
    /* Byte count len is minimum of buffer size and maxlen */
    int len = KSIZE < maxlen ? KSIZE : maxlen;
    memcpy(user_dest, kbuf, len);
    return len;
}
```

```
#define MSIZE 528

void getstuff() {
    char mybuf[MSIZE];
    copy_from_kernel(mybuf, MSIZE);
    printf("%s\n", mybuf);
}
```



Malicious Usage

```
/* Declaration of library function memcpy */
void *memcpy(void *dest, void *src, size_t n);
```

```
/* Kernel memory region holding user-accessible data */
#define KSIZE 1024
char kbuf[KSIZE];

/* Copy at most maxlen bytes from kernel region to user buffer */
int copy_from_kernel(void *user_dest, int maxlen) {
    /* Byte count len is minimum of buffer size and maxlen */
    int len = KSIZE < maxlen ? KSIZE : maxlen;
    memcpy(user_dest, kbuf, len);
    return len;
}
```

```
#define MSIZE 528

void getstuff() {
    char mybuf[MSIZE];
    copy_from_kernel(mybuf, -MSIZE);
    . . .
}
```

Summary



- Sign and unsigned variables in C
 - Bit pattern remains the same, just *interpreted* differently
 - Strange things can happen with our arithmetic when we convert/cast between sign and unsigned numbers
 - Type of variables affects behavior of operators (shifting, comparison)
- We can only represent so many numbers in w bits
 - When we exceed the limits, *arithmetic overflow* occurs
 - *Sign extension* tries to preserve value when expanding
- Shifting is a useful bitwise operator
 - Right shifting can be arithmetic (sign) or logical (0)
 - Can be used in multiplication with constant or bit masking

Q&A

