


File Structures

Ch03. A Disks, Raids and SSDs

2020. Spring
Instructor: Joonho Kwon
jhwon@pusan.ac.kr
Data Science Lab @ PNU



References



- **Operating Systems: Three Easy Pieces**
 - <http://pages.cs.wisc.edu/~remzi/OSTEP/>
 - Ch37. Hard Disk Drives
 - Ch38. Redundant Disk Arrays (RAID)
 - Ch44. Flash-based SSDs
- **Database Implementation**

Why Study Disks?




- Good Design
 - Responsive to the constraints of the medium and to the environment
- If files were stored just in memory
 - Data structures are enough
 - No need for file structures
- Secondary Storage
 - Very different from memory
 - Take much more time
 - Not all accesses are equal

Contents




- Ch37. Hard Disk Drives
- Ch38. RAID
- Ch39. SSD

Hard Drive Hardware



A Multi-Platter Disk




Addressing and Geometry



- Externally, hard drives expose a large number of **sectors** (blocks)
 - Typically 512 or 4096 bytes
 - Individual sector writes are **atomic**
 - Multiple sectors writes may be interrupted (**torn write**)
- Drive geometry
 - Sectors arranged into **tracks**
 - A **cylinder** is a particular track on multiple platters
 - Tracks arranged in concentric circles on **platters**
 - A disk may have multiple, double-sided platters
- Drive motor spins the platters at a constant rate
 - Measured in revolutions per minute (RPM)

Geometry Example




Common Disk Interfaces



- ST-506 → ATA → IDE → SATA
 - Ancient standard
 - Commands (read/write) and addresses in cylinder/head/sector format placed in device registers
 - Recent versions support **Logical Block Addresses** (LBA)
- SCSI (Small Computer Systems Interface)
 - Packet based, like TCP/IP
 - Device translates LBA to internal format (e.g. c/h/s)
 - Transport independent
 - USB drives, CD/DVD/Bluray, Firewire
 - iSCSI is SCSI over TCP/IP and Ethernet



Types of Delay With Disks



Track skew: offset sectors so that sequential reads across tracks incorporate seek delay

Three types of delay

1. Rotational Delay

- Time to rotate the desired sector to the read head
- Related to RPM

2. Seek delay

- Time to move the read head to a different track

3. Transfer time

- Time to read or write bytes

Access Time



- Data blocks can only be read and written if disk heads and platters are positioned accordingly.
- This design has implications on the **access time** to read/write a given block:

Access Time

1. Move disk arms to desired track (**seek time** t_s)
 2. Disk controller waits for desired block to rotate under disk head (**rotational delay** t_r)
 3. Read/write data (**transfer time** t_{tr})
- ➔ **Access Time:** $t = t_s + t_r + t_{tr}$

Example: Seagate Cheetah 15K.7



Cheetah 15K.7	
Capacity	600 GB
RPM	15000
Avg. Seek	3.4 ms
Max Transfer	163 MB/s

- Performance characteristics:
 - 4 disks, 8 heads, avg. 512 kB/track, 600 GB capacity
 - rotational speed: 15000 rpm (revolutions per minute)
 - average seek time: 3.4 ms
 - transfer rate 163 MB/s
- Q: What is the access time to read an 8KB block?
 - Average seek time: $t_s = 3.40 \text{ ms}$
 - Average rotational delay: $t_r = \frac{1}{2} \cdot \frac{1}{15000 \text{ min}^{-1}} = 2.00 \text{ ms}$
 - Transfer time for 8KB: $t_{tr} = \frac{8 \text{ KB}}{163 \text{ MB/s}} = 0.05 \text{ ms}$

Access Time for an 8KB block: $t = t_s + t_r + t_{tr} = 5.45 \text{ ms}$

Sequential vs. Random Access



- Q: Read 1,000 blocks of size 8 kB

- Random access: easy

- $t_{rnd} = 1,000 * 5.45 \text{ ms} = 5.45 \text{ s}$

- Sequential access

- $t_{seq} = t_s + t_r + 1000 \cdot t_{tr} + 16 \cdot t_{s,track-to-track}$
 - $= 3.4 \text{ ms} + 2.0 \text{ ms} + 50 \text{ ms} + 3.2 \approx 58.6 \text{ ms}$
 - Cheetah 15K.7 stores an average of 512 kB per track, with a 0.2 ms track-to-track seek time;
 - 8 kB blocks are spread across 16 tracks

Cheetah 15K.7	
Capacity	600 GB
RPM	15000
Avg. Seek t_s	3.4 ms
Max Transfer	163 MB/s
t_r	2.0 ms
t_{tr}	0.05 ms

Sequential vs. Random Access



- Comparison
 - Random I/O: **5.45 s**
 - Sequential I/O: **58.6 ms**
- Guidelines
 - Sequential I/O is **much** faster than random I/O
 - **Avoid random I/O** whenever possible
 - As soon as we need at least 58.6 ms $5,450 \text{ ms} = 1.07\%$ of a file, we better read the **entire** file sequentially

Random I/O results in very poor disk performance!

Evolution of Hard Disk Technology



- Disk seek and rotational latencies have only marginally improved over the last years (10% per year)
- **But:**
 - Throughput (i.e., transfer rates) improve by 50% per year
 - Hard disk capacity grows by 50% every year
- **Therefore:**
 - Random access cost hurts even more as time progresses

Ways to Improve I/O Performance



- The latency penalty is hard to avoid
- **But:**
 - Throughput can be increased rather easily by exploiting **parallelism**
- **Idea:**
 - Use multiple disks and access them in parallel, try to hide latency

Contents



- Ch37. Hard Disk Drives
- Ch38. RAID
- Ch39. SSD

Beyond Single Disks



- Hard drives are great devices
 - Relatively fast, persistent storage
- Shortcomings:
 - How to cope with disk failure?
 - Mechanical parts break over time
 - Sectors may become silently corrupted
 - Capacity is limited
 - Managing files across multiple physical devices is cumbersome
 - Can we make 10x 1 TB drives look like a 10 TB drive?

Redundant Array of Inexpensive Disks (1/2)




- RAID
 - use multiple disks to create the illusion of a large, faster, more reliable disk
- Externally, RAID looks like a single disk
 - i.e. RAID is **transparent**
 - Data blocks are read/written as usual
 - No need for software to explicitly manage multiple disks or perform error checking/recovery

Redundant Array of Inexpensive Disks (2/2)



- RAID
 - use multiple disks to create the illusion of a large, faster, more reliable disk
- Internally, RAID is a complex computer system
 - Disks managed by a dedicated CPU + software
 - RAM and non-volatile memory
 - Many different configuration options (**RAID levels**)


Example RAID Controller



RAID 0: Striping




- Key idea: present an **array** of disks as a single large disk
- Maximize parallelism by **striping** data cross all N disks



Addressing Blocks




- How do you access specific data blocks?
 - Disk = $\text{logical_block_number \% number_of_disks}$
 - Offset = $\text{logical_block_number / number_of_disks}$
- Example: read block 11
 - $11 \% 4 = \text{Disk 3}$
 - $11 / 4 = \text{Physical Block 2 (starting from 0)}$






Chunk Sizing



Chunk size = 1 block

- Chunk size impacts array performance
 - Smaller chunks → greater parallelism
 - Big chunks → reduced seek times
- Typical arrays use 64KB chunks



Chunk size = 2 block

Measuring RAID Performance (1/2)



- As usual, we focus on **sequential** and **random** workloads
- Assume disks in the array have **sequential** access time S
 - 10 MB transfer
 - $S = \text{transfer_size} / \text{time_to_access}$
 - $10 \text{ MB} / (7 \text{ ms} + 3 \text{ ms} + 10 \text{ MB} / 50 \text{ MB/s}) = 47.62 \text{ MB/s}$



Average seek time	7 ms
Average rotational delay	3 ms
Transfer rate	50 MB/s

Measuring RAID Performance (2/2)



- As usual, we focus on **sequential** and **random** workloads
- Assume disks in the array have **random** access time R
 - 10 KB transfer
 - $R = \text{transfer_size} / \text{time_to_access}$
 - $10 \text{ KB} / (7 \text{ ms} + 3 \text{ ms} + 10 \text{ KB} / 50 \text{ MB/s}) = 0.98 \text{ MB/s}$



Average seek time	7 ms
Average rotational delay	3 ms
Transfer rate	50 MB/s

Analysis of RAID 0




- Capacity: N
 - All space on all drives can be filled with data
- Reliability: 0
 - If any drive fails, data is permanently lost
- Sequential read and write: $N * S$
 - Full parallelization across drives
- Random read and write: $N * R$
 - Full parallelization across all drives




RAID 1: Mirroring



- RAID 0 offers high performance, but zero error recovery
- Key idea: make two copies of all data



RAID 0+1 and 1+0 Examples




- Combines striping and mirroring
- Superseded by RAID 4, 5, and 6

Analysis of RAID 1 (1/3)




- Capacity: $N / 2$
 - Two copies of all data, thus half capacity
- Reliability: 1 drive can fail, sometime more
 - If you are lucky, $N / 2$ drives can fail without data loss



Analysis of RAID 1 (2/3)




- Sequential write: $(N / 2) * S$
 - Two copies of all data, thus half throughput
- Sequential read: $(N / 2) * S$
 - Half of the read blocks are wasted, thus halving throughput



Analysis of RAID 1 (3/3)




- Random read: $N * R$
 - Best case scenario for RAID 1
 - Reads can parallelize across all disks
- Random write: $(N / 2) * R$
 - Two copies of all data, thus half throughput



The Consistent Update Problem



- Mirrored writes should be **atomic**
 - All copies are written, or none are written
- However, this is difficult to guarantee
 - Example: power failure
- Many RAID controllers include a **write-ahead log**
 - Battery backed, non-volatile storage of pending writes




Decreasing the Cost of Reliability



- RAID 1 offers highly reliable data storage
- But, it uses $N / 2$ of the array capacity
- Can we achieve the same level of reliability without wasting so much capacity?
 - Yes!
 - Use information coding techniques to build light-weight error recovery mechanisms

RAID 4: Parity Drive



Disk N only stores parity information for the other $N-1$ disks


Disk 0	Disk 1	Disk 2	Disk 3	Disk 4
0	0	1	1	$0 \wedge 0 \wedge 1 \wedge 1 = 0$
0	1	0	0	$0 \wedge 1 \wedge 0 \wedge 0 = 1$
1	1	1	1	$1 \wedge 1 \wedge 1 \wedge 1 = 0$
0	1	1	1	$0 \wedge 1 \wedge 1 \wedge 1 = 1$

Parity calculated using XOR


Updating Parity on Write




- How is parity updated when blocks are written?
 1. Additive parity



2. Subtractive parity



Random Writes and RAID 4



- Random writes in RAID 4
 1. Read the target block and the parity block
 2. Use subtraction to calculate the new parity block
 3. Write the target block and the parity block
- RAID 4 has terrible write performance
 - Bottlenecked by the parity drive


Analysis of RAID 4



- Capacity: $N - 1$
 - Space on the parity drive is lost
- Reliability: 1 drive can fail
- Sequential Read and write: $(N - 1) * S$
 - Parallelization across all non-parity blocks
- Random Read: $(N - 1) * R$
 - Reads parallelize over all but the parity drive
- Random Write: $R / 2$
 - Writes serialize due to the parity drive
 - Each write requires 1 read and 1 write of the parity drive, thus $R / 2$




RAID 5: Rotating Parity



Disk 0	Disk 1	Disk 2	Disk 3	Disk 4
0	0	1	1	$0 \wedge 0 \wedge 1 \wedge 1 = 0$
1	0	0	$0 \wedge 1 \wedge 0 \wedge 0 = 1$	0
1	1	$1 \wedge 1 \wedge 1 \wedge 1 = 0$	1	1
1	$0 \wedge 1 \wedge 1 \wedge 1 = 1$	0	1	1

Random Writes and RAID 5



Unlike RAID 4,
writes are spread
roughly evenly
across all drives

- Random writes in RAID 5
 1. Read the target block and the parity block
 2. Use subtraction to calculate the new parity block
 3. Write the target block and the parity block
- Thus, 4 total operations (2 reads, 2 writes)
 - Distributed across all drives

Analysis of Raid 5



- Capacity: $N - 1$ [same as RAID 4]
- Reliability: 1 drive can fail [same as RAID 4]
- Sequential Read and write: $(N - 1) * S$ [same]
 - Parallelization across all non-parity blocks
- Random Read: $N * R$ [vs. $(N - 1) * R$]
 - Unlike RAID 4, reads parallelize over all drives
- Random Write: $N / 4 * R$ [vs. $R / 2$ for RAID 4]
 - Unlike RAID 4, writes parallelize over all drives
 - Each write requires 2 reads and 2 write, hence $N / 4$





Comparison of RAID Levels

- N – number of drives
- S – sequential access speed
- R – random access speed
- D – latency to access a single disk

		RAID 0	RAID 1	RAID 4	RAID 5
Throughput	Capacity	N	$N/2$	$N-1$	$N-1$
	Reliability	0	1 (maybe $N/2$)	1	1
	Sequential Read	$N * S$	$(N/2) * S$	$(N-1) * S$	$(N-1) * S$
	Sequential Write	$N * S$	$(N/2) * S$	$(N-1) * S$	$(N-1) * S$
	Random Read	$N * R$	$N * R$	$(N-1) * R$	$N * R$
Latency	Random Write	$N * R$	$(N/2) * R$	$R/2$	$(N/4) * R$
	Read	D	D	D	D
	Write	D	D	$2 * D$	$2 * D$

RAID 6



- Any two drives can fail
- $N - 2$ usable capacity
- No overhead on read, significant overhead on write
- Typically implemented using Reed-Solomon codes

Choosing a RAID Level




- Best performance and most capacity?
 - RAID 0
- Greatest error recovery?
 - RAID 1 (1+0 or 0+1) or RAID 6
- Balance between space, performance, and recoverability?
 - RAID 5

Other Considerations



- Many RAID systems include a **hot spare**
 - An idle, unused disk installed in the system
 - If a drive fails, the array is immediately rebuilt using the hot spare
- RAID can be implemented in hardware or software
 - Hardware is faster and more reliable...
 - But, migrating a hardware RAID array to a different hardware controller almost never works
 - Software arrays are simpler to migrate and cheaper, but have worse performance and weaker reliability
 - Due to the **consistent update** problem

Contents




- Ch37. Hard Disk Drives
- Ch38. RAID
- Ch39. SSD

Beyond Spinning Disks



- Hard drives have been around since 1956
 - The cheapest way to store large amounts of data
 - Sizes are still increasing rapidly
- However, hard drives are typically the slowest component in most computers
 - CPU and RAM operate at GHz
 - PCI-X and Ethernet are GB/s
- Hard drives are not suitable for mobile devices
 - Fragile mechanical components can break
 - The disk motor is extremely power hungry

Solid State Drives (SSD) (1/2)




- SSDs provide **very low-latency random read access** (< 0.01 ms)
- **Random writes** however, are significantly **slower** than on traditional magnetic drives:
 - (Blocks of) Pages have to be **erased** before they can be updated
 - Once pages have been erased, sequentially writing them is almost as fast as reading

Solid State Drives (SSD) (2/2)



- NAND flash memory-based drives
 - High voltage is able to change the configuration of a floating-gate transistor
 - State of the transistor interpreted as binary data



Advantages of SSDs




- More resilient against physical damage
 - No sensitive read head or moving parts
 - Immune to changes in temperature
- Greatly reduced power consumption
 - No mechanical, moving parts
- Much faster than hard drives
 - >500 MB/s vs ~200 MB/s for hard drives
 - No penalty for random access
 - Each flash cell can be addressed directly
 - No need to rotate or seek
 - Extremely high throughput
 - Although each flash chip is slow, they are RAIDed

SSDs: Page-Level Writes, Block-Level Deletes



- Typical **page size**: 128 kB
- SSDs erase **blocks of pages**:
 - block \approx 64 pages (8 MB)
- Example
 - Perform block-level delete to accomodate new data pages



Example: Seagate Pulsar.2



- Performance characteristics:
 - NAND flash memory, 800 GB capacity
 - standard 2.5" enclosure, no moving/rotating parts
 - data read/written in pages of 128 kB size
 - transfer rate 370 MB/s
- Q: What is the access time to read an 8KB block?
 - Average seek time: $t_s = 0.00 \text{ ms}$
 - Average rotational delay: $t_r = 0.00 \text{ ms}$
 - Transfer time for 8KB: $t_{tr} = \frac{8 \text{ KB}}{370 \text{ MB/s}} = 0.30 \text{ ms}$

Access Time for an 8KB block: $t = t_s + t_r + t_{tr} = 0.3 \text{ ms}$



Sequential vs. Random Access with SSDs



Pulsar.2	
Avg. Seek t_s	0.0 ms
t_r	0.0 ms
t_{tr}	0.30 ms

- Q: Read 1,000 blocks of size 8 kB
- Random access: easy
 - $t_{rnd} = 1,000 * 0.3 \text{ ms} = 0.3 \text{ s}$
- Sequential access
 - $t_{seq} = \left\lceil \frac{1000 \cdot 8 \text{ kB}}{128 \text{ kB}} \right\rceil \cdot t_{tr} \approx 18.9 \text{ ms}$
 - Pulsar.2 (sequentially) reads data in 128 kB chunks.

=> Sequential I/O still beats random I/O (but random I/O is more feasible again)

- Adapting database technology to these characteristics is a current research topic

Challenges with Flash





- Flash memory is written in pages, but erased in blocks
 - Pages: 4 – 16 KB, Blocks: 128 – 256 KB
 - Thus, flash memory can become fragmented
 - Leads to the **write amplification** problem
- Flash memory can only be written a fixed number of times
 - Typically 3000 – 5000 cycles for MLC
 - SSDs use **wear leveling** to evenly distribute writes across all flash cells

Write Amplification



G moved to new block by the garbage collector

Cleaned block can now be rewritten



- Once all pages have been written, valid pages must be consolidated to free up space
- Write amplification:** a write triggers garbage collection/compaction
 - One or more blocks must be read, erased, and rewritten before the write can proceed

Garbage Collection



- Garbage collection (GC) is vital for the performance of SSDs
- Older SSDs had fast writes up until all pages were written once
 - Even if the drive has lots of “free space,” each write is amplified, thus reducing performance
- Many SSDs over-provision to help the GC
 - 240 GB SSDs actually have 256 GB of memory
- Modern SSDs implement background GC
 - However, this doesn’t always work correctly

The Ambiguity of Delete




- Goal: the SSD wants to perform background GC
 - But this assumes the SSD knows which pages are invalid
- Problem: most file systems don't actually delete data
 - On Linux, the “delete” function is unlink()
 - Removes the file meta-data, but not the file itself



Delete Example

File metadata
(inode, name,
etc.)



Metadata is
overwritten,
but the file
remains


1. File is written to SSD
2. File is deleted
3. The GC executes
 - 9 pages look valid to the SSD
 - The OS knows only 2 pages are valid

- Lack of explicit delete means the GC wastes effort copying useless pages
- Hard drives are not GCed, so this was never a problem

TRIM



- New SATA command TRIM (SCSI – UNMAP)
 - Allows the OS to tell the SSD that specific LBAs are invalid, may be GCed



- OS support for TRIM
 - Win 7, OSX Snow Leopard, Linux 2.6.33, Android 4.3
- Must be supported by the SSD firmware

Wear Leveling



- Recall: each flash cell wears out after several thousand writes
- SSDs use **wear leveling** to spread writes across all cells
 - Typical consumer SSDs should last ~5 years





Dynamic Wear Leveling

Wear Leveling



If the GC runs now, page G must be copied

Wait as long as possible before garbage collecting



Static Wear Leveling

Blocks with long lived data receive less wear



SSD controller periodically swap long lived data to different blocks

SSD Controllers



- SSDs are extremely complicated internally
- All operations handled by the SSD controller
 - Maps LBAs to physical pages
 - Keeps track of free pages, controls the GC
 - May implement background GC
 - Performs wear leveling via data rotation
- Controller performance is crucial for overall SSD performance

Flavors of NAND Flash Memory



Multi-Level Cell (MLC)

- Multiple bits per flash cell
 - For two-level: 00, 01, 10, 11
 - 2, 3, and 4-bit MLC is available
- Higher capacity and cheaper than SLC flash
- Lower throughput due to the need for error correction
- 3000 – 5000 write cycles
- Consumes more power

Single-Level Cell (SLC)

- One bit per flash cell
 - 0 or 1
- Lower capacity and more expensive than MLC flash
- Higher throughput than MLC
- 10000 – 100000 write cycles

Expensive, enterprise drives

Consumer-grade drives

Q&A

