

File Structures

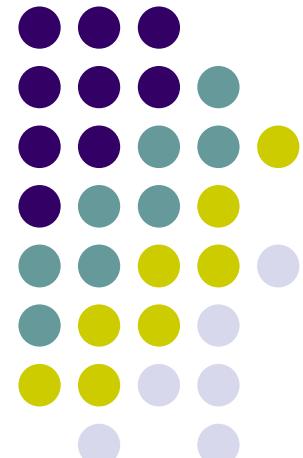
06. Organizing Files for Performance

2020. Spring

Instructor: Joonho Kwon

jhwon@pusan.ac.kr

Data Science Lab @ PNU



Outline



- 6.1 Data compression
 - skipped
- 6.2 Reclaiming space in files
- 6.3 Finding things quickly
 - An Introduction to internal sorting and binary searching
- 6.4 Keysorting

Modifying a file



- Scenario
 - modify a record in a variable-length record file, where the new record is longer than the original record
 - should keep RID intact
- Method1
 - append the new record to the end of the file, and then put a pointer from the original record space to the extension of the record
 - slower than it was originally
- Method2
 - rewrite the whole record at the end of the file, leaving a hole at the original location of the record
 - the file contains waste space

Space reclaiming



- Is the file organization deteriorated?
 - Record addition
 - No deterioration
 - Record updating and deleting
 - Yes!
- Space reclaiming is needed when
 - record updating
 - can be treated as a record deletion followed by a record addition
 - record deletion (we focus on this)
 - want to reuse the space

Record deletion and storage compaction (1/2)



- Storage compaction
 - Look for places in a file where there is no data at all and recover this space
 - Record deletion => empty spaces occur in files
- How to recognize records as deleted?
 - **to place a special mark in each deleted record**

Ames|John|123 Maple|Stillwater|OK|74075|...
Morrison|Sebastian|9035 South Hillcrest|Forest Village|OK|78420|
Brown|Martha|625 Kimbark|Des Moines|IA|50311|...

Ames|John|123 Maple|Stillwater|OK|74075|...
***|Sebastian|9035 South Hillcrest|Forest Village|OK|78420|**
Brown|Martha|625 Kimbark|Des Moines|IA|50311|...

After **compaction**

Ames|John|123 Maple|Stillwater|OK|74075|...
Brown|Martha|625 Kimbark|Des Moines|IA|50311|...

Record deletion and storage compaction (2/2)



- how to reuse the space from the deleted record
 - Storage compaction approach does not reuse the space for a while
 - Programs must include logic for ignoring records with deleted marks
 - Benefit: easily undelete a record
 - the reclamation of the space from the deleted records at once
 - a simple compaction
 - a file copy program that skips over deleted records
 - how often to run the storage compaction program
 - based on # of deleted records
 - based on the calendar

Dynamic Storage reclamation



- Storage compaction
 - The simplest and most widely used of the storage reclamation methods
 - In some applications wants to reuse the space from deleted records as soon as possible
- dynamic storage reclamation
 - to make record reuse happen more quickly
 - know immediately if there are empty slots in the file
 - jump directly to one of those slots if they exist

Case 1 . Deleting fixed-length records (1/2)



- How to reuse the space from deleted records
 - 1. deleted records must be marked in special way
 - 2. we could find the space which deleted records occupied
- Much slower approach when working with fixed-length records
 - Search sequentially through a file before adding a file
 - If deleted record is found → **reuse the space**
 - If not, the new record can be **appended at the end**

Case 2 . Deleting fixed-length records (1/2)




- To make record reuse quickly, we need
 - (1) a way to know immediately if there are empty slots in the file
 - (2) a way to jump directly to one of those slots if they exist
- Solution
 - Linked lists or Stacks for avail list
 - avail list : a list that is made up of deleted records

Detour: Linked lists and Stacks



● Linked Lists


- the available space list made up of deleted records



● Stacks

- handle a list as a stack

- All insertions and removals of nodes take place at one end of the list



10

Linking and Stacking deleted records (1/4)



- Placing the related records on stack
 - Meet two criteria for rapid access to reusable space
 - know immediately if there are empty slots in the file : The top of stacks contains the end-of-list value → no empty slots
 - jump directly to one of those slots if it exists : the stack top contains a valid node reference → reuse it
- Where do we keep the stack
 - A separate list (perhaps in a separate file)
 - Embedded within the data file

Linking and Stacking deleted records (2/4)



- Embedded within the data file
 - arranging and rearranging the links used to make one available record slot point to the next
 - deleted records are marked by replacing the first field with an asterisk
 - second field of deleted record points to next record
 - pointing is not done with pointers(byte offset), but RRNs. Why?
 - The reason: working with fixed-length records in a disk file

Linking and Stacking deleted records (3/4)



● Example

- Files with seven records (RRNs 0-6)

List head (first available record) => -1

0	1	2	3	4	5	6
Edwards...	Bates...	Wills...	Park....	Maters...	Browns...	Chavez

- Deleting record 3

List head (first available record) => 3

0	1	2	3	4	5	6
Edwards...	Bates...	Wills...	*-1	Maters...	Browns...	Chavez

- Deleting record 5

List head (first available record) => 5

0	1	2	3	4	5	6
Edwards...	Bates...	Wills...	*-1	Maters...	*3	Chavez

Linking and Stacking deleted records (4/4)



- Example (cont'd)

- Deleting record 1

List head (first available record) => **1**

0	1	2	3	4	5	6
Edwards...	*5	Wills...	*-1	Maters...	*3	Chavez

- After adding new records

- 1st on slot RRN1, 2nd on Slot RRN5, 3rd on Slot RRN3

List head (first available record) => **-1**

0	1	2	3	4	5	6
Edwards...	1 st new rec	Wills...	3 rd new rec	Maters...	2 nd new rec	Chavez

Implementing fixed-length record deletion



- Place deleted records on a linked avail list
 - Treat avail list as a stack
 - need a suitable place to keep the RRN of the first available record on the avail list
 - place the first available record in a header record at the start of the file

Case 2. Deleting variable-length records



- Supporting record reuse through an avail list
 - link the deleted records together into a list
 - an algorithm for adding newly deleted records to the avail list
 - an algorithm for finding and removing records from the avail list

An avail list of variable-length records (1/2)




- What kind of file structure do we need?
 - the length of each record by placing a byte count of the record contents at the beginning of each record
 - E.g) VariableLengthBuffer
- How to implement deletion
 - place a single asterisk in the first field, followed by a link field pointing to the next deleted record on the avail list
 - a link field: contain the byte offset
 - cannot use RRN for links

An avail list of variable-length records (2/2)




- Example

- A file with three records



- After deleting the second records





Adding and removing avail list records



- Avail list

- Problem: cannot access the avail list as a stack
 - The reason: the record slots on the avail list differ in size
- need to search through the avail list for a record slot that is the right size
 - define right size as "big enough"
 - a record that requires 55 bytes is to be added
 - 72(big enough) → remove it from the avail list




Internal fragmentation (fixed-length record)



- Internal fragmentation (in fixed-length record)

- waste space within a record

Unused space: Internal fragmentation




- variable-length records
 - minimize wasted space by doing away with internal fragmentation


External Fragmentation (variable length records) (1/4)



- Simple approach
 - After deleting the second records




- Adding new record:
 - keeping the 64-byte record slot intact → fragmentation




External Fragmentation (variable length records) (2/4)



- Another approach
 - After deleting the second records



- Adding new record:
 - Break the 64-byte record two parts: new rec and avail list




External Fragmentation (variable length records) (3/4)



- Another approach (cont'd)

- Adding new record:

Lee|Ed|Rt 2|Ada|OK|74820|



External Fragmentation (variable length records) (4/4)



- External fragmentation
 - too small space that the probability to reuse it is close to zero
 - Too fragmented to be reused
 - three possible solutions
 - storage compaction
 - regenerate the file when external fragmentation becomes intolerable
 - coalescing the holes
 - Combining adjacent record slots to make a single, larger record slot
 - minimizing fragmentation by adopting placement strategy

Placement strategies (1/2)



- ways to add and remove variable-length records from an avail list
- need to remove a record slot from the avail list (to add a record to the file)

Placement strategies (2/2)



- find a record slot that is big enough
 - First-fit
 - select **the first available record slot**
 - suitable when lost space is due to internal fragmentation
 - Best-fit
 - select **the available record slot closest in size**
 - avail list in ascending order
 - suitable when lost space is due to internal fragmentation
 - Worst-fit
 - select **the largest record slot**
 - avail list in descending order
 - suitable when lost space is due to external fragmentation



Outline

- 6.1 Data compression
 - skipped
- 6.2 Reclaiming space in files
- 6.3 Finding things quickly
 - An Introduction to internal sorting and binary searching
- 6.4 Keysorting

Turning point



- Searching and Sorting data in a memory
 - Searching and sorting algorithms → making many comparisons
 - Fast, if the size of data < the size of memory
- search a file for a particular piece of information and sorting a file
 - Goal: minimize the number of disk accesses

Finding things in simple field and record files



- File has fixed length records
 - RRN lets us to compute the **record's byte offset** and jump to it using **direct access**
- Access by key
 - Implies a sequential search
 - Have to look through the entire file

Search by guessing: Binary search



- Assumption: a file is sorted
- an algorithm for binary searching on a file of fixed-sized records

```
int BinarySearch (FixedRecordFile &file, RecType &obj, KeyType &key)
{
    int low = 0;
    int high = file.NumRecs() - 1;
    while (low <= high)
    {
        int guess = (high - low) / 2;
        file.ReadByRRN (obj, guess);
        if (obj.key() == key) return 1; // record found
        if (obj.key() < key) high = guess - 1; // search before guest
        else low = guess + 1; // search after guest
    }
    return 0;
}
```

Classes for binary search



```
Class KeyType
{
public:
    int operator == (KeyType &);
    int operator < (KeyType &);
};

class RecType
{
public:
    KeyType key();
};

class FixedRecordFile
{
public:
    int NumRecs();
    int ReadByRRN (RecType &record,
                   int RRN);
};
```

- **KeyType and RecType**
 - the minimum definition to allow a successful completion of BinarySearch
- **A class FixedRecordFile**
 - **NumRecs()**
 - return # of records in the FixedRecordFile
 - **ReadByRRN()**
 - read the record at a specific RRN and unpack it into a RecType object

Binary search versus sequential search




- to make use of binary searching, we need to sort a list on the basis of a key
 - a binary search: $O(\log_2 n)$
 - a sequential search: $O(n)$
- sorting is a very important part of file processing
 - internal sorting
 - external sorting

Sorting a disk file in RAM



- If applying sorting directory to data stored on a disk
 - A lot of seeking, and rereading of data
- Alternative
 - read the entire file from disk to memory
 - use internal sort (=sort in memory)
 - UNIX sort utility uses internal sort



The limitations of sort then binary search (1/3)



- problem1: binary searching requires more than one or two accesses
 - a binary search requires $\lceil \log n \rceil$ comparisons
 - $O(\log_2 1,000) = 9.5$ accesses per request
 - $O(\log_2 100,000) = 16.5$ accesses per request
 - Not a negligible cost
 - Alternatives
 - access records by RRN rather than by key → single access
 - the use of index structures

The limitations of sort then binary search (2/3)



- problem2: keeping a file sorted is very expensive
 - encounter difficulty when we add a record
 - Tradeoff : search benefit v.s. insert overhead
 - accumulate record additions in a transaction file and made in a batch merge → merge sort
- Better solutions:
 - not involve reordering of the records in the file when a new record is added
 - the use of simple indexes, can involve hashing
 - or, associated with data structures that allow for more rapid, efficient reordering of the file
 - use of tree structures, B-tree
 - to keep the file in order

The limitations of sort then binary search (3/3)



- problem3: an internal sort works only on small files
 - if the file is large, then need external sorting
 - develop a variation on internal sorting called a **keysort**

Outline



- 6.1 Data compression
 - skipped
- 6.2 Reclaiming space in files
- 6.3 Finding things quickly
 - An Introduction to internal sorting and binary searching
- 6.4 Keysorting

Keysorting



- Keysort

- a variation on internal sorting
 - to avoid the sorting of records in a file
- sort a file in RAM
 - **sorted thing is “key” only** , called tag sort
 - do not need to read the whole file into RAM during the sorting process

- Sorting procedure

- 1. Read only the keys into memory
- 2. Sort the keys
- 3. Rearrange the records in file by the sorted keys

Descriptions of the method



- Assume that we are dealing with a fixed-length record file

```
class FixedRecordFile
{public:
    int NumRecs();
    int ReadByRRN (RecType &record, int RRN);
    //additional methods required for keysort
    int Create (char *fileName);
    int Append (RecType &record);
};

class KeyRRN
{public:
    KeyType KEY; int RRN;
    KeyRRN();
    KeyRRN (KeyType key, int rrn);
};

int Sort (KeyRRN[], int numKeys);
```


Conceptual view of keysorting (1/2)



- before sorting

Memory


KEYNODES[]	
Key	RRN
Harrison Susan	1
Kellog Bill	2
Harris Margaret	3
...	
Bell Robert	k



- after sorting in memory

Memory

KEYNODES[]	
Key	RRN
Bell Robert	k
Harris Margaret	3
Harrison Susan	1
...	
Kellog Bill	2



Conceptual view of keysorting (2/2)




- Writing to a new file in the order of KEYNODES[]

Memory

KEYNODES[]	
Key	RRN
Bell Robert	k
Harris Margaret	3
Harrison Susan	1
...	
Kellog Bill	2

RecType

Bell Robert 8912 Hill....
Harris Margaret 4343 West....
Harrison Susan 387 Eastern....
...
Kellog Bill 17 Maple....



Algorithm for keysort (1/2)



- Read the (key, RRN) pairs into an array of KeyRRN objects
- Sort the KEYNODES[] array according to the KEY field
- reorganize the file according to new ordering

Algorithm for keysort (2/2)



```
int KeySort(FixedRecordFile & inFile, char * outFileName)
{
    RecType obj;
    KeyRRN * KEYNODES = new KeyRRN[inFile.NumRecs()];
    // read file and load keys
    for (int I = 0; I < inFile.NumRecs(); I++)
    {
        inFile.ReadByRRN(obj, I); //read record I
        KEYNODES[I] = keyRRN(obj.Key(), I); // put key and RRN into KEYNodes
    }
    sort (KEYNODES, inFile.NumRecs()); // sort Keys
    FixedRecordFile outFile; // file to hold records in key order
    outFile.Create(outFileName);
    // write new file in key order
    for (int j = 0; j < inFile.NumRecs(); j++)
    {
        inFile.ReadByRRN(obj, KEYNODES[j].RRN); // read in key order
        outFile.Append(obj); // write in key order
    }
    return 1;
}
```

Limitations of the keysort method




- never desirable to read entire records twice
 - One read for building array
 - Another read for writing out the new sorted file
- not reading through the input file sequentially
 - Seek in the input file to the record whose KEYNODES[i].RRN
 - require many random seeks into the input file
 - the getting-something-for-nothing aspect of keysort

Another solution: why bother to write the file back?



- the problem of rearranging all the records in the file
 - what if we skip to write out a sorted version of the file?
 - simply write out a copy of the array of canonical key nodes
 - **create a second kind of file, an index file**
 - output an index to the original file
- One of our favorite categories of solutions to computer science problems
 - If looking like a bottleneck, consider skipping it altogether



Pinned records (1/2)




- An avail list of deleted record slots
 - Created by linking all of the available slots together
 - linking is done by writing a link field into each deleted record that points to the next deleted record
 - a link field: have a pointer, physical locations
- Pinned records
 - Records that are referenced to physical location of themselves by other records

Pinned records (2/2)



- A File with pinned records



- to use an index file
 - to keep the sorted order of the records, while keeping the actual data file in its original order
 - need to take a close look at the use of indexes

Q&A

