

File Structures

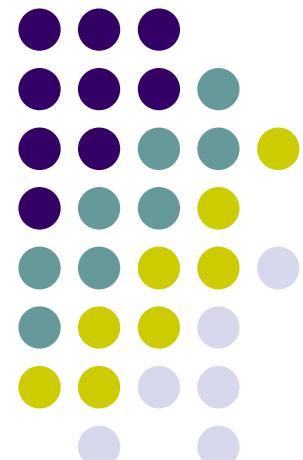
Ch09. A. Multilevel Indexing and B-Trees

2020. Spring

Instructor: Joonho Kwon

jhwon@pusan.ac.kr

Data Science Lab @ PNU





Outline

- 9.1 The invention of the B-tree
- 9.2 Statement of the Problem
- 9.3 Indexing with Binary Search Trees
- 9.4 Multilevel Index

The invention of the B-tree



- B-tree
 - de facto, **the standard organization for indexes** in a database system
 - to solve **how to access and efficiently maintain an index that is too large to hold in memory**
 - the **notion of a paged index** with retrieval time proportional to $\log_k I$ (I : index size, k : page size)
 - why the name B-tree?
 - B : balanced, broad, bushy, Bayer, Boeing

Statement of the Problem



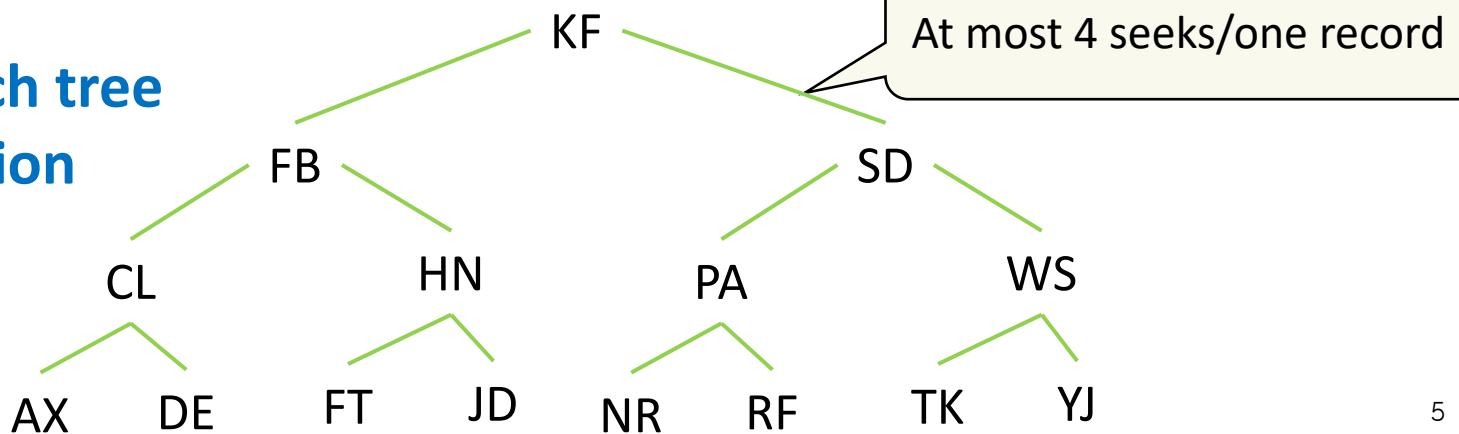
- Two problems !!!
 - searching the index must be faster than binary searching
 - Binary searching requires many seeks
 - searching for a key on a disk involves seeking to different disk tracks
 - insertion and deletion must be as fast as search
 - Inserting a key into an index involves moving a large # of other keys in the index
 - Why? Keep the index in sort order
 - → index maintenance is very impractical on disk
 - need to find a way to make insertions and deletions that have only local effects in the index

Indexing with Binary Search Trees (1/4)



- Binary search tree
 - BSTs keep their keys in sorted order
 - lookup and other operations can use the principle of [binary search](#)
 - Example: Sorted list of keys
 - AX, CL, DE, FB, FT, HN, JD, KF, NR, PA, RF, SD, TK, VI

Binary search tree representation



Indexing with Binary Search Trees (2/4)



- Internal Representation of Binary Tree
 - With RRN(fixed length record) or pointer

ROOT ➔ 9

		<i>key</i>	<i>left</i>	<i>right</i>
0	FB	10	8	
1	JD			
2	RF			
3	SD	6	13	
4	AX			
5	YJ			
6	PA	11	2	
7	FT			

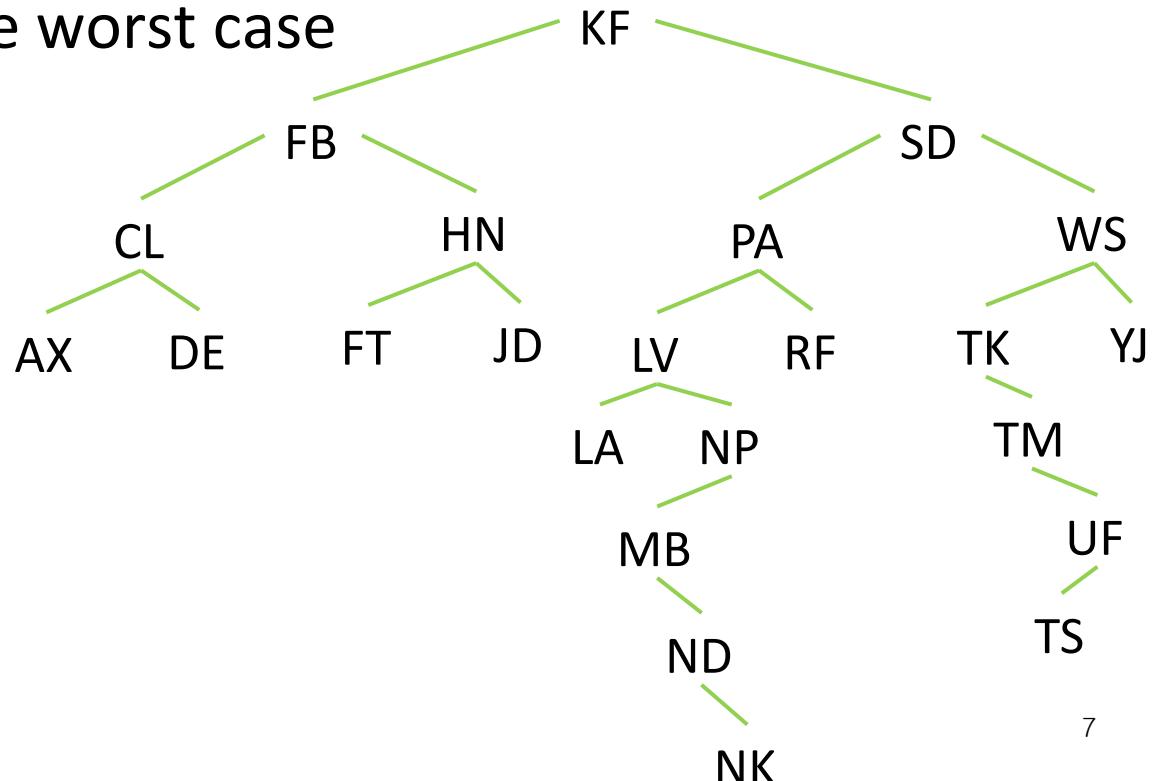
		<i>key</i>	<i>left</i>	<i>right</i>
8	HN	7	1	
9	KF	0	3	
10	CL	4	12	
11	NR			
12	DE			
13	WS	14	5	
14	TK			

Indexing with Binary Search Trees (3/4)



- Unbalanced Binary Tree

- require 7, 8, or 9 seeks for retrieval
- Binary search on a sorted list of 24 keys
- => 5 seeks in the worst case



Indexing with Binary Search Trees (4/4)



- Problems

- A BST is not fast enough for disk resident indexing
- lack of an effective strategy of balancing the tree

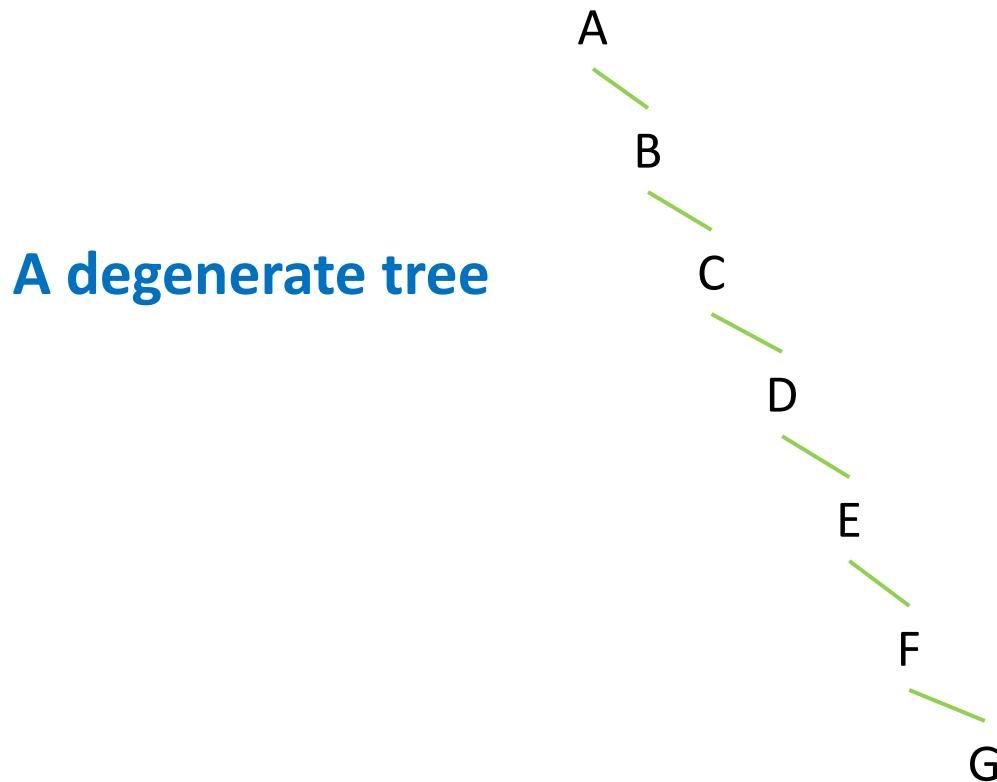
- Alternatives

- AVL tree
- Paged Binary Tree

AVL Trees (1/5)



- undesirable tree organizations
 - alphabetical order produces a degenerate tree
→ need to reorganize the nodes of the tree

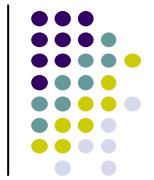


AVL Trees (2/5)

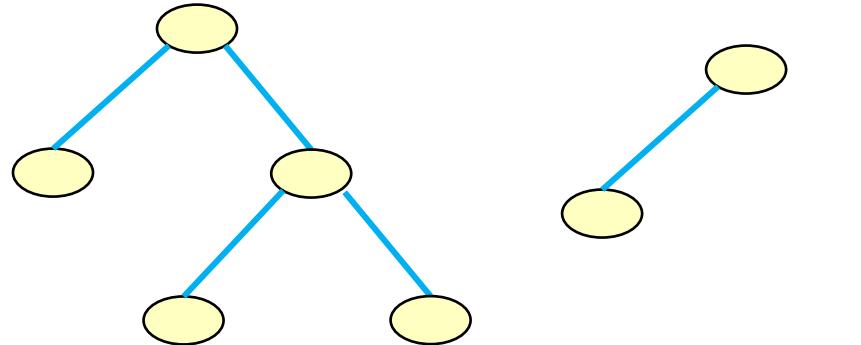


- AVL trees
 - defined by Russian mathematicians, G. M. Adel'son-Vel'skii and E.M. Landis
 - a near optimal tree structure
 - height-balanced 1-tree (HB(1) tree)
 - height-balanced trees with the maximum difference is 1
 - no two subtrees of any root differ by more than one level
 - not directly applicable to file structure problems
 - AVL trees have too many levels
 - guarantees that search performance approximates that of a completely balanced tree

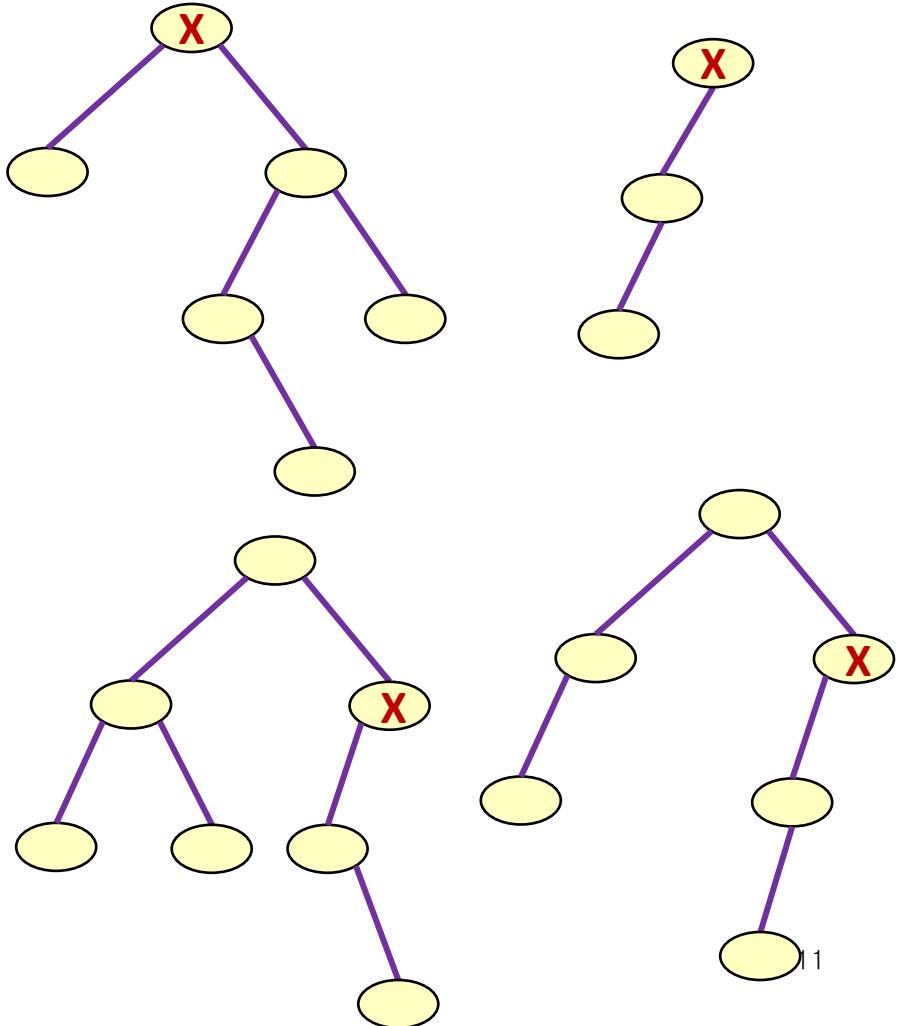
AVL Trees (3/5)



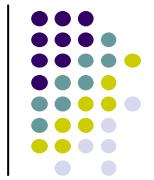
- AVL Trees



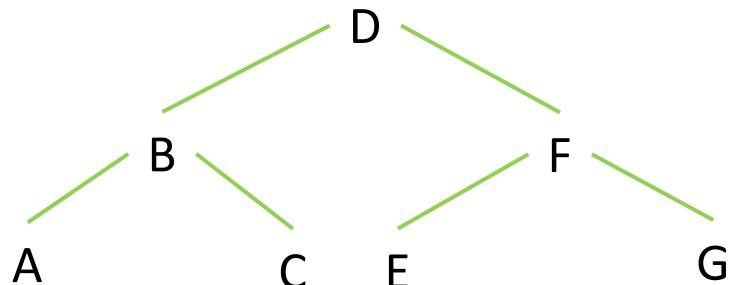
- Non AVL Trees



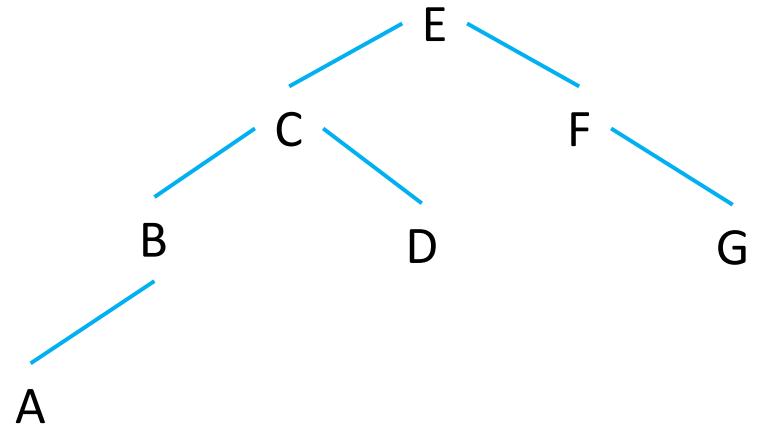
AVL Trees (4/5)



- Completed balanced BST from B,C,G,E,F,D,A



- AVL tree from B,C,G,E,F,D,A



AVL Trees (5/5)



- Complexity comparison (e.g., 1,000,000 keys)
 - (1) completely balanced tree
 - the worst-case search to find a key: $\log_2(N + 1)$
 - require seeking 20 levels
 - (2) AVL trees
 - the worst-case search to find a key: $1.44 \log_2(N + 2)$
 - require seeking 29 levels
- Again, two problems
 - **Binary searching requires too many seeks**
 - keeping an index in sorted order is expensive
 - binary trees provide an acceptable solution

Paged Binary Tree (1/3)

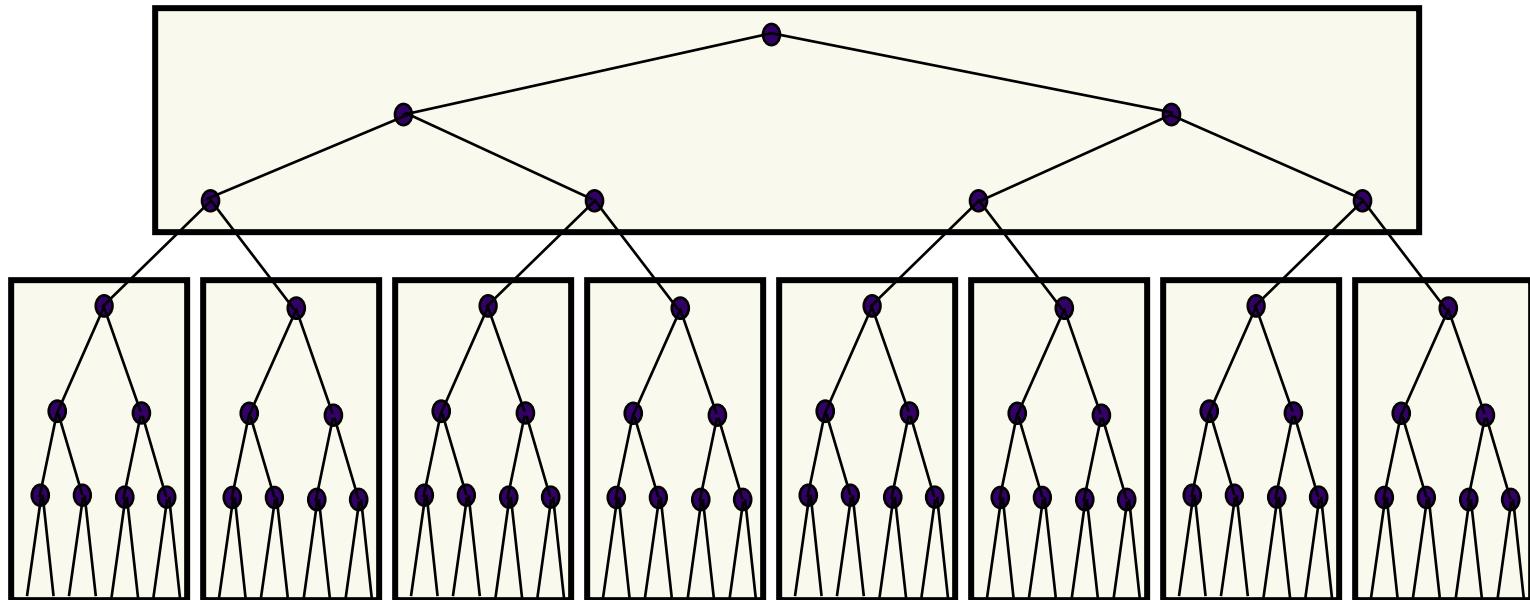


- Page
 - A unit of disk I/O for handling seek and transfer of disk data
 - Typically, 4k, 8k, 16k ...
- Paged Binary Tree
 - Divide a binary tree into pages and then store each page in a block of contiguous locations on disk.

Paged Binary Tree (2/3)



- An example of Paged binary tree
 - If every page holds 7 keys, 511 nodes(keys) in only three seeks



Paged Binary Tree (3/3)



- A typical example
 - page size : 8 KB for $k=511$ key/reference field pairs
 - N (# of keys) = 134,217,727
- ① Completely full, balanced binary tree
 - $\log_2(N + 1) = 27$ seeks
- ② Paged version of a completely full, balanced binary tree
 - $\log_{k+1}(N + 1) = 3$ seeks (k : # of keys held in a single page)
 - => can reduce disk seeks

Problems with Paged Trees (1/3)

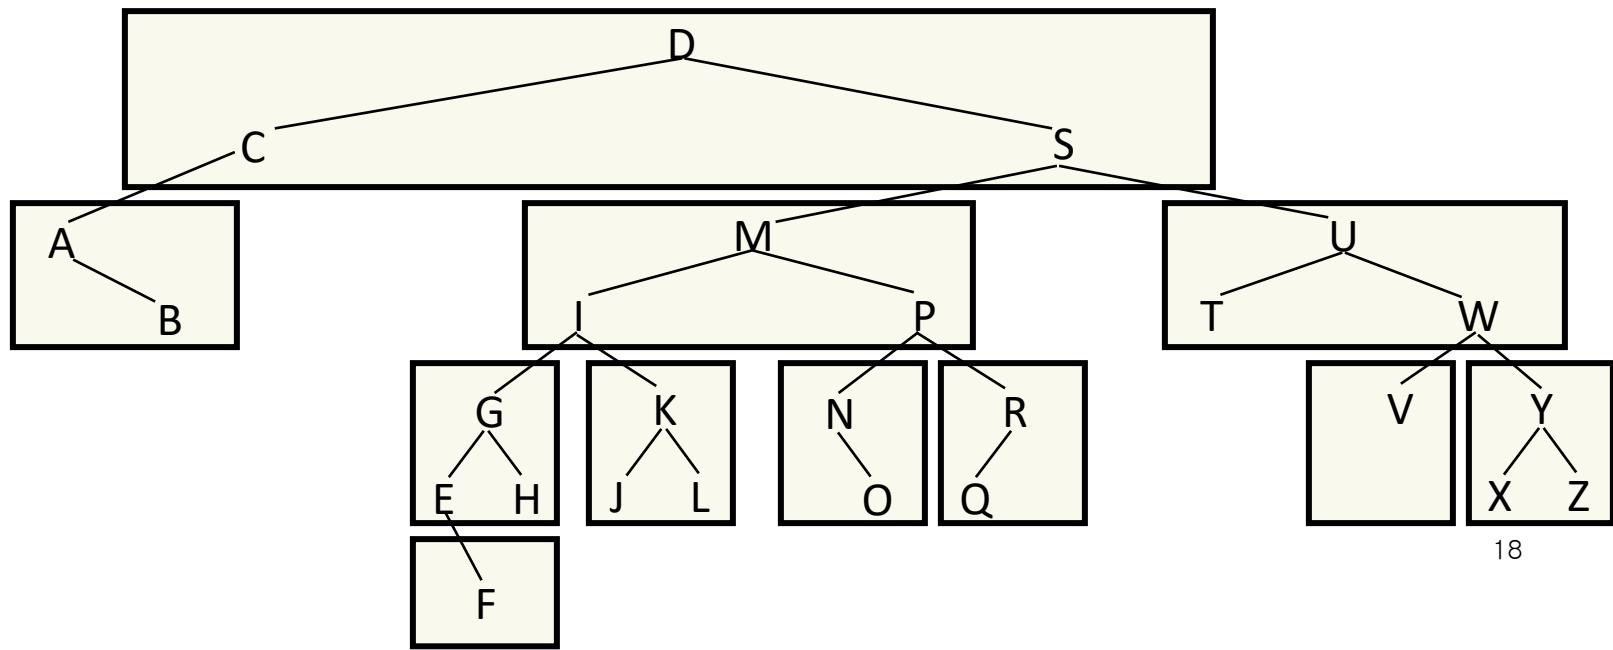


- Only valid when we have the entire set of keys in hand before the tree is built
- Problems due to out of balance
 - How to select a good separator
 - How to group keys
 - How to guarantee the maximum loading
- B-tree provides a solution for above problems!

Problems with Paged Trees (2/3)



- Construction of the paged binary trees
 - from the following sequence of keys:
 - C S D T A M P I B W N G U R K E H O L J Y Q Z F X V
 - contain a maximum of three keys per pages
 - rotate keys within a page to keep each page as balanced as possible



Problems with Paged Trees (3/3)



- Three unsolved questions
 - (1) the keys in the root turn out to be good separator keys ?
 - divide up the set of other keys more or less evenly
 - (2) how do we avoid grouping keys that should not share a page ?
 - C, D, S should not share the same page
 - (3) each of the pages contains at least some minimum number of keys?
 - maintain a lower bound
 - => **B-trees** : build trees **upward from bottom** instead of downward from the top

Multilevel indexing (1/5)

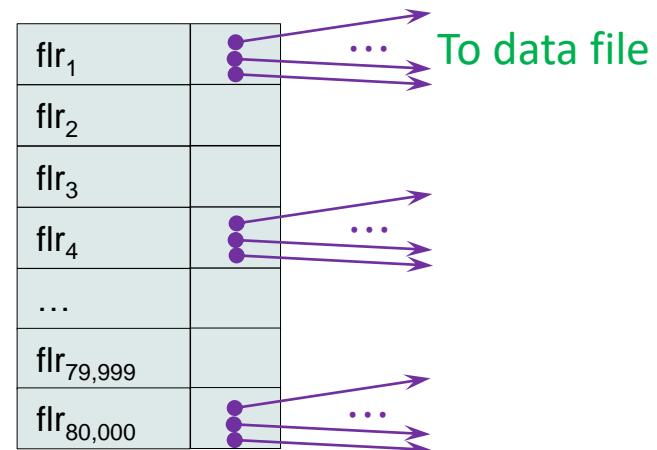


- Extension of the single record indexing
 - Approach as simple index record
 - limited on the number of keys allowed
 - Approach as multirecord index
 - consists of a sequence of simple index records
 - binary search is too expensive
 - Approach as multilevel index
 - reduced the number of records to be searched
 - speed up the search



Multilevel indexing (2/5)

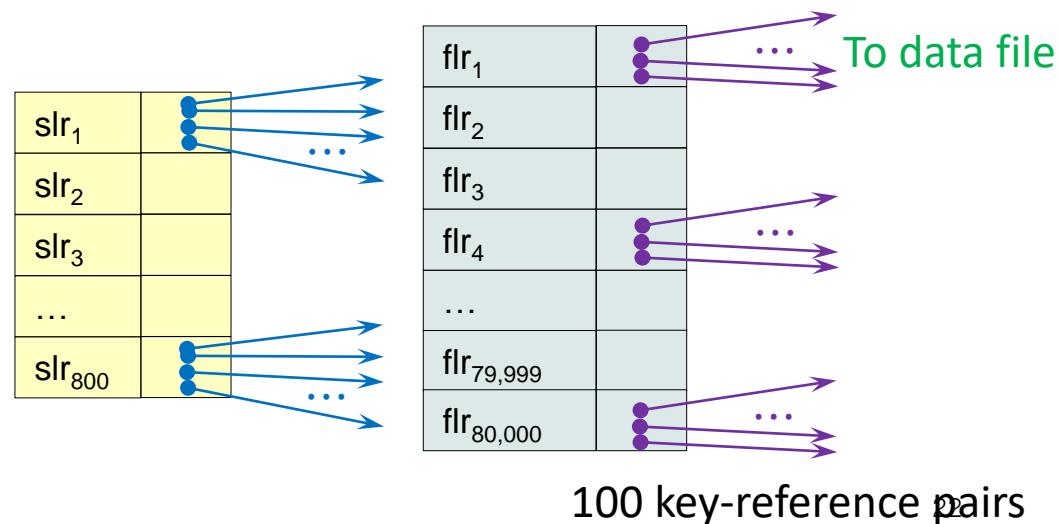
- 800 MB file of 8,000,000 records
 - 100 bytes each, with 10 byte keys
 - Index has 8,000,000 key-reference pairs
 - Each index record has 100 key-reference pairs
- (1) first level index
 - 80,000 ($=8,000,000/100$) records for 8,000,000 keys
 - Index to the data file



Multilevel indexing (3/5)



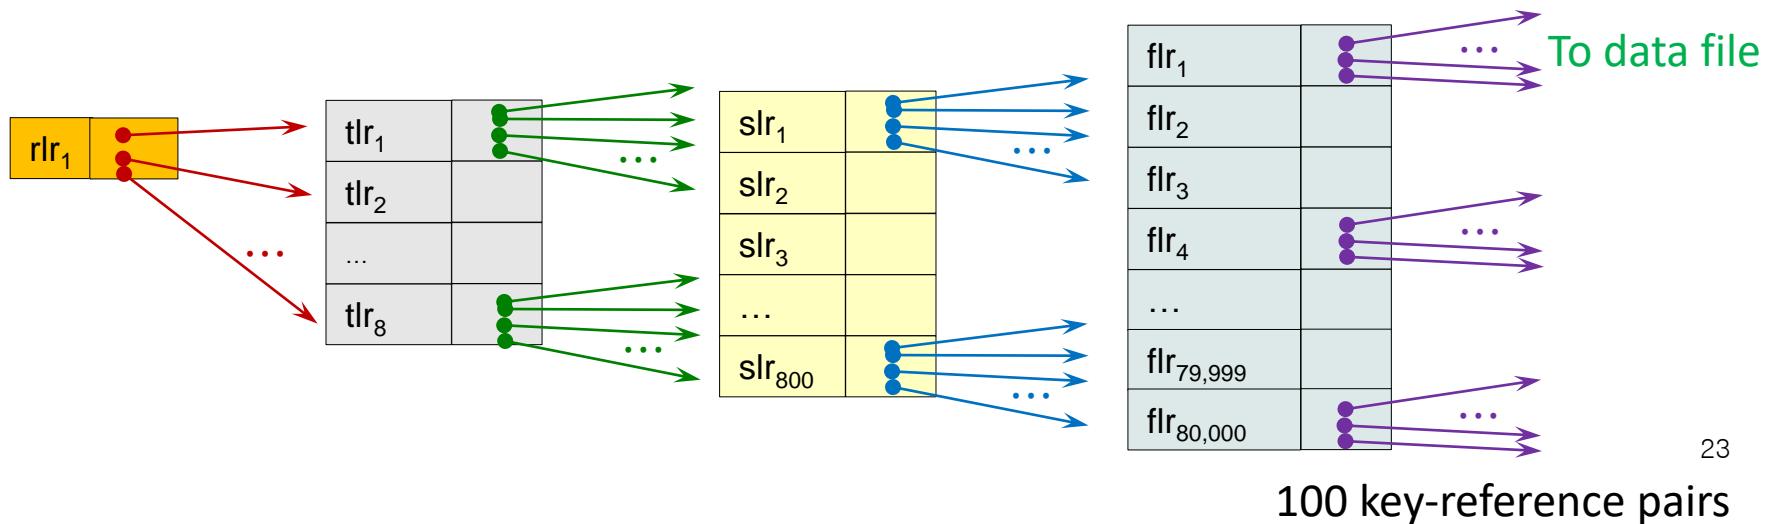
- (2) second level index
 - 800 index records for 80,000 keys
 - choose one of the keys (the largest) in each index record
 - use reference fields for index record addresses



Multilevel indexing (4/5)



- (3) third level index / (4) fourth level index
 - 8 index records for 800 keys / 1 index records for 8 keys
- Total
 - ① total 80,809 index records
 - ② average, min., max. # of disk access = 4 (\because 4 levels)



Multilevel indexing (5/5)



- How can we insert new keys into the multilevel index?
 - The index records in some level might be full
 - The several levels of indexes might be rebuilt
 - Overflow chain may be helpful, but still ugly
- Multi-level index structure is not strong in dynamic data processing applications
- B-tree will give you the right solution!



Outline

- 9.5 B-Trees
- 9.10 Formal Definition of B-Tree Properties
- 9.6 Example of Creating a B-Tree

B-Trees: Working up from the bottom



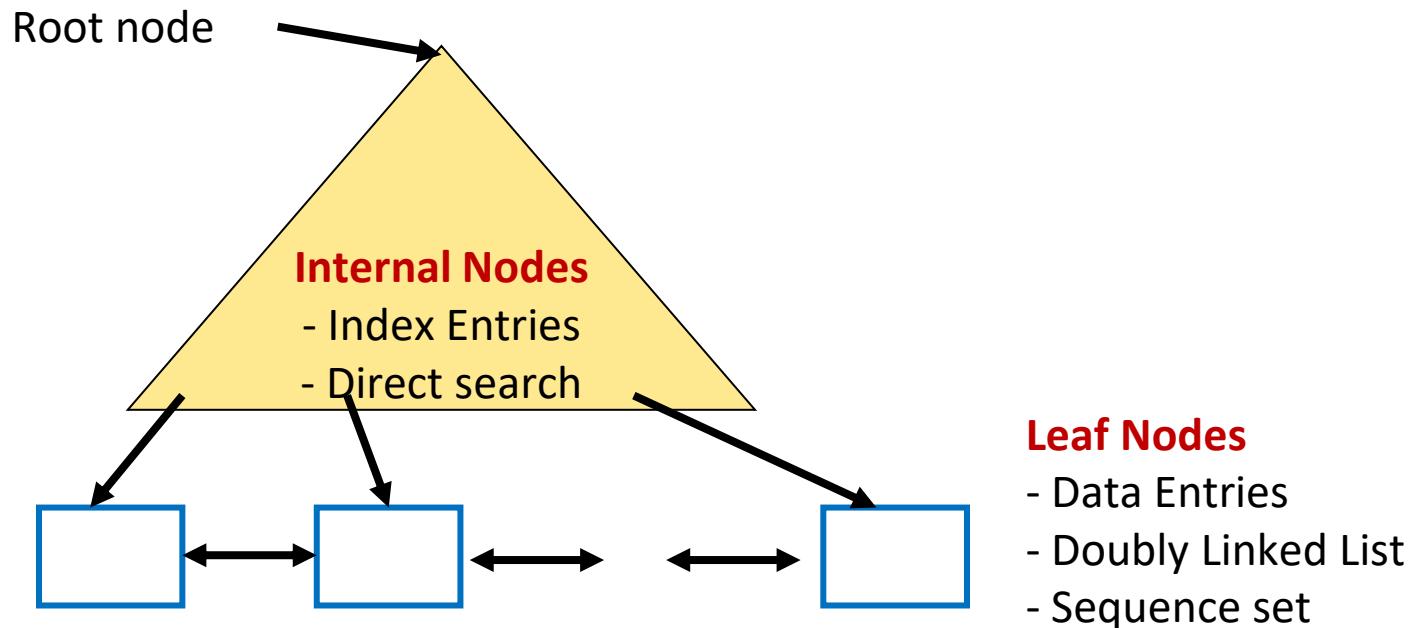
- B-trees
 - multilevel indexes defined by Bayer and McCreight
 - built **upward from the bottom**(i.e., leaves) instead of downward from the top
 - do not require that the index records be full
 - do not shift the overflow keys to the next record
 - **split an overfull record into two records**, each half full
 - deletion takes a similar strategy of **merging** two records into a single record when necessary

B-Trees

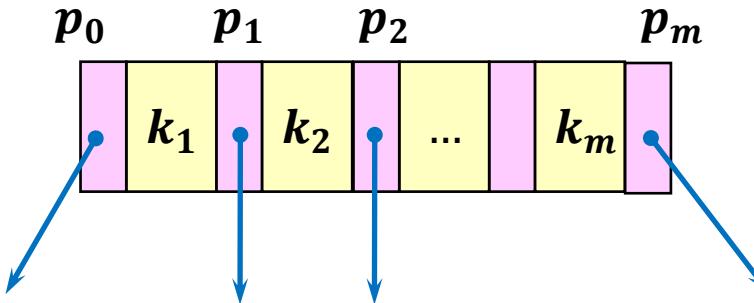


- Measures

- **Order:** the (maximum) number of indexing field values at each node
- Height: the number of indexing levels from root to any leaf

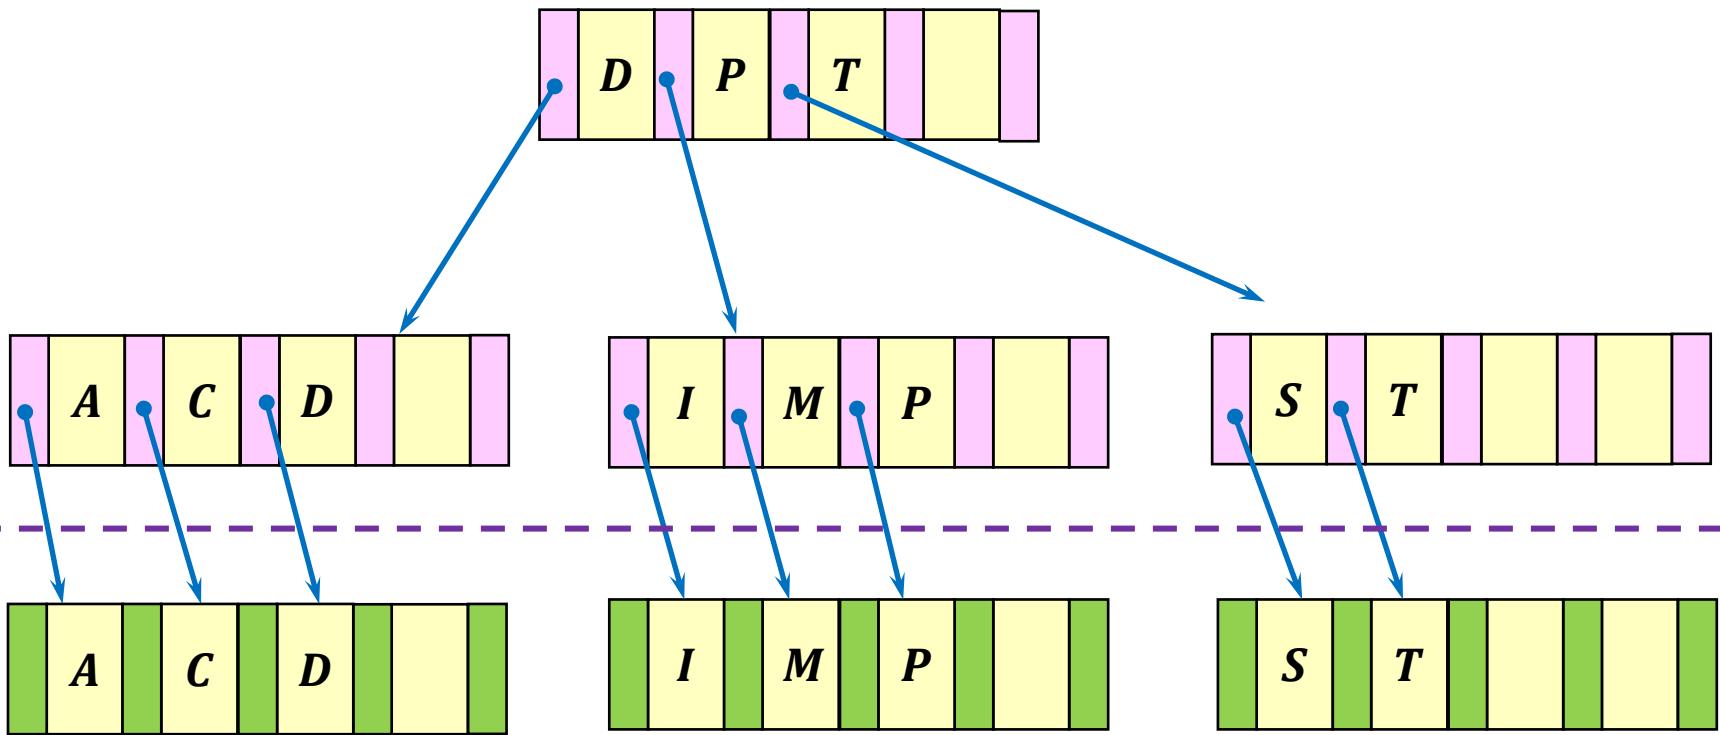


B-Tree Properties



- The properties of a B-tree of order n
 - 1. Every page has a **maximum of n** descendants
 - 2. Every page, except for the root and the leaves, has **at least ceiling of $(n/2)$** descendants
 - $n/2 \leq m \leq n$ (**half-full**), $n =$ order of tree (**fanout**)
 - 3. The root has **at least two** descendants (unless it is a leaf)
 - 4. All the leaves appear **on the same level**
 - 5. The leaf level forms a complete, ordered index of the associated data file

Sample B-Tree



Record Insertion(1/2)



- When a record is inserted in the data file, the B+-tree must be changed accordingly:
 - simple case
 - leaf not full: just insert (key, pointer-to-record)
 - leaf overflow
 - Internal node overflow
 - new root

Record Insertion(2/2)



- Algorithm: (from DB book, slightly different)
 - Find correct leaf L .
 - Put data entry onto L .
 - If L has enough space, *done!*
 - Else, must split L (*into L and a new node $L2$*)
 - Redistribute entries evenly, copy up middle key.
 - parent node may overflow
 - but then: push up middle key. Splits “grow” tree; root split increases height.

Splitting & Promoting (1/3)

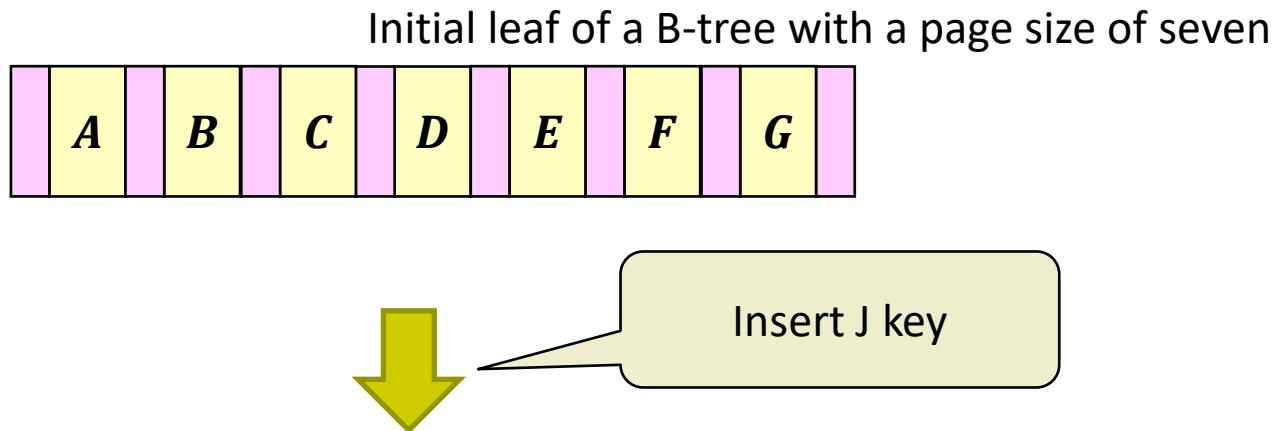


- Splitting
 - Creation of two nodes out of one because the original node becomes overfull
 - Result in the need to promote a key to a higher-level node to provide an index separating the two new nodes
- Promotion of a key
 - Movement of a largest key from one node into a higher-level node when split occurs

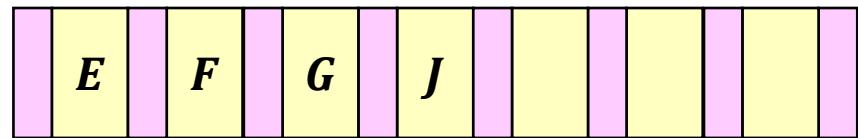
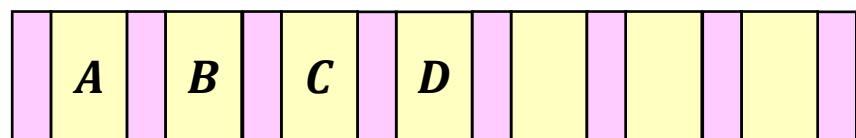
Splitting & Promoting (2/3)



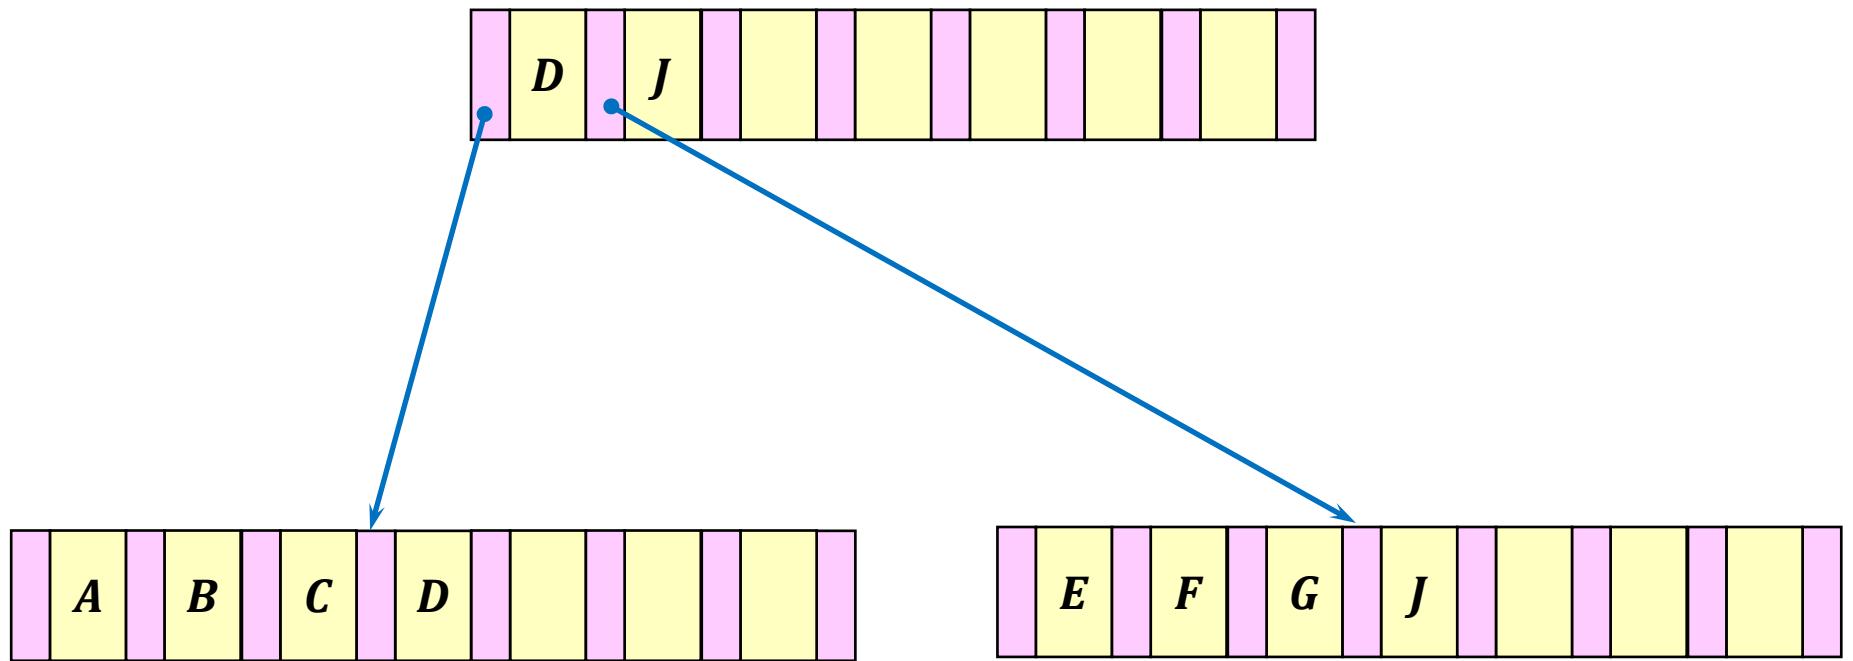
- Assume: order = 7



Splitting the leaf to accommodate the new J key



Splitting & Promoting (3/3)



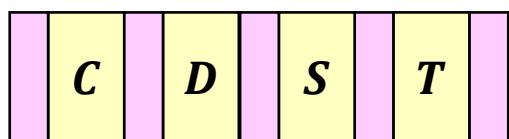
Promotion of the D, J keys into a root node (copy up)

Creating a B-tree (1/10)

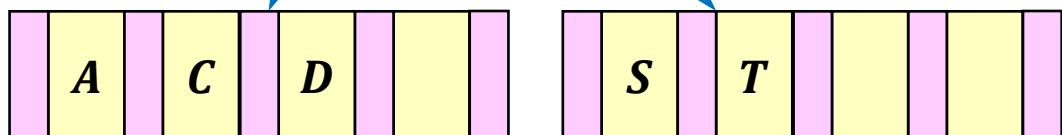
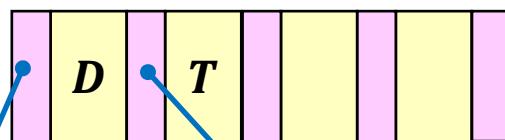


- Assume: Order of a B-tree = 4
- Input Sequence
 - C S D T A M P I B W N G U R K E H O L J Y Q Z F X V

Insertion of C, S, D, T
into the initial page



Insertion of A
causes node to **split**

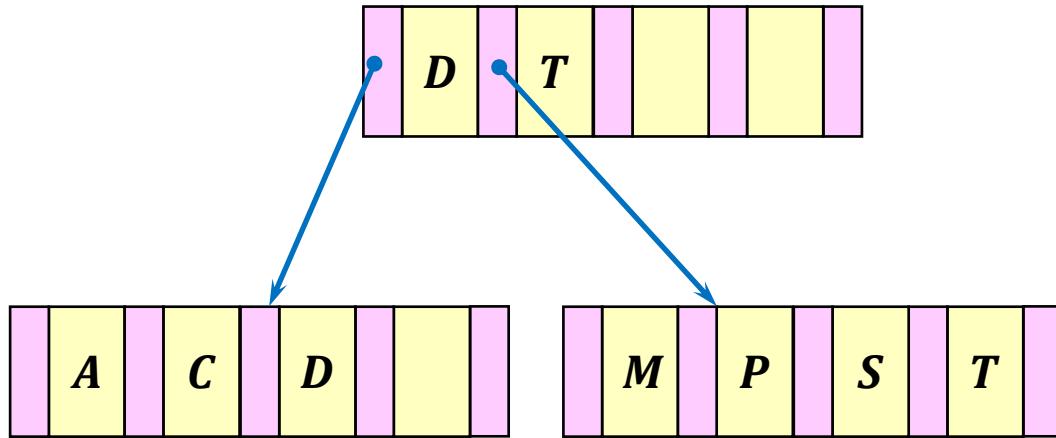


the largest key in each leaf node (D and T) to be placed in the root node

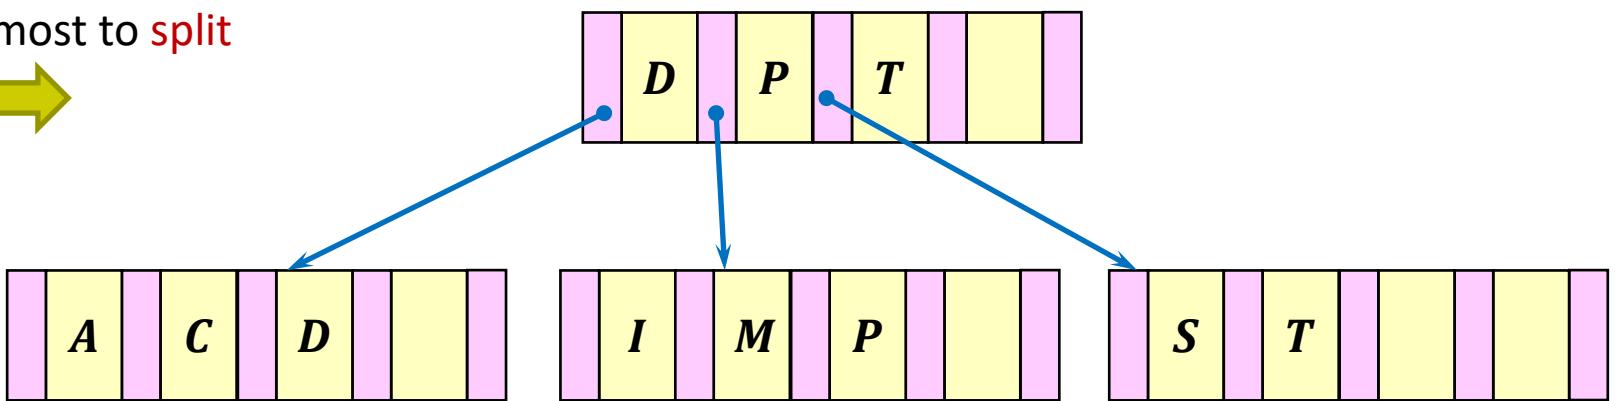
Creating a B-tree (2/10)



Insertion of M and P.
→ Put into the right
most leaf node



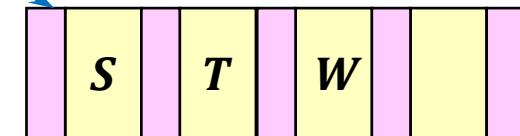
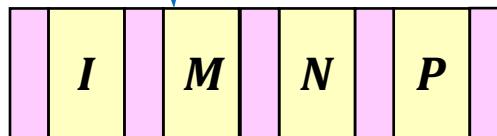
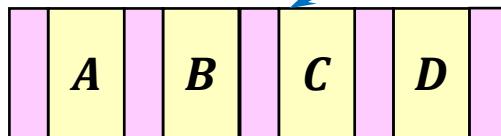
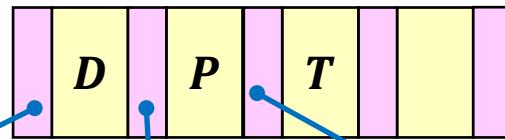
Insertion of I causes
the right most to **split**



Creating a B-tree (3/10)



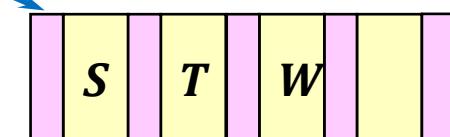
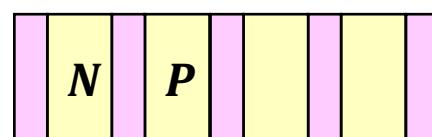
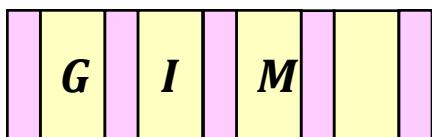
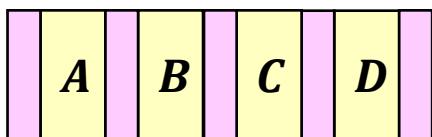
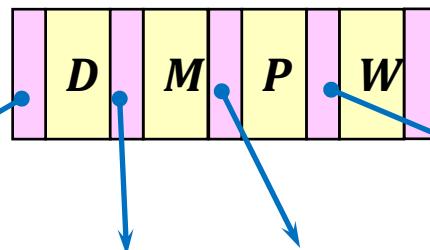
Insertion of B, W and N



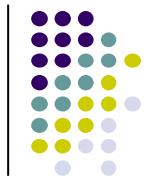
Insertion of G causes
another **split**



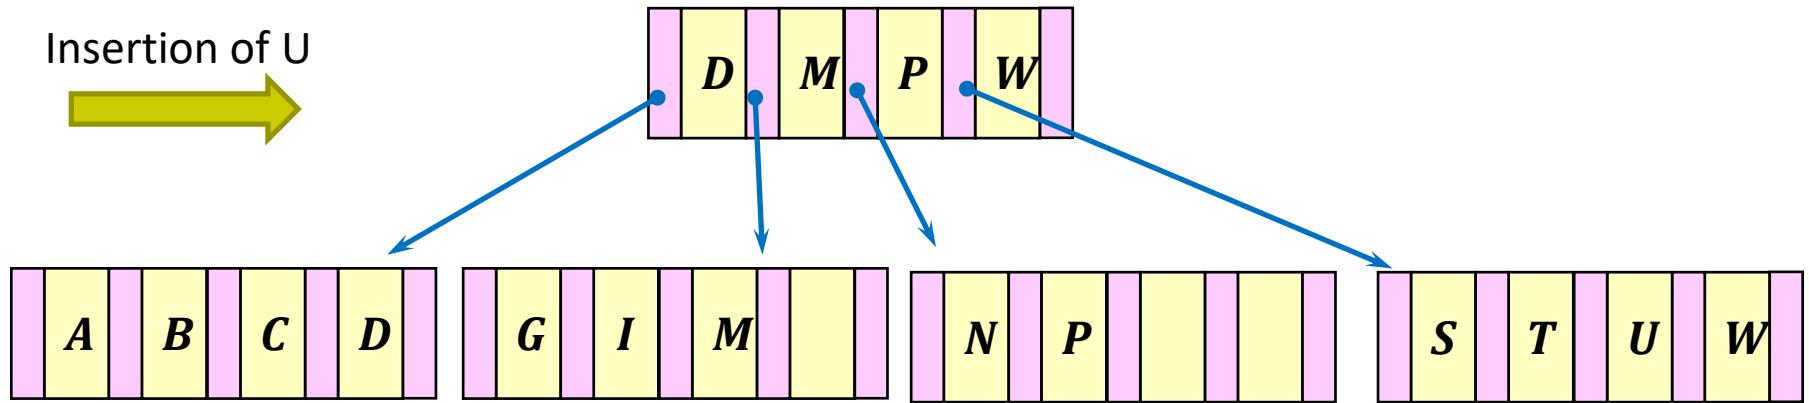
Root node is full



Creating a B-tree (4/10)



Insertion of U

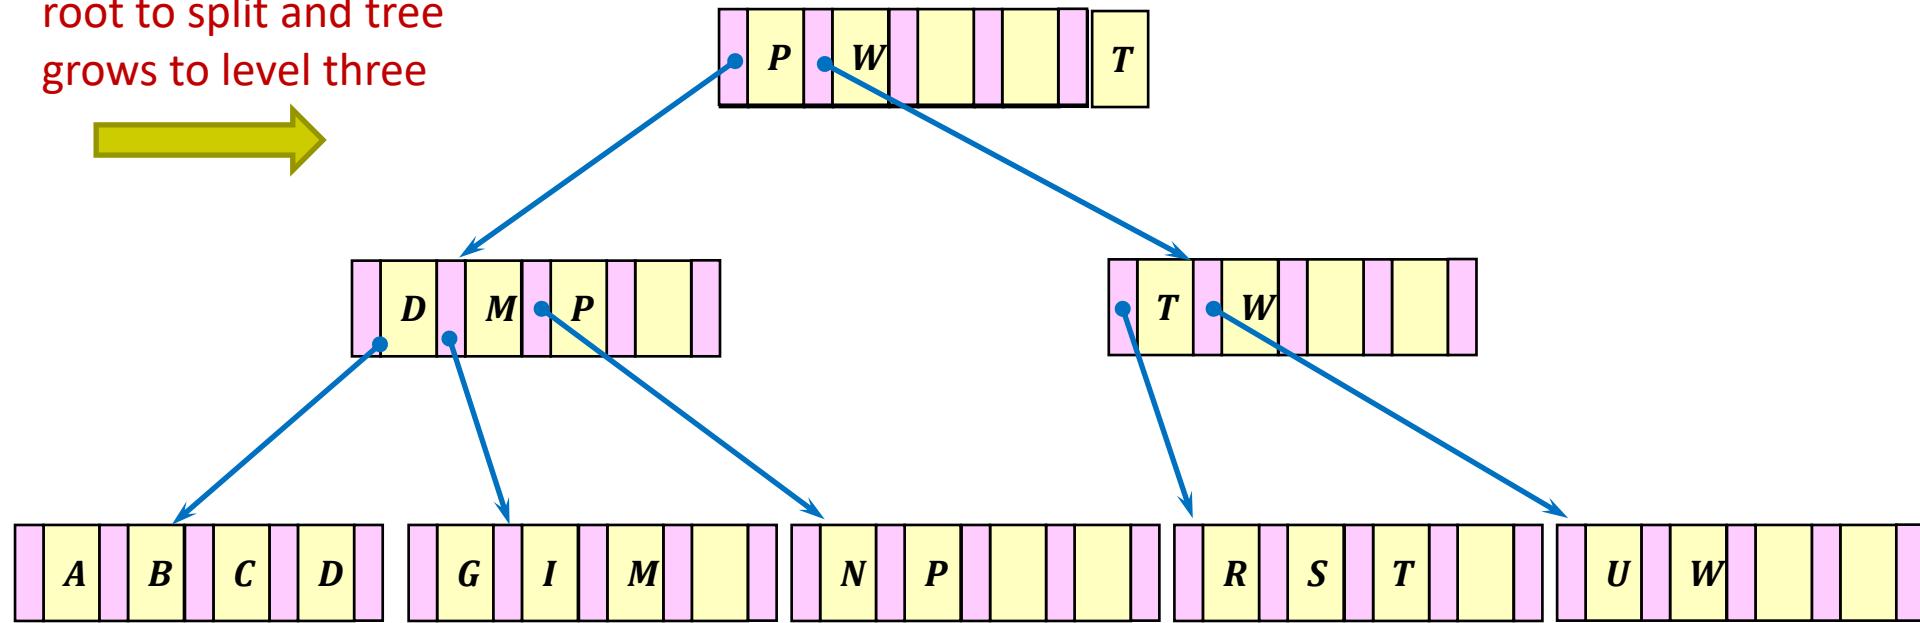


Creating a B-tree (5/10)

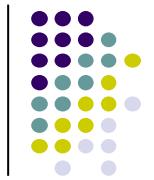


Insertion of R into the rightmost leaf to split,

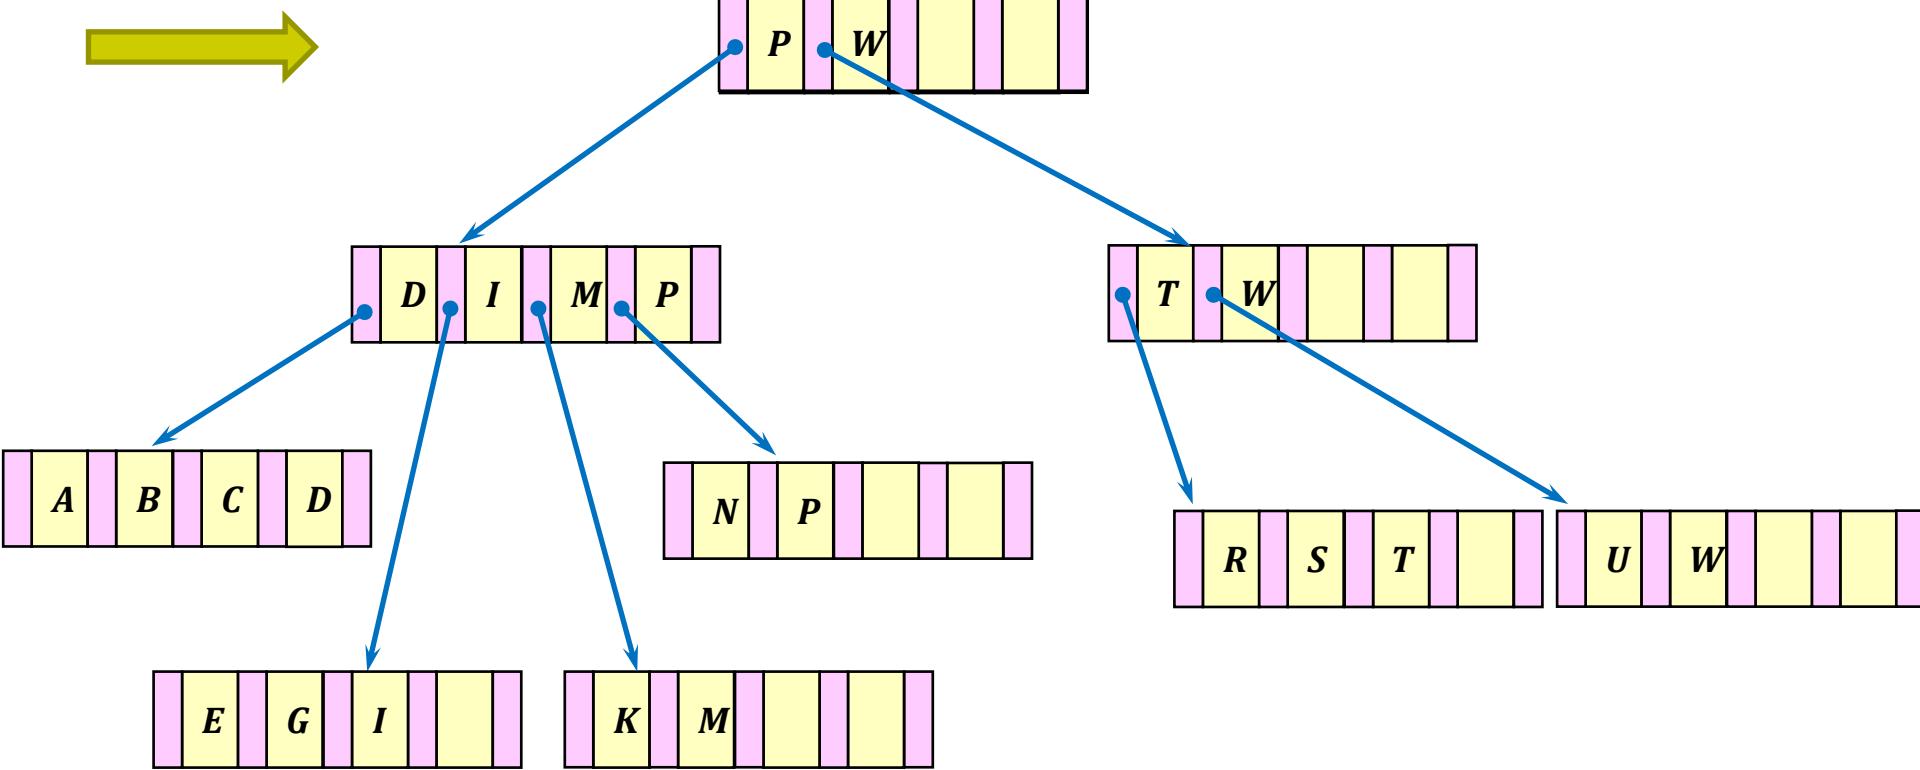
Insertion into the root to split and tree grows to level three



Creating a B-tree (6/10)



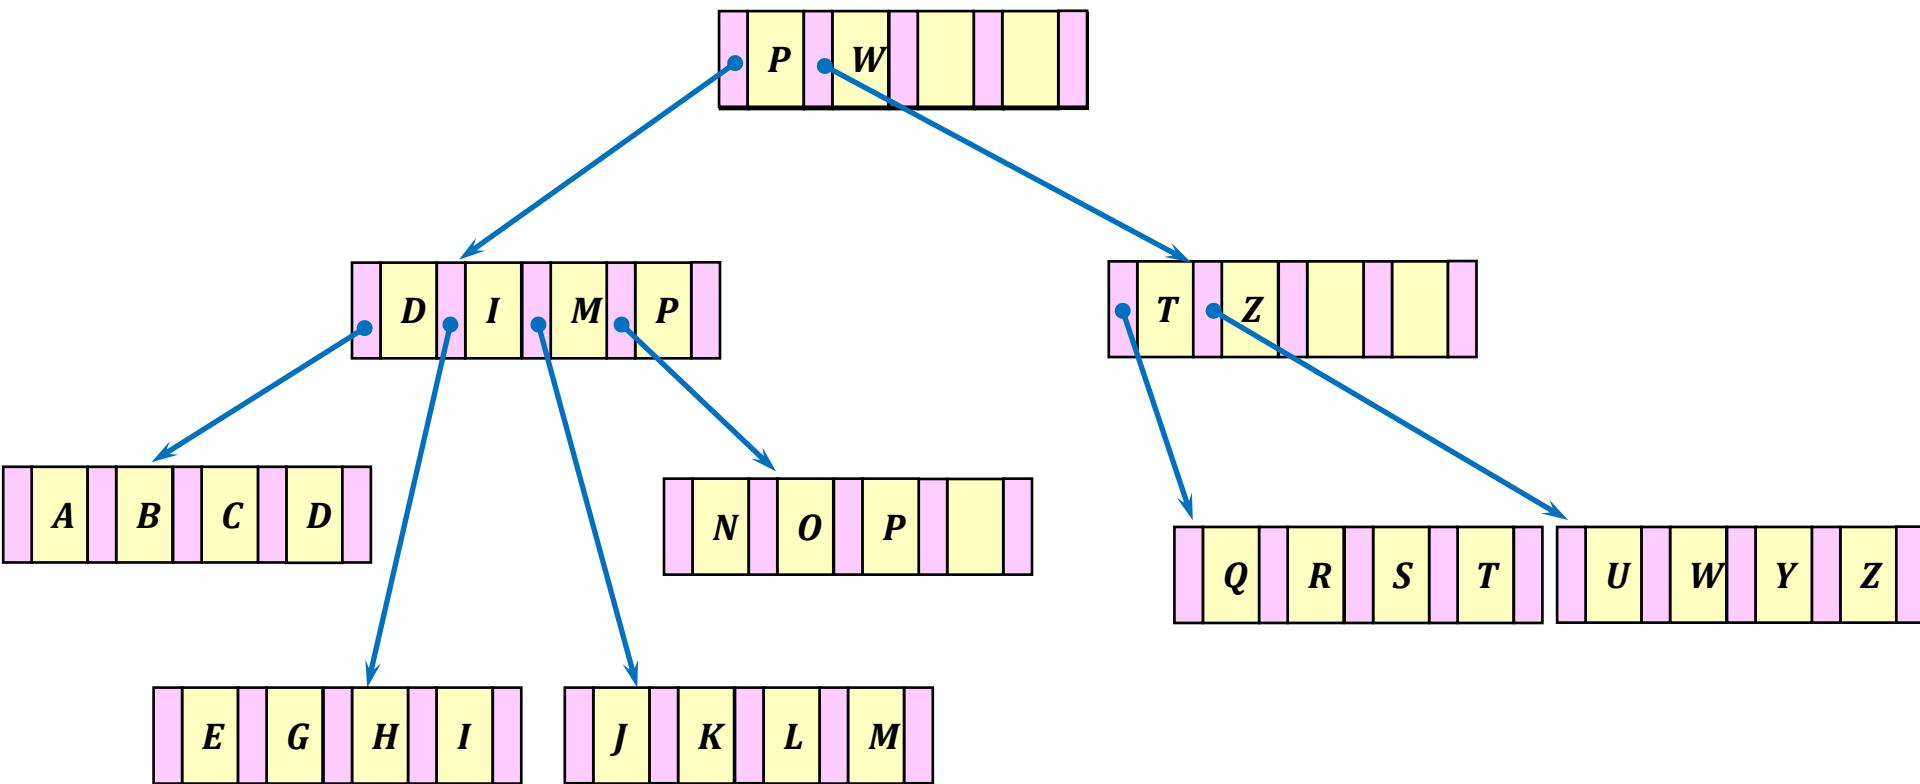
Insertion of K, E →
the second node to
split → I is promoted



Creating a B-tree (7/10)



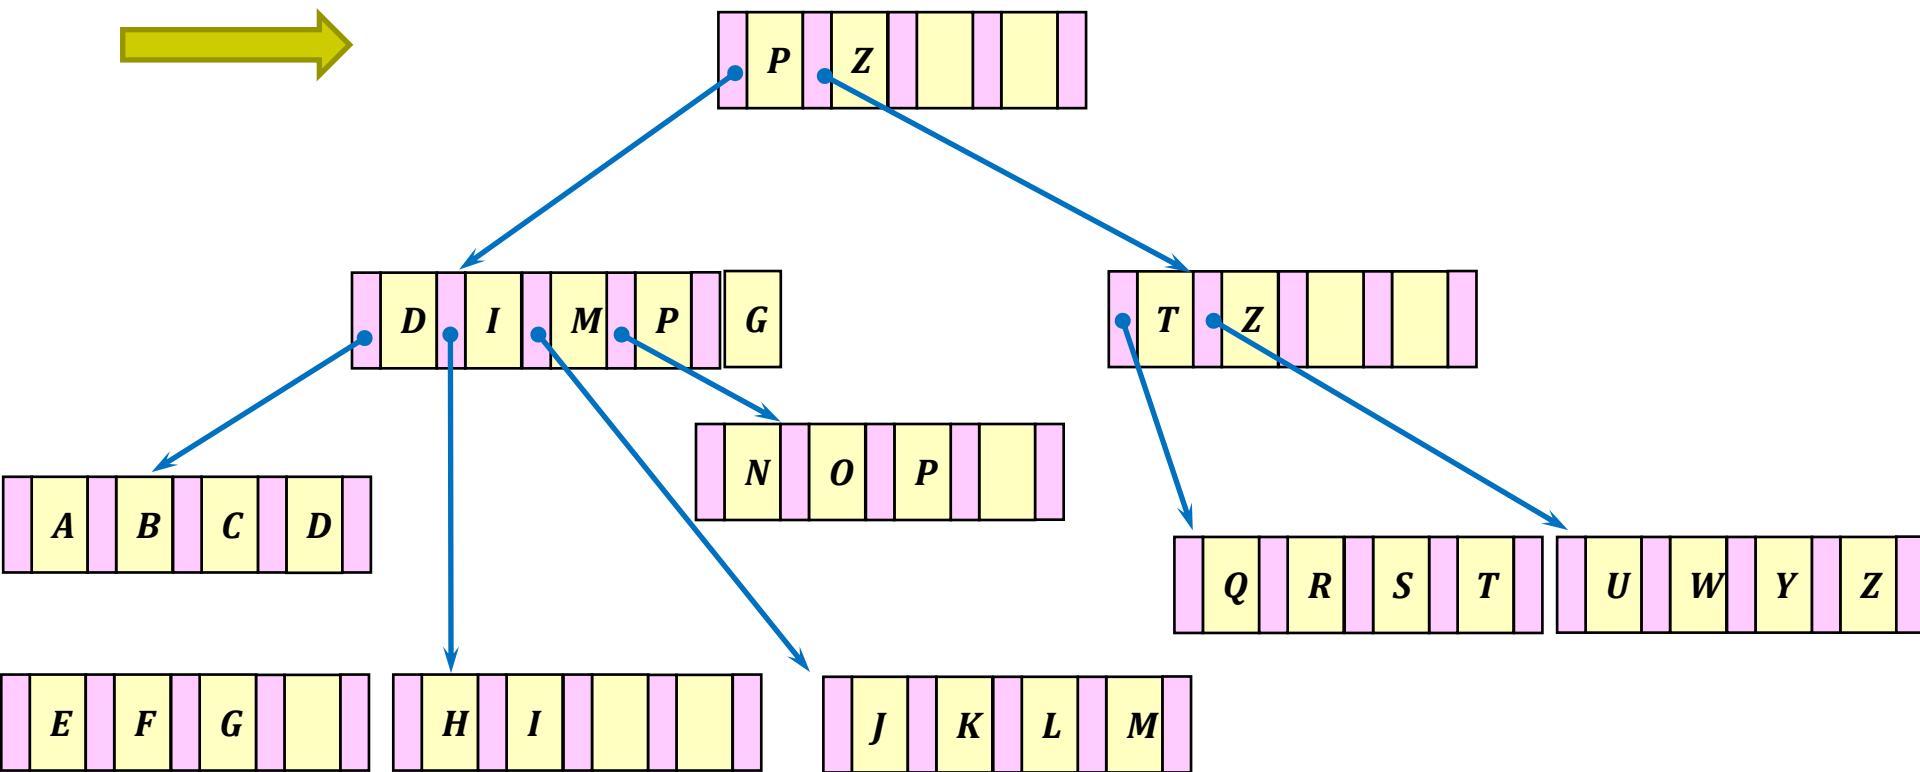
Insertion of H, O, L, J, Y, Q and Z



Creating a B-tree (8/10)



Insertion of F → split
Of the second leaf node
→ G is promoted



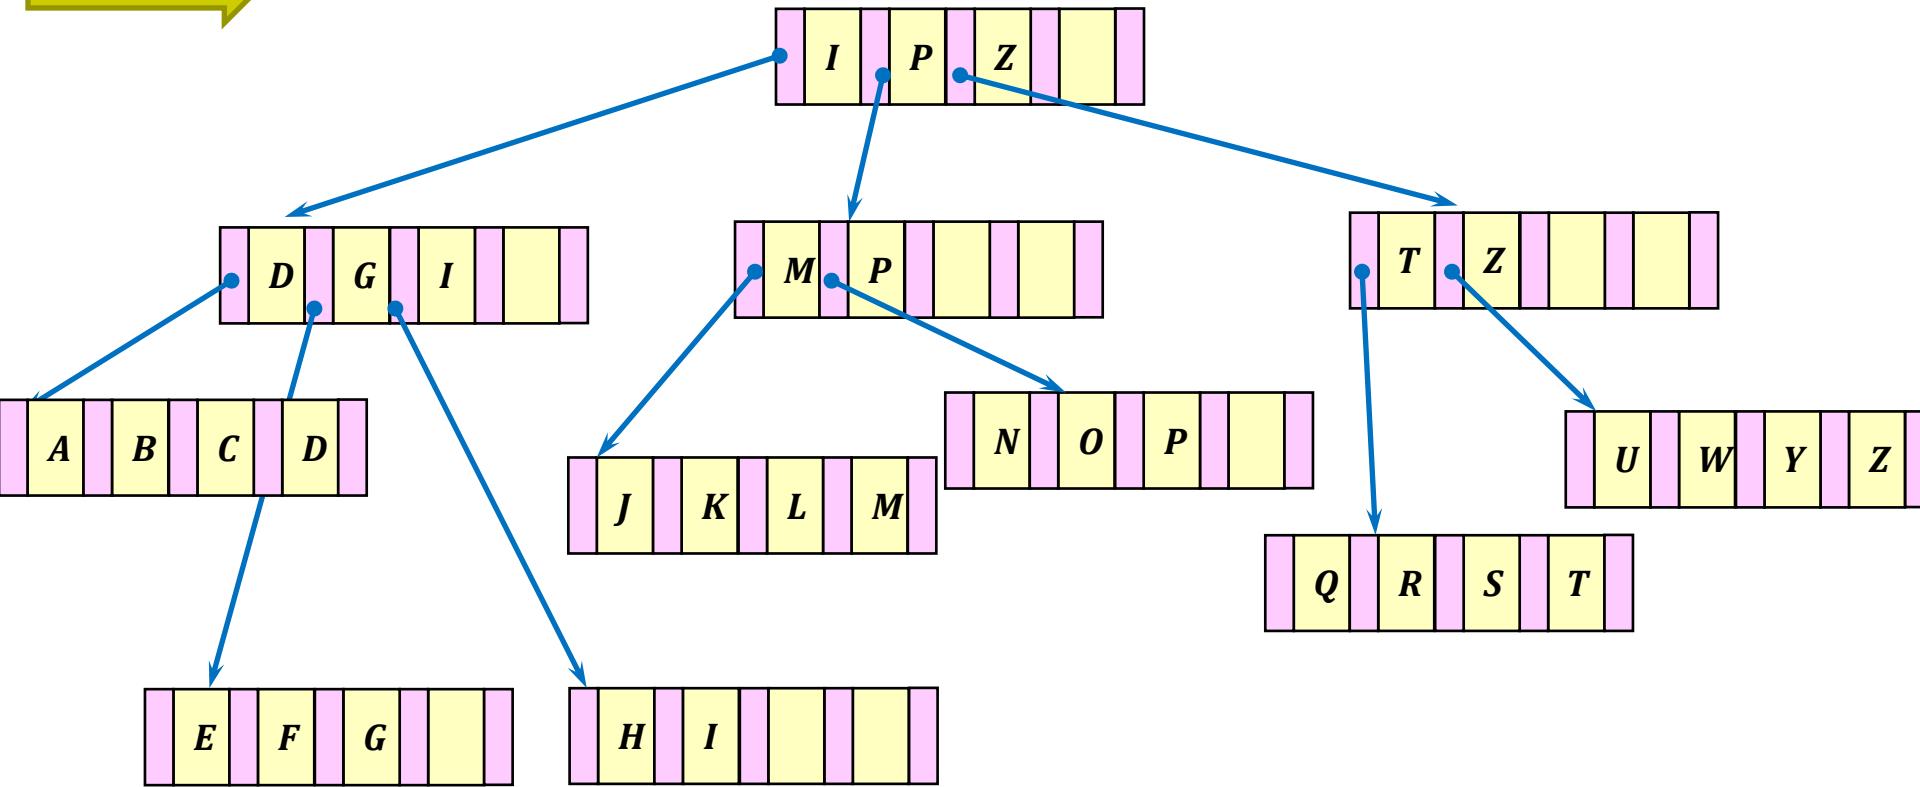
Creating a B-tree (9/10)



Insertion of F → split the second leaf node

→ G is promoted

→ Cascade split and I is promoted

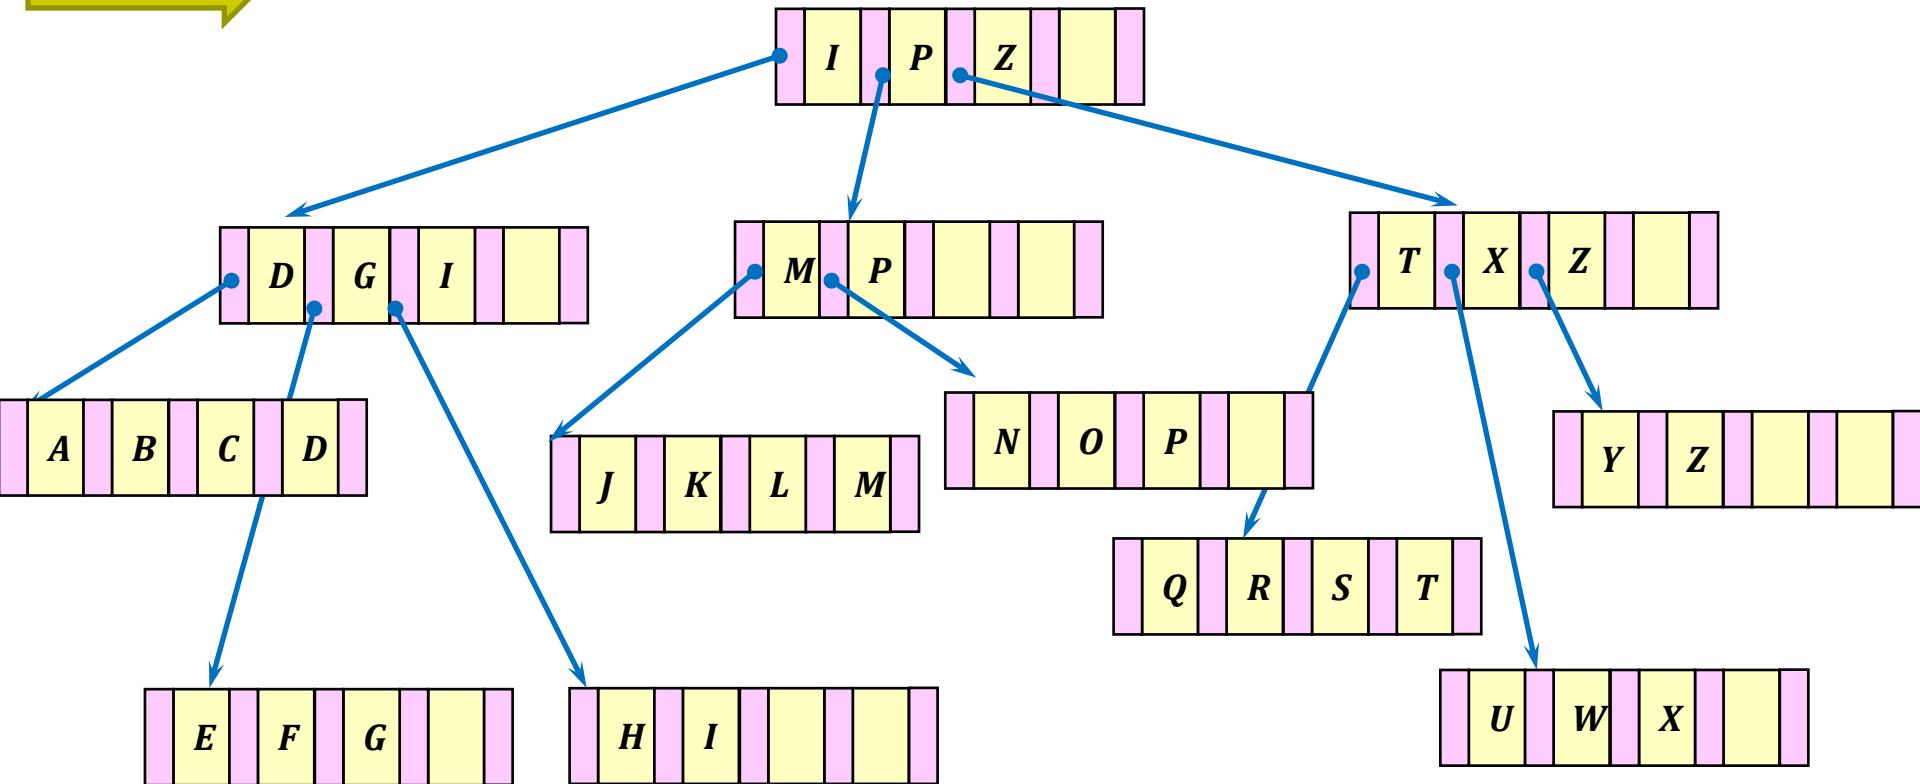


Creating a B-tree (10/10)



Insertion of X → split the rightmost node

→ X is promoted



Recap: Insertion in B-trees



- Major components of insertion
 - Split the node
 - Promote the middle key
 - Increase the height of the B-tree
- Insertion may touch no more than 2 nodes per level
- Insertion cost is strictly linear in the height of the tree

Q&A

