

# File Structures

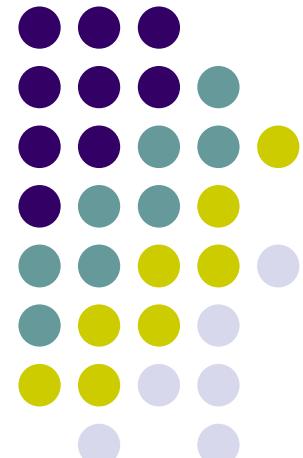
## Ch09. B. B-Tree Search, Insert

2020. Spring

Instructor: Joonho Kwon

[jhwon@pusan.ac.kr](mailto:jhwon@pusan.ac.kr)

Data Science Lab @ PNU



# Outline




- 9.7 An Object-Oriented Representation of B-Trees
- 9.8 B-Tree Methods Search, Insert, and Others
- 9.10. Formal Definition of B-tree properties
- 9.11 Worst-case search depth
- 9.9 B-Tree Nomenclature

# An OO representation of B-Trees



- OO representation
  - B-tree is an index file associated with a data file
  - Most of operations on B-trees, insertion and deletion, are applied to the B-tree nodes in memory
  - The template BTreenode class based on the SimpleIndex template class



# Recap: SimpleIndex



## ● SimpleIndex Class

```
template <class keyType>
class SimpleIndex
{
public:
    SimpleIndex (int maxKeys = 100, int unique = 1);
    ~SimpleIndex ();
    void Clear () // remove all keys from index
    int Insert (const keyType key, int recAddr);
    int Remove (const keyType key, const int recAddr = -1);
    int Search (const keyType key, const int recAddr = -1,
                const int exact = 1) const;
    ...
protected:
    int MaxKeys;
    int NumKeys;
    keyType * Keys;
    int * RecAddrs;
    int Find (const keyType key, const int recAddr = -1,
              const int exact = 1) const;
    ...
};
```

# Recap: SimpleIndex



- **Insert()**

```
template <class keyType>
int SimpleIndex<keyType>::Insert (const keyType key, int recAddr)
{
    int i;
    int index = Find (key);
    if (Unique && index >= 0) return 0; // key already in
    if (NumKeys == MaxKeys) return 0; //no room for another key
    for (i = NumKeys-1; i >= 0; i--)
    {
        if (key > Keys[i]) break; // insert into location i+1
        Keys[i+1] = Keys[i];
        RecAddrs[i+1] = RecAddrs[i];
    }
    Keys[i+1] = key;
    RecAddrs[i+1] = recAddr;
    NumKeys++;
    return 1;
}
```

# Class BTreenode (1/2)



- **Representing B-Tree Nodes in memory**
- **Public methods:**
  - insert : simply calls SimpleIndex::Insert and then check for overflow
  - remove a key, split and merge nodes
  - Search
    - inherited from SimpleIndex class(works perfectly well)
  - pack/unpack
    - manage the difference between the memory and the disk representation of BTreenode objects
- **Protected member**
  - store the file address of the node and the minimum and maximum number of keys

# Class BTreeNode (2/2)



```
// this is the in-memory version of the BTreeNode
template <class keyType>
class BTreeNode: public SimpleIndex <keyType>
{
public:
    BTreeNode(int maxKeys, int unique = 1);
    int Insert (const keyType key, int recAddr);
    int Remove (const keyType key, int recAddr = -1);
    int LargestKey (); // returns value of largest key
    int Split (BTreeNode<keyType> * newNode); // move keys into newNode
    int Merge (BTreeNode<keyType> * fromNode); // move keys from fromNode
    int UpdateKey (keyType oldKey, keyType newKey, int recAddr = -1);
    int Pack (IOBuffer& buffer) const;
    int Unpack (IOBuffer& buffer);
    static int InitBuffer (FixedFieldBuffer & buffer,
                          int maxKeys, int keySize = sizeof(keyType));
protected:
    int MaxBKeys; // maximum number of keys in a node
    friend class BTree<keyType>;
};
```

# BTreeNode constructor



- # of keys in the BTreeNode
  - Actually one more than the order of the tree
  - The SimpleIndex constructor creates an index record with maxKeys + 1 elements

```
template <class keyType>
BTreeNode<keyType>::BTreeNode(int maxKeys, int unique)
:SimpleIndex<keyType>(maxKeys+1, unique)
{   Init ();}
```

# BTreeNode::Insert()



- Method Insert simply calls SimpleIndex::Insert
  - Making the index record larger allows to create an overfull node
    - Respond to the overflow in an appropriate fashion
    - 1 for success, 0 for failure and -1 for overflow

```
template <class keyType>
int BTreeNode<keyType>::Insert (const keyType key, int recAddr)
{
    int result;
    result = SimpleIndex<keyType>::Insert (key, recAddr);
    if (!result) return 0; // insert failed
    // this is modified
    // accessing protected members of superclass in C++ with templates
    if (this->NumKeys >= this->MaxKeys) return -1; // node overflow
    //if (NumKeys >= MaxKeys) return -1; // node overflow
    return 1;
}
```

# BTreeNode::Remove()



- Can create an underfull node
  - 1 for success, 0 for failure and -1 for underflow

```
template <class keyType>
int BTreeNode<keyType>::Remove (const keyType key, int recAddr)
{
    int result;
    result = SimpleIndex<keyType>::Remove (key, recAddr);
    if (!result) return 0; // remove failed
    // modified
    // accessing protected members of superclass in C++ with templates
    if (this->NumKeys < MinKeys) return -1; // node underflow
    //if (NumKeys < MinKeys) return -1; // node underflow
    return 1;
}
```

# Class BTree (1/3)



- **Supporting Files of B-Tree Nodes**
  - Uses in-memory BTNode objects
  - adds the file access portion
  - enforces the consistent size of the nodes
- Methods
  - Create, Open, Close a B-Tree
  - Search, Insert, Remove key-reference pairs
- Protected area
  - **Fetch: transfer nodes from disk to memory**
  - **Store: transfer nodes back to disk**
  - root node, height of the tree, file of index records
  - BTNode \*\*Node
    - used to keep a **collection of tree nodes in memory** and reduce disk access

# Class BTree (2/3)





```
// this is the full version of the BTree
template <class keyType>
class BTree
{
public:
    BTree(int order, int keySize = sizeof(keyType), int unique = 1);
    ~BTree();
    //int Open (char * name, int mode);
    int Open (char * name, ios_base::openmode mode);
    //int BTree<keyType>::Open (char * name, ios_base::openmode mode)
    int Create (char * name, ios_base::openmode mode);
    //int Create (char * name, int mode);
    int Close ();
    int Insert (const keyType key, const int recAddr);
    int Remove (const keyType key, const int recAddr = -1);
    int Search (const keyType key, const int recAddr = -1);
    void Print (ostream & );
    void Print (ostream &, int nodeAddr, int level);
protected:
```

# Class BTree (3/3)



```
protected:  
    typedef BTNode<keyType> BTNode; // useful shorthand  
    BTNode * FindLeaf (const keyType key);  
  
    // load a branch into memory down to the leaf with key  
    BTNode * NewNode ();  
    BTNode * Fetch(const int recaddr);  
    int Store (BTNode *);  
    BTNode Root;  
    int Height; // height of tree  
    int Order; // order of tree  
    int PoolSize;  
    BTNode ** Nodes; // pool of available nodes  
    // Nodes[1] is level 1, etc. (see FindLeaf)  
    // Nodes[Height-1] is leaf  
    FixedFieldBuffer Buffer;  
    RecordFile<BTNode> BTreeFile;  
};
```

# Page Structure



content of PAGE 2, 3

	KEYCOUNT	KEY array	CHILD array
Page 2	4	<b>D M P W</b>	0 3 8 5
Page 3	3	<b>G I M</b>	Nil Nil Nil Nil

# Outline



- 9.7 An Object-Oriented Representation of B-Trees
- 9.8 B-Tree Methods Search, Insert, and Others
- 9.10. Formal Definition of B-tree properties
- 9.11 Worst-case search depth
- 9.9 B-Tree Nomenclature

# Algorithm for Search (1/4)



- Searching procedure
  - iterative
  - work in two stages
    - operating alternatively on entire pages ([Class BTree](#))
      - **BTree::Search(key, recAddr);**
    - and then within pages ([Class BTreeNode](#))
      - **SimpleIndex::Search(key, recAddr);**
  - Step1: Loading a page into memory
  - Step2: Searching through a page, looking for the key along the tree until it reaches the leaf level

# Algorithm for Search (2/4)



- Specifications of **Search** and **FindLeaf** methods

```
Template <class keyType>
int BTTree<keyType>::Search(const keyType key, const int recAddr)

template <class keyType>
BTTreeNode<keyType>* BTTree<keyType>::FindLeaf(const keyType key)
```

## Search method

```
recAddr = btree.Search('L')
call FindLeaf('L');
Search key in the leaf node, and then
  if key exists, return the data file address of record with key 'L'
  otherwise, return -1
```

## FindLeaf method

```
Search down to leafNode, beginning of the root
return the address of leafNode
```

# Algorithm for Search (3/4)



## ● Example

```
RecAddr = btree.Search('L');
==>leafNode = FindLeaf(key);
==> recAddr = Nodes[level-1]->Search(key, -1, 0)
==> Nodes[level] = Fetch(recAddr);
```

```
template <class keyType>
int BTTree<keyType>::Search (const keyType key, const int recAddr)
{
    BTNode * leafNode;
    leafNode = FindLeaf (key);
    return leafNode -> Search (key, recAddr);
}
```

# Algorithm for Search (4/4)



- FindLeaf

```
// load a branch into memory down to the leaf with key
template <class keyType>
BTreeNode<keyType> * BTree<keyType>::FindLeaf (const keyType key)
{
    int recAddr, level;
    for (level = 1; level < Height; level++) {
        //inexact search
        recAddr = Nodes[level-1]->Search(key, -1, 0);
        Nodes[level]=Fetch(recAddr);
    }
    return Nodes[level-1];
}
```

- Fetch

```
// load this node from File into a new BTreeNode
template <class keyType>
BTreeNode<keyType> * BTree<keyType>::Fetch(const int recaddr)
{
    int result;
    BTNode * newNode = new BTNode(Order);
    result = BTTreeFile.Read (*newNode, recaddr);
    if (result == -1) return NULL;
    newNode -> RecAddr = result;
    return newNode;
}
```

# Algorithm for Insertion (1/3)



- Observations of Insertion, Splitting, and Promotion
  - proceed all the way down to the leaf level
  - after finding the insertion location at the leaf level, the work proceeds upward from the bottom
- Iterative procedure as having three phases
  - Search to the leaf level, using **FindLeaf** method
  - Insertion, overflow detection, and splitting on the upward path
  - Creation of a new root node, if the current root was split

# Algorithm for Insertion (2/3)



- With no redistribution
- (Step 1) Locate node on bottom most level in which to insert record
  - Location is determined by key search.
- (Step 2) If vacant record slot is available
  - insert the record so that key sequencing is maintained.
  - Then, update the pointer associated with the record (Pointer is null for level 0 records). Then Stop!
- (Step 3) If no vacant record slot exists
  - identify median record.
  - All records and pointers to the left of the median records are stored in one node (the original) and those to the right are stored in another node(the new node).

# Algorithm for Insertion (3/3)




- (Step 4) If the topmost node was split
  - create a new topmost node which contains the median record identified in Step 3, filled with pointers to the original and split nodes.
  - Update the root node to point to the new topmost node. Then Stop!
- (Step 5) If topmost node was not split
  - prepare to insert median record identified in Step 3 and a pointer to the new node (created in Step 3).
  - Then Goto Step 2.
- Note : Step 4 makes B-tree increase in height by 1 level
- B-trees have 70% occupancy(like B+-trees) on an average


# Example of Inserting R with codes (1/7)



- Current State at creating a B-tree (4/10)
  - Try to Insert R



# Example of Inserting R with codes (2/7)



- (1) Search the leaf node for key R using FindLeaf

```
thisNode = FindLeaf (key);
```

- (2) Insert R into the leaf node

```
result = thisNode -> Insert (key, recAddr);
```

- (3) Detect an overflow

- The node must be split into two nodes

```
newNode = NewNode();  
thisNode->Split(newNode);  
Store(thisNode); Store(newNode);
```


# Example of Inserting R with codes (3/7)



- BTTreeNode::Split
  - Distribute the keys between the original page and the new page

```
template <class keyType>
int BTTreeNode<keyType>::Split (BTTreeNode<keyType> * newNode)
{
    // check for sufficient number of keys
    if (this->NumKeys < this->MaxKeys) return 0;
    // find the first Key to be moved into the new node
    int midpt = (this->NumKeys+1)/2;
    int numNewKeys = this->NumKeys - midpt;
    // check that number of keys for newNode is ok
    if (numNewKeys > newNode -> MaxBKeys || numNewKeys < newNode -> MinKeys)
        return 0;
    // move the keys and recaddrs from this to newNode
    for (int i = midpt; i< this->NumKeys; i++) {
        newNode->Keys[i-midpt] = this->Keys[i];
        newNode->RecAddrs[i-midpt] = this->RecAddrs[i];
    }
    newNode->NumKeys = numNewKeys; // set number of keys in the two Nodes
    this->NumKeys = midpt; // Link the nodes together
    return 1;
}
```



# Example of Inserting R with codes (4/7)



- (4) Update the parent node
  - The largest key in thisNode has changed
  - `thisNode->LargestKey() = 'T'`

```
level--; // go up to parent level
if (level < 0) break;
// insert newNode into parent of thisNode
parentNode = Nodes[level];
result = parentNode->UpdateKey(largestKey, thisNode->LargestKey());
```

# Example of Inserting R with codes (5/7)



- (5) Insert the largest value(**W**) in the new node into the root

```
result = parentNode->Insert (newNode->LargestKey(), newNode->RecAddr);  
thisNode=parentNode;
```


- Result = -1
- Promote the key **W** ( **thisNode** = **parentNode**)
  - cause the root to overflow with five keys
  - a new root node is created and the keys **P** and **W** are inserted : **(D,M,P)**, **(T,W)**

# Example of Inserting R with codes (6/7)




- (6) Create a new root node, and insert the keys P & W into it

```
// put previous root into file
int newAddr = BTreeFile.Append(Root);
// insert 2 keys in new root node
Root.Keys[0]=thisNode->LargestKey();
Root.RecAddrs[0]=newAddr;
Root.Keys[1]=newNode->LargestKey();
Root.RecAddrs[1]=newNode->RecAddr;
Root.NumKeys=2;
Height++;
```



# Example of Inserting R with codes (7/7)



# Create, Open, and Close (1/2)



- Create()
  - Write the empty root node into the file BTreerFile
- Open()
  - Open BTreerFile and load the root node into memory from the first record in the file
- Close()
  - Store the root node into BTreerFile and close it

# Create, Open, and Close (2/2)



```
template <class keyType>
int BTTree<keyType>::Create (char * name, ios_base::openmode mode)
{
    int result;
    result = BTTreeFile.Create(name, mode);
    if (!result) return result;
    // append root node
    result = BTTreeFile.Write(Root);
    Root.RecAddr=result;
    return result != -1;
}
```

```
template <class keyType>
int BTTree<keyType>::Close ()
{
    int result;
    result = BTTreeFile.Rewind();
    if (!result) return result;
    result = BTTreeFile.Write(Root);
    if (result== -1) return 0;
    return BTTreeFile.Close();
}
```

# Testing the B-Tree



- Full code of program to test creation and insertion of a B-tree

```
const char * keys="CSDTAMPIBWNGURKEHOLJYQZFXV";
const int BTreesize = 4;
int main (int argc, char ** argv)
{
    int result, i;
    BTrees <char> bt (BTreesize);
    result = bt.Create ("testbt.dat",ios::in|ios::out);
    if (!result){cout<<"Please delete testbt.dat"<<endl;return 0;}
    for (i = 0; i<26; i++)
    {
        cout<<"Inserting "<<keys[i]<<endl;
        result = bt.Insert(keys[i],i);
        bt.Print(cout);
    }
    bt.Search(1,1);
    return 1;
}
```

# Outline



- 9.7 An Object-Oriented Representation of B-Trees
- 9.8 B-Tree Methods Search, Insert, and Others
- 9.10. Formal Definition of B-tree properties
- 9.11 Worst-case search depth
- 9.9 B-Tree Nomenclature

# Motivation



- Assume
  - We store 1,000,000 keys using a B-tree of order 512 (maximum of 511 keys per page)
- Question

In the worst case, what will be **the maximum number of disk accesses** required to locate a key in the tree?



Same question


How deep the B-tree will be?

# How to answer the worst case search depth?



- Observe the formal definition of B-tree properties
  - To calculate the minimum number of descendants with a B-tree of some given order
- The worst case occurs
  - when every page of the tree has only the minimum # of descendants  
→ A maximal height with a minimum breadth

# Recap: Formal definition of B-Tree Properties



- The properties of a B-tree of order  $m$ 
  - 1. Every page has a **maximum of  $m$**  descendants
  - 2. Every page, except for the root and the leaves, has **at least ceiling of  $m/2$  ( $=\lceil m/2 \rceil$ )** descendants
  - 3. The root has **at least two** descendants (unless it is a leaf)
  - 4. All the leaves appear **on the same level**
  - 5. The leaf level forms a complete, ordered index of the associated data file

# Extract general pattern (1/3)



- A B-tree of order  $m$

- Level 1 (root)

- Minimum number of descendants from the root page is 2

→ 2 descendants

3. The root has **at least two** descendants

- Level 2:

- The second level of the tree has 2 pages
    - Each page of these pages, in turn, has at least  $[m/2]$  descendants

→  $2 \cdot [m/2]$

2. Every page,..., has **at least ceiling of  $m/2$  ( $=[m/2]$ )** descendants

# Extract general pattern (2/3)



- A B-tree of order  $m$ 
  - Level 3
    - The second level of the tree has  $2 \cdot [m/2]$  pages
    - Each page of these pages, in turn, has at least  $[m/2]$  descendants
- $2 \cdot [m/2]^2$
  
- Level  $d$ 
  - $2 \cdot [m/2]^{d-1}$



# Extract general pattern (3/3)

- The general pattern of the relation
  - between depth and the minimum number of descendants

Level	Minimum number of keys (children)
1 (root)	2
2	$2 \cdot [m/2]$
3	$2 \cdot [m/2]^2$
4	$2 \cdot [m/2]^3$
...	...
d	$2 \cdot [m/2]^{d-1}$

# Worst case search depth (1/2)



- For a tree with  $N$  keys in its leaves
  - The relationship between keys and the minimum height  $d$ 
    - $N \geq 2 \cdot \lceil m/2 \rceil^{d-1}$
- Solving the equation for  $d$ 
  - $d \leq 1 + \log_{\lceil m/2 \rceil}(\frac{N}{2})$

# Worst case search depth (2/2)



- Solving the equation for d
  - $d \leq 1 + \log_{\lceil m/2 \rceil}(\frac{N}{2})$
- Example
  - B-tree order d= 512, N=1,000,000 keys
    - $d \leq 1 + \log_{256}(500,000) = 3.37$
    - Dept of no more than level 3

# Outline



- 9.7 An Object-Oriented Representation of B-Trees
- 9.8 B-Tree Methods Search, Insert, and Others
- 9.10. Formal Definition of B-tree properties
- 9.11 Worst-case search depth
- **9.9 B-Tree Nomenclature**

# B-Tree Nomenclature





- Be aware that terms are not uniform in the literature
  - Definitions are also quite different
  - In fact, there are a number of B-tree variations
- In this text book
  - “B tree” for B+ tree by other books
  - “B+ tree” is B+ tree with a linked list of sorted data blocks

# Difference between B-tree and B+-tree




- In a B-tree
  - pointers to data records exist at all levels of the tree
- In a B+-tree
  - all pointers to data records exists **only at the leaf-level nodes**
  - A B+-tree can have less levels (or higher capacity of search values) than the corresponding B-tree


# B-tree in other books



# B+-tree in other books



Other Book	Our Book
B+-Tree	B-Tree





# Another aspect (node structures)



- B-Tree in other text
  - Homogeneous trees
    - leaf nodes and interior nodes have same structures;
    - Each contains both data pointers and tree pointers
  - Average search length less for homogeneous trees
    - because some searches may conclude before reaching a leaf node
- B+-Tree in other text
  - Heterogeneous trees
  - leaf nodes and interior nodes have different structures

# Recap: B-tree and B+ tree



# Comparison of B-Tree and B+-Tree in other text



Topic	B-tree	B+ -Tree
Algorithm complexity for insertion	Rather complexity	More simple
Retrieval efficiency	Less efficiency (B-tree is tall & spindle)	More efficient B+-tree is short & bushy
Storage efficiency	Slightly more efficient (is less space)	Less efficient (is more space)
1-pass structure creation algorithms	Rather complex	Simple

# Q&A

