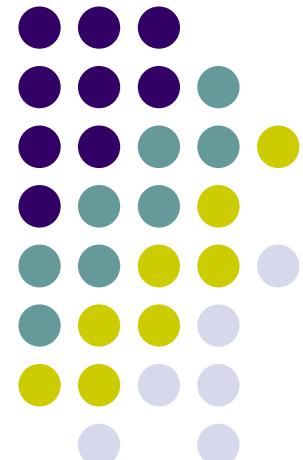


# File Structures

## Ch09. C. B-Tree Delete

---

2020. Spring  
Instructor: Joonho Kwon  
[jhwon@pusan.ac.kr](mailto:jhwon@pusan.ac.kr)  
Data Science Lab @ PNU



# Outline



- 9.12 Deletion, Merging, and Redistribution
- 9.13 Redistribution During Insertion
- 9.14 B\*-trees
- Skipped
  - 9.15 Buffering of Pages
  - 9.16 Variable-Length Records and Keys

# B-tree properties

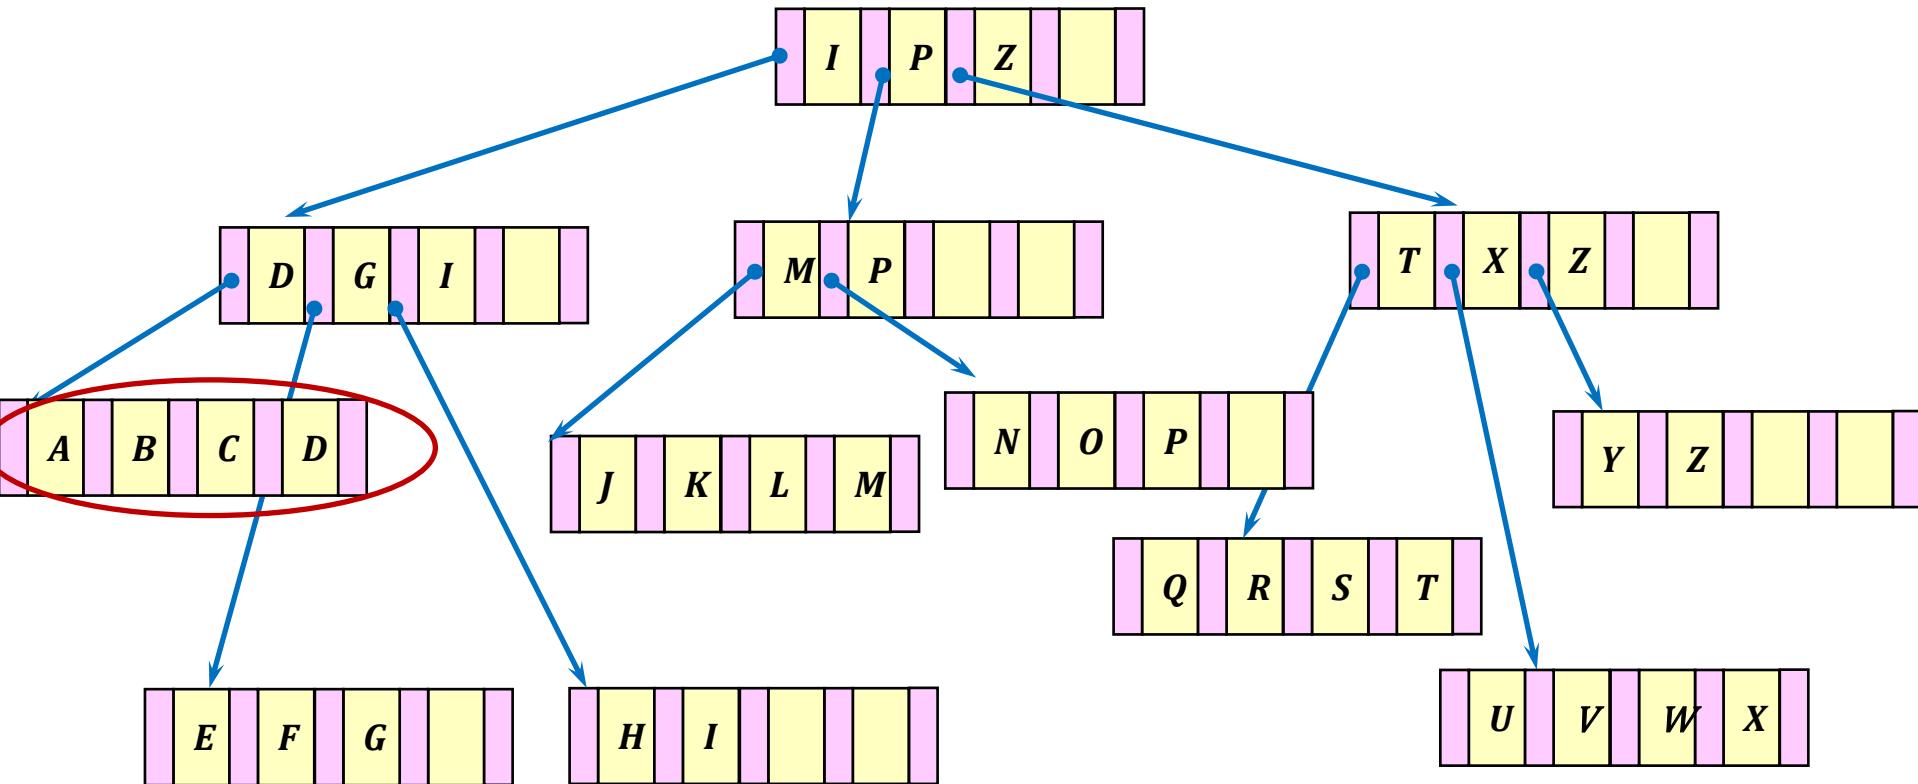


- Ensure that the B-tree properties
  - 1. Every page, except for the root and the leaves, has at least ceiling of  $m/2$  ( $=\lceil m/2 \rceil$ ) descendants
  - 2. A page contains at least  $\lceil m/2 \rceil$  keys and no more than  $m$  keys
- During insertions
  - The process of page splitting guarantees that these properties are maintained
- During deletions
  - ??

# Deletion example: case1 (1/2)



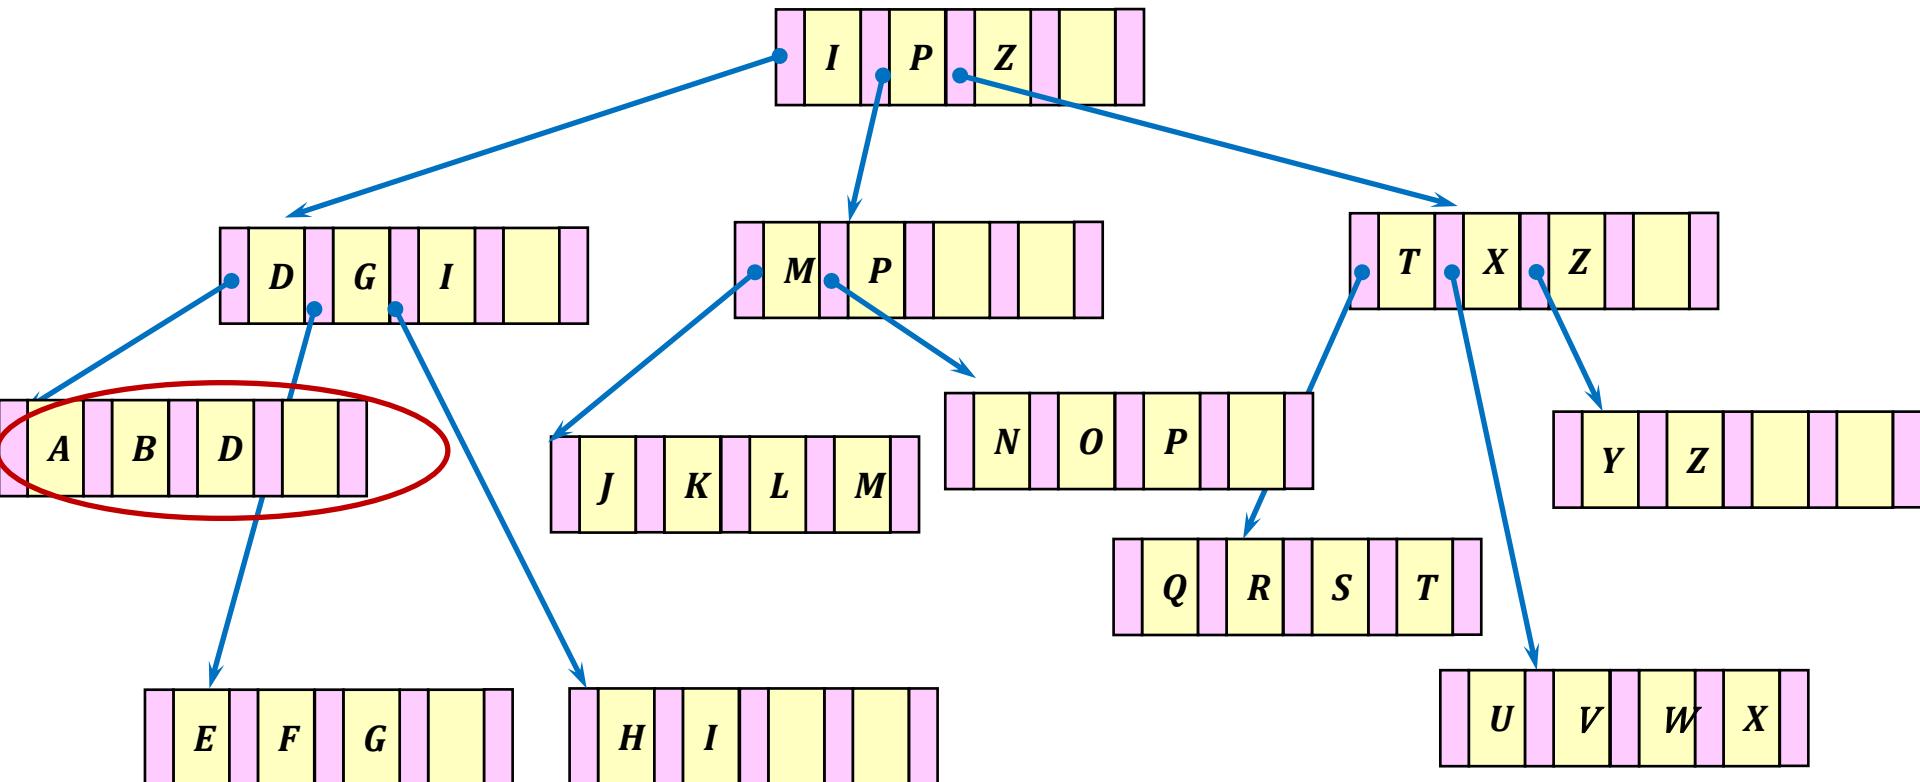
Removal of key C



# Deletion example: case1 (2/2)



- Change occurs only in leaf node

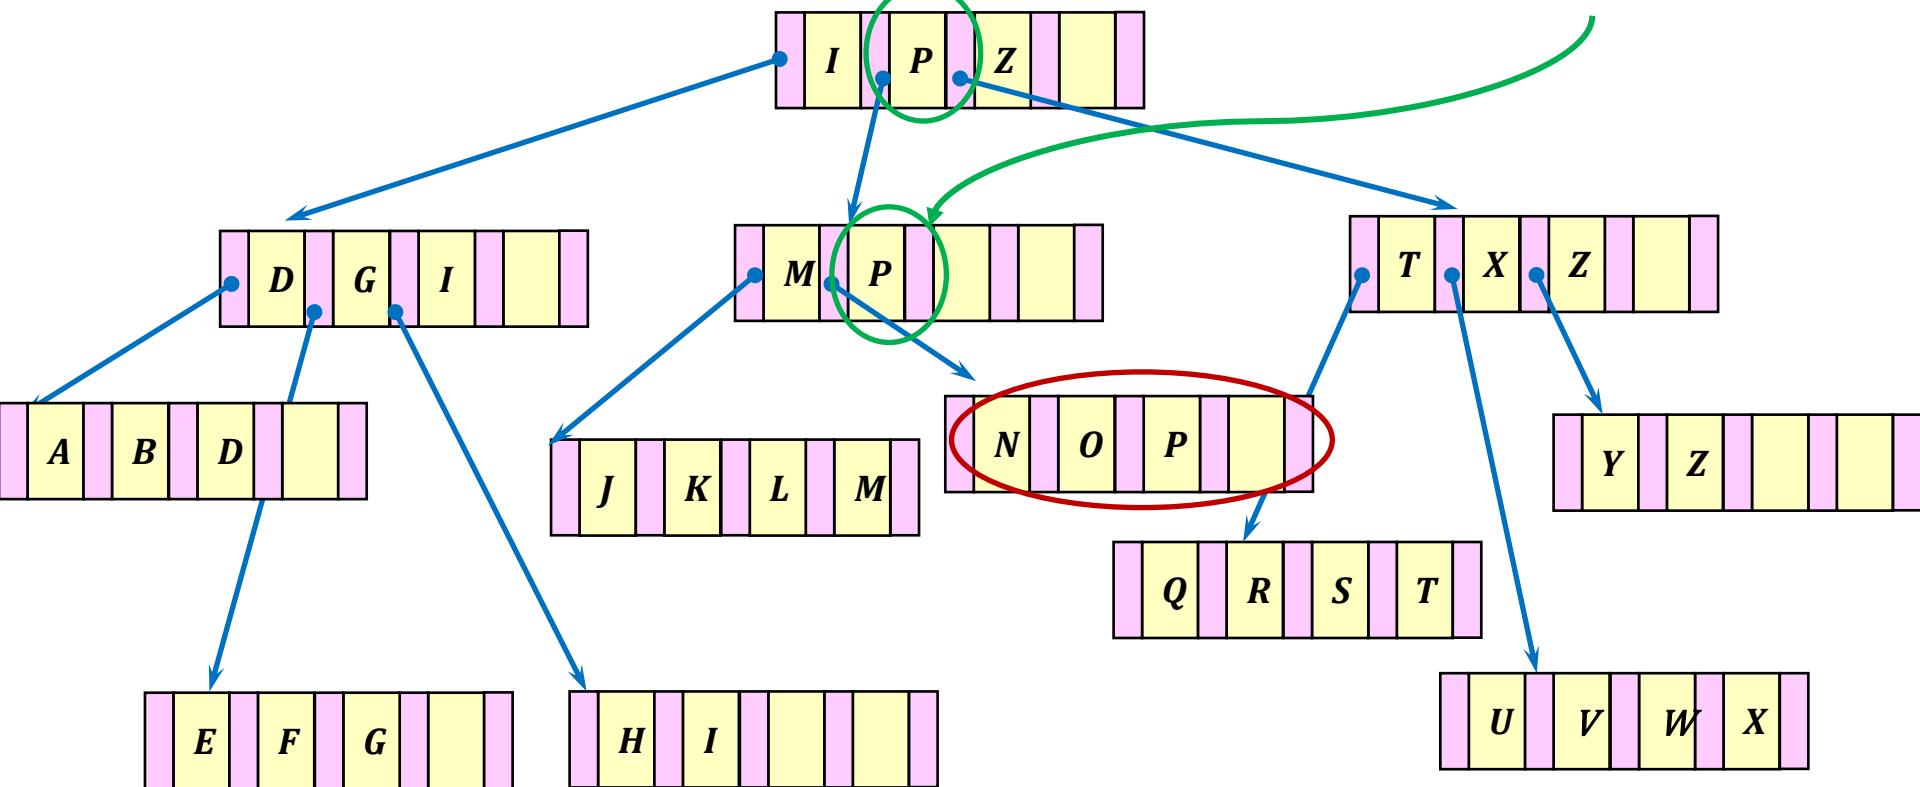


# Deletion example: case2 (1/4)



Removal of key P

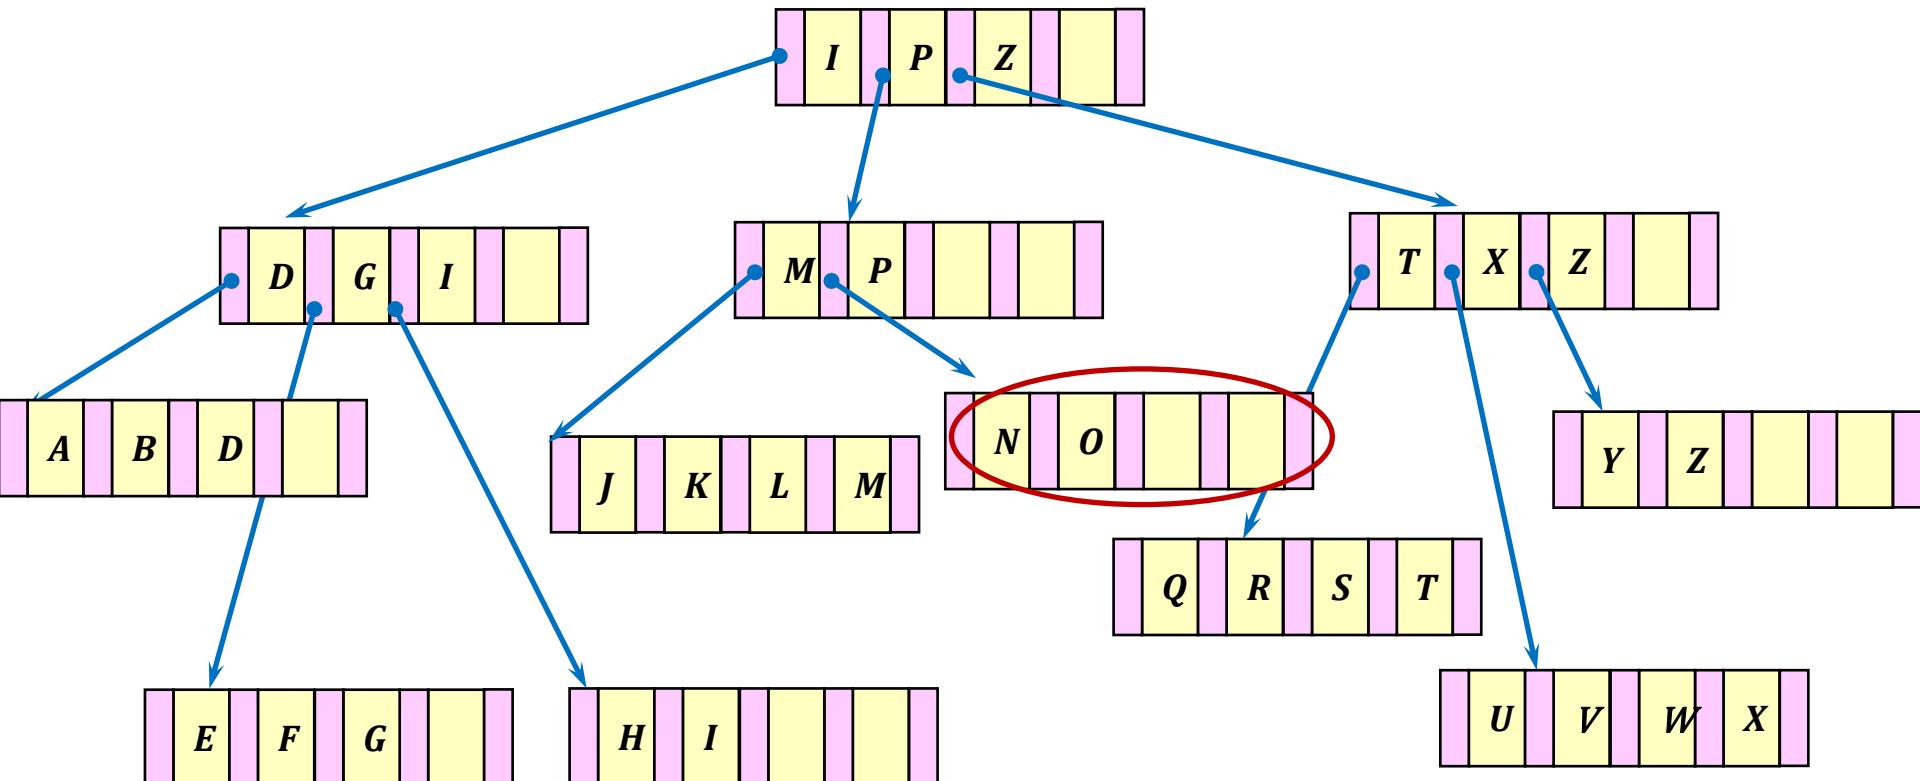
Does change the largest key in the node -> affects to the ancestor nodes



# Deletion example: case2 (2/4)



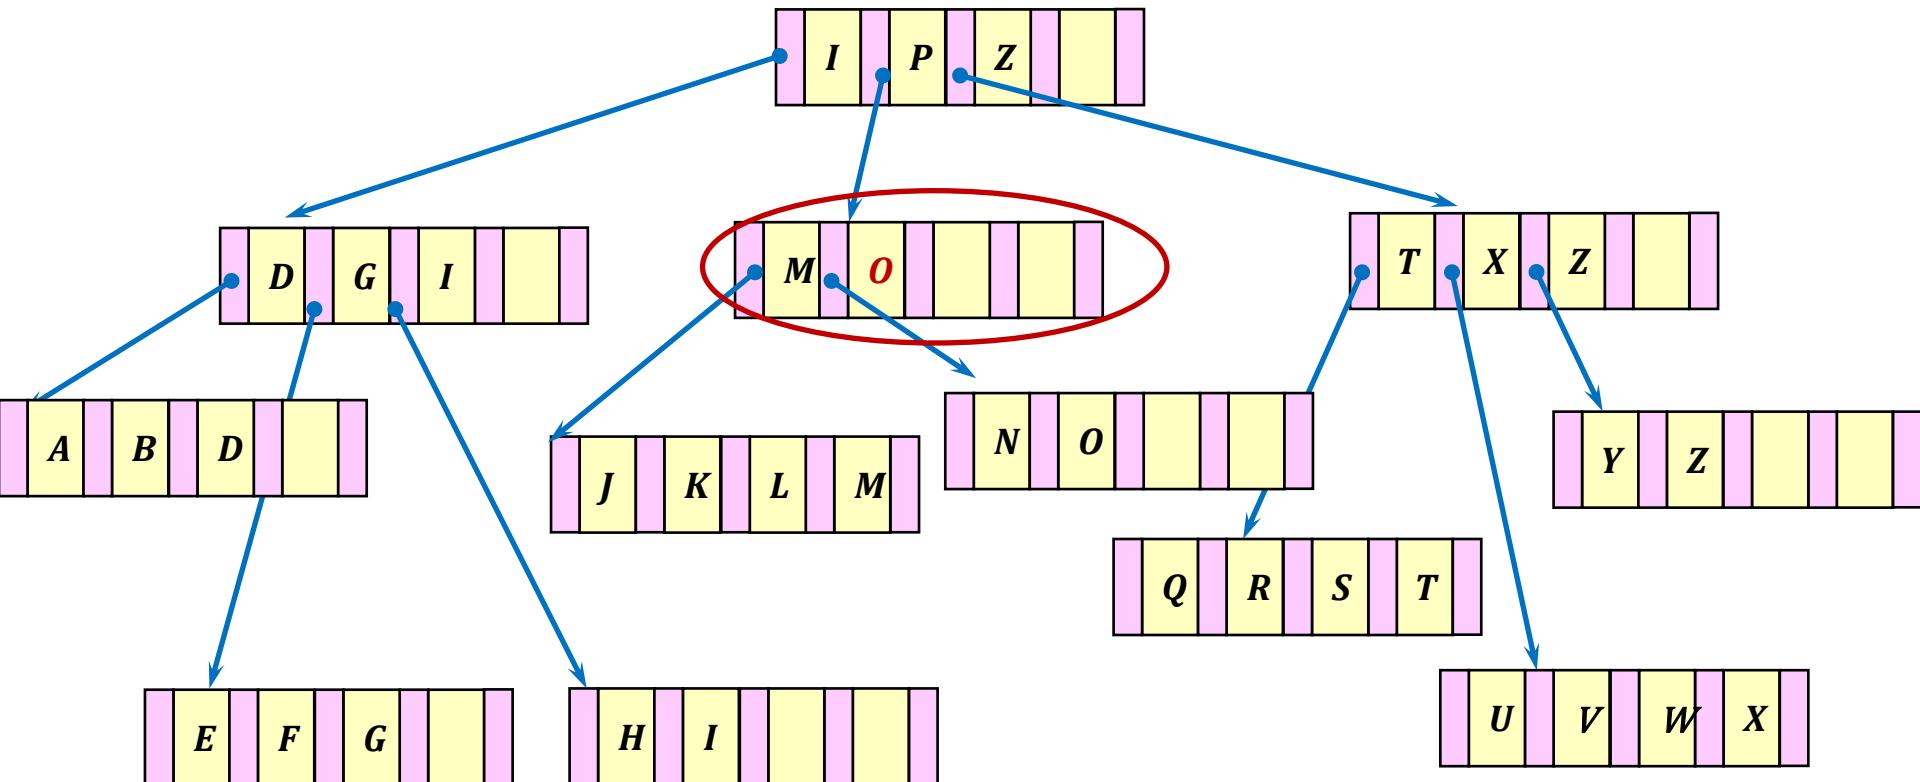
- P from the second leaf node
  - Change the largest key in the node



# Deletion example: case2 (3/4)



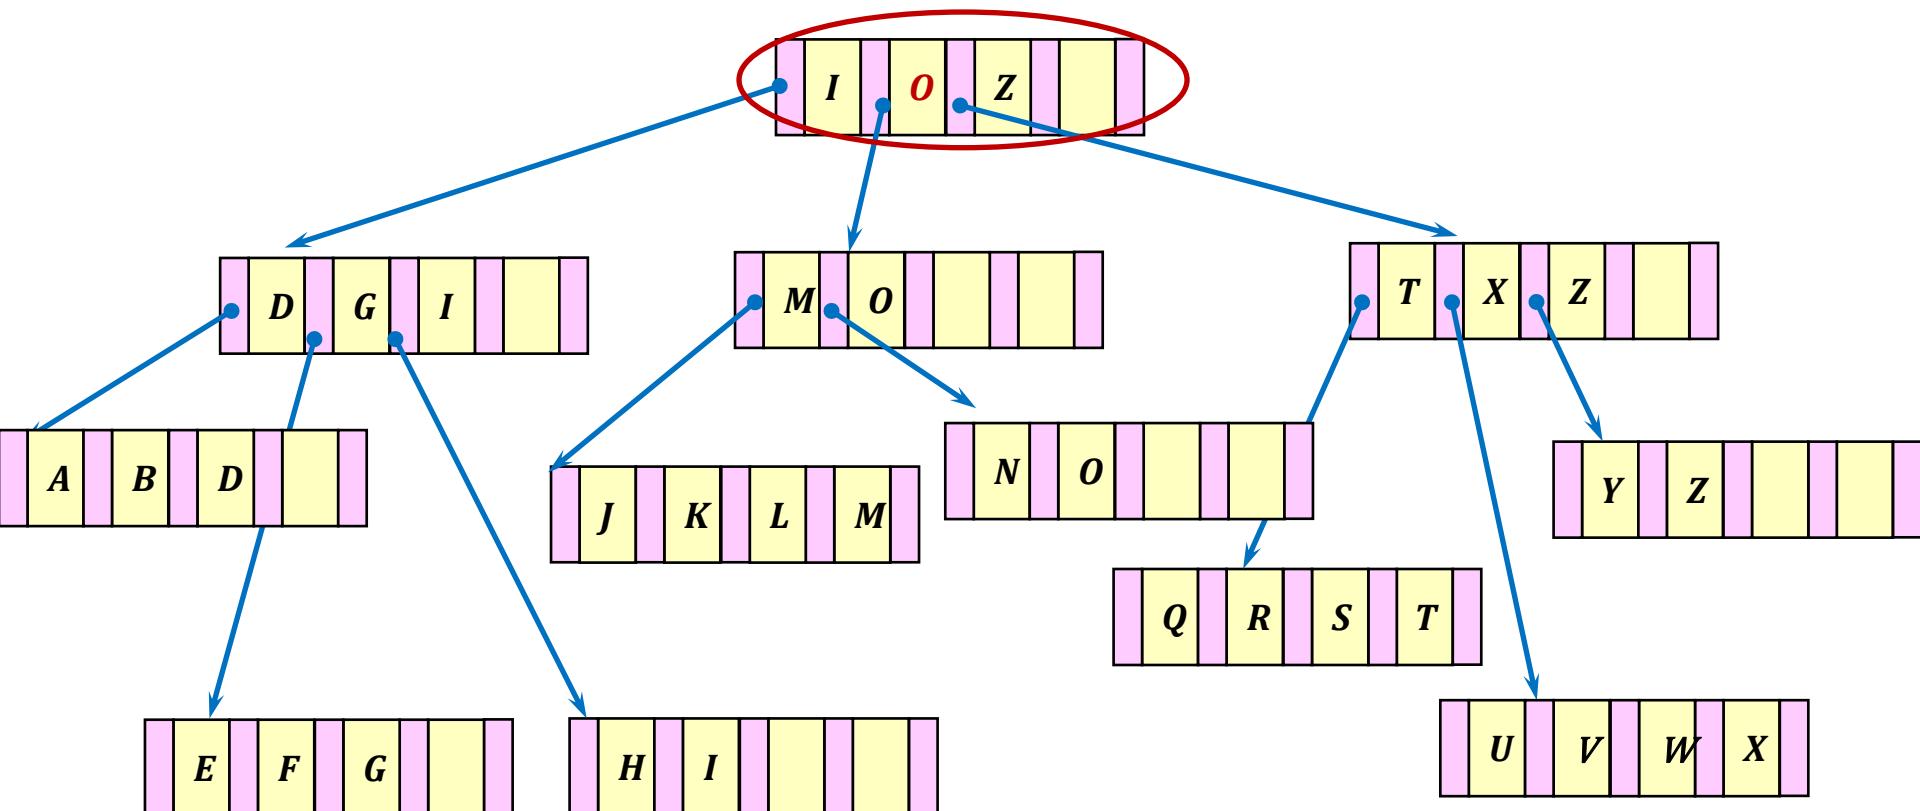
- The second-level node is modified
  - It contains O instead of P



# Deletion example: case2 (4/4)



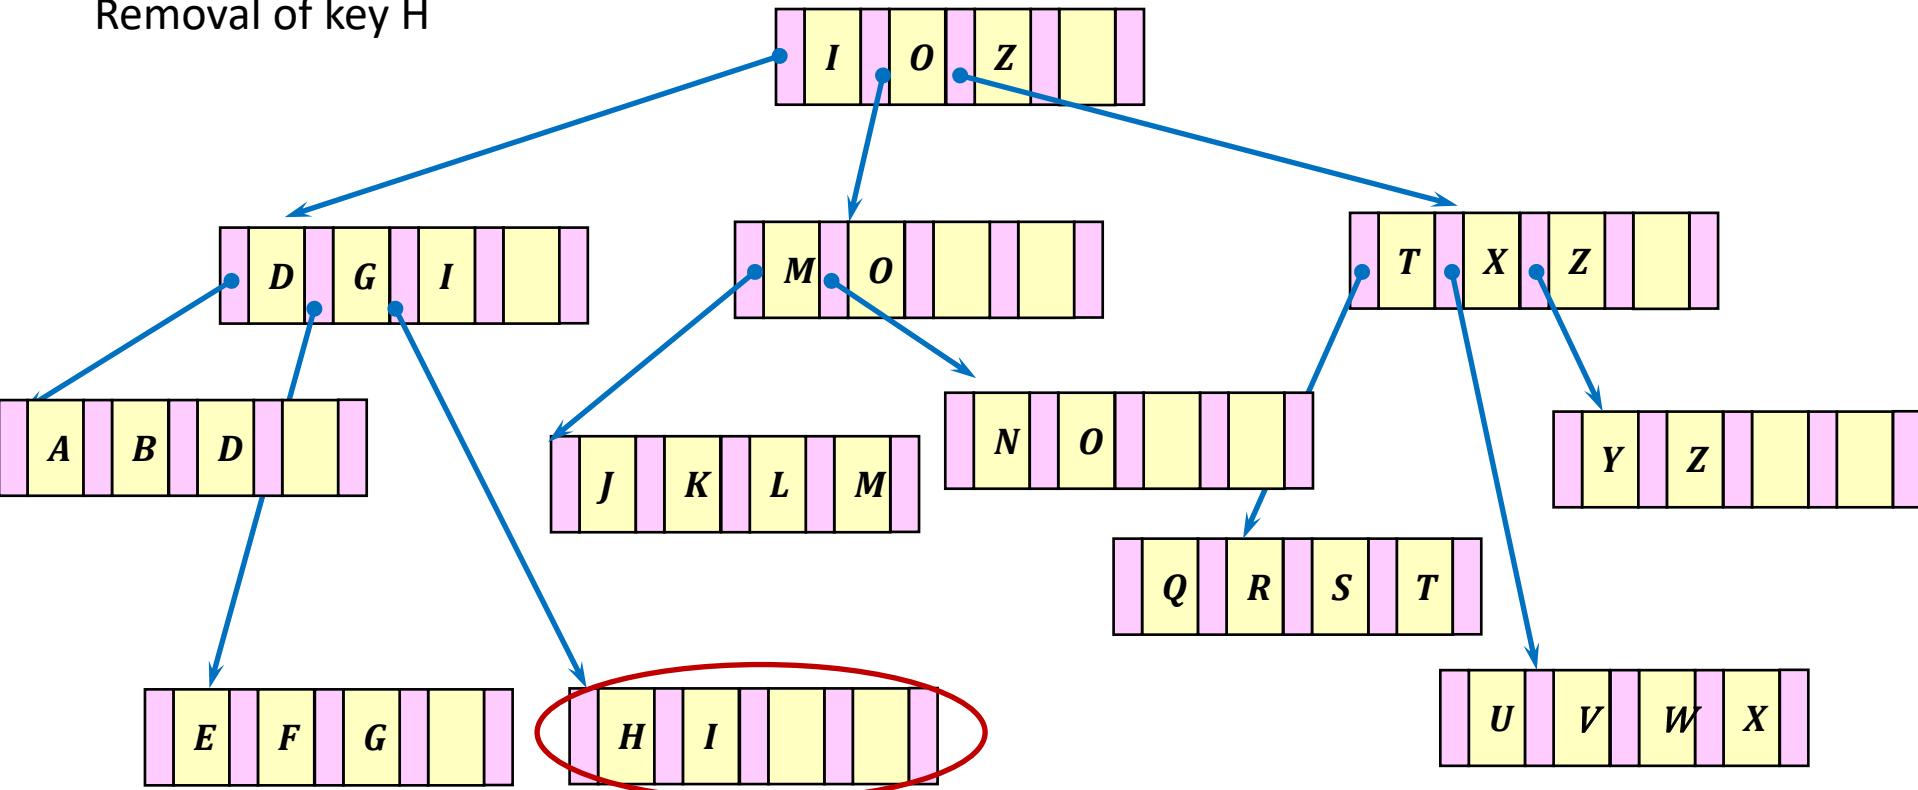
- The root is also modified
  - Key P is replaced by Key O



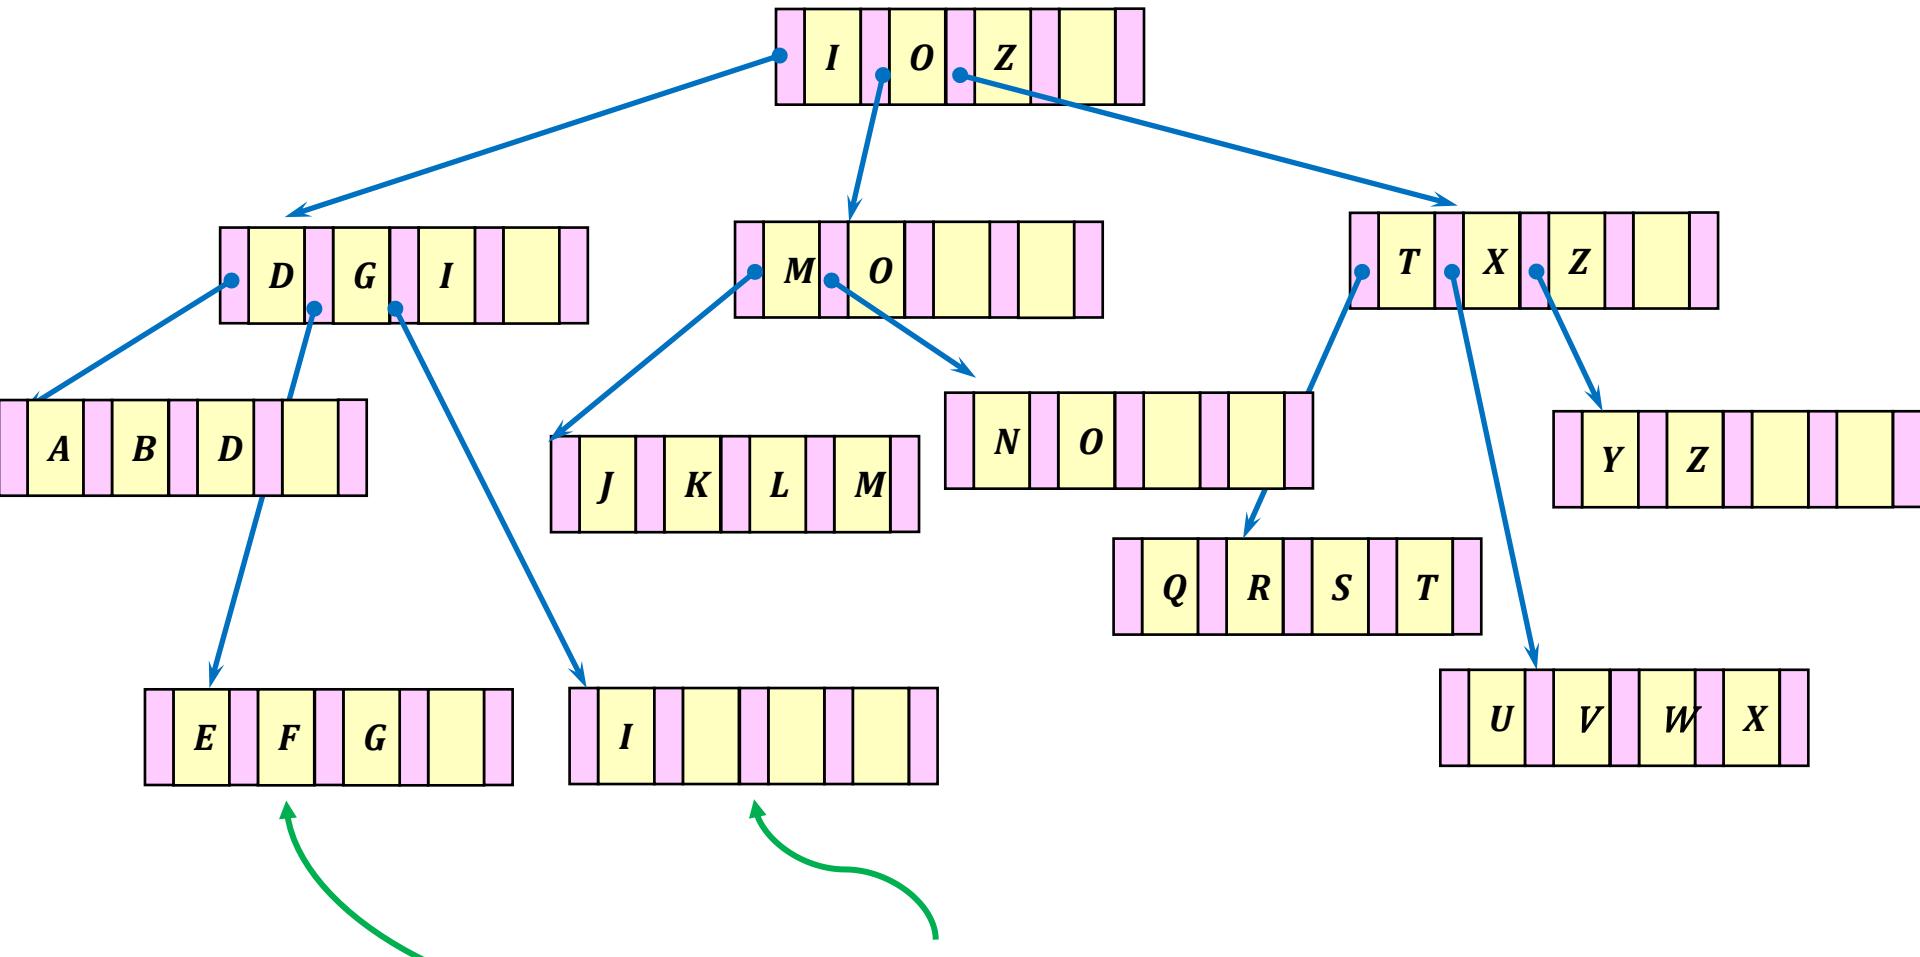
# Deletion example: case3 (1/3)



Removal of key H



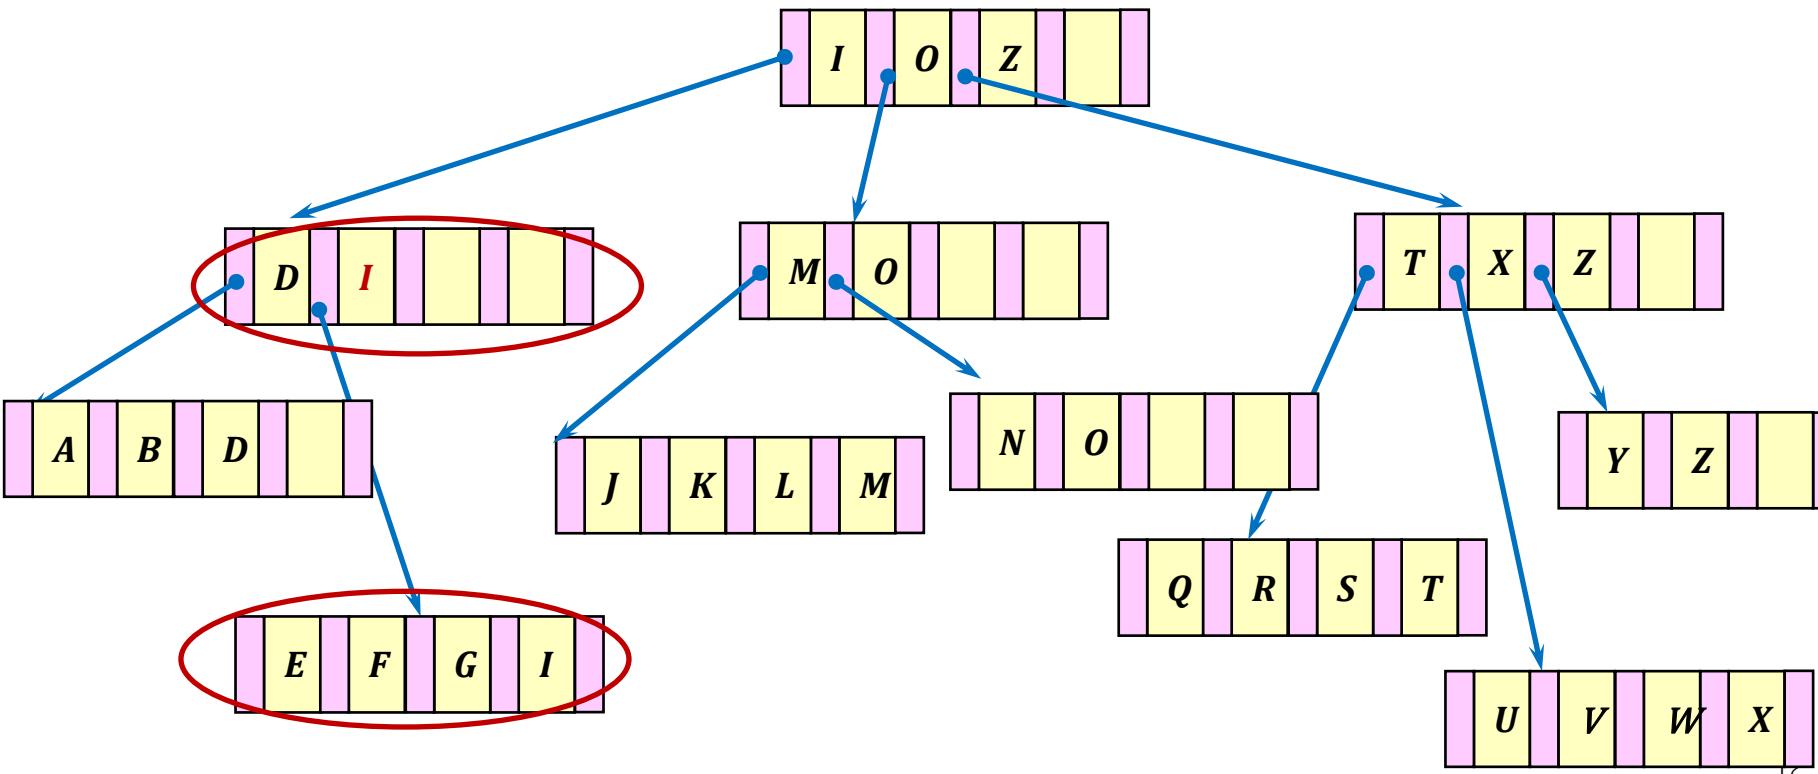
# Deletion example: case3 (2/3)



# Deletion example: case3 (3/3)



- After the merger
  - Second level node is modified to reflect the current status of the leaf nodes



# Deletion, Merging, and Redistribution (1/2)



- Rules for deleting a key  $k$  from a node  $n$  (order  $m$ )
  - 1. # of keys in  $n > m/2$ ,  $k$ : not the largest in  $n$ 
    - simply delete  $k$  from  $n$
  - 2. # of keys in  $n > m/2$ ,  $k$ : the largest in  $n$ 
    - delete  $k$
    - modify the higher level indexes to reflect the new largest key in  $n$
  - 3. # of keys in  $n = m/2$ , one of the siblings of  $n$  has few enough keys
    - merge  $n$  with its sibling
    - delete a key from the parent node

# Deletion, Merging, and Redistribution (2/2)

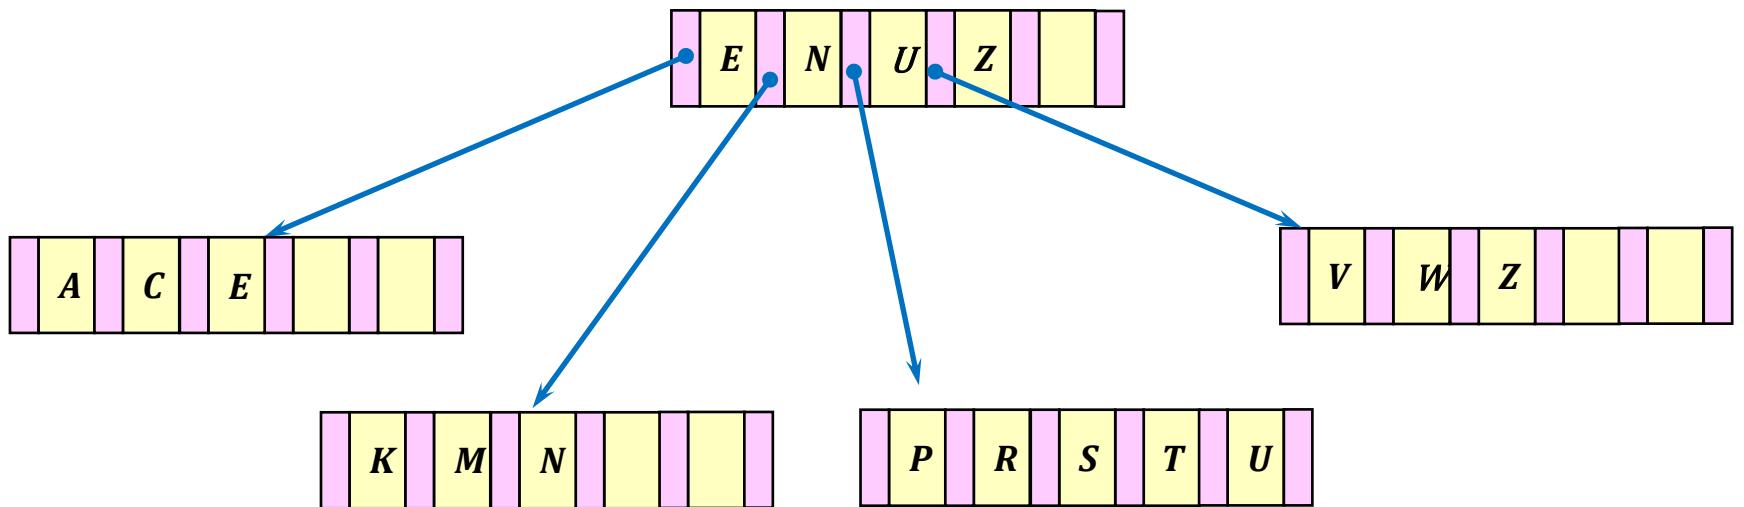


- Rules for deleting a key  $k$  from a node  $n$  (order  $m$ )
  - 4. # of keys in  $n = m/2$  and one of the siblings of  $n$  has extra keys
    - redistribute by moving some keys from a sibling to  $n$ , and
    - **modify the higher level indexes** to reflect the new largest keys in the affected nodes

# Another Example (1/4)



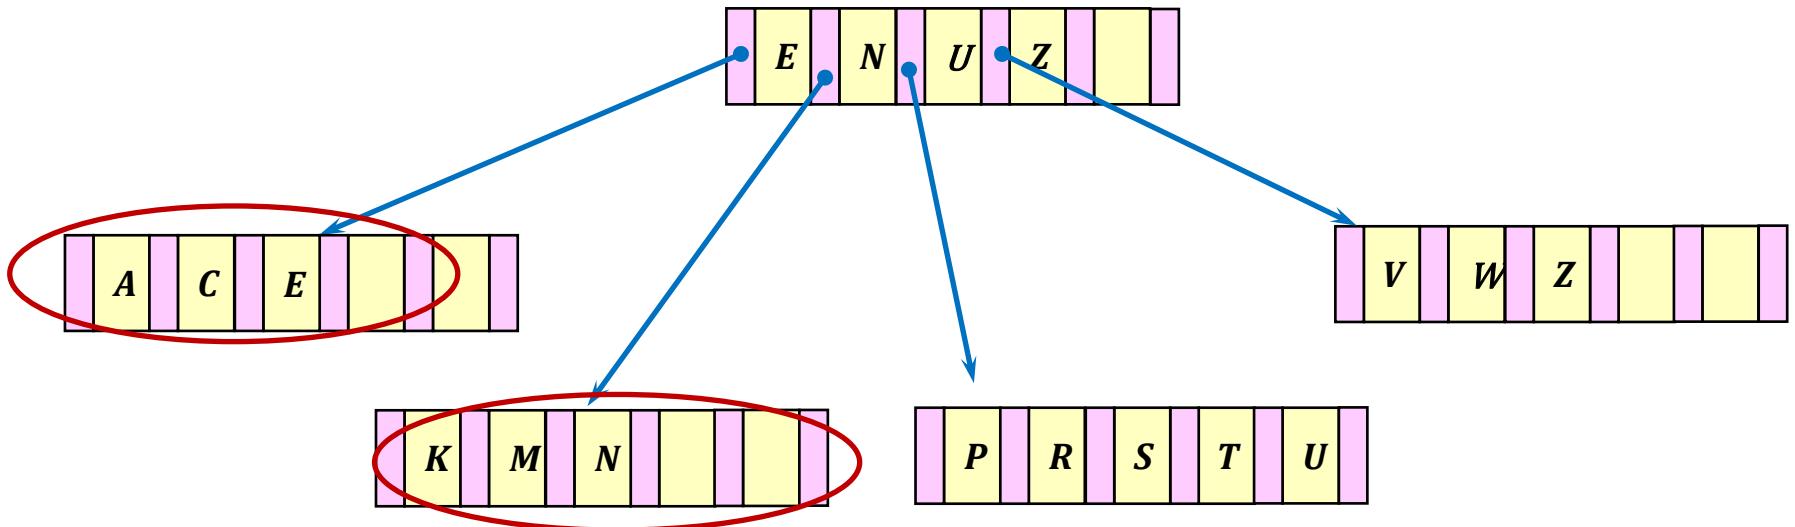
- A B-tree of order 6( $m=6$ )
  - 3 is minimum number of keys for each node ( $\lceil 6/2 \rceil = 3$ )



# Another Example (2/4)



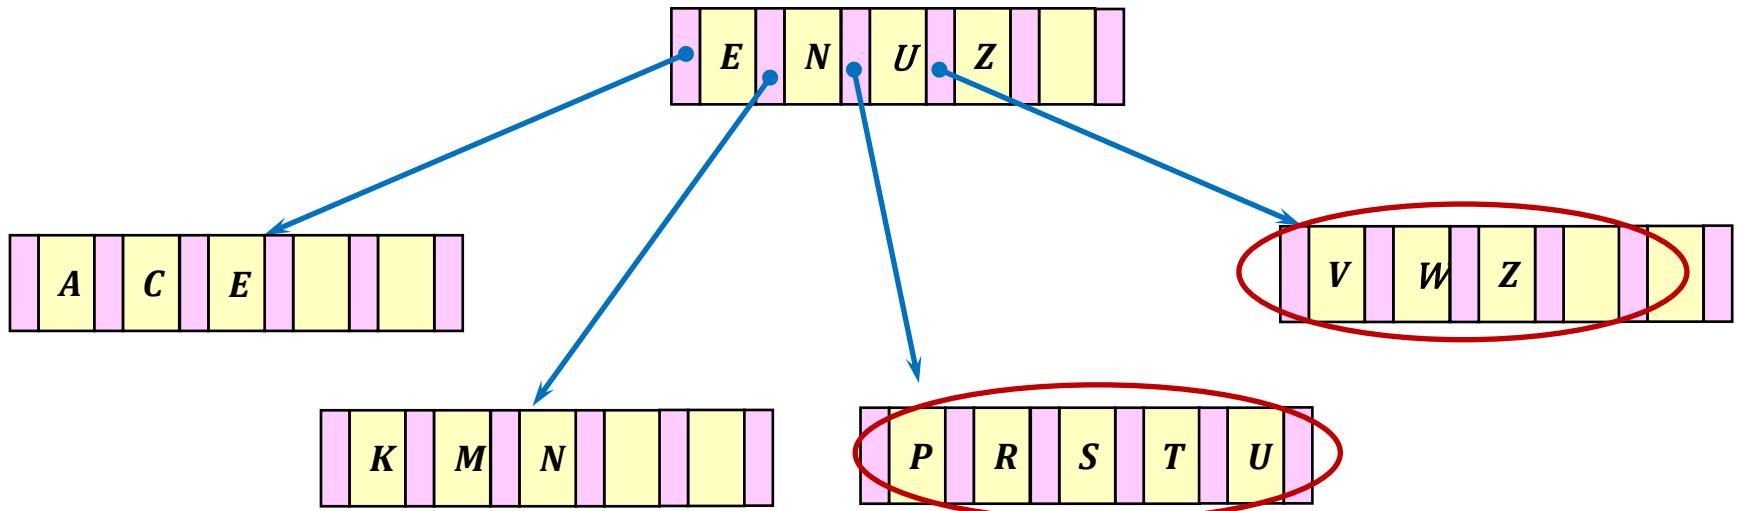
- Delete C
- Merging two sibling nodes



# Another Example (3/4)



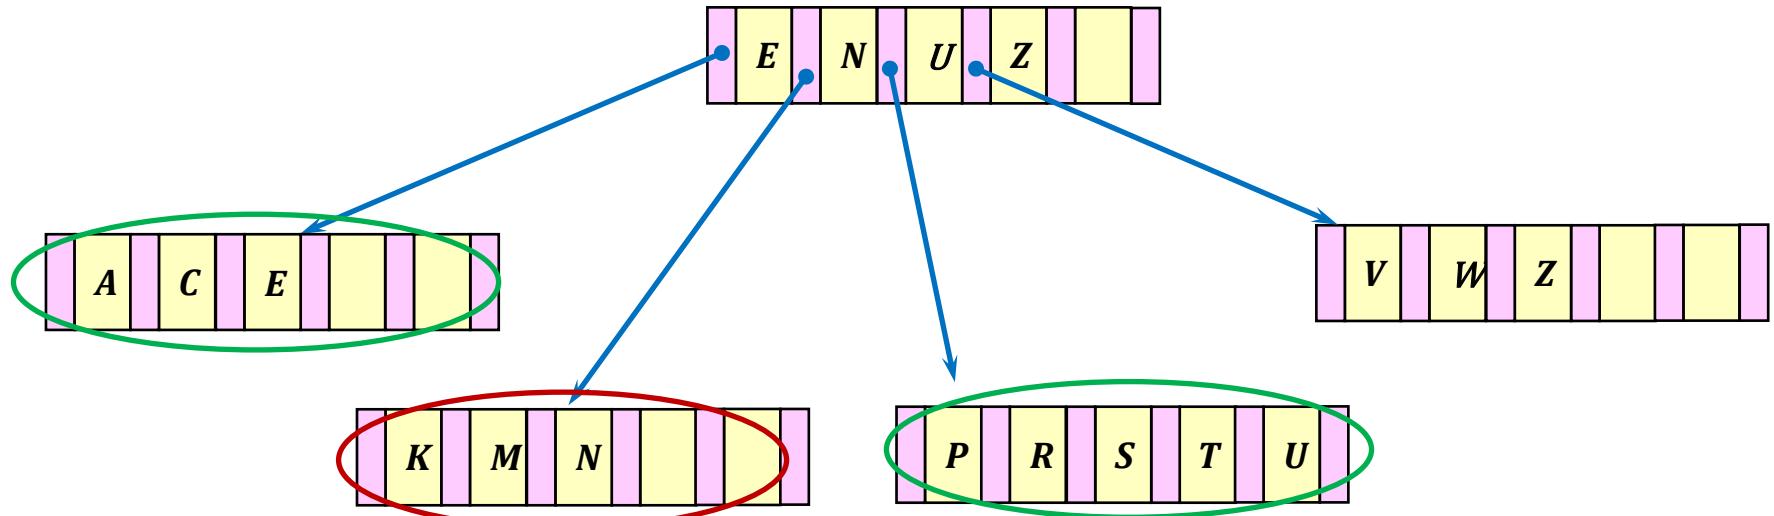
- Delete W
  - The only sibling has 5 keys, total 7 keys
    - redistribute



# Another Example (4/4)



- Delete M
  - Two options
    - Merge with the left sibling
    - Redistribute keys in the right sibling



# Deletion algorithm (1/2)



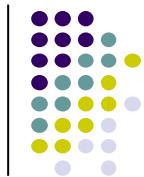
- Algorithm (with redistribution and merging)
  - 1. If the key to be deleted is not in a leaf
    - swap it with its immediate successor, which is in a leaf (might be redistributed or concatenated!)
  - 2. Delete the key

# Deletion algorithm (2/2)



- Algorithm (with redistribution and merging)
  - 3. If underflow occurs
    - (the leaf now contains one too few keys),
    - 3.1 If the left or right sibling has more than the minimum number of keys , redistribute
    - 3.2 Otherwise, concatenate the two leaves and the median key from the parent into one leaf
    - 3.3 Apply above step 3 to the parent as if it were deleted

# Redistribution



- Occur when a sibling has more than the minimum # of keys
- Idea: Move keys between siblings
- Result in a change in the key in the parent page
- Does not propagate : strictly local effects
- How many keys should be moved?
  - Not necessarily fixed
  - Even distribution is desired

# Merge (Concatenation)



- Occur in case of underflow
- Combining the two pages and the key from the parent page ==> make a single full page
- Reverse the splitting
- Concatenation must involve demotion of keys
  - *may cause underflow in the parent page*
- The effects propagate upward

# Redistribution During Insertion



- A way to improve storage utilization
- A way of avoiding the creation of new pages
- Tend to make an efficient B-tree in terms of space utilization
  - Worst case : around 50%
  - Average case : 67 ~ 69%
  - With redistribution during insertion : over 85%

# Outline



- 9.12 Deletion, Merging, and Redistribution
- 9.13 Redistribution During Insertion
- **9.14 B\*-trees**
- Skipped
  - 9.15 Buffering of Pages
  - 9.16 Variable-Length Records and Keys

# B\* Trees (1/2)



- Knuth, 1973, Addison-Wesley
- Use redistribution operation during insertion
- Perform two-to-three split
  - When split, the page has at least one sibling that is also full
  - After split, the pages are about 2/3 full
  - The page with at least  $(\text{ceiling of } (2m - 1)/3)$  keys  
c.f. remember  $(\text{ceiling of } (m/2)) - 1$  keys

# B\* Trees (2/2)



- Properties of a B\*-tree of order m
  - extend the notion of **redistribution during insertion** to include new rules for splitting
  - 1. Every page has a maximum of m descendants
  - 2. Every page except for the root has at least  $(2m-1)/3$  descendants (v.s.  $m/2$ )
  - 3. The root has at least two descendants (unless it is a leaf)
  - 4. All the leaves appear on the same level

# Q&A

