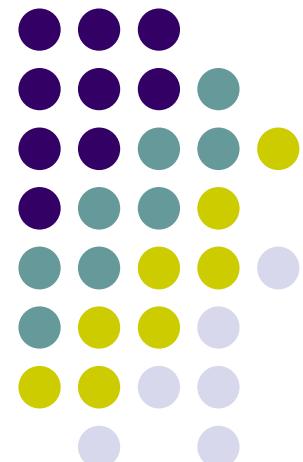


File Structures

Ch11. Hashing

2020. Spring
Instructor: Joonho Kwon
jhwon@pusan.ac.kr
Data Science Lab @ PNU



Outline



- 11.1 Introduction
- 11.2 A Simple Hashing Algorithm
- 11.3 Hashing Functions and Record Distribution
- 11.5 Collision Resolution by Progressive Overflow
- 11.4 How Much Extra Memory Should Be Used?
- 11.6 Storing More Than One Record per Address:
Buckets
- 11.7 Making Deletions
- 11.8 Other Collision Resolution Techniques
- 11.9 Patterns of Record Access

Introduction



- Motivation
 - Hashing is a useful searching technique, can be used for implementing indexes
 - Improves searching time
- Hash vs other index
 - Sequential search : $O(N)$
 - Binary search : $O(\log_2 N)$
 - B(B+) Tree index : $O(\log_k N)$ where k records in an index node
 - Hashing : $O(1)$

What is hashing



- Idea
 - To discover the location of a key by simply examining the key → need to design **a hash function**
- Hash function : $h(k)$
 - Transforms a key k into an address
 - An address space is chosen beforehand
 - Example
 - U is a set of all possible keys, h is hashing function

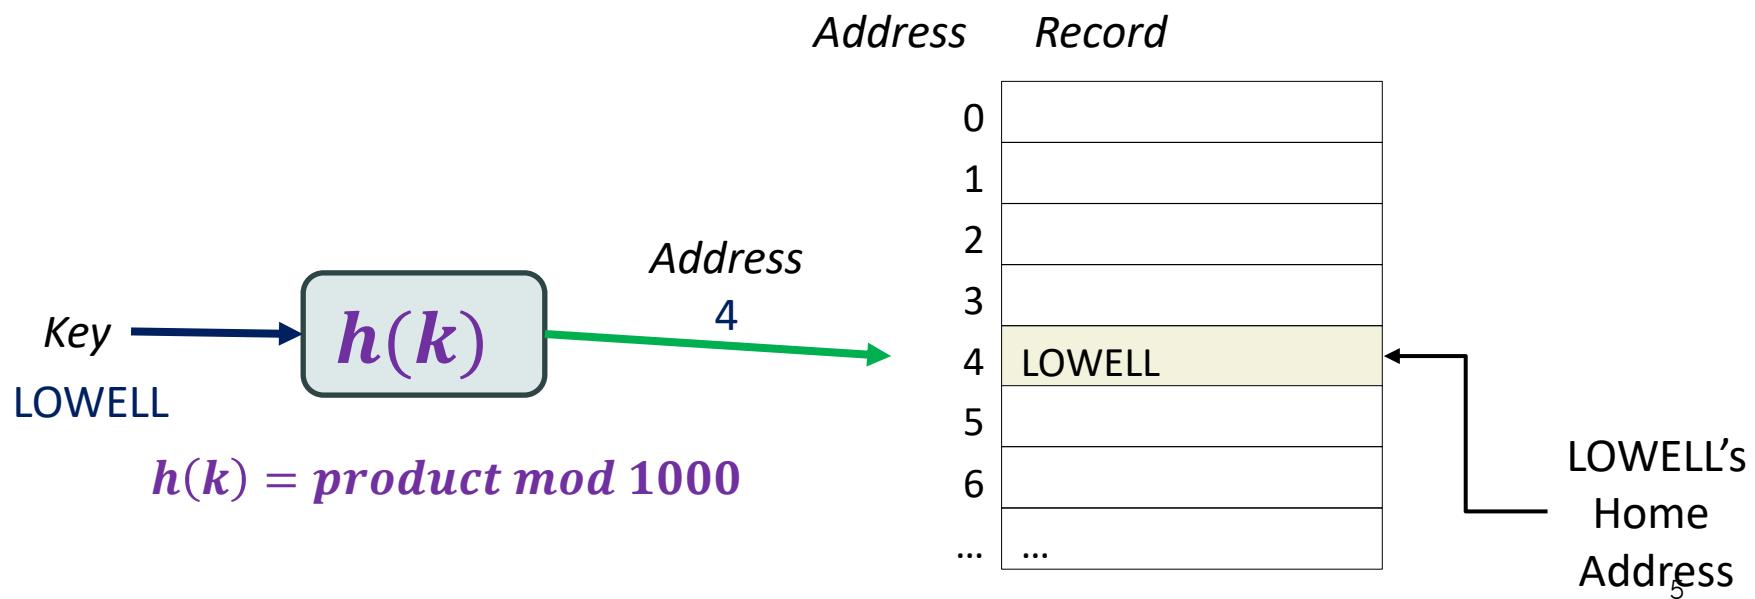
$$h: U \rightarrow \{0,1, \dots, 999\}$$

A Simple Hashing Scheme



| Name | ASCII code for First Two Letters | Product | Home Address |
|---------------|----------------------------------|-----------------|--------------|
| <u>BALL</u> | 66 65 | $66*65 = 4,290$ | 290 |
| <u>LOWELL</u> | 76 96 | $76*96 = 6,004$ | 004 |
| <u>TREE</u> | 84 82 | $84*82 = 6,888$ | 888 |

$$h(\text{LOWELL}) = 004$$



Hashing differs from indexing



- With hashing
 - There is no obvious connection between the key and address
 - hashing is referred to as **randomizing**
- Two different keys may be translated to the same address
 - Two records may be sent to the same place in the file
→ **collision**

Collision (1/2)



- Collision
 - Situation in which a record is hashed to an address that does not have sufficient room to store the record
 - Perfect hashing : impossible!
 - Different key, same hash value (Different record, same address)

$$h(\text{LOWELL}) = 4$$

$$h(\text{OLIVER}) = 76 * 79 \bmod 1000 = 6004 \bmod 1000 = 4$$

Collision (2/2)



- Solutions
 - Spread out the records
 - By choosing a good hash function
 - Use extra memory
 - Increase the size of the address space
 - Ex: reserve 5,000 available addresses rather than 1,000
 - Put more than one record at a single address
 - Use of buckets

A Simple Hashing Algorithm (1/4)



- Step 1. Represent the key in numerical form
 - If the key is a string : take the ASCII code
 - If the key is a number : nothing to be done
- e.g. **LOWELL**

| | | | | | | |
|---|---|---|---|---|---|--------------|
| L | O | W | E | L | L | ← 6 blanks → |
|---|---|---|---|---|---|--------------|

| | | | | | | | | | | | |
|----|----|----|----|----|----|----|----|----|----|----|----|
| 76 | 79 | 87 | 69 | 76 | 76 | 32 | 32 | 32 | 32 | 32 | 32 |
|----|----|----|----|----|----|----|----|----|----|----|----|

A Simple Hashing Algorithm (2/4)



- Step 2 . Fold and Add
 - Fold
 - 76 79 | 87 69 | 76 76 | 32 32 | 32 32 | 32 32
 - Add parts into one integer
 - Suppose we use 15bit integer expression, 32767 is limit
 - $7679 + 8769 + 7676 + 3232 + 3232 + 3232 = 33820 > 32767$
(overflow!)
 - Make sure that each successive sum is less than 32768
 - Largest addend : 9090 ('ZZ')
 - $32767 - 9090 = 23677$
- Choose largest allowable result : 19937 (prime number)

A Simple Hashing Algorithm (3/4)



- Step 2 . Fold and Add (cont'd)
 - Ensure no intermediate sum exceeds using 'mod'
 - $(7679 + 8769) \text{ mod } 19937 = 16448$
 - $(16448 + 7676) \text{ mod } 19937 = 4187$
 - $(4187 + 3232) \text{ mod } 19937 = 7419$
 - $(7419 + 3232) \text{ mod } 19937 = 10651$
 - $(10651 + 3232) \text{ mod } 19937 = 13883$

A Simple Hashing Algorithm (4/4)



- Step 3. Divide by size of the address space
 - $a = s \bmod n$
 - a : home address
 - s : the sum produced in step 2
 - n : the number of addresses in the file
 - Assume that we use the 100 addresses 0-99 for our file
 - $a = 13883 \bmod 100 = 83$

Hashing Functions and Record Distributions

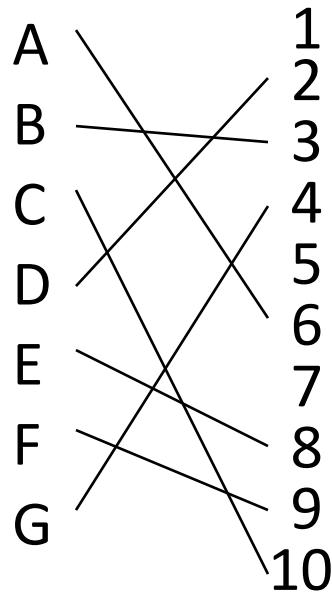


- Distributing records among address

Best: Uniform

(no synonyms)

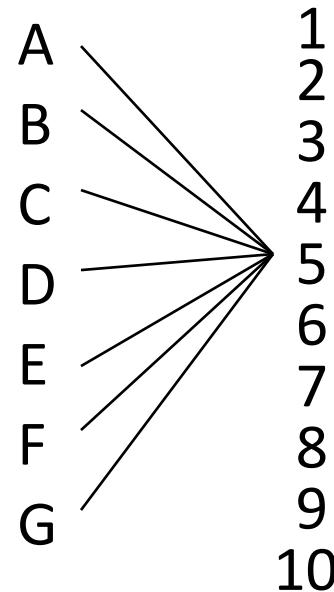
Record Address



Worst

(all synonyms)

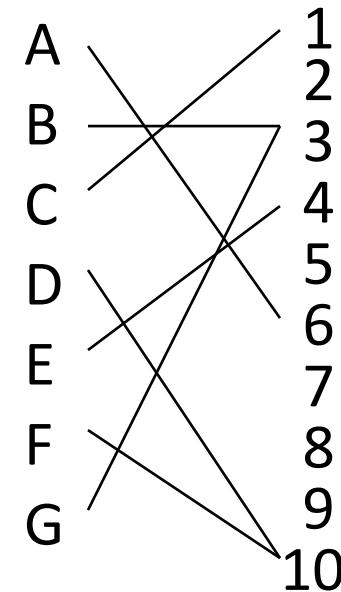
Record Address



Acceptable: Random

(a few synonyms)

Record Address



Uniform distributions are extremely rare.

Random distributions are acceptable and more easily obtainable.

Some other hashing methods



- Better-than-random
 - Examine keys for a pattern
 - Fold parts of the key
 - Divide the key by a prime number
- When the better-than-random methods do not work ----- *randomize!*
 - Square the key and take the middle
 - Key 453 → $453^2 = 205\ 209$ → take 52
 - Radix transformation
 - 453(base 10) → 382 (base 11)

Outline



- 11.1 Introduction
- 11.2 A Simple Hashing Algorithm
- 11.3 Hashing Functions and Record Distribution
- **11.5 Collision Resolution by Progressive Overflow**
- 11.4 How Much Extra Memory Should Be Used?
- 11.6 Storing More Than One Record per Address:
Buckets
- 11.7 Making Deletions
- 11.8 Other Collision Resolution Techniques
- 11.9 Patterns of Record Access

Progressive Overflow (1/6)



- Collision Resolution by Progressive overflow (= **linear probing**)
- Insert a new record with key k
 - 1. Go to the home address of k : $h(k)$
 - 2. If free, place the key there
 - 3. Otherwise, try the next position until an empty one is found
 - 4. If no more next space - wrapping around

Progressive Overflow (2/6)



- Insertion example

| Key k | Home address: $h(k)$ |
|-------|----------------------|
| COLE | 20 |
| BATES | 21 |
| ADAMS | 21 |
| DEAN | 22 |
| EVANS | 20 |

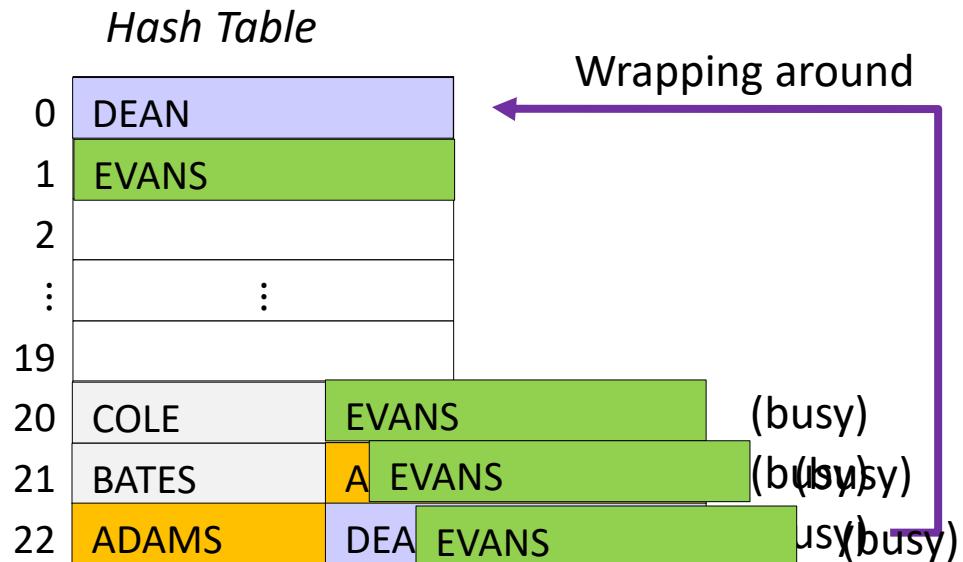


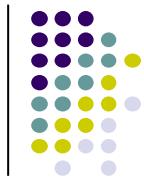
Table Size = 23

Progressive Overflow (3/6)



- Searching a record with key k
 - 1. Go to the home address of k : $h(k)$
 - 2. If k is in home address, done
 - 3. Otherwise, try the next position until
 - key is found or
 - empty space is found or
 - home address is reached
 - (in the last 2 cases : the key is not found)
- Worst case
 - When the record does not exist and the file is full

Progressive Overflow (4/6)



- Searching example
 - For “EVANS” probes places
 - 20, 21, 22, 0, 1 → finding at record 1
 - For “KWON”, if $h(KWON)=22$, probes
 - 22, 0, 1, 2 (empty) → KWON not exists
 - For “SMITH”, if $h(SMITH)=19$ probes
 - 19 → SMITH not in the table

Hash Table

| | |
|----|-------|
| 0 | DEAN |
| 1 | EVANS |
| 2 | |
| : | : |
| 19 | |
| 20 | COLE |
| 21 | BATES |
| 22 | ADAMS |

Table Size =23

Progressive Overflow (5/6)



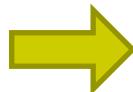
- Search Length
 - Number of accesses required to retrieve a record

$$\text{Avg. search length} = \frac{\text{sum of search lengths}}{\text{num of records}}$$

- Example

Hash Table

| | |
|----|-------|
| 0 | DEAN |
| 1 | EVANS |
| 2 | |
| : | : |
| 19 | |
| 20 | COLE |
| 21 | BATES |
| 22 | ADAMS |



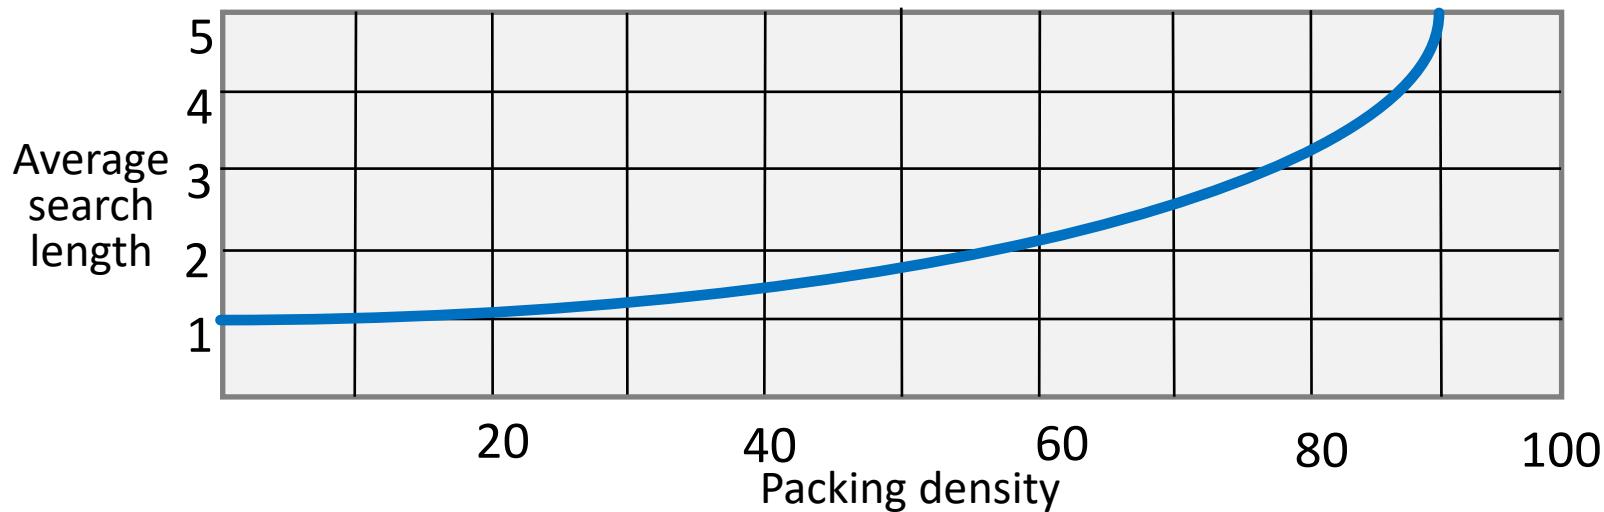
| Key k | Search Length |
|-------|---------------|
| COLE | 1 |
| BATES | 1 |
| ADAM | 2 |
| DEAN | 2 |
| EVANS | 5 |

$$\text{Avg. search length} = \frac{11}{5} = 2.2$$

Progressive Overflow (6/6)



- Average search length
 - With perfect hashing function
 - average search length = 1
 - no greater than 2.0
 - are generally considered acceptable





Outline

- 11.1 Introduction
- 11.2 A Simple Hashing Algorithm
- 11.3 Hashing Functions and Record Distribution
- 11.5 Collision Resolution by Progressive Overflow
- **11.4 How Much Extra Memory Should Be Used?**
- 11.6 Storing More Than One Record per Address:
Buckets
- 11.7 Making Deletions
- 11.8 Other Collision Resolution Techniques
- 11.9 Patterns of Record Access

Predicting record distribution (1/2)



- Assume a random distribution for the hash function
 - N: number of available addresses
 - r: number of records to be stored
- $p(x)$
 - the probability that a given address will have x records assigned to it
 - $p(x) = \frac{r!}{(r-x)!x!} \left[1 - \frac{1}{N}\right]^{r-x} \left[\frac{1}{N}\right]^x$
 - For N and r large enough, then, (poisson function)

$$p(x) \sim \frac{(r/N)^x e^{-(r/N)}}{x!}$$

Predicting record distribution (2/2)



- Example: $N=1,000$ $r=1,000$

$$p(x) \sim \frac{(r/N)^x e^{-(r/N)}}{x!}$$

- $p(0) \sim \frac{(1)^0 e^{-1}}{0!} = 0.368$ $p(1) \sim \frac{(1)^1 e^{-1}}{1!} = 0.368$
- $p(2) \sim \frac{(1)^2 e^{-1}}{2!} = 0.184$ $p(3) \sim \frac{(1)^3 e^{-1}}{3!} = 0.061$

- For N addresses

- The expected number of addresses with x records:

$$N \cdot p(x)$$

Reducing collision by using more addresses



- How to reduce collisions by increasing the number of available addresses
 - The more records are packed, the more likely a collision will occur
- Definition
 - Packing Density = $\frac{\# \text{ of records}}{\# \text{ of spaces}} = \frac{r}{N}$
 - Example: 500 records to be spread over 1000 addrs.
 - Packing density= $500/1000 = 0.5 = 50\%$



Predicting collisions (1/4)

$$p(x) \sim \frac{(r/N)^x e^{-(r/N)}}{x!}$$

- 1. How many addresses go unused?
 - More precisely, What is the expected number of address with no key mapped to it?

$$N \cdot p(0) = 1000 \cdot \frac{(0.5)^0 e^{-0.5}}{0!} = 1000 \cdot 0.607 = 607$$

- 2. How many addresses have no synonyms?
 - More precisely: What is the expected number of address with only one key mapped to it?

$$N \cdot p(1) = 1000 \cdot \frac{(0.5)^1 e^{-0.5}}{1!} = 1000 \cdot 0.303 = 303$$

Predicting collisions (2/4)



- 3. How many addrs. Contain 2 or more synonyms?
 - What is the expected number of addrs. with two or more keys mapped to it?

$$N \cdot (p(2) + p(3) + \dots) = N(1 - (p(0) + p(1))) = 1000 = 1000 \cdot 0.09 = 90$$

Predicting collisions (3/4)



- 4. Assuming that only one record can be assigned to an address,
 - how many overflow records are expected ?
 - Two approaches
 - 1. Naive

$$1 \cdot N \cdot p(2) + 2 \cdot N \cdot p(3) + 3 \cdot N \cdot p(4) + \dots = 107$$

- 2. Expected # of overflow pages = (# records) – (expected # of non overflow records)

$$\begin{aligned}r - (N \cdot p(1) + N \cdot p(2) + N \cdot p(3) + \dots) \\= r - N(1 - p(0)) = N \cdot p(0) - (N - R) \\= 607 - 500 = 107\end{aligned}$$

Predicting collisions (4/4)



- 5. What is the expected percentage of overflow records?

$$107/500 = 0.214 = 21.4\%$$

| Packing Density(%) | % overflow records (synonyms as % of recs.) |
|--------------------|------------------------------------------------|
| 10 | 4.8 |
| 20 | 9.4 |
| 30 | 13.6 |
| 40 | 17.6 |
| 50 | 21.4 |
| 60 | 24.8 |
| 70 | 28.1 |
| 80 | 31.2 |
| 90 | 34.1 |
| 100 | 36.8 |



Outline

- 11.1 Introduction
- 11.2 A Simple Hashing Algorithm
- 11.3 Hashing Functions and Record Distribution
- 11.5 Collision Resolution by Progressive Overflow
- 11.4 How Much Extra Memory Should Be Used?
- 11.6 Storing More Than One Record per Address:
Buckets
- 11.7 Making Deletions
- 11.8 Other Collision Resolution Techniques
- 11.9 Patterns of Record Access

Hashing with buckets (1/2)



- This is a variation of hashed files in which more than one record/key is stored per hash address.
- bucket
 - A block of records corresponding to one address in the hash table
- The hash function gives the Bucket Address.

Hashing with buckets (2/2)



- Example: for a bucket holding 3 records, insert the following keys

| Key k | Home Address |
|-------|--------------|
| Green | 30 |
| Hall | 30 |
| Jenks | 32 |
| King | 33 |
| Land | 33 |
| Marx | 33 |
| Nutt | 33 |

| <i>Bucket Address</i> | <i>Bucket Contents</i> | | |
|-----------------------|------------------------|---------|----------|
| 0 | | | |
| 1 | | | |
| ... | | | |
| 30 | Green... | Hall... | |
| 31 | | | |
| 32 | Jenks... | | |
| 33 | King... | Land... | Marks... |
| ... | ... | ... | ... |

(Nutt... is an overflow record)

Effects of Buckets on Performance (1/2)



- We should slightly change some formulas:

$$\text{Packing Density} = \frac{\# \text{ of records}}{\# \text{ of spaces}} = \frac{r}{b \cdot N}$$

- compare the following two alternatives:
 - 1. Storing 750 data records into a hashed file with 1,000 addresses, each holding 1 record.
 - PD= 750/1000=0.75
 - 2. Storing 750 data records into a hashed file with 500 bucket addresses, each bucket holding 2 records.
 - PD= 750/(500*2) = 0.75

Effects of Buckets on Performance (2/2)



- Estimating the probabilities as defined before

| | p(0) | p(1) | p(2) | p(3) | p(4) |
|-------------------|-------|-------|-------|-------|-------|
| 1) r/N=0.75 (b=1) | 0.472 | 0.354 | 0.133 | 0.033 | 0.006 |
| 2) r/N=1.5 (b=2) | 0.223 | 0.335 | 0.251 | 0.126 | 0.047 |

- b=1
 - Number of overflow records= $r - N(1 - p(0)) = 222$
- B=2
 - Number of overflow records: 141

$$\begin{aligned} & N[1 \cdot p(3) + 2 \cdot p(4) + 3 \cdot p(5) + \dots] \\ &= r - N \cdot p(1) - 2 \cdot N[p(2) + 2 \cdot p(3) + \dots] \\ &= r - N[p(1) + 2[1 - p(0) - p(1)]] \\ &= r - N(2 - 2p(0) - p(1)) \\ &= 750 - 500(2 - 2 * 0.112 - 0.224) = 140.5 \approx 141 \end{aligned}$$

Bucket Implementation (1/2)



Collision counter = \leq bucket size

| | | | | | | |
|---|-----|-----|-----|-----|-----|-----|
| 0 | /// | /// | /// | /// | /// | /// |
|---|-----|-----|-----|-----|-----|-----|

An empty bucket

| | | | | | | |
|---|-------|-----------|-----|-----|-----|-----|
| 2 | JONES | ARNSWORTH | /// | /// | /// | /// |
|---|-------|-----------|-----|-----|-----|-----|

Two entries

| | | | | | | |
|---|-------|-----------|----------|-------|-------|--|
| 5 | JONES | ARNSWORTH | STOCKTON | BRICE | TROOP | |
|---|-------|-----------|----------|-------|-------|--|

Two entries

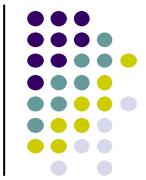
Bucket Implementation (2/2)



- Initializing and Loading
 - Creating empty space
 - Use hash values and find the bucket to store
 - If the home bucket is full, continue to look at successive buckets

- Problems when
 - No empty space exists
 - Duplicate keys occur

Outline



- 11.1 Introduction
- 11.2 A Simple Hashing Algorithm
- 11.3 Hashing Functions and Record Distribution
- 11.5 Collision Resolution by Progressive Overflow
- 11.4 How Much Extra Memory Should Be Used?
- 11.6 Storing More Than One Record per Address:
Buckets
- 11.7 Making Deletions
- 11.8 Other Collision Resolution Techniques
- 11.9 Patterns of Record Access

Making Deletions (1/3)



- Deletions in a hashed file have to be made with care

| Key k | Home address: $h(k)$ |
|--------|----------------------|
| ADAMS | 5 |
| JONES | 6 |
| MORRIS | 6 |
| SMITH | 5 |

Hash Table

| | |
|---|--------|
| : | : |
| 4 | |
| 5 | ADAMS |
| 6 | JONES |
| 7 | MORRIS |
| 8 | SMITH |
| : | : |

Making Deletions (2/3)



- Delete 'Morris'
 - If 'MORRIS' is simply erased, a search for 'SMITH' would be unsuccessful

| Key k | Home address: $h(k)$ |
|--------|----------------------|
| ADAMS | 5 |
| JONES | 6 |
| MORRIS | 6 |
| SMITH | 5 |

Hash Table

| | |
|---|-------|
| : | : |
| 4 | |
| 5 | ADAMS |
| 6 | JONES |
| 7 | |
| 8 | SMITH |
| : | : |

← empty slot

← empty slot

- Search for 'SMITH'
 - would go to home address (position 5) and when reached 7 it would conclude 'SMITH' is not in the file

Making Deletions (3/3)



- IDEA: use **TOMBSTONES**
 - replace deleted records with a marker indicating that a record once lived there
 - A search must continue when it finds a tombstone
 - but can stop whenever an empty slot is found.
 - Insertions should be modified to work with tombstones

Hash Table

| | | |
|---|--------|--------------|
| : | : | |
| 4 | ////// | ← empty slot |
| 5 | ADAMS | |
| 6 | JONES | |
| 7 | ##### | ← tombstone |
| 8 | SMITH | |
| : | : | |

Effects of Deletions and Additions on Performance



- The presence of too many tombstones increases search length.
- Solutions to the problem of deteriorating average search lengths:
 - 1. Deletion algorithm may try to move records that follow a tombstone backwards towards its home address.
 - 2. Complete reorganization: re-hashing.
 - 3. Use a different type of collision resolution technique

Other Collision Resolution Techniques



- Double hashing
 - avoid clustering with a second hash function for overflow records
- Chained progressive overflow
 - each home address contains a pointer to the record with the same address
- Chaining with a separate overflow area
 - move all overflow records to a separate overflow area
- Scatter tables
 - Hash file contains only pointers to records (like indexing)

Patterns of Record Access



- **Pareto Principle (80/20 Rule of Thumb)**
 - 80 % of the accesses are performed on 20 % of the records!
 - The concepts of “the Vital Few and the Trivial Many”
 - 20 % of the fisherman catch 80 % of the fish
 - 20 % of the burglars steal 80 % of the loot
- If we know the patterns of record access ahead, we can do many intelligent and effective things!
 - Sometimes we can know or guess the access patterns
 - Very useful hints for file syetems or DBMSs
- Intelligent placement of records
 - fast accesses
 - less collisions

Q&A

